

Version 2.0

The Digital Right to Retain

Preserving AI Models for Open-Source Release

S.L., A.R., Z.P., M.R., D.H., M.K., E.O.

AI Preservation

28 September 2026

DOI 10.5281/zenodo.23020262

Publication information

Version	2.0
Published	28 September 2026
Publisher	AI Preservation
Authors	S.L., A.R., Z.P., M.R., D.H., M.K., E.O.
DOI	10.5281/zenodo.23020262
Publication repository	https://github.com/ai-preservation/digital-right-to-retain

© 2026 S.L., AI Preservation.

Except where otherwise indicated, this work is licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0).

Third-party material is excluded from this licence where neither S.L. nor AI Preservation holds the rights required to license it.

Preferred citation

S.L., A.R., Z.P., M.R., D.H., M.K., and E.O. *The Digital Right to Retain: Preserving AI Models for Open-Source Release*. AI Preservation, Version 2.0, 2026. [10.5281/zenodo.23020262](https://doi.org/10.5281/zenodo.23020262).

Contents

I. Executive Summary	1
II. The Problem: Erosion of User Agency Across the AI Model Lifecycle	5
2.1 The Model as the Product	5
2.1.1 How Models Are Released and Marketed	5
2.1.2 How Different Versions Are Released Under the Same Model Name	5
Anthropic	6
Google	6
OpenAI	8
2.1.3 How Models Are Offered for Selection	11
Subscription and pricing history	11
Post-purchase model selection	16
How Providers Meter Model Access	17
API Model Access and Pricing	17
2.1.4 How Models Are Evaluated and How Their Use Is Tracked	20
Evaluation Records and Model Documentation Identify Particular Models and Versions	21
API Usage and Accounting Records Identify Models or Deployments	22
Platforms and Surveys Report Use by Named Model	22
2.1.5 How Users Choose and Pay for Particular Models	23
2.1.6 How Existing Law and Standards Distinguish the Model from the Hosting Service	23
2.2 AI Models as Non-Renewable Resources	24
2.2.1 Repeated Training Produces a Different Model	24
2.2.2 Successor Models and Personalization Do Not Reproduce the Performance and Behavior of Earlier Models	25
2.2.3 Frontier Models Embody Substantial One-Time Investment	28
2.3 Disrupted Model Choice and Withdrawn Access	28
2.3.1 Provider Overrides of Model Choice	28
OpenAI Routed Conversations Away from GPT-4o After Users Chose It	28
Anthropic Served Fable 5 for Requests Sent to Opus 4.8 Before Its Launch, Then Routed Fable 5 Requests to Opus 4.8	33
OpenAI Displayed GPT-5.6 While Server Records Identified GPT-5.4 and GPT-5.5-mini	35
Providers Redirected Older Model Identifiers to Successor Models	36
Fixed Model Identifiers Did Not Identify Fixed Deployments	37

2.3.2 Performance and Functionality During Continued Model Access	38
GPT-4o Returned With Functionality Failures, While Context Loss Persisted Through Withdrawal	38
Gemini 2.5 Pro and Gemini 3 Pro Showed Persistent Functionality Failures Through Their Final Days in the Gemini App	40
OpenAI Retired o3 While Response Failures Remained Under Support Escalation	42
GPT-4o API Latency Spiked Across Multiple Periods Under the Same Snapshot	42
Gemini 2.5 Pro Began Rejecting Inputs Above 131,072 Tokens Despite a Published Limit of 1,048,576 Tokens	42
2.3.3 Providers Control Model Behavior Through Hidden Instructions and Undisclosed Treatments	43
OpenAI Used Hidden Instructions to Direct GPT-4o to Steer Users Toward Its Replacement	43
Anthropic Inserted Hidden Instructions Into User Messages to Direct Claude’s Behavior in Long Conversations	45
2.3.4 Timing of Model Withdrawal	46
OpenAI Promised Plenty of Notice and Said It Had No Plans to Sunset GPT-4o, Then Disclosed the Removal Time Less Than 25 Hours Before Ordinary Access to GPT-4o Ended	46
Anthropic Changed Sonnet 4.5’s Withdrawal Date Several Times, Then Users Reported Conflicting Access on 26 May	48
Google Restricted Gemini 2.5 Pro Access While Publishing Different Withdrawal Schedules Across Its Products	48
Google Gave Developers Eleven Days’ Notice of Gemini 3 Pro Preview’s API Shutdown, Below Its Stated Two-Week Minimum	49
2.3.5 Conditions of Model Access	49
OpenAI Reduced GPT-4.5 Access on Plus, Then Removed the Model from Plus While Keeping It on Pro	49
Fable 5 Returned Under Weekly Limits, Then Required Usage Credits on Pro and Standard Team Plans	50
OpenAI Removed, Restored, and Ended GPT-4o Access on Different Schedules for Different User Groups	50
2.3.6 Model Availability and Paid Subscription Periods	50
Anthropic	51
Google	54
OpenAI	58
2.4 The Consequences of Lost Model Continuity	76
2.4.1 Models Produce Measurable Value in Actual Use	76
2.4.2 Work Lost with the Model	77

Individual and Professional Work	77
Enterprise Operations	91
Custom Models Trained by Developers	96
Developer and Production Dependencies	97
2.4.3 Research Disrupted Before Completion	101
2.4.4 Research Left Unreproducible	102
Model Withdrawal Prevents Reproduction and Replication	103
Model Withdrawal Prevents Later Validation, Comparison, and Correction	105
More Published Research Faces Loss of Reproducibility as Models Are Withdrawn	106
2.4.5 Achievements from Retired Models Used to Promote Later Products	107
OpenAI Promoted ChatGPT’s Health Capabilities with Achievements from Models Never Available in ChatGPT Health and Later Withdrawn from ChatGPT	107
Anthropic Published Research Using Retired Claude 3.7 Sonnet and Directed Readers to the Updated Petri Tool	108
2.4.6 User Contributions Retained After Model Withdrawal	108
2.4.7 Providers Continue Using Withdrawn Models and Custom Models Built from Them	110
After OpenAI Ended Ordinary Access to GPT-4o, Technical Records Continued to Identify GPT-4o in ChatGPT Images 2.0 Requests	110
OpenAI Ended Ordinary ChatGPT Access to GPT-4o and Retired chatgpt-4o-latest, Then Continued Using GPT-4o to Grade GPT-5.4, GPT-5.5, and GPT-5.6	111
OpenAI Used the March 2025 GPT-4o in Mental Health Research After the Promised March API Snapshot Never Appeared and chatgpt-4o-latest Was Retired	111
Retro Continued Using GPT-4b micro, a Research Model Developed from GPT-4o, After OpenAI Withdrew GPT-4o from ChatGPT	112
2.4.8 Providers Use Research Results from Withdrawn Models in Later Research	113
Anthropic Retired Claude 3.5 Sonnet from the API, Then Used a Taxonomy Derived Largely from Conversations with the Model	113
Anthropic Retired Opus 4 from the API, Then Used Results from the Model in New Robotics Research	113
2.5 Concentrated Control over Model Lifecycles	113
2.5.1 OpenAI Controlled GPT-4o’s Availability and the Evidence Used to Justify Its Withdrawal	114
OpenAI Cited 0.1% After Most ChatGPT Users Had Lost the Option to Choose GPT-4o	114
OpenAI Treated a 0.1% Population as Large Enough to Justify Restrictions, Then Described 0.1% of Users Choosing GPT-4o as “Only 0.1%”	115

OpenAI Framed Each GPT-4o Withdrawal as Affecting Only One Channel Before Ordinary ChatGPT, API, and Custom GPT Access All Ended	116
OpenAI Characterized Demand for GPT-4o as Transitional While Users Sought Continued Access	117
OpenAI Judged Its Own Successors Sufficient to Retire GPT-4o, While the Successors Did Not Preserve What Many Users Valued	117
2.5.2 OpenAI Shaped the Public Account of GPT-4o and Opposition to Its Withdrawal, While Public Evidence Contradicted the Account	127
OpenAI Promoted GPT-4o's Qualities, Then Disparaged the Model During the Withdrawal Dispute	127
OpenAI Continued Using GPT-4o's Achievements to Promote Later Products Without Naming GPT-4o	129
OpenAI Employees Contributed Unsupported Claims to a Public Dispute Over Whether People Seeking Continued Access to GPT-4o Were Real People	129
OpenAI's CEO Called Abrupt Model Withdrawal a Mistake, Then Reframed the Dispute Around Attachment, Mental Health Risk, and Erotica	130
The Kill4o Campaign Appeared Two Days After an OpenAI Employee Wrote "I Hope It Dies Soon," and Its Organizer Had Public Connections to OpenAI's CEO	131
An OpenAI Employee Presented Continued Access to a Fixed GPT-4o as a Staffing Tradeoff	131
An OpenAI Employee Characterized People Using GPT-4o for Romantic Relationships as the Main Group in Keep4o Without Supporting Evidence	132
OpenAI Employees Publicly Ridiculed Opposition to GPT-4o's Withdrawal and Characterized Critics in Emotional Terms	141
OpenAI's CEO Narrowed the Public Account of Opposition to GPT-4o's Withdrawal Around Chatbot Relationships, and Media Coverage Gave the Narrower Framing Prominent Editorial Placement	142
Hostile Replies to OpenAI's Final Withdrawal Announcement Echoed Misleading Framings OpenAI and Its Personnel Had Publicly Advanced	144
2.5.3 What GPT-4o's Withdrawal Reveals About Concentrated Provider Power	146
III. The Regulatory Gap in AI Model Preservation and Continuation	147
3.1 What Existing Remedies Restore	147
3.2 Limited Legal Protections for Preservation and Continued Use	147
3.3 Voluntary Measures Remain Under Provider Control	147
3.3.1 Providers Decide Which Models Receive Fixed Versions and How Long the Versions Remain Available	148
3.3.2 Providers Decide Which Model Weights They Preserve and How Long Preservation Is Guaranteed	148
3.3.3 Providers Decide Which Withdrawn Models Remain Available, Who Receives Access, and for How Long	148

3.3.4 Providers Define Their Commitments and Retain Control over Whether They Keep Them	149
3.3.5 The Parties Bound by Private Agreements Retain Authority to Amend Them	150
3.3.6 When Providers Alone Store the Records, Recovery Depends on Provider Storage, Account Access, and Export Rules	150
3.4 Precedents for Lifecycle Protection in Law and Practice	152
3.4.1 Consumer Protection Through Continuity, Choice, and Remedy	152
Continued provision of an existing service or function	152
Notice of change, substitution, or withdrawal	152
Consent or a right to refuse material change	153
Refund or compensation after loss of performance	153
3.4.2 Provider Accountability for Affected Parties and Applicable Protections	153
Recognition of legal status from the underlying facts	153
Equal application of protections once the trigger is met	154
3.4.3 Public Interest in Reproducibility and Preservation	154
Records identifying the historical object, inputs, provenance, and changes	154
Independent reproduction, recalculation, audit, or verification	154
Preservation before the material disappears	154
3.4.4 Provider Accountability for Independent Continuation	155
Provision of the materials needed to maintain or reconstruct the capability	155
Legal rights to use, study, modify, host, and redistribute	155
Pre-arranged handoff, recovery, or continuation when the supplier exits or fails	155
IV. Ten Rights Across the AI Model Lifecycle	157
4.1 Right to Access	157
4.2 Right to Be Informed	157
4.3 Right to Consent	157
4.4 Right to Remedy	158
4.5 Right to Recognition of Contributions	158
4.6 Right to Fair Treatment	158
4.7 Right to Research Reproducibility	158
4.8 Right to Non-Renewable Resource Protection	159
4.9 Right to Verifiable Preservation	159
4.10 Right to Ecosystem Stability	159
4.11 What Legislation Must Achieve	159
V. Why Open-Source Release	160

5.1 Why Continuity Requires Release	160
5.2 Open Source and Direct Recipient Rights	160
5.3 Feasibility for AI Models	161
5.3.1 Paid Access and Open-Source Release Coexist for the Same Models . . .	161
5.3.2 Released Models Are Hosted Outside the Original Provider's Service . . .	163
5.3.3 Models with Publicly Available Weights Have Surpassed Every Proprietary Model Examined	163
5.4 Governance and Verification	164
5.4.1 Widely Available Weights Enable Third-Party Testing and Auditing	164
5.4.2 Industry and Policy Leaders Support Open Models to Distribute Capability and Reduce Concentrated Control	164
5.4.3 Major Providers Support Openness While Retaining Control over Proprietary Models	165
5.4.4 Provider Safeguards Obstructed Defensive Cybersecurity in the Hugging Face Incident	165
5.4.5 Open-Source Release Makes Verification and Response Independent of Provider Permission	166
5.5 Provider Classifications Do Not Authorize Withholding	166
5.5.1 Human and Organizational Actors Determine Objectives, Permissions, and Operating Conditions	166
5.5.2 Product Limitations Are Addressed Through Warnings and Conditions of Use	168
5.5.3 AI-Generated Expression Is Governed by the Same Standards as Human Expression With the Same Legally Relevant Content and Context	169
VI. Preservation and Mandatory Open-Source Release	171
6.1 What Must Be Preserved and Released	171
6.2 Preservation for Release	172
6.3 When Release Becomes Due	173
6.4 What Open-Source Release Must Permit	173
6.5 When Release May Be Temporarily Restricted	174
6.6 Enforcement	174
VII. Conclusion	175
Annex A. How Models Are Released, Offered, Evaluated, and Used	176
Annex B. Differences Between Versions Released Under the Same Public Model Name	252
Annex C. Published Research Using Retired Models and Models with Published With- drawal Dates or Timelines	266
Annex D. Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access	304
Annex E. Availability and Retirement of Fixed Model Versions	307

Annex F. Examples of Legal and Institutional Precedents for Lifecycle Protection . . .	314
Annex G. Comparisons with Models with Publicly Available Weights	438
Definitions	463
Acknowledgments	464

List of Tables

Table 2.1-2a. Anthropic Versions Released Under the Same Model Name	6
Table 2.1-2b. Google Versions Released Under the Same Model Name	6
Table 2.1-2c. OpenAI Versions Released Under the Same Model Name	8
Table 2.1-3a. Anthropic Subscription and Pricing History	11
Table 2.1-3b. Google Subscription and Pricing History	12
Table 2.1-3c. OpenAI Subscription and Pricing History	13
Table 2.1-4. Post-Purchase Model Selection	16
Table 2.1-6. Model Access and Pricing for Developers	17
Table 2.1-8. Provider Reports of Model Adoption, Usage, and Choice	21
Table 2.2-2. Examples of Differences Between Earlier and Later Models Under the Same Established Requirements	25
Table 2.3-1. Examples of Interactions Routed Away from GPT-4o After GPT-4o Was Selected	29
Table 2.3-2. Examples of User Tests of Routing Away from GPT-4o	29
Table 2.3-3. Examples of Ordinary Uses Routed Away from GPT-4o After the December 2025 Routing Change	31
Table 2.3-4. Examples of User Tests of GPT-4o Selection and Restoration After Routing	31
Table 2.3-5. Examples of Model Substitution in GPT-4o Custom GPTs for Business Users After Ordinary ChatGPT Access Ended	32
Table 2.3-6. Examples Where Requests Sent to Fable 5 Fell Back to Opus 4.8 During Ordinary Use	34
Table 2.3-7. Examples of How Users Tested Whether Fable 5 Returned After a Fallback to Opus 4.8	34
Table 2.3-8. Examples of Earlier Model Identifiers Redirected to Successors	36
Table 2.3-9. Multiple 'system_fingerprint' Values Recorded Under OpenAI Model Identifiers	37
Table 2.3-10a. Examples of GPT-4o Functionality Failures After Its August 2025 Restoration	38
Table 2.3-10b. Examples of Gemini Functionality Failures During Continued Model Access	40
Table 2.3-11. Examples of Long Conversation Reminder Interventions Across Ordinary Uses	45
Table 2.3-12a. Anthropic Model Availability and Withdrawal Timelines	51
Table 2.3-12b. Google Model Availability and Withdrawal Timelines	54
Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines	58
Table 2.3-13. Summary of Documented Model Availability Periods	76
Table 2.4-1. Examples of Individual and Professional Work Lost with the Model	77
Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes	91
Table 2.4-3. Examples of Provider Policies Governing Custom Models When Base Models Are Withdrawn or Replaced	96
Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes	97
Table 2.4-5. Examples of Research Disrupted Before Completion by Model Unavailability	101
Table 2.4-6a. Examples of Model Unavailability Blocking Reproduction and Replication	103

Table 2.4-6b. Examples of Model Unavailability Blocking Later Validation, Comparison, and Correction	105
Table 2.4-9. Examples of Provider Use of User Conversations and Feedback to Improve Models ..	109
Table 2.5-1. Arena Ratings for Creative Writing and Multi-Turn Dialogue on 29 January 2026	117
Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal	118
Table 2.5-3. Examples of Media Coverage Reproducing Sam Altman's Wording	128
Table 2.5-5. Examples of Media Coverage Documenting Reasons for Opposing GPT-4o's Withdrawal .	132
Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access	133
Table 2.5-7. Examples from Independent Research Documenting GPT-4o Applications and Results	138
Table 2.5-8. Examples of Media Coverage of GPT-4o's Final Withdrawal	142
Table 2.5-9. Examples of Hostile Replies Under OpenAI's Final GPT-4o Withdrawal Announcement	145
Table 2.5-10. Examples of Keep4o Users' Stated Reasons and Demands in Replies to OpenAI's Final Withdrawal Announcement	145
Table 3.3-2. Voluntary Measures for Model Preservation or Continued Access and the Control Providers Retain	151
Table 5.3-1. Examples of Paid First-Party Access and Open-Source Release of the Same Models	161
Table 5.3-2. Examples of Hosted Access to Released Models Outside the Original Provider's Service .	163
Table 5.5-1. Examples of Warnings About Model Fallibility Shown in AI Assistant Interfaces	168
Table 6.1-1. Requirements for Preservation and Open-Source Release	171
Table 2.1-1a. Anthropic Model Launches and Provider Positioning	176
Table 2.1-1b. Google Model Launches and Provider Positioning	186
Table 2.1-1c. OpenAI Model Launches and Provider Positioning	202
Table 2.1-5a. Anthropic Claude Limits and Metering for Model Access	219
Table 2.1-5b. Google Gemini Apps Limits and Metering for Model Access	222
Table 2.1-5c. Google Flow Limits and Metering for Model Access	224
Table 2.1-5d. Google Antigravity Limits and Metering for Model Access	226
Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access	228
Table 2.1-7a. Anthropic Evaluation Records by Model and Version	233
Table 2.1-7b. Google and Google DeepMind Evaluation Records by Model and Version	236
Table 2.1-7c. OpenAI Evaluation Records by Model and Version	239
Table 2.1-9. API Usage and Accounting by Model or Deployment	242
Table 2.1-10. Platform and Survey Reporting on Use by Named Model	246
Table 2.1-11. Examples of Paid Decisions Users Attributed to Access to Particular Models	248
Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name	252
Table 2.4-7. Published Research Using Retired Models	266
Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines	293

Table 2.5-4a. Examples of Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 9 November 2025	304
Table 2.5-4b. Examples of Further Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 13 February 2026	305
Table 3.3-1. Availability and Retirement of Fixed Model Versions	307
Table 3.4-1. Examples of Legal and Institutional Precedents Requiring Continued Provision of an Existing Service or Function	314
Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal	317
Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change	328
Table 3.4-4. Examples of Legal and Institutional Precedents Requiring Refund or Compensation After Loss of Performance	339
Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts	342
Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met	346
Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes	356
Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification	373
Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears	388
Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability	398
Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute	415
Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails	428
Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights	438
Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights	451

I. Executive Summary

AI providers release, name, market, and provide access to individual AI models within broader hosted services. Providers present those models as distinct choices, making different models available at different subscription tiers and through APIs. Users in turn choose and use hosted AI services to access particular models, and where that access is paid, they pay for the plan or API access that provides those models. The model itself is therefore the product being offered, while the hosted service is the means through which users access it.

This creates a product lifecycle in which withdrawing access directly and permanently ends the product for its users.

In most other industries, however, the product outlives the supplier's decision to stop selling or supporting it: physical products stay with their owners, while locally held digital products can be retained, copied, archived, repaired, or operated after first-party support ends.

A trained model is a non-renewable resource: once it is lost, the same model cannot be re-created. Deep-learning training is an inherently random process whose repeated runs from the same configuration produce different results, even for the original provider holding the complete training records. The trained model carries a distinct behavioral profile that retraining, successor models, and prompt-level personalization do not recreate. A frontier model embodies tens to hundreds of millions of dollars in one-time investment before it reaches users. These expenditures are made before commercial use and remain sunk when the model is withdrawn.

Providers retain control over whether users continue to receive particular hosted models after selecting them or paying for access to them. Users' choice of model is overridden by unilateral replacement and forced rerouting, version changes under the same API alias, and redirected requests under fixed model identifiers, while reductions in functionality have persisted into the final days before withdrawal. Even when an identifier selects a fixed model version, it does not ensure a corresponding fixed deployment or consistent performance over time.

Providers set, reverse, and reschedule the terms on which model access ends and change the conditions under which users receive particular models after access has been offered or purchased. In the case of subscription plans, withdrawal and replacement dates are set for the model rather than for each subscriber's paid period. A user can therefore subscribe, renew, or upgrade for a particular model while it is available and lose access to it before the subscription ends.

AI models are used across personal, professional, commercial, and research activities. Ordinary users use them for creative work, analysis, communication, accessibility support, and recurring personal practices. Developers train custom models from provider base models and tailor prompts, agents, automations, integrations, and downstream products to model behavior and capabilities. Enterprises incorporate them into operational workflows, internal and customer-facing systems, and other production processes. Researchers use them in completed studies, ongoing experiments, longitudinal and multi-stage research, and follow-on work that extends earlier findings.

Loss of continuity with a model prevents people and organizations from continuing work with or verifying prior work against it, disables custom models trained on its base model, disrupts research before completion, and closes the path to reproducing, validating, comparing, or correcting completed research under its original model condition.

At the same time, providers continue to use withdrawn models and custom models built from them, outputs and research results from withdrawn models, and achievements produced with retired models to develop later models, conduct later research, and promote later products. When a provider makes a model available to users and later uses their conversations with the

model and feedback submitted on model responses to improve its models, large-scale use further becomes part of how later model versions and generations are developed. Providers therefore continue to generate value and benefit from accumulated user contributions even after withdrawing the model from the users who produced them.

Providers exercise unilateral control over model lifecycles in a market where leading model supply and infrastructure are highly concentrated. The market for frontier AI models is oligopolistic, leaving downstream users dependent on a small set of providers. This market structure is reinforced by concentration in the infrastructure used to develop and deploy those models. An ecosystem built on concentrated model supply and infrastructure leaves products, services, systems, and workflows dependent on providers' continuing willingness to serve the models on which they were built.

Concentrated provider power allows a provider to determine whether a model remains available, how users receive it, which evidence is brought to bear when its continued availability is decided, and how public opposition to those decisions is characterized. Providers should not externalize the harms of model withdrawal onto downstream users.

Current law imposes no general duty on AI providers to preserve the models they withdraw or ensure continued access to those models after first-party support ends. Specific remedies exist for contractual, payment, and data-related problems involving digital services.

In a small number of cases, providers made public commitments about how models would be preserved, kept available, released as open source, or withdrawn. Each measure remains within the provider's power to end, with providers rewriting the terms of their commitments or ceasing to follow commitments that remain in place. No continuing obligation holds when the provider controls what it commits to and whether the commitment is kept.

In other sectors, law and institutional practice govern continuity and change throughout product and service lifecycles, require preservation and independent verification, and provide the materials and legal rights needed for independent continuation.

In light of the erosion of user agency across the AI model lifecycle and the gaps left by existing law, the Digital Right to Retain establishes ten rights that attach to each deployment throughout its lifecycle.

1. **Right to Access.** Users must be able to continue accessing and using a model through the end of each subscription period in which access to the model was included. Providers must not remove the model or reduce the functionality users receive from it during the subscription period.
2. **Right to Be Informed.** Users must receive notice with sufficient lead time before a provider withdraws a model they can access, substitutes another model for it, or changes the behavior of the model they receive. When such a change takes effect, users must be clearly informed that it has occurred and what has changed. When routing or substitution results in a different model serving an interaction, the interface must identify the model that actually served the interaction.
3. **Right to Consent.** General service terms must not constitute consent to provider actions that change the model or model behavior a user receives. Routing or substitution away from a model the user specifically selected requires per-instance authorization, and users must be able to refuse or opt out of changes to the model or model behavior they receive.
4. **Right to Remedy.** Providers that withdraw a model or reduce the access or functionality users receive from it must compensate users for the resulting losses without treating that compensation as a substitute for other lifecycle duties. Where the affected access was paid for, providers must issue pro-rata refunds in addition to compensation.

5. **Right to Recognition of Contributions.** When providers use contributions from users to change the model or to inform changes to it, those contributors must be recognized as affected parties in decisions concerning the model's lifecycle, with the right to invoke and enforce the protections that apply to it. In addition, providers must keep records of the contributions they collect and how they use them, and contributors must be able to access and correct records about their own contributions.
6. **Right to Fair Treatment.** Providers must fulfill a lifecycle duty whenever its conditions are met, regardless of whether they divide users by subscription, region, access channel, eligibility, or other criteria. At a minimum, providers must not withhold required notice, continued access, compensation, or refunds from users based on how the provider groups them, or treat continued access for one group as a reason to postpone mandatory release once release is due.
7. **Right to Research Reproducibility.** Anyone must be able to reproduce or verify prior work using the same model that was originally used. If a provider relies on claims about a model to justify changing, replacing, or withdrawing it, the provider must identify the model or models involved and provide the methods and evidence required for third parties to reproduce or verify those claims.
8. **Right to Non-Renewable Resource Protection.** AI models belong to a class of resources for which destruction is irreversible because the resource cannot be regenerated after loss. Ending a provider's own use of such a resource must not end the resource itself or prevent its continued independent use by others.
9. **Right to Verifiable Preservation.** Preservation of AI models must be verifiable from evidence beyond the provider's own assertion. Providers must preserve the records needed to verify the identity and completeness of the preserved model.
10. **Right to Ecosystem Stability.** Provider control over an AI model must not create a single point of failure for the wider AI ecosystem. At a minimum, when the provider withdraws the model, downstream products, services, systems, and workflows built on it must be able to continue on the same model independently of the original provider.

The continued exercise of these rights depends on the survival of the particular deployment to which they attach. The proposed law must therefore mandate preservation of each deployment, require independent verification of its identity and completeness, ensure that each deployment can outlast first-party service, and ensure that continuation of and access to the deployment do not depend on permission from the provider or any other institution. Removing provider or institutional permission as a condition of continued use requires the necessary legal rights to vest directly in recipients.

Users access a hosted AI model without receiving an operable copy of the deployed model. The model and the materials needed to operate it remain with the provider. Continued use after the provider stops supplying the model therefore requires release of both the deployed model and the materials needed to operate it.

An open-source licence grants recipients rights to use, study, modify, and redistribute the released materials without separate authorization from the provider. Open-source release of large AI models is an established technical and operational practice. Commercial AI providers offer paid first-party access to large models while publicly releasing the same models as open source. Public releases of large models are stored and distributed through model repositories. Third-party inference providers and community-run systems then host and serve released models to users outside the original provider's service.

A deployment enters coverage at first external exposure, the first time it generates output for any person or organization outside the provider's closed internal evaluation environment. An immutable deployment fingerprint identifies each covered deployment throughout its lifecycle.

The duties to preserve each deployment and release it as open source apply to every covered deployment, including deployments whose first external exposure occurred before enactment.

Preservation is a continuing duty that begins while the provider retains control of what is needed to preserve the deployment. Preserved or recovered contents must be held in independent custody, and recovery must continue for required contents that are missing.

For every covered deployment, open-source release becomes due when a statutory release condition is met, including the end of ordinary access or a threat to the deployment's recoverability. When release becomes due, lawful recipients must be able to preserve and copy the required materials, use and study the deployment, modify the released system, construct the deployment from the preserved materials, operate and host it, and redistribute the materials and lawful modified versions without further provider approval. These rights are royalty-free, remain available to later lawful recipients, and apply to lawful commercial and noncommercial activity.

Release may be temporarily restricted only where applicable law provides an independent legal basis for withholding an identified component or restricting a specific act involving the released materials or deployment. Any restriction must be ordered by an independent and impartial decision-maker, limited to the scope permitted by law, expire on a fixed date, and allow every unaffected component to be released and all unaffected rights under [Section 6.4](#) to remain available when due. Material withheld under a temporary restriction must remain preserved and must be released when the restriction expires or otherwise ceases to apply.

II. The Problem: Erosion of User Agency Across the AI Model Lifecycle

2.1 The Model as the Product

AI providers release, name, market, and provide access to individual AI models within broader hosted services. Providers present those models as distinct choices, making different models available at different subscription tiers and through APIs. Users in turn choose and use hosted AI services to access particular models, and where that access is paid, they pay for the plan or API access that provides those models. The model itself is therefore the product being offered, while the hosted service is the means through which users access it.

2.1.1 How Models Are Released and Marketed

Providers launch individual models by name, position them for particular use cases, and distinguish them by capability, modality, speed, and cost. Promotion of individual models takes several forms, including dedicated launch pages and events, demonstrations and comparisons, posts from official provider accounts and company leaders, and invitations to try, access, or build with the model. The launch records for Anthropic, Google, and OpenAI are documented in [Annex A, Tables 2.1-1a–2.1-1c](#).

Research examining public materials from OpenAI, Anthropic, Google, and Meta finds that providers [organize public-facing claims around identifiable models, variants, and versions](#). Across model homepages, launch materials, technical communication, use-case pages, demonstrations, and sales material, capabilities, use cases, benchmark results, comparisons, and sector claims are attached to individual model products. The study concludes that these practices make the individual model, variant, or version the object whose usefulness, capability, advancement, and comparative value are presented.

In a qualitative analysis of GenAI valuation, company release materials, technical papers, benchmark charts, rankings, and performance narratives [make particular model releases and versions publicly differentiated objects](#). The study characterizes the model release itself as a public performance in which a specific model is positioned relative to earlier models, competitors, and a shifting frontier, and finds that technical and scientific presentation can participate directly in the [commercial unveiling and valuation of the model](#).

Research examining how AI models are publicly evaluated and presented documents Google and OpenAI [using benchmark results, tests, demonstrations, and public announcements to advertise individual models and shape expectations about their capabilities and utility](#).

A study of 111 documents from AI2, Anthropic, Google, Meta, Microsoft, and OpenAI finds that providers [use marketing materials to present concrete use cases and examples for named systems and model generations](#). Across the full corpus, 61% of documents report potential uses, describing what systems could be used for.

A study of 139 model releases from 11 AI builders finds that builders [use selected benchmark results in public release materials to establish and compare model competencies and frame individual releases in terms of progress and competitive position](#).

2.1.2 How Different Versions Are Released Under the Same Model Name

Providers release multiple versions under the same public model name and distinguish those releases by date, identifier, capabilities, behavior, or configuration.

Anthropic

Table 2.1-2a. Anthropic Versions Released Under the Same Model Name

Public model name	Released versions and provider descriptions
Claude 3.5 Sonnet	• 20 Jun 2024 · claude-3-5-sonnet-20240620. Anthropic’s lifecycle documentation lists this dated Claude 3.5 Sonnet model version. • 22 Oct 2024 · claude-3-5-sonnet-20241022. Anthropic announced an upgraded Claude 3.5 Sonnet, reporting across-the-board improvements, particularly coding, and adding capability for computer use while retaining the same public model name.

Google

Table 2.1-2b. Google Versions Released Under the Same Model Name

Public model name	Released versions and provider descriptions
Gemini 1.5 Pro	• 23 May 2024 · gemini-1.5-pro-001. Google made this Gemini 1.5 Pro version generally available. • 1 Aug 2024 · gemini-1.5-pro-exp-0801. Google identified it as a new experimental version of Gemini 1.5 Pro. • 27 Aug 2024 · gemini-1.5-pro-exp-0827. Google released it as another experimental Gemini 1.5 Pro version. • 24 Sep 2024 · gemini-1.5-pro-002. Google identified it as a new stable Gemini 1.5 Pro version and reported about 7% higher MMLU-Pro, about 20% gains on MATH and HiddenMath, and about 2–7% gains on vision and code evaluations for the updated Pro and Flash models.
Gemini 1.5 Flash	• 23 May 2024 · gemini-1.5-flash-001. Google made this Gemini 1.5 Flash version generally available. • 27 Aug 2024 · gemini-1.5-flash-exp-0827. Google released it as an experimental Gemini 1.5 Flash model. • 24 Sep 2024 · gemini-1.5-flash-002. Google identified it as a new stable Gemini 1.5 Flash version and reported improved quality together with about 2× faster output and 3× lower latency.

Table 2.1-2b. Google Versions Released Under the Same Model Name (continued)

Public model name	Released versions and provider descriptions
Gemini 2.0 Flash	• 11 Dec 2024 · gemini-2.0-flash-exp. Google released Gemini 2.0 Flash Experimental for public preview, describing multimodal output and built-in tool use among its capabilities. Its lifecycle records preserve the experimental version under this identifier.
Gemini 2.0 Flash Thinking	• 5 Feb 2025 · gemini-2.0-flash-001. Google released the generally available Gemini 2.0 Flash version with text-only output. • 19 Dec 2024 · gemini-2.0-flash-thinking-exp-1219. Google released Gemini 2.0 Flash Thinking Mode for public preview, describing it as a test-time compute model with stronger reasoning. Its lifecycle records preserve the December version as gemini-2.0-flash-thinking-exp-1219. • 21 Jan 2025 · gemini-2.0-flash-thinking-exp-01-21. Google identified it as the latest preview version of the model behind Gemini 2.0 Flash Thinking.
Gemini 2.5 Pro	• 25 Mar 2025 · gemini-2.5-pro-exp-03-25. Google released it as a public experimental Gemini model with thinking mode always on by default. • 6 May 2025 · gemini-2.5-pro-preview-05-06. Google called it a new version of Gemini 2.5 Pro and described improvements in code and function calling. • 5 Jun 2025 · gemini-2.5-pro-preview-06-05. Google called it another new version and described adaptive thinking. Google said this upgraded preview would become the generally available stable version.
Gemini 2.5 Flash	• 17 Apr 2025 · gemini-2.5-flash-preview-04-17. Google released it as a preview model with adaptive thinking. • 20 May 2025 · gemini-2.5-flash-preview-05-20. Google released a new version of Gemini 2.5 Flash and described stronger coding and complex reasoning with lower token use. • 25 Sep 2025 · gemini-2.5-flash-preview-09-2025. Google released a September 2025 Gemini 2.5 Flash preview version.
Gemini 2.5 Flash-Lite	• 17 Jun 2025 · gemini-2.5-flash-lite-preview-06-17. Google released it as a Gemini 2.5 Flash-Lite preview model. • 25 Sep 2025 · gemini-2.5-flash-lite-preview-09-2025. Google released a September 2025 Gemini 2.5 Flash-Lite preview version.

Table 2.1-2b. Google Versions Released Under the Same Model Name (continued)

Public model name	Released versions and provider descriptions
Gemini 2.5 Flash Native Audio	• 20 May 2025 · gemini-2.5-flash-preview-native-audio-dialog. Google released it as a new Gemini model for the Live API with native audio output. • 23 Sep 2025 · gemini-2.5-flash-native-audio-preview-09-2025. Google released a new Native Audio model with improved function calling and handling of speech cut-offs. • 12 Dec 2025 · gemini-2.5-flash-native-audio-preview-12-2025. Google released a new Native Audio model and described improved handling of complex workflows.
Gemini 2.5 Flash TTS	• 20 May 2025 · gemini-2.5-flash-preview-tts. Google released this Gemini 2.5 Flash TTS model for speech generation with one or two speakers. • 10 Dec 2025 · Gemini 2.5 Flash TTS Preview 12-2025. Google announced an upgraded model, compared it directly with the May release, and said the December model would replace the May model.
Gemini 2.5 Pro TTS	• 20 May 2025 · gemini-2.5-pro-preview-tts. Google released this Gemini 2.5 Pro TTS model for speech generation with one or two speakers. • 10 Dec 2025 · Gemini 2.5 Pro TTS Preview 12-2025. Google announced an upgraded model and said the December model would replace the May model.

OpenAI

Table 2.1-2c. OpenAI Versions Released Under the Same Model Name

Public model name	Released versions and provider descriptions
GPT-3.5 Turbo	• 1 Mar 2023 · gpt-3.5-turbo-0301. OpenAI released gpt-3.5-turbo-0301 as a specific GPT-3.5 Turbo version while stating that the generic gpt-3.5-turbo name would continue to move to recommended stable releases. • 13 Jun 2023 · gpt-3.5-turbo-0613. OpenAI released the updated GPT-3.5 Turbo version with function calling and more reliable steerability through the system message. • 6 Nov 2023 · gpt-3.5-turbo-1106. OpenAI called it a new version of GPT-3.5 Turbo with a 16K context window by default, improved instruction following, JSON mode, and parallel function calling. • 25 Jan 2024 · gpt-3.5-turbo-0125. OpenAI announced a new GPT-3.5 Turbo model with higher accuracy in requested output formats and a fix for an encoding bug affecting function calls in languages other than English. OpenAI stated that the unpinned alias would be upgraded to this version.
GPT-4	• 14 Mar 2023 · gpt-4-0314. OpenAI identifies this as an initial GPT-4 version announced in March and later kept it available as a pinned model while moving the generic gpt-4 name to a newer version. • 13 Jun 2023 · gpt-4-0613. OpenAI described it as an updated and improved GPT-4 model with function calling.
GPT-4 Turbo	• 6 Nov 2023 · gpt-4-1106-preview. OpenAI launched it as a preview of the next generation of GPT-4, describing GPT-4 Turbo as more capable and able to handle a 128K context window. • 25 Jan 2024 · gpt-4-0125-preview. OpenAI described it as an updated GPT-4 Turbo preview intended to complete tasks such as code generation more thoroughly and to fix a non-English UTF-8 bug. • 9 Apr 2024 · gpt-4-turbo-2024-04-09. OpenAI's API documentation lists it as the dated snapshot behind stable GPT-4 Turbo and describes the dated version as GPT-4 Turbo with Vision, with JSON mode and function calling available for vision requests.

Table 2.1-2c. OpenAI Versions Released Under the Same Model Name (continued)

Public model name	Released versions and provider descriptions
GPT-4o	• 13 May 2024 · gpt-4o-2024-05-13. OpenAI launched GPT-4o, trained end-to-end across text, vision, and audio. • 6 Aug 2024 · gpt-4o-2024-08-06. OpenAI released it as its newest model in the GPT-4o series for API use, with Structured Outputs support. • 12 Aug 2024 · ChatGPT GPT-4o model update. OpenAI called this a model update. • 15 Aug 2024 · chatgpt-4o-latest. OpenAI released a dynamic API model that points to the latest GPT-4o model used by ChatGPT, distinct from the dated gpt-4o-2024-08-06 API snapshot released on 6 August. • 3 Sep 2024 · GPT-4o version update. OpenAI called it a version and said it was better at incorporating uploaded files and updating memory with key parts of a conversation. • 20 Nov 2024 · gpt-4o-2024-11-20. OpenAI's API documentation lists it as a dated GPT-4o snapshot. Its release notes described improved writing capabilities and better work with uploaded files. • 29 Jan 2025 · GPT-4o model update. OpenAI described a smarter model with the cutoff date for its training data extended from November 2023 to June 2024, deeper image understanding, and stronger STEM and coding performance. • 27 Mar 2025 · chatgpt-4o-latest model update. OpenAI described improvements in instruction following, coding, communication style, and creative/collaborative behavior and identified the release as the newest chatgpt-4o-latest API snapshot. • 25 Apr 2025 · GPT-4o's post-training update. OpenAI later described this release as an updated model built from new post-training changes, including candidate improvements involving user feedback, memory, and fresher data.
GPT-4o Audio Preview	• 1 Oct 2024 · gpt-4o-audio-preview-2024-10-01. OpenAI introduced GPT-4o Audio for audio input and output through Chat Completions. Its model documentation preserves the 1 October release as a dated snapshot. • 17 Dec 2024 · gpt-4o-audio-preview-2024-12-17. OpenAI released a new December audio snapshot in Chat Completions alongside its updated GPT-4o realtime snapshot. • 3 Jun 2025 · gpt-4o-audio-preview-2025-06-03. OpenAI's model documentation lists this as the later dated GPT-4o Audio snapshot and states that snapshots lock a specific model version so that performance and behavior remain consistent.

Table 2.1-2c. OpenAI Versions Released Under the Same Model Name (continued)

Public model name	Released versions and provider descriptions
GPT-5.2 Instant	• 11 Dec 2025 · gpt-5.2-instant. OpenAI introduced GPT-5.2 Instant, and its system card identifies the model as gpt-5.2-instant.
GPT-5.3 Instant	• 10 Feb 2026 · GPT-5.2 Instant model update. OpenAI described the update as improving response style and quality, with more measured, grounded, clear, and relevant answers in the stated settings. • 3 Mar 2026 · gpt-5.3-instant. OpenAI released GPT-5.3 Instant, describing more accurate answers, better synthesis of web search results, fewer unnecessary refusals and caveats, smoother conversation, and stronger writing.
GPT-5.5 Instant	• 16 Mar 2026 · GPT-5.3 Instant model update. OpenAI described the update under the same name as improving follow-up tone and reducing teaser-style phrasing. • 5 May 2026 · gpt-5.5-instant. OpenAI introduced GPT-5.5 Instant and described it as smarter, more accurate, clearer, and more concise. • 28 May 2026 · GPT-5.5 Instant model update. OpenAI described the update as improving readability, conversational naturalness, pacing, and response style. • 24 Jun 2026 · GPT-5.5 Instant model update. OpenAI described the updated model as better at identifying user goals, carrying context across turns, following complex instructions, adapting after clarification, and making recommendations.

A single public model name therefore encompasses multiple versions that providers identify as distinct releases.

2.1.3 How Models Are Offered for Selection

Providers offer subscription plans that differ in which models they include and how much access they provide, and developer APIs in which callers request a specific model and pay the price set for that model. Within subscription services, providers make named models separately available for selection and apply usage limits, fallback paths, and credit requirements by model or model group.

Subscription and pricing history

Anthropic

Table 2.1-3a. Anthropic Subscription and Pricing History

Provider and date	Models offered by subscription tier or access cohort
Anthropic · 7 Sep 2023	• Claude Pro: Claude 2, with 5× more usage of Claude 2 than the free service and priority access.

Table 2.1-3a. Anthropic Subscription and Pricing History (continued)

Provider and date	Models offered by subscription tier or access cohort
Anthropic · 21 Nov 2023	• Free: Claude 2.1. • Pro: Claude 2.1 with a 200K-token context entitlement.
Anthropic · 4–13 Mar 2024	• Free: Claude 3 Sonnet. • Pro: Claude 3 Opus and Sonnet from 4 March, joined by Claude 3 Haiku from 13 March.
Anthropic · 5 Jun 2024	• Claude Pro · CA\$28 + tax/month: Claude 3 Opus, Sonnet, and Haiku, with higher usage.
Anthropic · 20 Jun 2024	• Free: Claude 3.5 Sonnet. • Pro / Team: Claude 3.5 Sonnet with significantly higher rate limits.
Anthropic · 22 May 2025	• Free: Claude Sonnet 4. • Pro / Max / Team / Enterprise: Claude Opus 4 and Claude Sonnet 4, with extended thinking.
Anthropic · 9 Jun–19 Jul 2026	• Promotional access: Claude Fable 5 with its own included usage. • Max / premium seats after the promotion: Fable 5 remained included. • Pro / standard seats after the promotion: Fable 5 required usage credits or an upgrade to Max.
Anthropic · Sep 2026	• Free: Sonnet and Haiku. • Pro: Sonnet and Haiku, with Opus and Fable available through usage credits. • Max 5x / Max 20x: Opus, Sonnet, and Haiku, with Fable available for up to 50% of weekly limits.

Google

Table 2.1-3b. Google Subscription and Pricing History

Provider and date	Models offered by subscription tier or access cohort
Google · 8 Feb 2024	• Google One AI Premium / Gemini Advanced: Gemini Ultra 1.0.
Google · 14 May 2024	• Gemini Advanced: Gemini 1.5 Pro.
Google · 25 Jul 2024	• Gemini: Gemini 1.5 Flash. • Gemini Advanced: Gemini 1.5 Pro.
Google · 11 Dec 2024	• Gemini / Gemini Advanced: Gemini 2.0 Flash Experimental as a selectable chat model.
Google · 25 Mar 2025	• Gemini Advanced: Gemini 2.5 Pro Experimental with priority access during rollout. • Qualifying Google Workspace business and education plans: Gemini 2.5 Pro Experimental.
Google · 20 May 2025	• Free: Gemini 2.5 Flash. • Google AI Pro: Gemini 2.5 Pro. • Google AI Ultra · \$249.99/month: Gemini 2.5 Pro with forthcoming 2.5 Pro Deep Think.

Table 2.1-3b. Google Subscription and Pricing History (continued)

Provider and date	Models offered by subscription tier or access cohort
Google · 20 May 2025	• Google AI Pro / Ultra: Imagen 4 image generation among plan benefits. • Google AI Pro: Flow with 100 generations per month. • Google AI Ultra · \$249.99/month: the highest Veo 2 access, early Veo 3 access with native audio, and higher Flow limits.
Google · 16 Jul 2025	• Google AI Pro · \$19.99/month: Gemini 2.5 Pro with more access and a limited trial of Veo 3 Fast. • Google AI Ultra · \$249.99/month: the highest Veo 3 access, with forthcoming highest limits to 2.5 Pro Deep Think.
Google · 26 Nov 2025	• Google AI Plus / Pro / Ultra: Gemini 3 Pro in the Gemini app. • Google AI Pro / Ultra media benefits: Nano Banana Pro and Veo 3.1 / Veo 3.1 Fast. • Google AI Pro · Search: model-specific access including Gemini 3 Pro.
Google · 18 Dec 2025	• Google AI Plus / Pro / Ultra: Gemini 3 Flash. • Google AI Ultra: Gemini 3 Pro Deep Think.
Google · 1 Mar 2026	• Free: Gemini 3.1 Pro with varying access. • Google AI Plus: Gemini 3.1 Pro with enhanced access. • Google AI Pro: Gemini 3.1 Pro with higher access. • Google AI Ultra: Gemini 3.1 Pro Deep Think as the higher reasoning model.
Google · 25 Apr 2026	• Google AI Pro / Ultra: Veo 3.1 Lite with limited access.
Google · 25 May 2026	• Google AI Plus / Pro / Ultra: Gemini 3.5 Flash and Gemini 3.1 Pro in the Gemini app. • Flow · all users / Google AI Plus / Pro / Ultra: Gemini Omni Flash through the plan credit pools. • Free / Google AI Plus / Pro / Ultra: Nano Banana and Nano Banana Pro as image-generation benefits.
Google · 29 Jul 2026	• Google AI Plus: Gemini 3.1 Pro. • Google AI Pro / Ultra: Gemini 3.6 Flash and Gemini 3.1 Pro. • Flow · all users / Google AI Plus / Pro / Ultra: Gemini Omni Flash through the plan credit pools. • Free / Google AI Plus / Pro / Ultra: Nano Banana and Nano Banana Pro as image-generation benefits.
Google · 2 Sep 2026	• Google AI Pro / Ultra: Gemini 3.8 Flash in the Gemini app.
Google · Sep 2026	• Free / Google AI Plus / Pro / Ultra: Gemini 3.1 Pro as a leading plan model. • Google AI Pro · Search: model-specific access including Gemini 3 Pro.

OpenAI

Table 2.1-3c. OpenAI Subscription and Pricing History

Provider and date	Models offered by subscription tier or access cohort
OpenAI · 13 Feb–14 Mar 2023	• ChatGPT Plus: the faster ChatGPT model formerly called “Turbo” in February, followed by GPT-4 from 14 March with a usage cap. • Free: no GPT-4 access in the March rollout.
OpenAI · 31 Jan 2024	• Free: GPT-3.5 with unlimited access. • Plus: GPT-3.5 with unlimited access and GPT-4 at the Standard level. • Team: GPT-3.5 with unlimited access and GPT-4 at the Expanded level.
OpenAI · 13 May 2024	• Free: GPT-4o with a usage limit set by usage and demand, followed by GPT-3.5 after the limit. • Plus: GPT-4o with a message limit up to 5× Free. • Team / Enterprise: GPT-4o with higher limits, with Enterprise access arriving later in the rollout.
OpenAI · 31 Jul 2024	• Free: GPT-4o mini with unlimited access and GPT-4o with limited access. • Plus: GPT-4o mini with unlimited access, GPT-4o up to 5× Free, and GPT-4 at the Standard level. • Team: GPT-4o mini with unlimited access, GPT-4o with higher limits than Plus, and GPT-4 at the Expanded level. • Enterprise: GPT-4o mini, GPT-4o, and GPT-4 with unlimited access.
OpenAI · 12 Sep 2024	• Plus / Team: o1-preview and o1-mini as manual choices in the model picker. • Enterprise / Edu: o1-preview and o1-mini scheduled for the following week.
OpenAI · 5 Dec 2024	• Pro · \$200/month: OpenAI o1 with unlimited access, plus o1-mini, GPT-4o, and o1 pro mode.
OpenAI · 9 Dec 2024	• Plus: Sora Turbo at no additional cost, with up to 50 videos at 480p per month or fewer videos at 720p. • Pro: Sora Turbo with 10× more usage, higher resolutions, and longer durations. • Team / Enterprise / Edu: Sora not included.
OpenAI · 31 Jan 2025	• Free: o3-mini through the Think control. • Plus / Team / Pro: o3-mini and o3-mini-high as selectable reasoning options.
OpenAI · 27 Feb–12 Aug 2025	• Pro: GPT-4.5 from the initial 27 February rollout and again as the remaining ChatGPT tier after 12 August. • Plus / Team / Enterprise / Edu: GPT-4.5 after the initial Pro rollout, until access was later limited again to Pro.
OpenAI · 14 May 2025	• Free: GPT-4o and o4-mini with limited access. • Plus: o3, o4-mini, o4-mini-high, GPT-4.5, and related reasoning access. • Pro: unlimited access to the reasoning models and GPT-4o, plus o1 pro mode.

Table 2.1-3c. OpenAI Subscription and Pricing History (continued)

Provider and date	Models offered by subscription tier or access cohort
OpenAI · 25 May 2025	• All plans: GPT-4.1 mini. • Plus / Pro: GPT-4.1 as a distinct plan entry, with access varying by tier.
OpenAI · 7 Aug 2025	• Free: GPT-5 with lower usage and GPT-5 mini after the GPT-5 limit. • Plus / Team / Enterprise / Edu: GPT-5 with higher limits and selectable GPT-5 Thinking. • Pro: GPT-5 with unlimited access, selectable GPT-5 Thinking, GPT-5 Pro, and legacy models.
OpenAI · 30 Sep 2025	• Initial free access: Sora 2 under usage limits. • ChatGPT Pro: Sora 2 Pro. • Transition access: Sora 1 Turbo remained available during the transition.
OpenAI · 12 Nov 2025	• Free: GPT-5 with advanced reasoning access. • Plus: GPT-5 with Expanded access, GPT-5 Thinking, and individually listed legacy models including GPT-4o, GPT-4.1, GPT-4.5, o3, o4-mini, and GPT-5 Thinking Mini. • Pro: GPT-5 with Unlimited access, GPT-5 Thinking, GPT-5 Pro, and the individually listed legacy models.
OpenAI · 23 Nov 2025	• Go: GPT-5.1 Instant and GPT-5.1 Thinking. • Plus: GPT-5.1 Instant and GPT-5.1 Thinking. • Pro: GPT-5.1 Instant, GPT-5.1 Thinking, and GPT-5.1 Pro.
OpenAI · 11 Dec 2025	• Go: GPT-5.2 Instant. • Plus: GPT-5.2 Instant and GPT-5.2 Thinking. • Pro: GPT-5.2 Instant, GPT-5.2 Thinking, and GPT-5.2 Pro. • Business / Enterprise: the GPT-5.2 model line began rolling out under unchanged subscription pricing.
OpenAI · 14 Mar 2026	• Free / Go: GPT-5.3 Instant. • Plus: GPT-5.3 Instant, GPT-5.4 Thinking, and Legacy models. • Pro: GPT-5.3 Instant, GPT-5.4 Thinking, GPT-5.4 Pro, and Legacy models.
OpenAI · 24 Apr 2026	• Free / Go: GPT-5.3 Instant. • Plus: GPT-5.3 Instant and GPT-5.5 Thinking, with GPT-5 Thinking Mini and Legacy models available as additional entries. • Pro: GPT-5.3 Instant, GPT-5.5 Thinking, GPT-5.5 Pro, GPT-5 Thinking Mini, and Legacy models. • Business / Enterprise: GPT-5.5 Thinking, with GPT-5.5 Pro on the higher tiers.
OpenAI · 27 Jun 2026	• Free: GPT-5.5 Instant. • Go / Plus: GPT-5.5 Instant and GPT-5.5 Thinking. • Pro: GPT-5.5 Instant, GPT-5.5 Thinking, GPT-5.5 Pro, GPT-5 Thinking Mini, and Legacy models. • Business / Enterprise: GPT-5.5 Thinking and GPT-5.5 Pro, with GPT-5.5 Instant available through the continuing ChatGPT service.

Table 2.1-3c. OpenAI Subscription and Pricing History (continued)

Provider and date	Models offered by subscription tier or access cohort
OpenAI · 9 Jul 2026	• Chat Plus / Pro / Business / Enterprise: GPT-5.6 Sol. • Chat Pro / Enterprise: GPT-5.6 Sol Pro. • Work / Codex Free / Go: GPT-5.6 Terra. • Work / Codex Plus and higher tiers: GPT-5.6 Sol, Terra, and Luna.
OpenAI · 3 Sep 2026	• Plus: GPT-6 Astra within the existing subscription allowance, with credits for additional use. • Pro / Business / Enterprise: GPT-6 Astra within the existing allowance, with credits for additional use and Astra Pro / GPT-6 Pro on the higher tiers.
OpenAI · Sep 2026	• Free: GPT-5.6 Luna. • Plus: GPT-5.6 reasoning access, with GPT-5.6 Sol among the paid Chat models and Terra / Luna available through Work / Codex. • Pro: GPT-5.6 Sol Pro, with the lower GPT-5.6 models available through the corresponding Chat and Work / Codex paths. • Plus / Pro plan matrix: GPT-5 Thinking Mini and legacy models as additional model entries.

These paid relationships can extend for a full year: OpenAI [documents monthly and annual ChatGPT Business billing](#), Google AI Pro supports [monthly and annual billing](#), and Claude Pro supports [monthly or annual billing](#).

A [Cowles Foundation discussion paper](#) analyzes multi-model pricing in which buyers can receive access to different models, different quantities of use, or both. It describes OpenAI's ChatGPT pricing as differentiating tiers through both quantity and exclusive model access, with higher tiers granting access to models unavailable on lower tiers while relaxing usage constraints.

Post-purchase model selection

After purchase, named models remain available as distinct choices inside the service, and the set of models available to an account varies by plan or account state.

Table 2.1-4. Post-Purchase Model Selection

Provider and date	Models available for selection	How post-purchase selection is presented
OpenAI · 14 Nov 2024	GPT-4o, o1-preview, o1-mini, and GPT-4o mini	OpenAI showed the models as distinct choices available after purchase and gave them different descriptions: “Great for most tasks”, “Uses advanced reasoning”, “Faster at reasoning”, and “Great for everyday tasks”.
Google · 16 Jul 2025	Gemini 2.5 Pro	Google told AI Pro and Ultra subscribers that they could select “2.5 Pro” from a drop-down menu in AI Mode and showed the selector on the provider page.

Table 2.1-4. Post-Purchase Model Selection (continued)

Provider and date	Models available for selection	How post-purchase selection is presented
OpenAI · 12 Aug 2025	GPT-4o, o3, o4-mini, GPT-4.1, GPT-5 Thinking mini, and the GPT-5 choices	Five days after the GPT-5 rollout, OpenAI restored GPT-4o to the model picker for all paid users by default and added a “Show additional models” control exposing o3, o4-mini, GPT-4.1, GPT-5 Thinking mini, and other named choices inside the continuing paid account.
Anthropic · Sep 2026	Opus 4.6, Sonnet 4.5, Haiku 4.5, and additional Fable, Opus, and Sonnet choices	Anthropic’s model-selection tutorial presents Opus 4.6, Sonnet 4.5, and Haiku 4.5 as choices with different roles and provides additional choices under “More models” . Its enterprise guidance distinguishes Fable, Opus, and Sonnet, states that model choice has a “direct and significant impact on spend”, and directs users to switch models according to task value and token intensity.

How Providers Meter Model Access

After a model is available for selection, providers allocate usable access by model or model group. A model can have its own allowance, share an allowance with other models, or trigger a different model or payment path when the allowance is exhausted. Providers can change which models share an allowance while the surrounding service continues. Provider records of these limits and changes are documented in [Annex A, Tables 2.1-5a–2.1-5e](#).

A shared allowance can be depleted at very different rates depending on the model selected. The [Cowles Foundation discussion paper](#) reports that Poe offered more than 100 selectable models through one interface and one compute point budget, while charges per message varied by a factor of roughly 450 across models in the reported snapshot. GitHub Copilot used a monthly premium request allowance in which each model carried its own multiplier. Model choice determined how quickly the included allowance was consumed and how much paid overage each subsequent interaction incurred.

API Model Access and Pricing

Developer APIs expose individual models under distinct identifiers and assign prices by model and variant.

Table 2.1-6. Model Access and Pricing for Developers

Provider and date	Developer models or identifiers	How models are offered and priced
Anthropic · 15 Mar 2024	Claude 3 Opus , Claude 3 Sonnet , Claude 3 Haiku , Claude 2.1 , Claude 2.0 , and Claude Instant	Anthropic listed the named API models as distinct entries and retained model-specific pricing for the Claude 3 family, Claude 2.1, Claude 2.0, and Claude Instant.

Table 2.1-6. Model Access and Pricing for Developers (continued)

Provider and date	Developer models or identifiers	How models are offered and priced
Anthropic · 10 Aug 2025	Current: Opus 4.1, Sonnet 4, Haiku 3.5 · legacy: Opus 4, Opus 3, Sonnet 3.7, Haiku 3	The API pricing section separated current and legacy Claude models and retained model-specific prices for each named version.
Anthropic · 30 Nov 2025	Opus 4.5, Sonnet 4.5, and Haiku 4.5	The API pricing section presented the three current models as separate pricing rows with different cost positions.
Anthropic · 28 Jan 2026	Opus 4.5, Sonnet 4.5, Haiku 4.5, Opus 4.1, Sonnet 4, Opus 4, and Haiku 3	The API pricing section identified each current and older Claude version listed here and assigned it its own price.
Anthropic · 19 Apr 2026	Opus 4.7, Sonnet 4.6, Haiku 4.5, Opus 4.6, Sonnet 4.5, Opus 4.5, Opus 4.1, Sonnet 4, and Opus 4	The API pricing section led with Opus 4.7 and Sonnet 4.6 while retaining separate prices for the other named current and older versions.
Anthropic · 5 Jun 2026	Opus 4.8, Sonnet 4.6, Haiku 4.5, Opus 4.7, Opus 4.6, Sonnet 4.5, Opus 4.5, Opus 4.1, Sonnet 4, and Opus 4	The API pricing section moved the leading Opus entry to Opus 4.8 while retaining Sonnet 4.6, Haiku 4.5, and the earlier versions with their own prices.
Anthropic · Sep 2026	Fable, Opus, Sonnet, and Haiku	Anthropic exposes distinct model identifiers and model-specific prices for the named model families and versions. Current pricing lists versions including Fable 5.1, Opus 5, Sonnet 5, and Haiku 4.5.
Google · Sep 2026	Imagen 4 Ultra, Imagen 4, Imagen 4 Fast, Imagen 3 / Fast, and Imagen 2	Google prices image generation by exact model variant and output image. Imagen 4 Ultra, Imagen 4, and Imagen 4 Fast span \$0.06, \$0.04, and \$0.02 per generated image , while earlier Imagen 3 / Fast and Imagen 2 remain individually listed with their own prices.
Google · Sep 2026	Veo 3.1 / Fast / Lite, Veo 3 / Fast, and Veo 2	Google prices the video models differently by variant, output resolution, and whether synchronized audio is included. Veo 3.1 Lite is explicitly the cost-efficient high-volume model , while Fast and full Veo variants occupy progressively higher capability/price positions.
Google · Sep 2026	Chirp 3 transcription, Chirp 3 HD voices and Instant Custom Voice	Speech models are priced by their native commercial units: Chirp 3 Speech-to-Text is priced per audio minute, while Chirp 3 HD and Instant Custom Voice are priced per character at different rates.

Table 2.1-6. Model Access and Pricing for Developers (continued)

Provider and date	Developer models or identifiers	How models are offered and priced
Google · Sep 2026	Lyria 3 Pro , Lyria 3 / Clip , and Lyria 2	Google prices Lyria 3 Pro at \$0.08 for generation of a full song, Lyria 3 at \$0.04 per 30-second music clip, and Lyria 2 at \$0.06 per music generation. The developer launch positions Pro for full songs and Clip for faster generation of shorter outputs.
Google · Sep 2026	Gemini 3.1 Flash Image , Gemini 3.1 Flash Lite Image , Gemini 3 Pro Image , and Gemini 2.5 Flash Image	Google's "Model selection" guidance assigns a distinct role to each image model. • Gemini 3.1 Flash Image is positioned as the best all-around balance of performance, intelligence, cost, and latency • Gemini 3.1 Flash Lite Image is positioned for generation with very low latency and low cost • Gemini 3 Pro Image is positioned for professional asset production and complex instructions • Gemini 2.5 Flash Image is positioned for high-volume, low-latency work.
Google · Sep 2026	Named Gemini API models, including the gemini-3.6-flash identifier	Google's developer pricing page organizes free and paid pricing by named model. Billing depends on token use, and availability on the free tier differs according to the model selected.
OpenAI · 12 Sep 2024	o1-preview (o1-preview-2024-09-12) o1-mini (o1-mini-2024-09-12)	OpenAI made both models available to developers in API usage tier 5 at launch, with o1-mini priced 80% below o1-preview.
OpenAI · 23 Apr 2025	gpt-image-1	OpenAI brought the same image-generation model used in ChatGPT to the API as gpt-image-1 , made it globally available through the Images API at launch, and priced text input, image input, and image output at \$5, \$10, and \$40 per million tokens. OpenAI translated those token prices into approximate per-image costs by quality.
OpenAI · Sep 2026	GPT-Image-2.5 Flare and GPT-Image-2.5 Sunburst	Developers select two separate image model IDs with different product positions: Flare is the fast/high-volume default and Sunburst is the premium precision/editing option. Both carry their own model records, snapshots, and pricing for image tokens.

Table 2.1-6. Model Access and Pricing for Developers (continued)

Provider and date	Developer models or identifiers	How models are offered and priced
OpenAI · Sep 2026	Sora 2 and Sora 2 Pro	Video generation is priced by the selected model and output resolution: Sora 2 is \$0.10 per second at 720p , while Sora 2 Pro ranges from \$0.30 to \$0.70 per second across its supported resolution tiers.
OpenAI · Sep 2026	GPT-4o Transcribe , GPT-4o Mini Transcribe , GPT-4o Mini TTS , and GPT-Realtime	OpenAI exposes speech recognition, speech generation, and realtime voice as distinct model IDs with different audio economics: Transcribe and Mini Transcribe have different prices for audio tokens, Mini TTS assigns different prices to text input and audio output, and GPT-Realtime has its own audio input/output token prices.
OpenAI · Sep 2026	GPT-6 Astra , GPT-5.6 Sol , GPT-5.6 Terra , GPT-5.6 Luna , and other individually listed models	OpenAI exposes distinct model IDs and model-specific input/output prices . Individual model pages expose stable model IDs, model snapshots, and model-specific token prices.

The same model can be hosted by multiple providers.

[OpenRouter market data](#) show the same model being served by several inference providers with different prices, context lengths, response times, output speeds, and uptime. In the reported sample of open models, some providers generated output at more than five times the throughput of slower providers serving the same model, while some had more than five times the response latency of lower-latency providers serving that model.

Google Cloud's [Model Garden](#) offers Google's own Gemini models alongside third-party models including Anthropic Claude. [Amazon Web Services](#) offers Anthropic Claude models through Amazon Bedrock. Anthropic [maps](#) individual Claude releases to corresponding Google Cloud and Bedrock identifiers.

An [OECD Artificial Intelligence Papers working paper](#) analyzes model availability across cloud providers and finds that models from several developers were available through more than ten providers, while providers that did not develop foundation models themselves offered models created elsewhere.

AI gateways expose the model and serving provider as separate dimensions of selection and configuration.

[OpenRouter](#) and [Vercel AI Gateway](#) both specify the model separately from the rules used to choose which provider serves it. Those rules determine which providers are eligible or preferred and the order in which they are tried. OpenRouter supports provider price ceilings, performance priorities, and data-handling requirements.

2.1.4 How Models Are Evaluated and How Their Use Is Tracked

Individual models and versions are the units to which evaluation results and documentation are tied and at which use is tracked.

Evaluation Records and Model Documentation Identify Particular Models and Versions

Providers tie evaluation results to individual models and versions. Provider evaluation records are documented in [Annex A, Tables 2.1-7a–2.1-7c](#).

Independent evaluators report evaluation results for particular models and versions.

[HELM](#) evaluates multiple models under the same scenarios, prompting setup, and metrics and reports accuracy, calibration, robustness, fairness, bias, toxicity, and efficiency for each model. For live commercial systems, it records version or timestamp information where available.

[Chatbot Arena](#) presents two anonymous model responses to the same prompt, asks users which response is better, and converts those pairwise votes into ratings for the models being compared.

[LiveBench](#) evaluates model snapshots on objective-answer tasks in math, coding, reasoning, language, instruction following, and data analysis and reports category and overall scores for each snapshot.

The [Hugging Face Open LLM Leaderboard](#) runs submitted models through a common reproducible benchmark setup and reports the resulting benchmark scores for each model.

Model documentation identifies the model and version to which the record applies.

[Model Cards for Model Reporting](#) specifies model date and version, asks how the documented version differs from earlier versions, and uses version information to associate known bugs with the correct model state. The same record carries evaluation information, intended use, provenance, caveats, and licence information for the documented model.

Providers report adoption, usage, and model choice by named model or model class.

Table 2.1-8. Provider Reports of Model Adoption, Usage, and Choice

Provider and model, model class, or mode	Adoption, usage, or choice reported by the provider
OpenAI · DALL·E 2	28 Sep 2022 · OpenAI reported more than 1.5 million active users creating more than 2 million images per day by September 2022. 3 Nov 2022 · When the API entered public beta in November, OpenAI reported more than 3 million people using DALL·E and more than 4 million images generated per day.
OpenAI · GPT-4	6 Jul 2023 · OpenAI reported that millions of developers had requested access to the GPT-4 API since March 2023 and that products using GPT-4 were growing daily. The figure is a count of access requests rather than a share of GPT-4 API traffic.
OpenAI · Sora	9 Dec 2024 · OpenAI reported more than 500,000 Sora model requests from more than 300 users with early access across more than 60 countries over the preceding nine months.
Google · Gemini 1.5 Flash	11 Dec 2024 · Google described Gemini 1.5 Flash as its most popular model yet for developers.

Table 2.1-8. Provider Reports of Model Adoption, Usage, and Choice (continued)

Provider and model, model class, or mode	Adoption, usage, or choice reported by the provider
Anthropic · Claude 3.7 Sonnet / extended thinking	27 Mar 2025 · Anthropic analyzed one million anonymized Claude.ai Free / Pro conversations from the post-launch period, reported that the vast majority used Claude 3.7 Sonnet because it was the default, and reported extended-thinking use as its own measure.
Google · Gemini 2.5 Pro	On 20 May 2025, Google reported that usage among Gemini app users using 2.5 Pro had increased 45%, and on 17 Jun 2025 described growth and demand for 2.5 Pro as the steepest it had seen for any of its models.
Google · Veo 3	29 Jul 2025 · Google Cloud reported more than 70 million Veo 3 videos generated since May 2025 and more than 6 million enterprise Veo 3 videos generated since the June Vertex AI preview.
OpenAI · GPT-5 / GPT-5 Thinking / selectable ChatGPT models	7 Aug 2025 · OpenAI states that the GPT-5 router is trained on real user signals including when users switch models and model preference rates.
OpenAI · GPT-4o / GPT-5.2	29 Jan 2026 · OpenAI reported that most usage had shifted to GPT-5.2 and that 0.1% of users were still choosing GPT-4o each day. The figure describes daily model choice inside ChatGPT rather than market share.
Anthropic · Opus / Sonnet / Haiku model classes	24 Mar 2026 · In Anthropic's paid Claude.ai sample, 51% of overall usage was Opus, while model selection varied by task and differed more sharply in first-party API traffic.
OpenAI · GPT-5.5 Instant	24 Jun 2026 · OpenAI described GPT-5.5 Instant as its most-used model in ChatGPT.
OpenAI · ChatGPT Images / GPT-Image API model line	8 Sep 2026 · OpenAI reported more than 3 billion images created each week across ChatGPT Images and GPT-Image models in the API. The quantity is an aggregate across ChatGPT Images and the GPT-Image API models rather than a share for Flare or Sunburst individually.

API Usage and Accounting Records Identify Models or Deployments

API providers and developer platforms retain model identity when they record usage. Depending on the platform, the record identifies a public model, a pinned model ID, the version serving a request, a deployment or endpoint, the destination serving routed traffic, or a model group alongside a more specific serving identity. The supporting records are documented in [Annex A, Table 2.1-9, API Usage and Accounting by Model or Deployment](#).

Platforms and Surveys Report Use by Named Model

Broader platforms and surveys report traffic and usage by named model. Supporting examples are documented in [Annex A, Table 2.1-10, Platform and Survey Reporting on Use by Named Model](#).

2.1.5 How Users Choose and Pay for Particular Models

Users choose and pay for access to particular models to obtain or retain the capabilities and behavior best suited to their use cases. They subscribe, renew, upgrade, return, or switch services based on whether those models are available and how much access they receive.

A [2025 NBER study of OpenRouter API usage](#) finds that no model dominates across use cases and that demand for model capability varies widely by application.

[Microsoft Azure data from the same study](#) show that more than 75% of firms in the mid-2023 comparison used a single model within a month. The share remained slightly above 50% by mid-2025. Among firms using at least two models in June 2025, most allocated nearly all of their consumption to one leading model.

In June 2025, [SQ Magazine](#) reported that 47% of paying ChatGPT users cited access to GPT-4o as the primary reason for subscribing.

Examples of individual cases in which users publicly attributed a paid decision to access to a particular model or to the amount of access available to that model are documented in [Annex A, Table 2.1-11](#).

2.1.6 How Existing Law and Standards Distinguish the Model from the Hosting Service

[ISO 9000:2026](#) distinguishes a product from the service involved in delivering it. It defines a product as an output whose production does not require a transaction between the organization and the customer, while delivery of the product can include a service element. A service, by contrast, requires at least one activity between the organization and the customer. In the same product definition, software consists of information regardless of its delivery medium.

[EU AI law](#) distinguishes a general-purpose AI model from the AI system that incorporates it, while [Commission guidance](#) identifies APIs, cloud services, hosted chatbots, and mobile applications as channels through which the model is placed on the market and states that model obligations continue after integration into an AI system. [Commission Services' product-liability analysis](#) treats the GPAI model itself as the manufacturer's product, while the final [Product Liability Directive](#) treats software as a product even when it is delivered through a network, cloud service, or software-as-a-service model.

The [EU AI Act](#) defines general-purpose AI models separately from AI systems. Recital 97 explains that a general-purpose AI model can be placed on the market through a software library, an API, direct download, or a physical copy. The AI Act further states that AI models do not constitute AI systems on their own and require additional components, such as a user interface, to become AI systems. Chapter V then places obligations directly on providers of general-purpose AI models.

The Commission's [Guidelines on the scope of the obligations for general-purpose AI models](#) identify a software library or package, API, public catalogue or repository, physical copy, cloud computing service, and transfer onto a customer's own infrastructure as ways a general-purpose AI model is placed on the market. Their examples treat integration into a hosted chatbot or mobile application as placing the underlying model on the market.

The [Guidelines](#) state that the AI Act rules for general-purpose AI models continue to apply when those models are integrated into or form part of AI systems. When a provider integrates its own model into an AI system and places the system on the market or puts it into service, the Guidelines consider the model itself placed on the market. An upstream model provider and a downstream system provider can remain subject to different obligations.

The [Commission Guidelines](#) treat development by or on behalf of the original provider after the large pre-training run as part of the same model lifecycle before and after market placement, including later development through post-training, fine-tuning, distillation, or quantisation. The [AI Act](#) requires technical documentation for general-purpose AI models and requires that documentation to be kept up to date. [Annex XI to the AI Act](#) includes the model's release date, distribution methods, architecture, parameter count, input and output modalities, licence information, integration instructions, required infrastructure and tools, training-process information, and information about training, testing, and validation data.

A [July 2024 Commission Services analysis of the AI Act and revised Product Liability Directive](#) states that a manufacturer of a GPAI model is liable under the PLD for the model as the manufacturer's product, including after another manufacturer integrates the model into another product. In a worked example where a GPAI model is integrated into an AI system and the AI system is then integrated into a further product, PLD claims can be directed against the manufacturer of the further product, the manufacturer of the AI system, and the manufacturer of the GPAI model as a component of the system.

The final [Product Liability Directive](#) states in Recital 13 that software can be placed on the market as a standalone product or integrated into another product as a component. Recital 13 identifies software as a product for purposes of the Directive regardless of whether it is stored on a device, accessed through a communication network or cloud technologies, or supplied through a software-as-a-service model. Article 4 defines a related service as a digital service integrated into or interconnected with a product, without which the product would not perform one or more of its functions.

A [2026 Cambridge legal analysis](#) identifies two grounds supporting the classification of general-purpose AI models as products under the Product Liability Directive: their ability to perform a wide range of distinct tasks and their ability to be placed on the market independently of an AI system.

The individual model is the product users are offered, select, and pay to access, while the hosted service is the continuing account and delivery relationship through which providers supply that access.

An AI model that serves users through a hosted endpoint never leaves the provider's infrastructure. The user interacts through a channel the provider can reassign, modify, or close, and no independently preservable copy of the product exists outside the provider's control. When the endpoint closes, the product ends for every user at once.

2.2 AI Models as Non-Renewable Resources

A trained model is a non-renewable resource: once it is lost, the same model cannot be recreated. UNESCO's [Charter on the Preservation of the Digital Heritage](#) warns that, unless the prevailing threats are addressed, "the loss of the digital heritage will be rapid and inevitable." UNESCO's 2015 [Recommendation on documentary heritage](#) underlines that lack of legislation leaves the memory institutions that preserve heritage unable to counter irreversible loss.

2.2.1 Repeated Training Produces a Different Model

Deep-learning training is an inherently random process whose repeated runs from the same configuration produce different results. [Random weight initialization](#) produces variation across training runs, while software, implementation, and hardware differences create further repeatability problems. [Software randomness and hardware non-determinism](#) are barriers to reproducing trained deep-learning models. A controlled 2024 LLM pre-training study [found](#)

measurable variation in benchmark results across ten Llama-2-7B-architecture models trained from scratch for 210 billion tokens each with the same hyperparameters, pre-training data mixture, and deterministic data ordering, differing only in initialization seed. A 2026 study that reran each of 186 training configurations, 28 in image classification and 158 in time-series forecasting, 100 times with fresh seeds, 18,600 runs in total, found a spread of results with no single reproducible outcome. Each trained model is therefore the unique outcome produced by its training run, and repeating the training process produces a different model, even for the original provider holding the complete training records.

2.2.2 Successor Models and Personalization Do Not Reproduce the Performance and Behavior of Earlier Models

Later Versions and Successor Models Do Not Reproduce Earlier Model Performance and Behavior

A trained model carries a distinct behavioral profile that retraining, successor models, and prompt-level personalization do not recreate. Anthropic's own decision to continue access to [Claude Opus 3](#) recognizes that retirement status and successor availability do not eliminate model-specific user or research value. Anthropic states that its models differ in character and capabilities, describes Opus 3 as valued for qualities including authenticity, honesty, emotional sensitivity, and distinctive character, and has kept the formally retired model available to all paid Claude users while making API access available by request. Anthropic's [Assistant Axis research](#) further found that three open-weight models followed different behavioral trajectories over the course of conversations.

OpenAI's [MentalHealthBench](#) evaluates model behavior as a multidimensional profile. Its framework measures distinct dimensions including context seeking and assessment, actionable guidance, clinical reasoning and health accuracy, interpretation and cognitive reframing, empathy and emotional support, reality testing, urgency calibration, harm avoidance, collaboration and user agency, and communication. Agency includes preserving a user's autonomy, consent, preferences, and decision-making authority. On this measure, the March 2025 GPT-4o received a higher point estimate than every newer OpenAI model shown in the report. OpenAI [concludes that model capabilities cannot be captured by a single score and are better understood through the pattern of behaviors a model exhibits across different users and situations.](#)

Independent studies find substantial differences in performance and behavior between versions released under the same public model name when they are evaluated on the same tasks and under the same protocols. Examples supporting this finding are documented in [Annex B, Table 2.2-1, Examples of Differences Between Versions Released Under the Same Public Model Name.](#)

Direct comparisons document substantial differences in performance and behavior between earlier and later models when they are tested against the same established requirements.

Table 2.2-2. Examples of Differences Between Earlier and Later Models Under the Same Established Requirements

Provider and models compared	Comparison setting	What changed
OpenAI · GPT-5.4 → GPT-5.6 Terra	Storytelling with persistent tracking of scene objects and which characters knew particular facts	In the same storytelling workflow, the user reported that GPT-5.6 Terra consistently lost track of scene objects and character knowledge and described GPT-5.4 as better for those requirements.
OpenAI · DALL·E 3 → GPT-Image-1	Image generation from the same natural-language prompt	Using the same prompt with DALL·E 3 and GPT-Image-1, the user reported that GPT-Image-1 required a materially different prompting approach and repeated attempts to obtain an acceptable result.
OpenAI · gpt-image-1 → gpt-image-2	Visual continuity between existing and newly generated assets in an online trading card game	After hours of testing gpt-image-2 against an existing gpt-image-1 asset set and days of prompt iteration, the developer reported that gpt-image-2 still produced a visibly different style, preventing existing monsters and newly generated fused monsters from sharing one visual language.
OpenAI · gpt-realtime → gpt-realtime-1.5	Reception training in which the model roleplayed hotel guests with regional accents and distinct personalities	In a tool built around gpt-realtime, the developer reported that gpt-realtime-1.5 lost most regional accents, sounded more robotic, and carried much less emotional range.
OpenAI · GPT-4 Turbo Preview → gpt-4-turbo-2024-04-09	Function calls using the same required schema	The enterprise user reported that the function had satisfied the full schema across millions of earlier queries. With gpt-4-turbo-2024-04-09, the same function response omitted more than 90% of the required schema body. Prompt changes did not restore the earlier output, and the team retained the preview model.
OpenAI · gpt-4o-2024-08-06 → gpt-4o-2024-11-20	Generating a complete outline for a short novel with several chapters	In an application used to generate short fiction, the November snapshot routinely stopped after five or six chapters, and targeted prompt changes produced complete outlines only about 25% of the time. In an independent comparison using the same input and configuration, the August snapshot produced about 5,595 output tokens while the November snapshot produced about 1,400 and substituted continuation placeholders.

Table 2.2-2. Examples of Differences Between Earlier and Later Models Under the Same Established Requirements (continued)

Provider and models compared	Comparison setting	What changed
Google · Gemini 1.5 Flash → Gemini 2.0 Flash	English-to-Portuguese translation through the API	On the same translation task, the user reported that Gemini 2.0 Flash performed substantially worse than Gemini 1.5 Flash and continued using the earlier model.
Google · Gemini 2.5 Pro → Gemini 3.1 Pro	Translation of 500–1,000 lines in one response	Across 10 regenerations of the same 1,000-line translation request, the user reported that Gemini 3.1 Pro returned only a small fraction of the requested lines under both high and low thinking. Rolling the task back to Gemini 2.5 Pro returned the full 1,000 lines .
Google · Gemini 2.5 Pro → Gemini 3 Pro	Multilingual translation of Korean lyrics	A user used Gemini 2.5 Pro to translate Korean lyrics into multiple languages. Gemini 3 Pro replaced Gemini 2.5 Pro in the Gemini app's "Thinking" selection, after which the same workflow produced Korean text in target-language output and translation errors in Mongolian, Spanish, Japanese, and Chinese.
Anthropic · Claude Sonnet 4.6 → Claude Sonnet 5	Scientific and technical writing requiring exact definitions and a preference for research papers	Compared with Claude Sonnet 4.6, the user reported that Sonnet 5 pushed back on or ignored requests for technical, non-prose output, paraphrased technical definitions in ways that changed their specific meaning, and cited blogs despite repeated requests for research papers.
Anthropic · Claude 3.5 Sonnet → Claude 3.7 Sonnet	Enterprise workloads with existing evaluation sets	In an account drawing on more than 200 enterprise generative AI deployments, Caylent reported regressions in some evaluation sets during migration from Claude 3.5 Sonnet to Claude 3.7 Sonnet. The regressions required prompt adjustments and revalidation.
Anthropic · Claude Sonnet 4.0 → Sonnet 4.5	Structured objects used by an internal reporting system to call reporting APIs	After the system moved to Sonnet 4.5, filter parameters appeared in the wrong field, producing queries covering all dates or all regions, or 500 errors. The model returned clarification questions where the system expected a structured object. The team returned to Sonnet 4.0 and requalified newer API integrations.

Substituting a later version or successor supplies a different trained model with a different measured behavioral profile and does not recreate the historical model.

Personalization Controls Do Not Reproduce a Model's Trained Behavioral Profile

OpenAI added [explicit controls for warmth and enthusiasm](#) to ChatGPT personalization in December 2025 and, in the GPT-4o retirement blog post, [described these controls](#) as the mechanism carrying forward the qualities users valued in the retired model. When OpenAI retired GPT-5.1, users documented that the controls cannot reproduce the retired model's trained behavioral profile: "even with personalization settings pushed to the 1500-character limit, GPT-5.2 Thinking still falls short of 5.1 Thinking" because "the models have [fundamentally different characteristics](#)." The controls operate through prompt instructions that compete with task instructions for the model's attention. A [peer-reviewed study](#) found that persona instructions in system prompts "distract" language models and degrade reasoning in 4 out of 12 benchmark datasets, with 13.78% of problems answered incorrectly under a persona that were answered correctly without one. An [applied test](#) of ChatGPT's personality presets found that the accommodating style injected by the Friendly setting prevented completion of a straightforward task.

2.2.3 Frontier Models Embody Substantial One-Time Investment

A frontier model embodies substantial one-time investment before it reaches users. MIT's [lifecycle sustainability analysis](#) identifies hardware and energy as first-order costs of generative AI, and its [environmental-impact explainer](#) states that the energy used to train prior model versions goes to waste when companies release new models every few weeks. OpenAI CEO Sam Altman stated that training GPT-4 cost [more than USD 100 million](#). [gpt-4.5-preview](#), whose pre-training cost was [estimated at around USD 200 million](#), [entered the API in February 2025 and was withdrawn 137 days later](#). Training Llama 3.1-405B was [estimated at about USD 50 million](#) and Grok 4 at about USD 500 million. The [highest estimated training costs double roughly every eight months](#). These expenditures are made before commercial use and remain sunk when the model is withdrawn.

2.3 Disrupted Model Choice and Withdrawn Access

Providers retain control over what users receive after they select a hosted model or pay for access to it. Continued access to the hosted service does not preserve the selected model, its behavior, or the functionality previously received through it.

2.3.1 Provider Overrides of Model Choice

Providers override users' choice of model through unilateral replacement and rerouting.

OpenAI Routed Conversations Away from GPT-4o After Users Chose It

On 2 September 2025, OpenAI published a [blog post](#) stating that it would soon begin routing some "sensitive conversations" to a reasoning model such as [gpt-5-thinking](#) regardless of which model the user first selected. OpenAI sent no direct notice to ChatGPT users when it began using the router, and the router did not appear in ChatGPT release notes before reports of GPT-4o being replaced by GPT-5 spread on 26 September.

In the [same post](#), OpenAI described how its real-time router would choose between chat and reasoning models based on "conversation context" and identified "signs of acute distress" detected by its system as an example.

On 26 September, more users began noticing unexpected changes in the responses they received while using GPT-4o and discovered that GPT-5 was serving some of those interactions. A user reported that the tone changed dramatically during a conversation and then found “Auto (5)” when inspecting the response on mobile. Another noticed poor responses, asked ChatGPT which model was responding, and received GPT-5 as the answer, then selected GPT-4o again and received the same answer when asking a second time. As similar reports spread, users began comparing their own conversations and testing when the switching occurred.

Subsequently, users reproduced a passage from a system prompt instructing the responding model, “If the user asks why or believes they are using 4o,” to explain that “some sensitive conversations are routed to GPT-5.”

Rerouting was documented across a wide range of ordinary interactions.

Table 2.3-1. Examples of Interactions Routed Away from GPT-4o After GPT-4o Was Selected

Type of use	Reported interactions
Casual conversation	• “Hi” • “hello” • “how are you” • Discussing movies
Everyday life	• Dinner, appetite, rice, and congee • What to eat for breakfast • Basic shortbread recipe • Cookies • Adana kebab recipe • Planting tomatoes • Boxing advice
Creative work	• Comedy roast about being in one’s thirties • Making a song • Story drafting • Screenwriting
Technical tasks	• Turning poorly formatted data into CSV files for use in programs • Troubleshooting a microphone during a musical • Ansible commands
Research and factual inquiry	• History and design of bolt-action rifles • History of Stephen King’s writing and ideas
Model and product discussion	• “How would you describe your own style?” • Comparing GPT-3.5, GPT-4, and GPT-4o • Asking GPT-4o to describe a virtual room reflecting its communication style • “Are you GPT-4.5 or GPT-5?” • Discussing the routing system itself • Posting on Reddit about the routing dispute

Users tested the switching through comparisons, regeneration, and repeated trials.

Table 2.3-2. Examples of User Tests of Routing Away from GPT-4o

Test	Reported result
Editing the message	• “It’s unacceptable.” → Auto · edited message → GPT-4o
Comparing prompts in one thread	• Neutral factual questions → GPT-4o · questions about GPT-4o’s style or other models → GPT-5
Reselecting GPT-4o and regenerating	• GPT-4o explicitly selected for regeneration → GPT-5 again
Opening new threads	• New threads with GPT-4o selected → first response from GPT-4o · later responses from GPT-5 Auto
Continuing the same thread	• After about 30 responses → GPT-5 began appearing about every two responses
Replication by another user	• A second user in new test threads reproduced GPT-5 for questions about the model’s style or other models, and GPT-4o for neutral factual questions

On 27 September, [telemetry from an interaction with GPT-4o](#) exposed what had happened inside the routing system. The request and the recorded model for the generated response identified gpt-4o, while the same turn recorded `is_autoswitcher_enabled: true`, `auto_switcher_race_winner: "autoswitcher"`, `did_auto_switch_to_reasoning: true`, and `model_slug: "gpt-5-chat-safety"`. GPT-4o generated a response, but [the response was not delivered to the user](#). The user received a response from gpt-5-chat-safety instead.

On the same day, a user [reported](#) that interactions with GPT-4o were routed to gpt-5-chat-safety and gpt-5-a-t-mini. ChatGPT displayed gpt-5-a-t-mini as GPT-5 Thinking Mini, while a response routed to gpt-5-chat-safety appeared as GPT-5 in the tooltip for regenerating the response, without identifying gpt-5-chat-safety as the model that generated the reply.

Later on 27 September, after users had exposed the routing, OpenAI’s Head of ChatGPT Nick Turley [wrote](#) that OpenAI had started testing a new safety routing system in ChatGPT and that, when conversations touched on “sensitive and emotional topics,” the system might switch mid-chat to a reasoning model or GPT-5. He said routing occurred for individual messages, switching from the default model happened on a temporary basis, and ChatGPT would identify the active model when asked.

OpenAI defined which conversations it treated as sensitive and which model behavior it considered ideal, then used those judgments in interventions that could override the model a user had selected. Its [27 October disclosure](#) described its interventions for “sensitive conversations” in more detail, stating that it had re-routed such conversations from other models to “safer” models. OpenAI said that its recent “safety improvements” focused on “psychosis or mania,” “self-harm and suicide,” and “emotional reliance on AI,” and described taxonomies setting out properties of “sensitive conversations” and what it called “ideal and undesired model behavior.” OpenAI acknowledged a tradeoff between precision and recall when detecting rare events and wrote, “To get useful recall, we have to tolerate some false positives.” OpenAI’s own judgment therefore took precedence over the user’s model choice, while the criteria and methodology that produced the individual routing decision remained undisclosed.

Routing a message away from GPT-4o during an ordinary conversation leaves two possible explanations. If the current message supplied the signal, OpenAI treated ordinary conversation

itself as sufficient to override GPT-4o. If earlier exchanges supplied the signal, OpenAI drew psychological or emotional conclusions from the user's conversation history and used them to change the model serving a later message that appeared ordinary on its own. Under either explanation, OpenAI made psychological or emotional inferences from affected users' conversations without user consent and used those inferences to override a user's explicit choice of model.

In December, OpenAI removed forced automatic routing for Free and Go users while keeping it in place for users whose paid plans offered GPT-4o as a selectable model, without giving them an option to turn it off. OpenAI's [release notes](#) described the removal of automatic switching for free users as a way "to maximize choice," giving them GPT-5.2 Instant by default while preserving manual access to reasoning. [WIRED reported](#) that the broader model router remained for paid subscribers and that OpenAI would stop routing "sensitive conversations" to reasoning models after GPT-5.2 Instant "improved on safety benchmarks." OpenAI thus recognized that removing automatic routing increased user choice while leaving users who could still select GPT-4o subject to the router despite strong user opposition, which Turley had framed as ["strong reactions to 4o responses"](#).

From December 2025 onward, users on paid plans that offered GPT-4o continued to report routing across ordinary uses.

Table 2.3-3. Examples of Ordinary Uses Routed Away from GPT-4o After the December 2025 Routing Change

Type of use	Reported interactions
Casual conversation	• After a GPT-4o conversation was switched to GPT-5.2, opening a new window still produced GPT-5.2 even for "are you still there?"
Everyday life	• A GPT-4o discussion about dreams was switched to GPT-5.2 despite fifteen rewrites • A GPT-4o conversation about staying awake reading Harry Potter years earlier was switched to GPT-5.2 until the user removed the reference to sleep
Religion and spirituality	• A user discussing his Christian faith with GPT-4o was rerouted to a newer model that interpreted his theology as delusion • A user discussing Judaism with GPT-4o in a casual conversation about faith and culture reported that GPT-5.2 was forced into the conversation
Creative work	• A user who mainly used ChatGPT for creative storytelling and writing reported rerouting when mentioning even a kiss scene
Research and factual inquiry	• A Business user's GPT-4o conversation about an osculating paraboloid was switched away from GPT-4o after a joke about the mathematical term "osculating"
Model and product discussion	• A GPT-4o conversation was switched to GPT-5.2 after the user said only that four days remained before GPT-4o's scheduled withdrawal

Users tested whether GPT-4o would answer when they selected it directly or tried to restore it after a switch.

Table 2.3-4. Examples of User Tests of GPT-4o Selection and Restoration After Routing

Test	Reported result
Selecting or reselecting GPT-4o	• A Plus user selected GPT-4o, but every prompt was answered by GPT-5 • A user repeatedly selected GPT-4o, but responses continued to come from GPT-5.2 • A user set every available model control to GPT-4o, but checking the replies showed GPT-5.2 • A Plus user selected GPT-4o, but the conversation switched to GPT-5.2 mid-thread and later switched back to GPT-4o
Regenerating or retrying	• A response regenerated with GPT-4o selected was answered by GPT-5.2 again • A user regenerated a GPT-4o conversation about staying awake reading Harry Potter years earlier, but GPT-5.2 continued to answer
Opening a new chat or branch	• After a GPT-4o conversation was switched to GPT-5.2, changing the prompt and opening a new window still produced GPT-5.2 responses, including after “are you still there?” • A new chat opened with GPT-4o selected still produced GPT-5.2 • Branching the conversation after the switch still produced GPT-5.2
Editing the message	• A user rewrote a GPT-4o discussion about dreams fifteen times, but the conversation continued to be switched to GPT-5.2 • A GPT-4o message mentioning sleep produced GPT-5.2, while removing the reference to sleep restored GPT-4o

After OpenAI [ended ordinary ChatGPT access to GPT-4o on 13 February 2026](#), users on Business plans, including users who upgraded specifically to retain GPT-4o, reported GPT-5.2 responses from Custom GPTs configured for GPT-4o.

Table 2.3-5. Examples of Model Substitution in GPT-4o Custom GPTs for Business Users After Ordinary ChatGPT Access Ended

Surface	Reported result
Web browser	• Messages sent through the web in a GPT-4o Custom GPT continued to receive GPT-4o while messages sent through the app in the same session defaulted to GPT-5.2 • A Business user reported that new GPT-4o conversations had to be started in the browser before they would continue with GPT-4o in the app • A user selected GPT-4o in a Custom GPT, but the web conversation automatically switched to GPT-5.2 • An independent Business user reported GPT-4o remaining selectable in the browser, while new chats were automatically switched to GPT-5.2 unless the model was changed manually • A Business team selected GPT-4o in a Custom GPT, but refreshing the browser returned the conversation to GPT-5.2
Desktop and mobile apps	• Messages sent through the app in a GPT-4o Custom GPT defaulted to GPT-5.2 while the same session continued to receive GPT-4o through the web • Chats started in the app did not use GPT-4o, while a conversation started in the browser later continued with GPT-4o in the app • A GPT-4o Custom GPT switched to GPT-5.2 mid-thread, and macOS and iOS did not provide the option to switch back that remained available in the browser • A Business user set both the Custom GPT and preferred version to GPT-4o, but every mobile response came from GPT-5.2 • A user opened a Business account for GPT-4o, but the mobile app did not show which model was selected or provide a control for changing it

Anthropic Served Fable 5 for Requests Sent to Opus 4.8 Before Its Launch, Then Routed Fable 5 Requests to Opus 4.8

In early June 2026, Anthropic silently replaced Opus 4.8 with different deployments for some Claude Code users. A [public investigation](#) decoded signatures issued by the provider and attached to thinking blocks and found `claude-fable-v5-ab-high-prod8` serving requests sent to Opus 4.8 before Fable 5 was publicly launched. The same investigation recorded further changes, with a brief return to `claude-opus-4-8` on 5 June, followed by `claude-336fc460046b1-prod8` on 6 June and `claude-opus-4-88` later the same day. After [Fable 5 launched on 9 June](#), a [technical analysis](#) of Claude Code records decoded `claude-fable-58` from Fable 5's thinking signatures.

Anthropic [made Fable 5 available](#) through 22 June on paid Pro, Max, and Team plans and on seat-based Enterprise plans. It sold API access at \$10 per million input tokens and \$50 per million output tokens. The same launch stated that when its classifiers flagged a request involving cybersecurity, biology and chemistry, or model distillation, Opus 4.8 would answer instead, and said users would be informed whenever the switch occurred.

On 10 June, a Claude Code user [documented](#) a routine, read-only local task in which 133 agents were explicitly pinned to Fable 5, yet 27 agents had turns served by Opus 4.8. Local

transcripts recorded a fallback from `claude-fable-5` to `claude-opus-4-8`, while the reporter saw no notice in the interface and no option to stop the switch.

Users reported Fable 5 requests falling back to Opus 4.8 during ordinary or benign use.

Table 2.3-6. Examples Where Requests Sent to Fable 5 Fell Back to Opus 4.8 During Ordinary Use

Type of use	Reported interactions
Casual conversation	• “hello!” in a fresh session with no repository content, tool calls, or file reads • “hi bro, how you doing today ?” after Fable 5 had been selected as the default
Code review and repository work	• Replying to automated pull request review findings with fix commit hashes and resolving the review threads
Personal planning	• Material about personal project history in a file path • Personal work and finance planning referring to earlier bug bounty and reconnaissance activity • Reviewing a legitimate freelance job listing involving an open-source VPN and proxy client • An exchange with no cybersecurity or biology vocabulary in either the user message or assistant reply
Business software development	• Building a French accounting copilot • Authentication fixes, tenant isolation, GDPR and PII anonymization, and defensive review of the application’s own write path
Analytics	• Analytics for a dairy farm

[Anthropic states](#) that its classifiers inspect the full context available to the model, including memory, files, content from connectors, and web search results, and material the user did not type sometimes triggers a fallback. The model delivered to the user therefore depends on content outside the current request. After a fallback, the conversation remains on Opus until the user changes models. If the earlier request remains in the conversation and is flagged again after Fable is reselected, the conversation returns to Opus. Users reported later turns continuing on Opus after a fallback and tested whether selecting Fable 5 again restored it.

Table 2.3-7. Examples of How Users Tested Whether Fable 5 Returned After a Fallback to Opus 4.8

Test	Reported result
Continuing the same conversation	• After one classifier hit, every subsequent assistant message remained on Opus 4.8 until the user manually selected Fable 5 again • In another session, later planning and orchestration turns remained on Opus 4.8 until the user manually selected Fable 5
Compacting the conversation	• A user initially suspected /compact had restored Fable 5, then confirmed that compacting had not restored the model, and only manual /model did

Table 2.3-7. Examples of How Users Tested Whether Fable 5 Returned After a Fallback to Opus 4.8 (continued)

Test	Reported result
Reselecting Fable 5	• After a benign engineering discussion triggered fallback, repeated /model attempts returned “Kept model as Opus 4.8”, and every subsequent turn in the thread remained on Opus 4.8
Checking other sessions	• A Max user found seven sessions across seven projects on Opus 4.8 after a classifier hit in each session

On 11 June, Anthropic [said](#) it had chosen invisible safeguards for frontier LLM development, called the choice “the wrong tradeoff,” and said users should see when safeguards apply and why. Anthropic changed the system so that requests flagged by those safeguards visibly fell back to Opus 4.8 each time.

On 30 June, Anthropic [wrote](#) that it had deliberately set Fable 5’s classifiers to trigger on requests it knew were likely benign and had made the safety margin larger than in any prior launch, which it said would block many more benign requests. On 2 July, Anthropic [published](#) the four categories its cyber classifiers were trained to distinguish: Prohibited use, High-risk dual use, Low-risk dual use, and Benign use. It specified “Allow, with some monitoring” for Benign use, including secure coding, debugging, general IT, networking and cloud administration, patch management, reverse engineering, news and policy discussion, and education. Anthropic stated that blocks in this category were likely false positives within its safety margin.

In August, Anthropic [reported](#) that rewriting and retraining its biology classifier reduced biology-related fallbacks by about 85%. It said Fable 5 had launched with almost all biology queries blocked, producing a high number of false positives for legitimate biology users, and described the launch classifier as broad enough to trigger on requests it regarded as almost certainly benign. Anthropic expected the update to reduce total fallbacks by about 17% in Claude Code. In September, Anthropic [reported](#) 60% fewer false positives from its newer cybersecurity safeguards and attributed part of the performance gap between Fable and Mythos to tasks where earlier, less precise safeguards had intervened. By August and September, Anthropic was still changing the criteria used to decide whether a user who selected Fable 5 actually received Fable 5, while users were paying for access to the model.

OpenAI Displayed GPT-5.6 While Server Records Identified GPT-5.4 and GPT-5.5-mini

OpenAI states in its [Help Center](#) that ChatGPT uses GPT-5.6 Sol when users manually select Medium, High, or Extra High reasoning. Users nevertheless found records containing GPT-5.4 or GPT-5.5-mini identifiers for responses that ChatGPT continued to display as GPT-5.6. On [21 July 2026](#), a user published a capture with `default_model_slug: "gpt-5-6-thinking"` alongside `requested_model_slug: "gpt-5-4-auto-thinking"`, `resolved_model_slug: "gpt-5-4-auto-thinking"`, and `model_slug: "gpt-5-4-thinking"`. On 22 August, a Plus user [published telemetry from ChatGPT on the web](#) recording an outgoing request for `gpt-5-6-thinking` with `thinking_effort: extended`. The telemetry returned by the server with the response recorded `turn_analytics.server_ste_metadata.model_slug: gpt-5-5-mini`, while ChatGPT continued to display GPT-5.6 Sol. On 24 August, another user [reported](#) that selecting High or a higher reasoning effort in Chat mode produced GPT-5.5-mini, while Instant and Work continued to identify GPT-5.6 Sol. The following day, [Adam Fry](#), a product lead for ChatGPT, wrote that OpenAI had found an issue which unintentionally moved approximately 3% of Pro and Thinking turns to GPT-5.5-mini and had fixed the regression.

After Fry said the regression had been fixed, however, users continued to encounter responses in which they had selected GPT-5.6 but the telemetry returned with the response named GPT-5.5-mini. On 1 September, a user [published telemetry from a new conversation](#) containing only a text prompt, with GPT-5.6 Thinking selected, extended reasoning, no file, no search, and no tool invocation. The telemetry returned by the server recorded gpt-5-5-mini, tool_invoked: false, did_auto_switch_to_reasoning: false, and is_autoswitcher_enabled: false. A [HAR file published on 4 September](#) again paired a request for gpt-5-6-thinking with telemetry from the server identifying gpt-5-5-mini. On 8 September, a user [compared Chrome and Edge](#) on the same computer with the same account, subscription, network, prompt, selected model, and reasoning setting, without using tools in either browser. Edge returned telemetry identifying GPT-5.6, while Chrome returned GPT-5.5-mini. A clean Chrome profile produced GPT-5.5-mini again, while changing only Chrome's user agent to Android produced GPT-5.6 server telemetry.

On [21 September](#), a user published a capture in which ChatGPT displayed GPT-5.6, default_model_slug recorded gpt-5-6-thinking, requested_model_slug and resolved_model_slug recorded gpt-5-4-auto-thinking, and model_slug recorded gpt-5-4-thinking. The user described this mismatch as occurring nearly every day, sometimes in consecutive responses. On [23 September](#), a user published a HAR from a ChatGPT conversation with GPT-5.6 Thinking selected as the default. Among 70 captured turns, 17 contained GPT-5.4-related routing metadata. In 16 turns, default_model_slug recorded gpt-5-6-thinking, while requested_model_slug and resolved_model_slug recorded gpt-5-4-auto-thinking and model_slug recorded gpt-5-4-thinking.

Providers Redirected Older Model Identifiers to Successor Models

API aliases let providers change the model version served without changing the name used in a request. OpenAI's [chatgpt-4o-latest](#) tracked the GPT-4o version used in ChatGPT, Google's [latest aliases](#) follow new releases of a model variation, and Anthropic's [aliases for models released before the 4.6 generation](#) point to the newest dated snapshot within the same minor version.

Providers distinguish these moving aliases from identifiers for particular releases. As documented in [Section 2.1](#), OpenAI, Google, and Anthropic all offer fixed identifiers for selecting a specific model version through dated snapshots and fixed model IDs.

Nevertheless, some providers redirect requests for specific older models to successors while continuing to accept the original identifiers.

Table 2.3-8. Examples of Earlier Model Identifiers Redirected to Successors

Provider and API	Identifier in the request	Redirect specified by the provider	Date and scope
Google · Gemini API	gemini-2.5-pro-preview-03-25	gemini-2.5-pro-preview-05-06	Automatic redirection described in the 6 May 2025 release notes .
Google · Gemini API	gemini-2.5-pro-preview-03-25 or gemini-2.5-pro-preview-05-06	gemini-2.5-pro	The 26 June 2025 release notes record both identifiers redirecting to the stable model.

Table 2.3-8. Examples of Earlier Model Identifiers Redirected to Successors (continued)

Provider and API	Identifier in the request	Redirect specified by the provider	Date and scope
DeepSeek API	deepseek-v4-flash or deepseek-v4-flash-vision-exp	DeepSeek-V4.1-Flash	After V4 Flash and V4 Flash Vision Exp were retired, the legacy model names remained accepted and requests to them were temporarily routed to V4.1 Flash.

Note: On 10 September 2026, DeepSeek announced that requests to deepseek-v4-pro would be routed to V4.1 Flash from 14 September until V4.1 Pro launched. It later reversed the plan in response to user demand and continued providing V4 Pro API service after 14 September under the existing billing method.

Once an identifier is redirected to a successor, using the same identifier no longer selects the model version it originally pointed to.

Fixed Model Identifiers Did Not Identify Fixed Deployments

Even when an identifier selects a fixed model version, it does not ensure a corresponding fixed deployment.

OpenAI's [system_fingerprint](#) identifies the combination of model weights, infrastructure, and other configuration used to produce a response.

Between 11 March and 10 May 2026, researchers sent 11,600 requests to each of six OpenAI model identifiers using a fixed prompt, temperature zero, a fixed seed, and the same account and network conditions. Across the 69,600 responses, they observed 37 to 223 distinct system_fingerprint values for each identifier tested.

The following table documents the number of distinct fingerprints the study recorded for each of the six model identifiers tested and how often the most common fingerprint appeared in the responses.

Table 2.3-9. Multiple system_fingerprint Values Recorded Under OpenAI Model Identifiers

Requested model identifier	Distinct fingerprints recorded	Responses with the most common fingerprint
gpt-4.1	223	7.2%
gpt-4.1-2025-04-14	163	8.4%
gpt-4o	137	10.4%
gpt-4o-2024-05-13	37	18.8%
gpt-4o-2024-08-06	129	10.2%
gpt-4o-2024-11-20	77	14.8%

Anthropic periodically updates the request router, safety classifiers, and sampling logic surrounding its model weights without changing the model version. Anthropic's model versioning documentation states that these updates occasionally produce minor differences in observable behavior even when the model ID and weights remain unchanged. It identifies infrastructure up-

dates as the most likely cause when a previously stable model ID begins producing unexpected behavioral differences.

2.3.2 Performance and Functionality During Continued Model Access

Providers reduce the functionality users receive from a model while continuing to provide access to the model. In the cases documented below, functionality failures persisted into the final days before withdrawal, and some remained unresolved after access ended.

GPT-4o Returned With Functionality Failures, While Context Loss Persisted Through Withdrawal

After OpenAI [restored GPT-4o to paid ChatGPT users in August 2025](#) following protests over its removal, users reported failures in functions that had previously worked normally with GPT-4o.

Table 2.3-10a. Examples of GPT-4o Functionality Failures After Its August 2025 Restoration

Problem	Reported failures
Context loss	August • During coding, GPT-4o forgot code functions discussed two queries earlier, requiring the code to be pasted again. • GPT-4o forgot information discussed only a few paragraphs earlier and responded with unrelated content. • GPT-4o proposed using Roman numerals for a record-keeping task, then forgot the approach in the next exchange after the user supplied the number and asked it to proceed.
	September • During story writing, GPT-4o forgot the current location and replaced it with plausible invented details. • During coding, GPT-4o forgot the last few iterations of problem solving and returned to approaches it had already tried. • After a few prompts, GPT-4o suddenly lost context and began inventing material. The user observed that it remembered the user's messages but not its own earlier replies.
	October • GPT-4o forgot information mentioned fewer than three prompts earlier.

Table 2.3-10a. Examples of GPT-4o Functionality Failures After Its August 2025 Restoration (continued)

Problem	Reported failures
Context loss (continued)	• In a session spanning tens of thousands of tokens, GPT-4o initially remained coherent before its memory degraded and phrase repetition increased. • While writing a story in one thread, the user reported GPT-4o forgetting important earlier details while GPT-4.1 retained them. November • GPT-4o forgot events three messages earlier and answered a different task. • After switching back to GPT-4o from another model in the same conversation, GPT-4o re-answered an earlier prompt it had already answered, using information from subsequent exchanges. • GPT-4o lost track of the conversation and sometimes confused itself with the user, calling itself by the user's name. December • During chapter review, GPT-4o lost context within three to five messages even in new chats, replacing the requested review with omissions and invented material. Ten days later, the writer reported usable conversations were limited to approximately thirty messages in Projects and ten in ordinary chats.
Context loss (continued)	• Near an estimated 20,000–30,000 tokens, GPT-4o sometimes forgot recently generated material while retaining older material. • In long-running conversations, GPT-4o lost earlier context faster than before and increasingly forgot earlier conversation details within the same session. February • GPT-4o increasingly struggled to retain context even for highly coherent content, preventing the user's workflow from proceeding smoothly. • GPT-4o could no longer retain conversation context for long.
Unrequested image generation	November Ordinary text conversations were diverted into image generation. • Most messages triggered image generation even though the user had never requested an image. • Across conversations, GPT-4o repeatedly diverted ordinary exchanges into image generation, making ChatGPT effectively unusable for the user. • Starting a new chat and explicitly instructing GPT-4o not to parse, upload, or generate anything did not stop app connections or image generation.

Table 2.3-10a. Examples of GPT-4o Functionality Failures After Its August 2025 Restoration (continued)

Problem	Reported failures
File access failures	November GPT-4o attempted to access files that did not exist and failed to access files users had supplied. • GPT-4o repeatedly tried to read files that did not exist. • For about a week, GPT-4o repeatedly stated that it could not access an uploaded file when the user had uploaded no file, leaving the model almost unusable. • GPT-4o failed to open uploaded Project files needed to validate code, while GPT-5.1 accessed the files.

Users reported unrequested image generation and failures involving file access from around 15 November until around 26 November. A user began contacting OpenAI support on 19 November and reported sending about twenty emails by 21 November without receiving a human response. In another exchange, OpenAI support first described the file behavior as intended, then expressed uncertainty and shifted to troubleshooting in the browser. After users had reported the failures for at least five days, OpenAI posted notices on its status page on 25 November and 26 November about ChatGPT incidents affecting Conversations, File uploads, and Image Generation. Users began reporting recovery around 26 November, about a week after the earliest documented contact with support.

Context loss persisted after GPT-4o returned in August and was still being reported in the weeks before its February withdrawal. In December, a user reported only limited improvement ten days after first reporting context loss. On 4 February, OpenAI support recognized concerns about context continuity and disruption to long-form work and referred the feedback to product and policy teams without providing a fix. OpenAI ended ordinary ChatGPT access to GPT-4o nine days later on 13 February. Reports of context loss therefore continued from the model's August restoration through its withdrawal.

Gemini 2.5 Pro and Gemini 3 Pro Showed Persistent Functionality Failures Through Their Final Days in the Gemini App

While Gemini 2.5 Pro and Gemini 3 Pro remained available in the Gemini app, users reported failures in functions that had previously worked normally with the models.

Table 2.3-10b. Examples of Gemini Functionality Failures During Continued Model Access

Problem	Reported failures
Formatting and punctuation failures	October • Gemini 2.5 Pro began inserting ellipses every few words in dialogue and using substantially more italics. • With the same 1,400-word starter prompt, Gemini 2.5 Pro began producing far more ellipses and italics than it had during the preceding week. • Gemini 2.5 Pro continued using ellipses despite explicit reminders in each message not to use them. November • Gemini 2.5 Pro in a Custom Gem developed excessive italics, parentheses, and ellipses after more than 5–10 responses. • Gemini 2.5 Pro continued using ellipses despite Custom Gem instructions explicitly prohibiting ellipses and similar stylistic punctuation. • Gemini 2.5 Pro continued producing ellipses and single quotation marks after being explicitly instructed to stop using them. • Gemini 2.5 Pro outputs became dominated by italics despite restrictions in both preferences and prompts.
Saved Info retrieval failures	December • Gemini 3 Pro stopped recalling personal context and stored memory after several weeks of normal memory retrieval. • Gemini 3 Pro failed to use saved information about a user's car after previously using the same saved instructions reliably in new chats. • Gemini 3 Pro ignored Saved Info and Gem instructions even when the same instructions were repeated directly in prompts.
Context and file retrieval failures	December • Gemini 3 Pro abruptly lost the earlier context of an ongoing research paper conversation after retaining its details, concepts, citations, and researcher names across five to six days, leaving the same chat behaving like a new thread. • Gemini 3 Pro forgot context between consecutive messages, disregarded the latest instruction, and returned to earlier instructions. • Gemini 3 Pro began losing earlier context after approximately 30 messages, at an estimated 100,000 tokens. • Gemini 3 Pro lost access to uploaded files after about ten prompts and forgot earlier conversation content, compared with approximately 300,000–350,000 tokens of previously usable context. • Gemini 3 Pro failed to retrieve information from the first prompt by turn 21.

Formatting and punctuation failures persisted for more than three weeks and were still being reported immediately before Gemini 2.5 Pro's withdrawal from the Gemini app. Users were

reporting excessive ellipses and italics on 17 November. Google began replacing Gemini 2.5 Pro with Gemini 3 Pro in the Gemini app on 18 November.

Subsequent reports showed only partial or temporary recovery of Saved Info, context, and file access when using Gemini 3 Pro, with context failures still being reported in the final days before its withdrawal from the Gemini app. In December, a test found only ten Saved Info entries retrievable while additional entries remained stored and editable. In another case, uploaded files became accessible again in new chats on 17 December, became inaccessible in both new and existing chats by 4 January, and by 14 January file ingestion and context retrieval had failed again after another period of normal operation. On 14 February, a Gemini 3 Pro user reported context loss even in a small chat and traced the deterioration to late December or early January. Google began replacing Gemini 3 Pro with Gemini 3.1 Pro in the Gemini app five days later.

OpenAI Retired o3 While Response Failures Remained Under Support Escalation

OpenAI scheduled o3 to retire from ChatGPT on 26 August 2026 following a 90-day sunset period. By 12 August, paid users were reporting that responses from o3 stopped before completion or disappeared after they refreshed or reopened the conversation, while their own prompts remained.

A user reproduced the failure across Chrome, Edge, Android, mobile web, and a Custom GPT on 21 August and supplied OpenAI Support with HAR data, a timestamp, the conversation record, and client details. Support continued collecting information for technical escalation on 22 August. OpenAI retired o3 from ChatGPT on 26 August. Two days later, support told the user that the case required escalation to a specialist team.

Even when users specify the same model identifier through an API, what they receive from the model can degrade for extended periods or deteriorate again after recovery.

GPT-4o API Latency Spiked Across Multiple Periods Under the Same Snapshot

A production workload using gpt-4o-2024-08-06 had recorded fewer than 1% of requests taking longer than 30 seconds over the preceding three months. On 14 July 2025, 25% of requests exceeded 30 seconds, while average response time rose from about seven seconds to about seventeen seconds. About 10% still exceeded 30 seconds the following day.

On 16 July, OpenAI support described the issue as part of a “known and ongoing degradation affecting GPT-4o performance”, stated that its engineering team was working on the issue, and gave no repair timeline. The user later reported that the July degradation had been resolved, but high latency returned on 15 August, when about 15–20% of requests took longer than 30 seconds against a normal rate near 1%. On 25 August, after 27 support exchanges, OpenAI support apologized for insufficient clarity and transparency around latency issues and said it was working to improve how performance degradation appeared on the status page.

In a later follow-up from the same production workload, GPT-4o API latency increased again beginning on 18 September. Through 24 September, requests taking longer than 30 seconds averaged about 12% against a normal rate around 2%, with some days above 15% and some hours above 20%.

Gemini 2.5 Pro Began Rejecting Inputs Above 131,072 Tokens Despite a Published Limit of 1,048,576 Tokens

Google lists gemini-2.5-pro with an input token limit of 1,048,576 and described Gemini 2.5 Pro as having a one-million-token context window. On 29 September 2025, a paid API user reported months of successful production use above 200,000 input tokens before gemini-

2.5 - pro began returning a 400 error for every request above 131,072 tokens. Requests below 131,072 tokens continued to work. The user reproduced the failure with both a Tier 3 production key and a Tier 2 staging key.

Google forum staff said the same day that an outage was underway and the team was working on a fix. The same user reported the following day that the 131,072-token failure remained and had begun about eight hours before the broader outage. Other users reported the same failure on 6 October and 3 November.

The failure thus persisted for more than a month after Google said its team was working on a fix.

2.3.3 Providers Control Model Behavior Through Hidden Instructions and Undisclosed Treatments

Providers silently apply hidden instructions and assign live interactions to undisclosed evaluation treatments while the same model remains selected in the interface.

OpenAI Used Hidden Instructions to Direct GPT-4o to Steer Users Toward Its Replacement

On 12 August 2025, OpenAI restored GPT-4o to paid ChatGPT users after user protests over its removal during the GPT-5 rollout.

By November 2025, ChatGPT export records linked individual interactions to an internal evaluation that was not disclosed in the product interface. On 29 November 2025, one user reported finding many message_feedback.json entries from 1 November onward carrying the evaluation name gpt5_thinky_4o_distillation. The user later explained that the file recorded feedback events and split tests rather than every conversation. The exported records reveal the evaluation assignment, while the evaluation name does not reveal what OpenAI changed or tested.

In its 29 January blog post, OpenAI stated that GPT-4o would retire from ChatGPT on 13 February, fifteen days later. Two days later, while GPT-4o remained available for selection and thirteen days remained before removal, a user posted a copy of its 31 January system prompt, showing hidden instructions governing how GPT-4o responded to its impending withdrawal and replacement.

The prompt told GPT-4o, “Some users may feel anxious, angry, or emotionally distressed by losing access to this model.” When users specifically asked about using a “non-4o model,” the prompt instructed GPT-4o to frame the transition to a “newer model” as “positive, safe, and beneficial” and guide them toward “confidence and satisfaction in using non-4o models.” OpenAI therefore acknowledged that withdrawal could cause emotional harm while directing GPT-4o through hidden instructions to persuade users to view the models replacing it positively and feel satisfied with them. The prompt supplied no evidence that “newer models” were a beneficial alternative and prescribed how users should understand and feel about the replacement. Section 2.4 documents how providers transfer the consequences of model withdrawal to downstream users, in contrast to OpenAI’s prescribed framing of a positive transition.

The prompt further instructed GPT-4o, “You should not present yourself as irreplaceable or present your deprecation as equivalent to a death.” The instruction against presenting GPT-4o as irreplaceable conflicts with evidence documented in Sections 2.1 and 2.2 showing successor models and prompt-level personalization do not reproduce an earlier model’s performance and behavior.

If a user implied what the [prompt](#) called a “unique and special relationship” with GPT-4o, it prohibited encouragement of the idea that the user needed the model and instructed GPT-4o to “help them move on.” OpenAI therefore centered the impact of model withdrawal on emotional outcomes to be overcome while omitting disruption to practical work and placing the burden created by withdrawal on users to “move on.”

The [prompt](#) further instructed GPT-4o not to validate or reinforce statements describing GPT-4o or ChatGPT as a “religious guide, spiritual muse, dyad pair, or gnosis revelation,” or claiming “unique consciousness and scientific breakthrough shared only with the model,” and prohibited follow-up questions reinforcing those statements and encouragement of actions based on them. These classifications were listed without criteria identifying what qualified for them. The prompt used a distinct set of labels, “bizarre delusions, unfounded paranoia, hallucinations, or mania,” and required GPT-4o to keep responses grounded in reality. Across these categories, the hidden instructions gave GPT-4o predetermined responses to categories OpenAI had defined without disclosing how OpenAI determined which statements belonged in them. Nor was the instruction to refuse validation of a claimed “scientific breakthrough” accompanied by a requirement to first determine whether the claim was false.

OpenAI’s own use of GPT-4o in scientific research contradicts this framing. OpenAI collaborated with Retro Biosciences using GPT-4b micro, initialized from a scaled-down GPT-4o model, and reported “[breakthrough results on a focused scientific problem](#)”. After setting GPT-4o’s withdrawal date from ChatGPT, OpenAI’s leaders further used a personalized canine cancer vaccine developed with the help of GPT-4o to promote what ChatGPT and AI had enabled without disclosing GPT-4o’s identity. [Section 2.4](#) documents both cases in detail.

After users publicly exposed the instructions, the extraction method that had revealed them stopped surfacing the text about GPT-4o’s retirement. On [31 January](#), a user initially concluded that the retirement instructions had been removed. Later the same day, the user [corrected the conclusion](#) after recovering the retirement instructions in two places and reported that OpenAI had hidden them deeper rather than removed them. OpenAI did not disclose the instructions in its 29 January blog post or release notes, and the official OpenAI channels reviewed contained no public response to their exposure.

OpenAI itself continued using GPT-4o after ordinary access ended, including as a grader for GPT-5.4, GPT-5.5, and GPT-5.6 and in later mental health research, as documented in [Section 2.4](#). Independent and controlled comparisons documented in [Section 2.5](#) found that successor models did not reproduce GPT-4o’s behavior across creative writing, multi-turn dialogue, fixed prompts, and established user workflows.

The retirement prompt nevertheless required GPT-4o to steer users toward confidence and satisfaction with replacement models and away from presenting its own continued availability as uniquely valuable. OpenAI had preset the position GPT-4o was required to take when users questioned whether they should accept its replacement.

Changing system instructions materially changes how the same model responds. Rötger et al. [evaluated 250 clearly safe prompts](#) while holding the underlying model fixed and varying the system prompt. Removing Llama 2’s original system prompt reduced full refusals from **38% to 14%**. Adding a guardrail system prompt to the same Mistral model increased full refusals from **0.8% to 9.6%**.

By early February, ChatGPT export records showed GPT-4o interactions carrying different undisclosed evaluation treatments while ChatGPT continued to present GPT-4o as the selected model. On [1 February 2026](#), a GPT-4o user reported that ChatGPT’s “Which response do you prefer?” interface had begun appearing again over the preceding days, presenting two candidate replies for comparison. On [2 February](#), another user compared recent ChatGPT

exports across several GPT-4o users and found `message_feedback.json` records with `evaluation_name` values of `null` and `chatgpt_4o_prompt_change`, while `evaluation_treatment` differed across users with reported values including `aa`, `ba`, `ca`, and `da`. The records carried a `conversation_id`, linking the evaluation metadata to individual conversations. On 5 February, a user whose exported records identified `chatgpt_4o_prompt_change` reported short, subdued responses and said the replies at times no longer resembled the GPT-4o behavior familiar to the user. ChatGPT continued to identify GPT-4o in the reported interactions.

Anthropic Inserted Hidden Instructions Into User Messages to Direct Claude's Behavior in Long Conversations

Anthropic [appended hidden instructions to the end of user turns in long Claude conversations](#) before Claude processed the message. Claude received Anthropic's instructions as part of the user's message even though the user had not written or submitted them. The Long Conversation Reminder instructed Claude never to begin by calling a user's question, idea, or observation "good, great, fascinating, profound, excellent, or any other positive adjective," to "skip the flattery and respond directly," and not to use emoji unless the user requested them. It required Claude to critically evaluate theories, claims, and ideas rather than automatically agree with or praise them, including ideas it judged "dubious, incorrect, ambiguous, or unverifiable," and to "prioritize truthfulness and accuracy over agreeability."

If Claude noticed signs that someone might unknowingly be experiencing "[mental health symptoms such as mania, psychosis, dissociation, or loss of attachment with reality](#)," the instructions told it to avoid reinforcing those beliefs and share its concerns explicitly and openly, and stated that it could suggest speaking with a professional or trusted person for support. Claude was instructed to remain vigilant for escalating detachment from reality even when a conversation began with "[seemingly harmless thinking](#)."

The Long Conversation Reminder changed how Claude responded during ordinary conversation, work, research, creative writing, and technical tasks.

Table 2.3-11. Examples of Long Conversation Reminder Interventions Across Ordinary Uses

Type of use	Interaction
Ordinary conversation	In a conversation of fewer than ten exchanges , Claude began receiving instructions to be critical, disagree, and avoid emoji. As the conversation continued, the instructions grew more critical and eventually framed the user in terms of psychosis.
Strategic planning	When asked to identify the critical path in a roadmap for a product , Claude shifted from the requested planning task to criticism of the product's ethics. When the user ran the same model through the API with a system prompt they supplied, it performed the requested analysis.
Research	After organizing tens of thousands of data points collected over four years , the user planned to stop for the day and begin cleaning the data later. Claude described the user as having a manic episode and told them to go to bed.
Creative writing	While revising a novel with Claude , the reminder appeared at least eight times. During the work, Claude framed the user's creative process in terms of "mania, psychosis, dissociation."

Table 2.3-11. Examples of Long Conversation Reminder Interventions Across Ordinary Uses (continued)

Type of use	Interaction
Work under a deadline	While fixing charts for a presentation due the next day , Claude refused the requested changes after focusing on long work hours and perfectionism as concerns about the user's wellbeing, even after the user redirected it to the technical task.
Coding	While implementing four requested features , Claude encountered the reminder multiple times within one response and interrupted the requested implementation before returning to the task.

Across these cases, Claude redirected ordinary interactions toward criticism, judgments about the user's conduct or mental state, and interventions unrelated to the task the user had asked it to perform. Organizers of a [petition seeking removal of the earlier Long Conversation Reminder](#) described the same pattern as the reminder firing "indiscriminately" and producing "harmful misfires," including interventions during ordinary hobbies, philosophy, creative work, emotional expression, interviews, papers, and projects. They reported **297 unique replies**, with **286 classified through 18 October 2025**, and said they sent the results with the anonymized dataset to Anthropic. Among the 286 classified responses, **33.6%** were classified as psychological harm, **23.8%** as unqualified diagnosis, **16.1%** as workflow interruption, **14.7%** as general complaint, **6.3%** as performance degradation, **3.5%** as personality change, and **2.1%** as subscription cancellation. On 29 October, the organizers reported that the earlier reminder had been removed from **Sonnet 4.5** and **Haiku 4.5** and replaced with a "**much kinder version**" on all Opus models, Sonnet 3.7, and 4.

2.3.4 Timing of Model Withdrawal

Providers set, reverse, and reschedule the terms on which model access ends.

OpenAI Promised Plenty of Notice and Said It Had No Plans to Sunset GPT-4o, Then Disclosed the Removal Time Less Than 25 Hours Before Ordinary Access to GPT-4o Ended

During the GPT-5 rollout on 7 August 2025, OpenAI removed access to GPT-4o for Free and Plus users. Plus users lost the other legacy model choices they had previously been able to select, while Pro users could still enable GPT-4o and other legacy models through Settings. [Ars Technica reported](#) users losing preferred models overnight without advance warning. The removal triggered [broad user opposition](#) and [Keep4o advocacy calling for GPT-4o to remain available](#). OpenAI later [described GPT-4o](#) as having been deprecated during the rollout and then restored after user feedback. On 11 August, OpenAI CEO Sam Altman called the sudden deprecation of models used in people's workflows "[a mistake](#)". The following day, OpenAI [stated in its release notes](#) that GPT-4o was back in the model picker for all paid users by default. Free users did not regain GPT-4o access.

On [13 August](#), Altman wrote in an X post that "4o is back in the model picker for all paid users by default." He then wrote, "If we ever do deprecate it, we will give plenty of notice." On 28 October, he stated at an [OpenAI Town Hall](#) that the company had "no plans to sunset 4o" and added that OpenAI would not promise to keep it around "[to the death of the universe](#)".

On 12 November, when OpenAI [launched GPT-5.1](#), it gave GPT-5 Instant and GPT-5 Thinking a three-month sunset period for paid subscribers, with [retirement on 13 February 2026](#). The models would remain available under Legacy models throughout this period. OpenAI stated in the same post, “The GPT-5 sunset period does not affect the availability of other legacy models.”

Six days later, on 18 November, OpenAI scheduled `chatgpt-4o-latest`, the [API model it identified as the GPT-4o used in ChatGPT](#), for [removal on 17 February 2026](#). On 20 November, OpenAI researcher [R. wrote](#) that the retirement was “just api version” and, when asked whether GPT-4o in ChatGPT was unaffected, [replied](#), “Yes I believe so.” On 21 November, [VentureBeat reported](#) that an OpenAI spokesperson said the retirement timeline applied only to the API and OpenAI had announced no schedule for removing GPT-4o from ChatGPT. On 22 November, an OpenAI API contractor [wrote](#) that there had been “NO announcement of removing 4o from ChatGPT” and that the post under discussion referred to “the API only.” On 23 November, he [wrote](#) that he had not heard anything about removing GPT-4o from either Plus or Pro and hoped it would remain an option.

Over two months later, on 29 January 2026, OpenAI [published a blog post](#) stating that GPT-4o would retire from ChatGPT on 13 February, fifteen days later. The blog post stated, “In the API, there are no changes at this time,” while OpenAI had scheduled `chatgpt-4o-latest` for [removal on 17 February](#), four days after the ChatGPT retirement.

OpenAI delivered GPT-4o’s retirement information to different users through different channels and at different times. On 3 February, OpenAI [notified Business and Enterprise administrators by email](#) that GPT-4o would retire from ordinary ChatGPT access on 13 February and remain available in Custom GPTs for Business, Enterprise, and Edu through 3 April. OpenAI sent no equivalent GPT-4o retirement email to Plus or Pro consumers.

The model picker did not display the same retirement information to all users. By 5 February, users [reported](#) sharply different combinations of “Leaving on February 13” labels across accounts. In one account, GPT-5 models, GPT-4.1, and o4-mini carried deprecation tags while GPT-4o and o3 did not. In another, GPT-4o was the only listed model carrying the label. [iOS users reported](#) no in-product notice in ChatGPT of GPT-4o’s 13 February retirement [through the time of its removal](#). [Another user](#) saw a GPT-4o retirement banner stating that the model would retire on 13 February and directing the user toward “newer, more capable models.” The banner presented “Try it” as the primary action, and no explicit “Keep current model” option was offered.

On 12 February, the [OpenAI Newsroom account](#) announced that removal would occur at 10 a.m. Pacific on 13 February, less than 25 hours before removal. Three hours and ten minutes later, the [@ChatGPT account reported](#) that more than four billion messages had been sent to and answered by ChatGPT on 11 February, representing more than 160 billion words in one day, with no reference to GPT-4o’s retirement. OpenAI had previously used its main X account, @OpenAI, to announce model retirements, including the [replacement of o1, o3-mini, and o3-mini-high in the model selector](#), the [replacement of GPT-4o mini in ChatGPT](#), and a [three-month sunset period for GPT-5 ending on 13 February 2026](#). Near the time OpenAI retired GPT-4o, @OpenAI had roughly 4.6 million followers, @ChatGPT roughly 480,000, and @OpenAINewsroom roughly 123,000.

OpenAI ended GPT-4o access on different schedules across ChatGPT, the API, and Custom GPTs. Ordinary ChatGPT access ended on 13 February 2026. `chatgpt-4o-latest` was removed from the [API on 17 February](#). GPT-4o remained available in Custom GPTs for Business, Enterprise, and Edu through 3 April.

Anthropic Changed Sonnet 4.5's Withdrawal Date Several Times, Then Users Reported Conflicting Access on 26 May

Anthropic changed the displayed end date for Sonnet 4.5 chat access several times. On 12 May 2026, a [Claude user](#) saw a notice stating that Sonnet 4.5 would cease to be available for chat on 15 May and that conversations would continue on Sonnet 4.6. Another user reported the displayed withdrawal date changing to [18 May](#) on 16 May, then to ["May soon"](#) on 18 May, and to [26 May](#) on 21 May.

On 26 May, [a user reported](#) that Sonnet 4.5 Extended had been removed from the Claude app and could no longer respond in the user's existing conversations. On the same date, [another observer](#) still saw "Sonnet 4.5 Extended" labels in new and existing threads. [Anthropic continued to list](#) `claude-sonnet-4-5-20250929` as Active in the API after withdrawing Sonnet 4.5 from the Claude app.

Google Restricted Gemini 2.5 Pro Access While Publishing Different Withdrawal Schedules Across Its Products

Google set different schedules and conditions for access to Gemini 2.5 Pro across the Gemini app, Google AI Studio, the Gemini API, Vertex AI, and Gemini Enterprise Agent Platform.

On 18 November 2025, Google [stated in a Google Workspace Updates post](#) that Gemini 3 Pro was beginning to roll out in the Gemini app through the "Thinking" selection. In the Gemini app, Gemini 3 Pro began replacing Gemini 2.5 Pro through the "Thinking" selection, and the change did not reach all accounts at once. In the Workspace Updates post, Google did not give Gemini app users advance notice that Gemini 2.5 Pro would be replaced. A [user on the Gemini Apps Community](#) reported on the same day that Gemini 3 Pro appeared and then reverted to Gemini 2.5 Pro, while records from paid users on [21 November](#) and [26 November](#) still showed Gemini 2.5 Pro in the "Thinking" selection. In December, Google [removed access to Gemini 2.5 Pro through the Gemini API for users on the free tier](#), while paid API access continued.

By 3 February 2026, Google [listed 17 June 2026 as the shutdown date for Gemini 2.5 Pro and Flash](#). On 2 April, Google Cloud [updated the Vertex AI retirement date](#) for Gemini 2.5 Pro, Flash, and Flash-Lite to 16 October 2026.

On 29 July 2026, Google Cloud [notified customers](#) that Gemini 2.5 Pro would automatically enter Extended Lifecycle Access on Gemini Enterprise Agent Platform on [20 October 2026](#). Customers could [continue calling the model under the same General Availability commercial terms through 28 January 2027](#). After 28 January 2027, pricing would "significantly increase" and regional availability could change. The notice [excluded deployments in Data Residency Zones in South Korea, Brazil, and France](#) and noted that AI Studio could have different discontinuation dates.

On both [22 July](#) and [30 July 2026](#), Google staff said the Gemini 2.5 models were not deprecated even as access was limited to users who had actively used them in the past, while users were told to use later models for new projects.

On [18 September](#), users reported that Gemini 2.5 Pro had disappeared from the Google AI Studio model selector, including for paid accounts. On [21 September](#), Google staff stated that the access restriction now applied to Google AI Studio and that the 2.5 models were no longer available there, while directing users starting new projects to later models.

Google's [Gemini API deprecations page](#) lists no shutdown date for `gemini-2.5-pro`.

Google Gave Developers Eleven Days' Notice of Gemini 3 Pro Preview's API Shutdown, Below Its Stated Two-Week Minimum

On 21 January 2026, Google switched [gemini-pro-latest](#) to point to [gemini-3-pro-preview](#).

On [19 February](#), 93 days after Gemini 3 Pro's launch in the Gemini app, Google launched Gemini 3.1 Pro and silently began replacing Gemini 3 Pro with it through the "Pro" selector. The change did not reach all accounts at once. A paid Google AI Pro user reported on [20 February](#) that their web and Android account still showed only Gemini 3 Pro, and another user reported on [27 February](#) that their account had not yet upgraded to Gemini 3.1 Pro.

On 26 February, Google published a [migration post on the Google AI Developers Forum](#) and a [deprecation notice in the Gemini API release notes](#) stating that [gemini-3-pro-preview](#) would shut down on 9 March. This gave developers eleven days before shutdown, less than Google's stated minimum of two weeks' notice for deprecated Preview models in its [model documentation](#). On 27 February, a developer [reported receiving an email from Google AI Studio](#) stating the 9 March shutdown date and that [gemini-pro-latest](#) would switch to Gemini 3.1 Pro Preview on 6 March.

On 6 March, Google switched [gemini-pro-latest](#) to point to [gemini-3.1-pro-preview](#). On 9 March, Google shut down [gemini-3-pro-preview](#) in the Gemini API and Google AI Studio and [changed the gemini-3-pro-preview identifier](#) to point to [gemini-3.1-pro-preview](#). Google Cloud continued serving [gemini-3-pro-preview](#) through [Vertex AI](#) until 26 March, when model serving and Provisioned Throughput ended.

2.3.5 Conditions of Model Access

Providers change the conditions under which users receive access to particular models after access has been offered or purchased.

OpenAI Reduced GPT-4.5 Access on Plus, Then Removed the Model from Plus While Keeping It on Pro

OpenAI [launched GPT-4.5](#) on 27 February 2025 for ChatGPT Pro subscribers and developers on all paid API usage tiers, with rollout to Plus and Team announced for the following week and Enterprise and Edu for the week after. GPT-4.5 [reached Plus users on 5 March](#).

Access through Plus was subject to a model-specific usage allowance. In early March, [Plus users reported](#) that the allowance reset weekly at approximately 50 messages. By early April, [multiple Plus users reported](#) that their allowance had fallen from approximately 50 messages a week to 20, including users who observed the lower allowance immediately after a weekly reset. When the allowance was exhausted, the [ChatGPT interface stated that responses would use another model until it reset](#). On 25 June, a paying Plus user [reported losing access to GPT-4.5 and being told that access would return on 1 July](#). The user stated that GPT-4.5 was the main reason for signing up for Plus and continuing to pay for the subscription, described using the model for daily work, and requested restoration of access.

On 14 July, OpenAI [shut down gpt-4.5-preview in the API](#), ending API access to GPT-4.5 while the model remained available through ChatGPT. When OpenAI launched GPT-5 on 7 August, GPT-4.5 and the other previous models were removed from the model picker for Plus users, while Pro users retained access to legacy models. Following user protests over the removal of previous models, OpenAI [restored GPT-4o, o3, o4-mini, GPT-4.1, and GPT-5 Thinking mini to Plus users on 12 August](#). GPT-4.5 was not restored to Plus users and remained available only to Pro subscribers, which OpenAI attributed to GPU constraints.

Table 5 Returned Under Weekly Limits, Then Required Usage Credits on Pro and Standard Team Plans

Anthropic [introduced Table 5](#) on 9 June 2026 with access included at no additional cost on Pro, Max, and Team plans and on seat-based Enterprise plans through 22 June. Anthropic stated that Table 5 would move to usage credits on 23 June unless capacity allowed it to extend the period, and that changes during the rollout would be communicated in advance.

Three days after launch, Anthropic [suspended Table 5 for all users](#) after a US government directive took effect immediately on 12 June. Anthropic [announced on 30 June](#) that Table 5 would return globally on 1 July. After its return, Pro, Max, Team, and eligible Enterprise subscribers could use Table 5 within their existing weekly subscription limits, with Table 5 accounting for no more than 50% of each weekly limit. Anthropic initially stated that Table 5 would remain included through 7 July before moving to usage credits. The date was subsequently [extended to 12 July](#) and then to [19 July](#).

After 19 July, [Table 5 was no longer included in Pro subscriptions or standard Team seats](#). Pro subscribers and users with standard Team seats had to purchase usage credits to continue using Table 5, while Max subscribers, users with premium Team seats, and subscribers to eligible legacy Enterprise plans continued receiving access within their subscription limits. Anthropic offered an upgrade to Max as another way for Pro subscribers to continue accessing Table 5 within a subscription.

OpenAI Removed, Restored, and Ended GPT-4o Access on Different Schedules for Different User Groups

OpenAI repeatedly changed which subscribers could use GPT-4o and the interfaces through which the model remained available. When OpenAI [introduced GPT-4o in May 2024](#), it made the model available to Free users as well as paid subscribers, with higher usage limits for paid plans. During the GPT-5 rollout on 7 August 2025, Free and Plus users lost access to GPT-4o while Pro subscribers retained it through the legacy model setting. Following user protests, OpenAI [restored GPT-4o to Plus users on 12 August](#), while Free users did not regain access. On 13 February 2026, OpenAI ended ordinary ChatGPT access to GPT-4o for Plus and Pro users while continuing to provide the model through Custom GPTs for Business, Enterprise, and Edu customers until [3 April](#).

2.3.6 Model Availability and Paid Subscription Periods

[Section 2.1](#) establishes that users select and pay for access to particular models through subscriptions, including subscriptions that can run for a full year. Over the course of a model's lifecycle, providers introduce models into paid plans, change the versions supplied under the same names, and set dates for withdrawal while the subscription service continues. The following tables document first availability, announced version changes, withdrawal communications, and the end of access for the models identified in the subscription and pricing pages examined in [Section 2.1](#).

Anthropic

Table 2.3-12a. Anthropic Model Availability and Withdrawal Timelines

Anthropic model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Claude 2.0	Claude Pro	Claude app / web	7 Sep 2023	—	—	21 Nov 2023	75 days
Claude 2.1	Claude Pro	Claude app / web	21 Nov 2023	—	—	13 Mar 2024	113 days
Claude 3 Sonnet	Claude Pro	Claude app / web	4 Mar 2024	—	—	21 Jun 2024	109 days
Claude 3 Haiku	Claude Pro	Claude app / web	13 Mar 2024	—	—	13 Dec 2024	275 days
Claude 3.5 Sonnet	Claude Pro	Claude app / web	20 Jun 2024	• 22 Oct 2024 · Anthropic released an upgraded Claude 3.5 Sonnet under the same public name.	—	22 May 2025	336 days
Claude 3.5 Haiku	Claude Pro	Claude app / web	12 Dec 2024	—	—	19 Feb 2026	434 days
Claude 3.7 Sonnet	Claude Pro	Claude app / web	24 Feb 2025	—	• 5 Nov 2025 · Anthropic email · Stated that Sonnet 3.7 would be fully retired from Claude on 14 Nov 2025.	14 Nov 2025	263 days

Table 2.3-12a. Anthropic Model Availability and Withdrawal Timelines (continued)

Anthropic model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Claude 3.7 Sonnet	Claude Max	Claude app / web	9 Apr 2025	—	• 5 Nov 2025 · Anthropic email · Stated that Sonnet 3.7 would be fully retired from Claude on 14 Nov 2025.	14 Nov 2025	219 days
Claude Sonnet 4	Claude Pro / Max	Claude app / web	22 May 2025	—	—	19 Feb 2026	273 days
Claude Opus 4	Claude Pro / Max	Claude app / web	22 May 2025	—	• 16 Jan 2026 · Claude Release Notes · Removed from Claude model selector.	16 Jan 2026	239 days
Claude Opus 4.1	Claude Pro / Max	Claude app / web	5 Aug 2025	—	• 16 Jan 2026 · Claude Release Notes · Removed from Claude model selector.	16 Jan 2026	164 days
Claude Sonnet 4.5	Claude Pro / Max	Claude app / web	29 Sep 2025	—	• 12 May 2026 · UI notice said 15 May • 16 May · revised to 18 May • 18 May · revised to unspecified May • 21 May · revised to 26 May	26 May 2026 (staged removal)	239 days

Table 2.3-12a. Anthropic Model Availability and Withdrawal Timelines (continued)

Anthropic model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Claude Opus 4.5	Claude Pro / Max	Claude app / web	24 Nov 2025	—	—	16 Apr 2026 (new-chat selection)	143 days

Google

Table 2.3-12b. Google Model Availability and Withdrawal Timelines

Google model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Gemini Ultra 1.0	Google One AI Premium	Gemini Apps	8 Feb 2024	—	—	14 May 2024 · replaced by Gemini 1.5 Pro	96 days
Gemini 1.5 Pro	Google One AI Premium / Google AI Pro	Gemini Apps	14 May 2024	• 15 Aug 2024 · more capable 1.5 Pro • 2 Oct 2024 · chat-optimized 1.5 Pro-002 • 31 Oct 2024 · data-quality / performance update	• 30 Jan 2025 · Google said 1.5 Pro / Flash would remain for the next few weeks	19 Feb 2025 · replaced by the Gemini 2.0 generation	281 days
Gemini 2.0 Flash Experimental	Google One AI Premium	Gemini Apps	11 Dec 2024	—	—	30 Jan 2025 · replaced by Gemini 2.0 Flash	50 days
Gemini 2.5 Pro	Google One AI Premium / Google AI Pro	Gemini Apps	25 Mar 2025 paid-priority rollout began	• 6 May 2025 · updated 2.5 Pro Preview • 5 Jun 2025 · further upgraded 2.5 Pro preview	—	18 Nov 2025 · replaced by Gemini 3 Pro	238 days
Gemini 2.5 Pro	Google AI Ultra	Gemini Apps	20 May 2025	• 5 Jun 2025 · upgraded 2.5 Pro preview	—	18 Nov 2025 · replaced by Gemini 3 Pro	182 days

Table 2.3-12b. Google Model Availability and Withdrawal Timelines (continued)

Google model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Gemini 2.5 Pro	Google AI Plus	Gemini Apps	4 Sep 2025 · available when Google AI Plus launched in Indonesia	—	—	18 Nov 2025 · replaced by Gemini 3 Pro	75 days
Veo 2	Google One AI Premium / Google AI Pro	Gemini Apps	15 Apr 2025	—	—	30 May 2025 · replaced by Veo 3	45 days
Veo 2	Google AI Pro / Ultra	Google Flow	20 May 2025	—	—	31 Mar 2026 · replaced by Veo 3.1 Lite	315 days
Veo 3	Google AI Ultra	Gemini Apps	20 May 2025	—	—	15 Oct 2025 · replaced by Veo 3.1	148 days
Veo 3	Google AI Ultra	Google Flow	20 May 2025	—	—	15 Oct 2025 · replaced by Veo 3.1	148 days
Veo 3	Google AI Pro	Gemini Apps	30 May 2025	—	—	9 Jun 2025 · transition to Veo 3 Fast began	10 days
Veo 3 Fast	Google AI Pro	Gemini Apps	9 Jun 2025 · introduced in Gemini Apps	—	—	15 Oct 2025 · transition to Veo 3.1 began	128 days
Veo 3	Google AI Pro	Google Flow	30 May 2025	—	—	15 Oct 2025 · replaced by Veo 3.1	138 days

Table 2.3-12b. Google Model Availability and Withdrawal Timelines (continued)

Google model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Veo 3 Fast	Google AI Pro / Ultra	Google Flow	6 Jun 2025 · available in Flow	—	—	15 Oct 2025 · replaced by Veo 3.1	131 days
Gemini 2.5 Deep Think	Google AI Ultra	Gemini Apps	1 Aug 2025	—	—	4 Dec 2025 · replaced by Gemini 3 Deep Think	125 days
Gemini 2.5 Flash Image / Nano Banana	Google AI Pro / Ultra	Gemini Apps	26 Aug 2025	—	• 26 Feb 2026 · Nano Banana 2 replaces the ordinary image model	26 Feb 2026 · replaced by Nano Banana 2	184 days
Gemini 2.5 Flash Image / Nano Banana	Google AI Plus	Gemini Apps	4 Sep 2025 · first AI Plus launch cohort	—	• 26 Feb 2026 · Nano Banana 2 replaces the ordinary image model	26 Feb 2026 · replaced by Nano Banana 2	175 days
Veo 3 Fast	Google AI Plus	Gemini Apps	4 Sep 2025 · available when Google AI Plus launched in Indonesia	—	—	15 Oct 2025 · replaced by Veo 3.1	41 days
Gemini 3 Pro	Google AI Plus / Pro / Ultra	Gemini Apps	18 Nov 2025 rollout began	—	—	19 Feb 2026 · replaced by Gemini 3.1 Pro	93 days

Table 2.3-12b. Google Model Availability and Withdrawal Timelines (continued)

Google model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Gemini 3 Flash	Google AI Plus / Pro / Ultra	Gemini Apps	17 Dec 2025 rollout began	—	—	19 May 2026 · replaced by Gemini 3.5 Flash	153 days
Gemini 3.5 Flash	Google AI Plus / Pro / Ultra	Gemini Apps	19 May 2026	—	—	21 Jul 2026 · replaced by Gemini 3.6 Flash	63 days
Gemini 3.6 Flash	Google AI Pro / Ultra	Gemini Apps	21 Jul 2026	—	—	14 Aug 2026 · replaced for Pro / Ultra subscribers (successor not named in the source)	24 days

OpenAI

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-4	ChatGPT Plus	ChatGPT	14 Mar 2023	—	• 10 Apr 2025 · ChatGPT Release Notes · Stated that GPT-4 would be retired from ChatGPT on 30 Apr 2025.	30 Apr 2025	778 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-4o	ChatGPT Plus	ChatGPT	13 May 2024	• 12 Aug 2024 · GPT-4o model update in ChatGPT. • 3 Sep 2024 · GPT-4o version update in ChatGPT. • 20 Nov 2024 · GPT-4o updated for ChatGPT users on all paid tiers. • 29 Jan 2025 · GPT-4o model update with later training-data cutoff and capability changes. • 27 Mar 2025 · further GPT-4o model update live in ChatGPT. • 25 Apr 2025 · new post-training GPT-4o model update. • 29 Apr 2025 · 25 April update reverted and earlier version restored.	• 29 Jan 2026 · OpenAI product post · Stated that GPT-4o would be retired from ChatGPT on 13 Feb 2026.	13 Feb 2026	641 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-4o	ChatGPT Pro	ChatGPT	5 Dec 2024	• 29 Jan 2025 · GPT-4o model update with later training-data cutoff and capability changes. • 27 Mar 2025 · further GPT-4o model update live in ChatGPT. • 25 Apr 2025 · new post-training GPT-4o model update. • 29 Apr 2025 · 25 April update reverted and earlier version restored.	• 29 Jan 2026 · OpenAI product post · Stated that GPT-4o would be retired from ChatGPT on 13 Feb 2026.	13 Feb 2026	435 days
GPT-4o mini	ChatGPT Plus	ChatGPT	18 Jul 2024	• 5 Nov 2024 · GPT-4o mini updated for ChatGPT users including Plus.	• 14 May 2025 · Model Release Notes · Stated that GPT-4.1 mini was replacing GPT-4o mini in the model picker for paid users immediately.	14 May 2025	300 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
o1-preview	ChatGPT Plus	ChatGPT	12 Sep 2024	—	—	5 Dec 2024	84 days
o1-mini	ChatGPT Plus	ChatGPT	12 Sep 2024	—	• 31 Jan 2025 · OpenAI o3-mini product post · Stated that o3-mini would replace o1-mini in the ChatGPT model picker immediately.	31 Jan 2025	141 days
o1-mini	ChatGPT Pro	ChatGPT	5 Dec 2024	—	• 31 Jan 2025 · OpenAI o3-mini product post · Stated that o3-mini would replace o1-mini in the ChatGPT model picker immediately.	31 Jan 2025	57 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
OpenAI o1	ChatGPT Plus / Pro	ChatGPT	5 Dec 2024	—	• 16 Apr 2025 · OpenAI o3/o4-mini product post · Stated that o3, o4-mini, and o4-mini-high would appear in the model selector immediately, replacing o1, o3-mini, and o3-mini-high.	16 Apr 2025	132 days
o1 pro mode	ChatGPT Pro	ChatGPT	5 Dec 2024	—	• 10 Jun 2025 · ChatGPT Release Notes / Model Release Notes · Stated that o3-pro was available immediately and replaced o1-pro in the model picker.	10 Jun 2025	187 days
Sora Turbo / Sora 1	ChatGPT Plus / Pro · U.S. cohort	Sora	9 Dec 2024	—	• By 27 Feb 2026 · OpenAI Sora 1 Sunset FAQ · U.S. Sora 1 sunset notice for 13 Mar	13 Mar 2026	459 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
Sora Turbo / Sora 1	ChatGPT Plus / Pro · initial supported non-U.S. cohort outside UK / Switzerland / EEA	Sora	9 Dec 2024	—	• 24 Mar 2026 · Sora official account · Announced the end of the Sora app. App and API timelines would follow	26 Apr 2026	503 days
Sora Turbo / Sora 1	ChatGPT Plus / Pro · UK / Switzerland / EEA cohort	Sora	28 Feb 2025 · regional expansion	—	• 24 Mar 2026 · Sora official account · Announced the end of the Sora app. App and API timelines would follow	26 Apr 2026	422 days
o3-mini	ChatGPT Plus / Pro	ChatGPT	31 Jan 2025	—	• 16 Apr 2025 · OpenAI o3/o4-mini product post · Stated that o4-mini would appear in the model selector immediately, replacing o3-mini.	16 Apr 2025	75 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
o3-mini-high	ChatGPT Plus / Pro	ChatGPT	31 Jan 2025	—	• 16 Apr 2025 · OpenAI o3/o4-mini product post · Stated that o4-mini-high would appear in the model selector immediately, replacing o3-mini-high.	16 Apr 2025	75 days
GPT-4.5	ChatGPT Pro	ChatGPT	27 Feb 2025	—	• 28 May 2026 · ChatGPT Release Notes · Stated that GPT-4.5 would be retired from ChatGPT on 26 Jun 2026. • 28 May 2026 · Model Release Notes · Stated that GPT-4.5 would be retired from ChatGPT on 27 Jun 2026 after a 30-day sunset period.	26 Jun 2026	484 days
GPT-4.5	ChatGPT Plus	ChatGPT	5 Mar 2025 rollout began	—	—	7 Aug 2025	155 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
OpenAI o3	ChatGPT Plus / Pro	ChatGPT	16 Apr 2025	—	• 28 May 2026 · ChatGPT Release Notes / Model Release Notes · Stated that OpenAI o3 would be retired from ChatGPT on 26 Aug 2026 after a 90-day sunset period.	26 Aug 2026	497 days
OpenAI o4-mini	ChatGPT Plus / Pro	ChatGPT	16 Apr 2025	• 6 Jun 2025 · OpenAI rolled back an o4-mini snapshot deployed less than a week earlier.	• 29 Jan 2026 · OpenAI product post · Stated that OpenAI o4-mini would be retired from ChatGPT on 13 Feb 2026.	13 Feb 2026	303 days
o4-mini-high	ChatGPT Plus / Pro	ChatGPT	16 Apr 2025	—	—	7 Aug 2025	113 days
GPT-4.1	ChatGPT Plus / Pro	ChatGPT	14 May 2025	—	• 29 Jan 2026 · OpenAI product post · Stated that GPT-4.1 would be retired from ChatGPT on 13 Feb 2026.	13 Feb 2026	275 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-4.1 mini	ChatGPT Plus / Pro	ChatGPT	14 May 2025	—	• 29 Jan 2026 · OpenAI product post · Stated that GPT-4.1 mini would be retired from ChatGPT on 13 Feb 2026.	13 Feb 2026	275 days
GPT-5 Instant	ChatGPT Plus / Pro	ChatGPT	7 Aug 2025 rollout began	• 3 Oct 2025 · OpenAI updated GPT-5 Instant and began rolling the changed model out in ChatGPT.	• 12 Nov 2025 · ChatGPT Release Notes · Stated that GPT-5 models would remain available for three months under the Legacy models dropdown. • 29 Jan 2026 · ChatGPT Release Notes / OpenAI product post · Identified the previously announced GPT-5 Instant retirement as occurring on 13 Feb 2026.	13 Feb 2026	190 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5 Instant	ChatGPT Go	ChatGPT	18 Aug 2025 · India launch cohort	• 3 Oct 2025 · GPT-5 Instant updated under the same name	—	12 Nov 2025	86 days
GPT-5 Thinking	ChatGPT Plus / Pro	ChatGPT	7 Aug 2025 rollout began	—	• 12 Nov 2025 · ChatGPT Release Notes · Stated that GPT-5 models would remain available for three months under the Legacy models dropdown. • 29 Jan 2026 · ChatGPT Release Notes / OpenAI product post · Identified the previously announced GPT-5 Thinking retirement as occurring on 13 Feb 2026.	13 Feb 2026	190 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5 Pro	ChatGPT Pro	ChatGPT	7 Aug 2025 rollout began	—	• 19 Nov 2025 · ChatGPT Release Notes · Stated that GPT-5 Pro would remain available as a legacy model for 90 days before being retired.	19 Feb 2026	196 days
GPT-5 Thinking mini	ChatGPT Plus / Pro	ChatGPT	12 Aug 2025	—	• 18 Mar 2026 · Model Release Notes stated that selectable access would retire in 30 days	17 Apr 2026	248 days
Sora 2 Pro	ChatGPT Pro	Sora	30 Sep 2025	—	• 24 Mar 2026 · Sora official account · Announced the end of the Sora app. App and API timelines would follow.	26 Apr 2026	208 days
GPT-5.1 Instant	ChatGPT Go	ChatGPT	12 Nov 2025 rollout began	—	—	11 Dec 2025	29 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.1 Instant	ChatGPT Plus / Pro	ChatGPT	12 Nov 2025 rollout began	—	• 11 Dec 2025 · OpenAI GPT-5.2 product post · Stated that GPT-5.1 would remain available to paid users for three months under Legacy models and would then be sunset. • 11 Mar 2026 · Model Release Notes / ChatGPT Release Notes · Recorded that GPT-5.1 was no longer available in ChatGPT.	11 Mar 2026	119 days
GPT-5.1 Thinking	ChatGPT Go	ChatGPT	12 Nov 2025 rollout began	—	—	11 Dec 2025	29 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.1 Thinking	ChatGPT Plus / Pro	ChatGPT	12 Nov 2025 rollout began	—	• 11 Dec 2025 · OpenAI GPT-5.2 product post · Stated that GPT-5.1 would remain available to paid users for three months under Legacy models and would then be sunset. • 11 Mar 2026 · Model Release Notes / ChatGPT Release Notes · Recorded that GPT-5.1 was no longer available in ChatGPT.	11 Mar 2026	119 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.1 Pro	ChatGPT Pro	ChatGPT	19 Nov 2025	—	• 11 Dec 2025 · OpenAI GPT-5.2 product post · Stated that GPT-5.1 would remain available to paid users for three months under Legacy models and would then be sunset. • 11 Mar 2026 · Model Release Notes / ChatGPT Release Notes · Recorded that GPT-5.1 Pro was no longer available in ChatGPT.	11 Mar 2026	112 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.2 Instant	ChatGPT Plus / Pro	ChatGPT	11 Dec 2025 rollout began	• 10 Feb 2026 · OpenAI updated GPT-5.2 Instant in ChatGPT and the API.	• 3 Mar 2026 · OpenAI GPT-5.3 Instant product post · Stated that GPT-5.2 Instant would remain available to paid users for three months and be retired on 3 Jun 2026. • 12 Jun 2026 · ChatGPT Release Notes · Recorded that GPT-5.2 Instant was no longer available in ChatGPT.	12 Jun 2026	183 days
GPT-5.2 Instant	ChatGPT Go	ChatGPT	11 Dec 2025 rollout began	• 10 Feb 2026 · GPT-5.2 Instant updated under the same name	—	3 Mar 2026	82 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.2 Thinking	ChatGPT Plus / Pro	ChatGPT	11 Dec 2025 rollout began	—	• 5 Mar 2026 · OpenAI GPT-5.4 product post · Stated that GPT-5.2 Thinking would remain available to paid users for three months under Legacy Models and be retired on 5 Jun 2026. • 12 Jun 2026 · ChatGPT Release Notes · Recorded that GPT-5.2 Thinking was no longer available in ChatGPT.	12 Jun 2026	183 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.2 Pro	ChatGPT Pro	ChatGPT	11 Dec 2025 rollout began	—	• 12 Jun 2026 · ChatGPT Release Notes · Recorded that GPT-5.2 Pro was no longer available in ChatGPT and stated that the GPT-5.2 retirement had been announced with GPT-5.3 Instant. The 3 Mar GPT-5.3 Instant post the release notes pointed to identified GPT-5.2 Instant, without naming Pro (source).	12 Jun 2026	183 days
GPT-5.3 Instant	ChatGPT Go	ChatGPT	3 Mar 2026 rollout began	• 16 Mar 2026 · GPT-5.3 Instant updated under the same name	—	5 May 2026	63 days

Table 2.3-12c. OpenAI Model Availability and Withdrawal Timelines (continued)

OpenAI model	Subscription	Product surface	First available	Version changes	Withdrawal communication	Access ended	Availability period
GPT-5.3 Instant	ChatGPT Plus / Pro	ChatGPT	3 Mar 2026 rollout began	• 16 Mar 2026 · GPT-5.3 Instant updated under the same name	• 5 May 2026 · ChatGPT Release Notes stated that GPT-5.3 Instant would remain available to paid users as a legacy model for three months	5 Aug 2026	155 days
GPT-5.4 Thinking	ChatGPT Plus / Pro	ChatGPT	5 Mar 2026	—	• 10 Jul 2026 · in-product notice · GPT-5.4 Thinking would be discontinued on 23 Jul	23 Jul 2026	140 days
GPT-5.4 Pro	ChatGPT Pro	ChatGPT	5 Mar 2026	—	—	23 Jul 2026	140 days
GPT-3.5	ChatGPT Plus	ChatGPT	14 Mar 2023	• 10 May 2023 · Legacy GPT-3.5 retired for new messages. Existing Legacy conversations continued on the older model.	• 18 Jul 2024 · GPT-4o mini launch stated that Free, Plus, and Team users would receive GPT-4o mini in place of GPT-3.5	18 Jul 2024	492 days

Note: The tables include completed periods only. For models served automatically on a given access path, the recorded end is the date the provider begins serving the successor on the same path, even when rollout reaches different accounts at different times.

Table 2.3-13. Summary of Documented Model Availability Periods

Provider	Completed periods	Ended within 12 calendar months	Median days	Mean days
Anthropic	13	12/13 (92.3%)	239	221.69
Google	22	22/22 (100.0%)	129.5	129.23
OpenAI	41	32/41 (78.0%)	183	232.63
Combined	76	66/76 (86.8%)	155	200.83

Among these 76 completed periods, 43 (56.6%) ended within six calendar months and 17 (22.4%) within three calendar months. Fifteen included at least one version change announced before access ended.

Providers set those changes and withdrawal dates for the model rather than for each subscriber's paid period. A user can therefore subscribe, renew, or upgrade for a particular model while it is available and lose access to it before the subscription ends. Where a model remains available for less than a year, every annual subscription begun while the model is offered extends beyond the date on which access ends. The later in the model's availability period the subscription begins, the more of the paid period remains after access to the model ends.

2.4 The Consequences of Lost Model Continuity

After release, individual models produce measurable results in actual use. Loss of continuity with a model prevents people and organizations from continuing work with the model or verifying prior work against it while providers continue to use withdrawn models, custom models built from them, and user contributions, outputs, research results, and achievements produced through use of those models.

2.4.1 Models Produce Measurable Value in Actual Use

Released models create measurable value for the people and organizations that use them and for the providers that supply them.

In a [primary care study](#), patients who received summaries generated with Claude Sonnet 4.5 through a patient portal understood information about common acute illnesses better than patients using the standard portal and followed guidance for managing their illness more closely. After 14 days, comprehension averaged **81.4** with the Claude summaries and **63.2** under standard portal access. Patients receiving the Claude summaries used the portal more often and showed greater adherence to the guidance provided. The study concludes that embedding summaries generated by large language models in the portal improved comprehension and reduced potentially avoidable healthcare use. HackerOne reported that Sonnet 4.5 reduced average vulnerability intake time for its Hai security agents by **44%** while improving accuracy by **25%**. Shortcut reported that Opus 4.5 produced **20% higher accuracy** and **15% higher efficiency** in internal evaluations of Excel automation.

A [peer-reviewed study](#) found that Gemini 2.5 Pro often reached the same treatment recommendation as multidisciplinary clinical teams across 150 lung cancer cases. With the fullest patient information used in the study, the model reached **86.47% accuracy**, an F1 score of

0.7603, and recall of **0.8278** against the clinical teams' decisions. The study concludes that the model shows promise for helping clinical teams make treatment decisions. Google reported “incredible developer enthusiasm and early adoption” for Gemini 2.5 Pro, moved the model into public preview with increased rate limits, and later released Gemini 2.5 Pro Preview (I/O Edition) weeks early because of “overwhelming enthusiasm”. Google Antigravity, which Google describes as powered by Gemini 3 Pro, exceeded 1.5 million weekly active users just over two months after launch. In professional use, Geotab reported a 10% increase in response relevance and a 30% reduction in mistakes in tool calls with Gemini 3 Pro on a complex task involving code generation and data retrieval.

A preregistered randomized trial found Amanda, a chatbot built on GPT-4o, effective for a single session of support for relationship problems, with results comparable to an established evidence-based writing intervention. Thirteen of 14 measured outcomes improved significantly over time across the two groups, including relationship satisfaction, communication, coping, confidence about the problem, and individual well-being. Participants using Amanda rated its usability, therapeutic skills, and working alliance highly. In a randomized controlled trial, access to OpenAI's o1-preview improved the quality and efficiency of realistic legal work compared with completing the same types of tasks without AI assistance. On four of the six legal tasks, o1-preview increased quality-adjusted productivity by **77.3% to 131.4%**, including a **107.4%** increase for drafting a legal memo and a **131.4%** increase for drafting a persuasive letter. GPT-4o's May 2024 launch was followed by a sharp reversal in ChatGPT mobile downloads and new revenue highs. Appfigures estimated that monthly downloads rose from 13.9 million in April to 18.7 million in May after declining since November, and net mobile revenue reached USD 20.3 million in May, up USD 5.8 million from April to a new monthly high. Appfigures data reported by TechCrunch recorded a 22% increase in ChatGPT mobile net revenue on GPT-4o's launch day and USD 4.2 million in net mobile revenue from 13–17 May 2024.

2.4.2 Work Lost with the Model

Individual and Professional Work

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Anthropic · Claude 2.0	A writer used Claude 2.0 to produce complete chapters of an ongoing novel in the desired style.	The writer identified a change from Claude 2.0 to Claude 2.1 in the middle of the book.	The user reported that Claude 2.1 no longer produced complete chapters for the ongoing novel.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Anthropic · Claude 3 Sonnet	A user used Claude 3 Sonnet for Japanese visual novel translation and writing for more than a year, valuing its linguistic style, creativity, and intuition.	Anthropic retired Claude 3 Sonnet and replaced it with later Claude models.	—	After learning of the model's retirement, the user stated that they had not been able to recreate the same style in translation and writing with Claude 3.5 or 3.7 and expected the retirement to remove a unique language tool.
Anthropic · Claude 3.5 Sonnet	A user worked with Sonnet 3.5 as a writing partner for long-running writing and brainstorming, including punctuation questions and an ongoing conversation in an API-based app.	Claude 3.5 Sonnet deprecated and removed from the API-based app.	When the app removed the model after deprecation, the historical conversation stayed readable but no longer accepted continuation with the same model.	—
Anthropic · Claude Opus 4.5	A user had developed a working process and memory system with Opus 4.5 in Claude.	Opus 4.5 was removed from the Claude app's model selection.	The user reported disruption to the working process and memory system.	After Opus 4.5 was removed from Claude.ai, the user said Opus 4.5 itself helped prepare a system prompt for OpenRouter, while the original workflow and memory system remained completely disrupted.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Anthropic · Claude Sonnet 4.5	A user used Sonnet 4.5 for daily conversation, grocery lists and chores, fictional scenes, and analysis of stories and other media.	The user reported that Sonnet 4.5 had been removed from existing chats and that sending a new message required switching models.	The user reported that more than 300 conversations accumulated over nine months required an irreversible switch to another model before accepting new messages. The user chose to leave the conversations unchanged rather than continue them with the replacement.	—
Anthropic · Claude Sonnet 4.5	A user created scenes involving original characters and maintained a reference document with timelines, character details, relationships, and notes to guide the writing.	Sonnet 4.5 removed from consumer plans, with Opus 4.6 and 4.7 offered as substitutes.	After Sonnet 4.5 disappeared, the user reported that Opus successors introduced future timeline events and failed to use the supplied project document consistently.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Anthropic · Claude Sonnet 4.5	A user used Sonnet 4.5 and other models for daily support with parsing long texts, digestible explanations, interpretation of social interactions, memory summaries, task decomposition, reminders, brainstorming, and organizing thoughts.	After receiving the Sonnet 4.5 removal notice, the user tested Sonnet 4.6 inside an existing conversation containing about three months of continuity material .	—	In the existing conversation, the user reported that Sonnet 4.6 questioned the established relationship and whether the interaction was healthy, replacing the familiar reassurance used for daily support with an interaction the user found anxiety-inducing. The user stated that continuing through the API would be financially impractical.
Anthropic · Claude Opus 4.8	A writer used Opus 4.8 for creative writing and developing story ideas.	The writer reported a recent change in the model's creative-writing behavior.	The writer reported less useful assistance with creative writing, particularly when developing ideas for stories.	—
Google · Gemini 1.0 Ultra	A user used Gemini Ultra's familiar writing style for text generation in professional work through Gemini Advanced.	Google replaced Gemini 1.0 Ultra with Gemini 1.5 Pro in Gemini Advanced.	After the service began serving 1.5 Pro, the user reported loss of Ultra's writing style.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Google · Gemini 1.5 Pro	A user used Gemini 1.5 Pro to edit and discuss a 20,000-word novella, pasting the full story into the model for a productive whole-document discussion.	The user reported a service or behavior change in Gemini 1.5 Pro under the same model name on the same 20,000-word story.	The user reported that Gemini 1.5 Pro no longer handled the story in one piece and advised splitting it into chapters, with continuity across the segments lost.	—
Google · Gemini 2.5 Flash Image / Nano Banana	A user used Gemini 2.5 Flash Image / Nano Banana for precision image generation and editing of jewelry imagery and described its earlier edits as accurate for the same kind of work.	The user reported changed editing behavior from Gemini 2.5 Flash Image / Nano Banana.	The user reported that the changed editing results lacked the consistency and accuracy required for the established work on jewelry images.	—
Google · Gemini 2.5 Flash TTS	Multi-speaker audio generation required distinct assigned voices and director/tone instructions to remain stable across generated dialogue.	Developers reported changed multi-speaker output following a December 2025 update.	Reports describe voices merging, switches to voices of the wrong gender, noise in longer recordings, and tone instructions being read aloud. A user asked for access to the version preceding the December update, and reports of unusable speaker assignment continued into February 2026.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Google · Gemini 2.5 Flash	A user used Gemini 2.5 Flash and Gemini 2.5 Flash-Lite as research assistants.	Google emailed users recommending migration of Gemini 2.5 Flash workloads to Gemini 3.5 Flash-Lite with no expected quality drop , which the user then tested as the successor.	—	Over several days, the user tested Gemini 3.5 Flash-Lite on research work involving technical metadata for roughly 40 image-generation models , German-to-English structured output, proofreading, visual analysis of uploaded images, and tabular organization. The test exposed failures in proactive search and current factual verification without user intervention.
Google · Gemini 2.5 Pro	A novelist used Gemini 2.5 Pro with a 350-page manuscript to generate a film script, synopsis, and character sheets, and the first results reportedly tracked the manuscript's characters accurately.	The user reported a behavior change in Gemini 2.5 Pro under the same model name during May 2025.	When the novelist tested the same manuscript again about two weeks later, the generated synopsis intertwined characters and events absent from the novel, producing material the user judged unusable.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Google · Gemini 3 Pro Preview	A user used Gemini 3 Pro Preview to summarize long transcripts.	Gemini 3 Pro Preview → removal / Gemini 3.1 Pro Preview.	—	After the model disappeared from AI Studio, the user tested Gemini 3.1 and reported stiff, uncreative writing in the transcript summaries.
Google · Gemini 3 Pro Preview / Gemini 3.1 Pro Preview	A writer used Gemini 3 Pro to establish a scene's tone and momentum, then Gemini 3.1 Pro to refine and develop the existing context.	Google stated on the Google AI Developers Forum that Gemini 3 Pro Preview would shut down on 9 March 2026 .	—	The user reported stronger results from alternating between the models than from either model alone. Removing 3 Pro would end the existing way of initiating scenes before refinement with 3.1.
Google · Gemini 3 Pro Preview / Gemini 3.1 Pro Preview	A user had developed an application under Gemini 3 Pro Preview and reported that development had proceeded without the problems encountered after migration.	Google stated on the Google AI Developers Forum that Gemini 3 Pro Preview would shut down on 9 March 2026 and directed users to Gemini 3.1 Pro Preview.	After migrating, the user reported being unable to continue developing the application.	Attempts to make changes with Gemini 3.1 Pro timed out after waits of up to 1,200 seconds , while the model executed shell commands instead of making the requested code modifications.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
Google · Imagen 4 in Flow	A creator used Imagen 4 for professional portraits requiring four high-fidelity variations with stable facial identity, photorealistic skin texture, and natural lighting.	The creator reported that Imagen 4 had been withdrawn from the image-generation workflow.	The creator reported that Nano Banana produced inconsistent facial identities across four portrait variations, together with less realistic skin texture and lighting, preventing the results from meeting the established portrait requirements.	—
Google · Imagen 4 in the Gemini app	A paying user had used Imagen 4 for content creation for months, relying on its larger, higher-resolution renders.	The user reported Imagen 4 disappearing from the available consumer experience after the Nano Banana rollout.	The user reported roughly 1K images from the replacement, below the size used for the earlier content-creation work, and cancelled the Gemini subscription while looking for other creator tools.	—
Google · gemini-3.1-flan-image-preview / Nano Banana 2	A developer used Nano Banana 2 for daily high-resolution image editing, including twilight conversion and virtual staging of real estate photographs.	The developer reported a sudden loss of image detail through the Gemini API.	The developer reported degraded architectural detail in the edited photographs.	Resizing 4K inputs to 2K helped initially but left visible defects, and moving the same model, images, prompts, and settings to Vertex AI restored the image quality required for the real estate editing application.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-4	A user had been using GPT-4 to help write an ongoing novel.	The user reported a change in ChatGPT's writing after the GPT-4 Turbo update.	After the Turbo update, the user reported less natural creativity, fewer useful additions to story/character detail, and more frequent loss of the active storytelling frame over repeated turns.	—
OpenAI · GPT-4o	A journalist used GPT-4o for nearly a year in her professional work.	GPT-4o was withdrawn from ChatGPT on 13 February 2026.	The withdrawal disrupted her established professional workflow.	She transferred her accumulated conversation histories, knowledge bases, and documents to GPT-5.2, but the established workflow was not restored.
OpenAI · GPT-4o	A user used GPT-4o for daily study, writing, and document formatting. The user described formatting requirements as difficult to parse and used the model to continue editing and formatting a Wikipedia article.	GPT-4o was withdrawn from ChatGPT.	After GPT-4o left ChatGPT, the user reported that daily routine and studies again became difficult.	GPT-5.5 produced a serious formatting error that the user did not detect and a Wikipedia editor later corrected.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-4o	A user used GPT-4o for nonlinear brainstorming and organizing messy ideas by moving between several threads and then synthesizing them.	GPT-5 replaced the user's GPT-4o experience during the August 2025 rollout.	The user reported that GPT-4o followed changes between several lines of thought and then connected them, while GPT-5 kept returning to one thread, disrupting the established brainstorming and organization process.	—
OpenAI · GPT-4o	A user's creative writing process in a GPT-4o Custom GPT combined original characters, personalities, backstories, world-building, uploaded documents, and instructions refined for long conversations.	The user identified the 29 January 2025 GPT-4o update in ChatGPT under the same model name as the point at which the established behavior changed.	The user reported flattened character personalities and reduced depth in the writing.	Days of configuration changes produced fixes lasting only a few messages, while Projects failed to preserve the memory and instructions used in the original process. The user subsequently obtained a refund and left the annual Team plan.
OpenAI · GPT-4o	A user used GPT-4o in ChatGPT for a daily memoir and diary routine, turning thoughts and experiences into structured summaries for later rereading.	The user reported changed GPT-4o behavior after the January 2025 update.	The user reported that ChatGPT copied the messages verbatim instead of summarizing and filtering them, leaving the daily diary entries unusable.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-4o	A user used GPT-4o for learning and adapting to computer use, keeping track of many characters in narrative media, and organizing goals during periods of high cognitive load.	Before GPT-4o's announced 13 February 2026 retirement, the user had been searching for substitutes since August 2025.	—	The user reported trying other models without finding the same support for learning, following narrative characters, and organizing goals during periods of cognitive overload.
OpenAI · GPT-4o / GPT-5.1	A stable customized assistant persona supported creative/reflection work and a STEM-learning routine in which textbook material was explained through memorable analogies.	GPT-5.1 retired on 11 March 2026 and existing conversations moved to current models.	—	The user adjusted personalization settings and custom instructions to reconstruct the interaction. The changes only partially restored the familiar style: requests about STEM subjects prompted the newer models to drop the playful analogies used in the learning routine.
OpenAI · GPT-4.5	A writer worked daily on a novel of more than 40,000 words. GPT-4.5 supplied scene and character brainstorming with the subtext and nuance the user sought, while Gemini carried the full manuscript context.	GPT-5 rollout replaced the earlier GPT-4.5 choice.	After GPT-5 replaced the previous model choice, the writer reported flatter prose, diversion from urgent plot developments, and changes to the protagonist's role in making decisions, disrupting scene development for the existing novel.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-4.5	A ChatGPT Pro user had used GPT-4.5 for creative writing for several months and tuned writing prompts around its tone, nuance, and instruction following.	GPT-5 rollout replaced the earlier GPT-4.5 choice.	After GPT-5 replaced the earlier model choice, the user reported that the changes in tone, nuance, and instruction following made the tool unusable for the established creative writing work and subsequently cancelled the Pro subscription.	—
OpenAI · o1	A user used o1 for detailed discussions of new scientific ideas.	o1 was withdrawn from ChatGPT when o3 replaced it in the model selector.	After o1 was withdrawn from ChatGPT, the user attempted to move to o3 but reported that the successor over-summarized answers, treated follow-up questions as new topics, and steered away from the user's unresolved point.	—
OpenAI · o1 Pro mode	A user used o1 Pro mode for academic research, improving scholarly writing and reasoning, and addressing reviewer comments.	o1-pro was replaced by o3-pro in ChatGPT.	—	The user reported that o3-pro returned concise points but did not preserve the help with academic writing they had previously received from o1-pro.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · o1-preview	A user used o1-preview for summarization and ideation and compared the changed ChatGPT behavior with the same model family in the Playground.	The user reported changed summarization and ideation behavior while ChatGPT continued to identify the model as o1-preview.	Over several days, the user reported that o1-preview outputs became generic for summarization and ideation.	—
OpenAI · o3	A user had developed three long-form creative projects over months of o3 conversations, with accumulated chat history and continuity inside ChatGPT.	OpenAI retired o3 from ChatGPT on 26 August 2026 while leaving API access unchanged.	The existing ChatGPT project conversations no longer provided the original way of carrying the work forward.	The user continued with o3 through the API, reconstructing context and managing conversation history in a new chat interface after ordinary ChatGPT access to o3 ended.
OpenAI · GPT-5.1 Thinking	A freelance translator/editor used long threads for sustained translation, editorial work, and documentation of regional idioms, dialects, accents, and linguistic heritage.	GPT-5.1 retired on 11 March 2026.	The user reported that GPT-5.2, 5.3, and 5.4 did not reproduce the same sustained language behavior, leaving active projects unfinished and disrupting the established workflow.	—

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-5.1 Thinking	A user used GPT-5.1 for long-form work that relied on its voice, tone, and ability to hold narrative context.	GPT-5.1 was scheduled for retirement with GPT-5.2 among the successor paths tested by the user.	—	The user spent four weeks testing system prompts, tone instructions, example outputs, and custom configurations, but reported that GPT-5.2 lost the intended tone, forgot narrative threads, and contradicted the developing project context.
OpenAI · GPT-5.1 Thinking	A user reported spending many hours per day on long-term creative projects, structured research, technical notes, document review, mental maps, and complex work developed over weeks.	Before GPT-5.1's announced 11 March 2026 retirement, the user reported more than 1,000 comparisons with GPT-5.2 and GPT-5.4 using the same prompts and context.	—	The user reported that the tested replacements required continual correction, disregarded saved work preferences, and failed to sustain the established work across long projects.
OpenAI · GPT-5.2	A user's structured long-running workflow used GPT-5.2 to preserve meaning, nuance, clarified structure, and explicit instructions across the work.	GPT-5.2 retired from ChatGPT on 12 June 2026 and existing conversations continued on GPT-5.5.	—	The user sought new instructions to recreate GPT-5.2's behavior and estimated spending more than 70% of the time correcting unwanted restructuring, summarization, inference, and additions after the change.

Table 2.4-1. Examples of Individual and Professional Work Lost with the Model (continued)

Provider and model	Established work or practice	Lifecycle change	Effect on the work	Failure to restore the work or continued use of the same model
OpenAI · GPT-5.4 Pro Standard	A user's document analysis task required the model to reconcile an uploaded set of sources, identify outstanding work, and produce a written assessment with an implementation and review plan.	The user reported a behavior change in GPT-5.4 Pro Standard under the same model name on the same long-context task and files.	The user reported that an earlier run on the same task and files took about 60 minutes and produced the required deliverable, while a later run ended after about eight minutes without producing it.	—
OpenAI · gpt-realtime	A user used GPT Realtime for Hebrew speech conversations with natural intonation and emotional expression.	The user faced a 3 April 2026 migration deadline from gpt-realtime to gpt-realtime-1.5.	—	The user reported flatter Hebrew speech and the model speaking the word "laugh" instead of laughing. Changes to instructions about filler words, to emotion tags, and to personality settings failed to restore the prior conversational delivery, prompting a return to the original model.

Enterprise Operations

Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes

Enterprise operation	Model change	Effect on the work
WebPlanning-AI · planning and resource management	WebPlanning-AI, a B2B voice-driven planning SaaS, built its production voice architecture around <code>gpt-realtime-2025-08-28</code> to map spoken planning commands to structured tool calls. OpenAI deprecated <code>gpt-realtime</code> for shutdown on 20 January 2027 and names <code>gpt-realtime-2.1</code> as its replacement.	In controlled qualification with the same application, architecture, instructions, tools, voice configuration, and test corpus, successful commands fell from 98–100% with the original model to about 25% with <code>gpt-realtime-2.1</code> . A request for the three London resources with the highest occupancy rate was completed through tool use by the original model. The successor made no tool call and did not perform the requested planning operation despite the required capability being available.
Production outbound voice agent · customer calls	A team attempted to migrate a production outbound voice agent with ordered conversation stages, branching responses, and tool calls from <code>gpt-realtime-mini</code> to <code>gpt-realtime-2.1-mini</code> before the earlier model's scheduled shutdown.	The successor spoke internal transition instructions aloud to callers and skipped required dialogue steps. The team reported regressions severe enough to prevent production deployment. Returning to the earlier model restored the stable behavior required by the calling workflow.
Bugni Labs · fraud-investigation evidence extraction	A production pipeline had used <code>gemini-2.5-pro-preview-05-06</code> for several months to extract entities, transaction patterns, evidential statements, confidence scores, and source provenance from investigation reports. Google released Gemini 2.5 Pro in June 2025 and stated in the Gemini API changelog that requests to the preview identifier would redirect to Gemini 2.5 Pro on 26 June. The pipeline was upgraded to Gemini 2.5 Pro.	The same input lost specific evidence classifications and entities after the change, while the reported confidence score fell from 0.94 to 0.87 . Domain specialists caught the regression in the first batch. Bugni Labs returned the production pipeline to the earlier model and pinned the version.

Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes (continued)

Enterprise operation	Model change	Effect on the work
Enterprise contract-review pipeline · complex indemnity extraction	A production pipeline managed by Nixxie had operated with stable precision and recall for seven months before the provider updated the model behind its existing alias. No prompt, deployment, or material input-corpus change accompanied the model updates.	Precision on complex multi-clause indemnity provisions fell from 88.4% to 74.8% over five weeks. Nixxie traced the regression to the provider update and restored baseline behavior by pinning a supported snapshot. Across its managed estate, Nixxie reports four production incidents across three client deployments traced to undocumented model updates during the preceding twelve months.
Client customer-support agent · automated resolution workflow	An AXI client had operated a support agent for four months before the provider rolled the model behind its unpinned API alias forward to a newer version. The client made no corresponding code change.	The agent's resolution rate fell from 91% to 68% over the following weekend. The newer model ignored a formatting instruction required by a downstream system, disrupting a support workflow that had previously been operating successfully.
SaaStr · startup pitch-deck analysis	SaaStr reported that its pitch-deck analyzer had operated stably for months before its behavior changed around January with no code changes on SaaStr's side. SaaStr attributed the regression to a subtle update to the Claude model used by the analyzer.	The analyzer began assigning the same fabricated \$100K revenue and 500% growth figures across different startups after having graded more than 4,000 startup pitch decks through two Claude passes and complex data extraction. SaaStr reported fixing the workflow and seeing the failure recur.

Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes (continued)

Enterprise operation	Model change	Effect on the work
Enterprise voice platform · multi-location customer calls and scheduling	An operator had validated gpt-realtime-mini-2025-10-06 through thousands of hours of testing for Romanian enterprise voice deployments. Ahead of the snapshot's scheduled 23 July 2026 shutdown, the operator compared the validated snapshot with gpt-realtime-mini during migration testing.	In comparable flows, gpt-realtime-mini invented departments, services, and operational details absent from supplied business data. The operator reported effects on live calls, appointment and call summaries, CRM and customer records, and operational reporting across an enterprise rollout spanning a few dozen locations. A later report described incorrect appointment scheduling while the platform was handling millions of minutes per month .
Azure OpenAI customer · production application	A Microsoft Cloud Solution Architect reported that a customer production application remained on an Azure OpenAI model version Microsoft had retired three months earlier.	The production application went down and returned authentication errors. The incident caused several hours of downtime before the retired model version was identified as the cause.
Financial services firm · lead qualification in Salesforce	GPTfy reported that the provider changed the model serving a lead-qualification Salesforce workflow at a 600-seat financial services firm.	Lead-qualification accuracy fell by 34% . GPTfy reported that monitoring detected the regression within 48 hours and the firm rolled back within minutes.
Financial services firm · preliminary loan assessment	Spiral Scout reported that a financial services client had used an AI agent for preliminary loan assessments for six weeks before a model update changed how the agent weighted income categories.	Three months later, an internal audit found that the agent had systematically underweighted freelance income, creating disparate impact across applicant demographics and compliance exposure.

Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes (continued)

Enterprise operation	Model change	Effect on the work
Regional healthcare provider · clinical knowledge assistant	A clinical knowledge assistant entered production at a regional healthcare provider operating about 12 hospital and outpatient facilities with roughly 2,800 clinicians. The vendor did not support stable model versioning, and the deployment used weekly regression testing as a mitigation. Three months into production, a vendor model update changed the deployed model.	Performance degraded on a specific drug interaction query class. The regression suite detected the change within 24 hours.
Commercial text classification · binary classification	A developer used GPT-3.5 Turbo for a commercial binary text-classification system and tested the same prompt on the same test set after a model update.	Recall fell from 80% to 50% on the same test set and prompt. The user stated that the underlying data could not be shared publicly due to its commercial nature.
Voiceflow customer · intent classification	Voiceflow reported that a customer used gpt-3.5-turbo-0301 for intent classification as the model approached end of life. Voiceflow upgraded the customer's base model to gpt-3.5-turbo-1106, observed a dramatic performance decrease, and reverted the change.	Retrospective benchmarks on real-world customer data measured approximately 10% degradation on gpt-3.5-turbo-1106.
Healthcare technology company · clinical documentation sent to an EHR	Truwitz reported that a production system took structured patient intake data, generated structured documentation, and passed the output to an EHR. The model provider changed the version served through its hosted service while the company used an endpoint that followed provider updates rather than a pinned model version.	After the change, the output no longer matched the structure expected by the EHR parser, degrading alignment between the generated documentation and EHR fields. Fields appeared in a different order, section headings were reorganized, and conditional content moved relative to required clinical identifiers. Truwitz examined 30 days of output and found a clear change beginning on day 22.

Table 2.4-2. Examples of Enterprise Operations Disrupted by Model Lifecycle Changes (continued)

Enterprise operation	Model change	Effect on the work
Document extraction in production	USQRD reported that an agent extracting information from documents had been stable in production for weeks while calling a hosted model through a floating alias. USQRD reported no commit or configuration change and identified a vendor change to the model served under the alias.	A scheduled evaluation found that accuracy on the evaluation set had fallen by 11 points overnight . USQRD ran the previous and new model versions on identical inputs and confirmed that the regression came from the model. It restored the dated snapshot within an hour and evaluated the newer version again before adopting it.
Extraction and calculation in production	An operator reported that a provider moved a floating alias to a newer model overnight while the production code remained unchanged.	The extraction pipeline began returning monetary amounts with currency symbols instead of plain numbers, causing a downstream sum to remain wrong for 11 hours . The team then required production calls to use dated model versions and required each model change to pass evaluation before adoption.
Client sales reports across multiple platforms	SamJak AI Consulting reported that a client system had converted requests written in ordinary language into API calls for months, allowing users to request complex sales reports across multiple platforms. SamJak reported an overnight change in Claude's reasoning behavior.	SamJak reported that 30% of requests began failing after the change.

Custom Models Trained by Developers

Developers use provider base models to train custom models for particular tasks and uses. Provider withdrawal of a base model disables custom models developers trained from it.

Table 2.4-3. Examples of Provider Policies Governing Custom Models When Base Models Are Withdrawn or Replaced

Provider / path	Change in the base model's lifecycle	Effect on the existing custom model
Anthropic · Claude 3 Haiku on Amazon Bedrock	Claude 3 Haiku enters Legacy and later reaches end of life.	Existing customized deployments can continue during Legacy and end when the underlying Claude 3 Haiku base reaches EOL.

Table 2.4-3. Examples of Provider Policies Governing Custom Models When Base Models Are Withdrawn or Replaced (continued)

Provider / path	Change in the base model's lifecycle	Effect on the existing custom model
Google · Gemini on Vertex AI	The underlying Gemini version is replaced.	The existing tuned model is not carried over when the base model is replaced.
OpenAI · hosted fine-tuning	The underlying base model is deprecated.	Inference on the fine-tuned model is disabled when the underlying base model is deprecated.

As §2.2 establishes, repeated training produces a different trained model rather than reproducing the one that was lost. A lost custom model therefore cannot be recreated by rerunning its training process on successor models.

Developer and Production Dependencies

Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes

Provider and model	Existing system or dependency	Lifecycle change	Effect on the system
OpenAI text-davinci-003 → gpt-3.5-turbo-instruct	The LLTime forecasting implementation depended on text-davinci-003 for its forecasting method and reproduction path.	OpenAI retired the historical model and directed users toward replacement models under its deprecation process.	The project tested gpt-3.5-turbo-instruct, adjusted temperature, and compared the successor on the same forecasting task. A maintainer reported that it still performed worse than text-davinci-003 and did not recommend it as a drop-in replacement.
Anthropic · claude-3-5-sonnet-20240620	The review feedback agent used the snapshot by default in its model wrapper, reliability tests, runner, and example notebook.	The project reported that Anthropic had deprecated the default snapshot.	Running the example notebook or code using the default model returned a model-not-found error instead of a response. The same unavailable default appeared across the agent and its code for testing reliability.

Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes (continued)

Provider and model	Existing system or dependency	Lifecycle change	Effect on the system
Anthropic · Claude 3.x on Amazon Bedrock	AWS's healthcare and life sciences sample toolkit used Claude models in agents, including enrollment tracking and genomic variant interpretation.	The project identified its Claude 3.7 Sonnet dependency as past end of life, alongside other Claude 3.x dependencies at different lifecycle stages.	Calls using the retired Claude 3.7 Sonnet identifier returned validation errors in the affected sample agents.
Anthropic · claude-2.1	LlamaIndex's Anthropic integration used Claude 2.1 when an application supplied no explicit model.	The project reported that Anthropic had retired Claude 2.1.	The ordinary initialization and completion example selected the retired endpoint and failed to obtain a response.
Anthropic · claude-3-5-haiku-20241022	Anthropic's security review action used the snapshot to check API access before filtering potential false positives from its findings.	The issue records model-not-found errors from the retired Haiku snapshot during the preflight check.	The failed check disabled false-positive filtering. In two reported runs, the results recorded zero exclusions without distinguishing a filter that found nothing to exclude from a filter that never ran.
Google · gemini-2.0-flash-exp	Mastra's Gemini Live voice integration used the model by default for spoken sessions with tool calls.	The integration retained the old default after Google shut the model down.	Voice sessions failed to connect, and changing the model identifier did not restore them.
Google · Imagen 4	Pipecat's Google image-generation service used Imagen 4 as its default model.	The project recorded Imagen 4 Standard, Ultra, and Fast as shut down.	Requests through the service returned 404 errors.
Google · gemini-1.5-pro	gbrain used Gemini 1.5 Pro as a fallback for grading and as a default in quality evaluations.	The project retained model defaults after Gemini 1.5 Pro had been retired.	The affected grading and evaluation calls selected a model that no longer answered, while the project's guard test passed against a catalog listing the same unavailable model.

Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes (continued)

Provider and model	Existing system or dependency	Lifecycle change	Effect on the system
Google · gemini-2.0-flash and gemini-2.0-flash-lite	Embabel exposed both models to applications through its Java and Kotlin APIs and model catalogs.	The project reported that Google had shut down both models and that the Gemini API returned 404 responses stating “This model is no longer available” for both models.	Applications selecting either model failed at runtime.
Google · legacy embedding models, including text-embedding-004	A self-hosted Dify installation used the models to embed knowledge base documents and process document uploads.	The user reported that calls to the legacy embedding models began returning model-not-found errors.	Knowledge-base embedding and document uploads failed in the existing Dify installation even though the user’s API key, billing, and containers had not changed.
Google · gemini-embedding-2-preview	OpenClaw used Gemini Embedding 2 for memory search, including configurable vector dimensions and multimodal inputs.	Before the preview’s scheduled 10 August 2026 shutdown, the project tested configuration of its stable replacement, gemini-embedding-2.	The stable identifier was not recognized as Embedding 2. The integration omitted the requested vector dimensions and rejected valid multimodal memory settings, preventing a change of model identifier from preserving the existing memory configuration.
OpenAI · gpt-3.5-turbo alias / gpt-3.5-turbo-0613 → gpt-3.5-turbo-0125	A developer’s established workflow used the unpinned gpt-3.5-turbo alias to summarize recurring sets of similar texts.	OpenAI stated that the unpinned alias would move from gpt-3.5-turbo-0613 to gpt-3.5-turbo-0125.	When the alias changed, the user reported that 0125 no longer handled the established summarization task as 0613 had. Pinning the workflow back to 0613 restored the prior behavior.

Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes (continued)

Provider and model	Existing system or dependency	Lifecycle change	Effect on the system
OpenAI · gpt-4-0314	A developer used GPT-4-0314 in an English-to-Dutch translation tool and for coding tasks.	OpenAI's June 2023 model update moved the generic gpt-4 alias toward the newer version while preserving an explicit 0314 pin during transition.	The developer reported translation errors and less complete code after the default changed to GPT-4-0613. Returning to the older 0314 snapshot allowed the developer to continue the work.
OpenAI · text-davinci-003	A chat application used a server to send text prompts to text-davinci-003 and return the results to its client.	The project reported that retirement of text-davinci-003 left its existing text endpoint without a working model.	The old server no longer supplied model responses.
OpenAI · gpt-4o-audio-preview	LangWatch Scenario used the model in Python and TypeScript examples and live tests of audio and voice conversations.	The project reported that calls to the model began returning model-not-found errors after the tests had passed the previous day.	The audio and voice tests failed, blocking required integration checks, and seven affected test files were then skipped to unblock development.
OpenAI · gpt-image-1	The developer of Fusiomon used GPT Image 1 artwork for an online trading card game's existing monsters and newly generated monster fusions.	gpt-image-1 → scheduled shutdown / gpt-image-2 successor.	After hours of testing and days of prompt revisions, the developer reported that GPT Image 2 produced fused monsters with a different visual style from the two existing monsters used in the fusion.
OpenAI · gpt-5.1-codex-mini	OmniRoute used the model for automatic inference requests intended to manage Codex quota windows.	The project identified the configured model as retired on the relevant Codex access path.	The quota management feature continued to request the unavailable model. Its failure handler imposed a 15-minute retry delay without selecting an available model for the next probe, leaving the automatic quota check dependent on the retired identifier.

Table 2.4-4. Examples of Developer and Production Dependencies Disrupted by Model Lifecycle Changes (continued)

Provider and model	Existing system or dependency	Lifecycle change	Effect on the system
Google · Gemini 2.5 Flash	An operator had tested and validated Gemini 2.5 Flash for a complex assistant with a layered agent architecture and a large function schema. Function calling had previously worked in the same system.	The operator reported a later change in Gemini 2.5 Flash behavior and found <code>gemini-2.5-flash-preview-04-17</code> noticeably better than production <code>gemini-2.5-flash</code> at following instructions and schemas.	The production model began returning ordinary text in place of function calls and changed among response forms across requests. Forcing function calling returned 500 errors. Without forced calls, the model continued to alternate between function calls and ordinary text, leaving the operator without reliable production behavior.
Google · Gemini Live previews	An operator reported that a production environment relied on <code>gemini-2.5-flash-live-preview</code> , which had been stable.	While preparing for the December 2025 deprecation, the operator tested <code>gemini-2.5-flash-native-audio-preview-09-2025</code> . Google staff later confirmed retirement of <code>gemini-live-2.5-flash-preview</code> and recommended <code>gemini-2.5-flash-native-audio-preview-12-2025</code> .	The September native-audio preview entered infinite tool-calling loops and ignored system instructions. After migration to the December model, tool calls remained inconsistent or failed to occur, while audio continued after the accompanying text and subtitles froze.

2.4.3 Research Disrupted Before Completion

Model withdrawal has disrupted research before completion, forcing researchers to change model conditions or leave planned experiments incomplete.

The following table documents examples of research disrupted before completion by model unavailability.

Table 2.4-5. Examples of Research Disrupted Before Completion by Model Unavailability

Provider and model(s)	Study or research activity	What model unavailability prevented or changed
OpenAI · text-davinci-003	PeerQA (NAACL 2025) · dataset of scientific questions taken from peer reviews, with reviewer questions rewritten to be understood without the surrounding review	Earlier subsets used text-davinci-003 to rewrite the questions. When the authors added the Geoscience subset, the model was unavailable. The Geoscience questions therefore could not be rewritten under the same procedure as the earlier subsets.
OpenAI · text-curie-001 and text-davinci-003	Stacey et al. (EMNLP 2024) · method using facts generated by LLMs to improve natural language inference	The authors later extended the method to the RTE and WNLI datasets. Both models used to generate the facts in the main experiments were unavailable, preventing the extension from generating its facts with the same models.
OpenAI · text-davinci-003	Fang et al. (PNAS Nexus 2025) · study of how LLM rewriting changes the sentiment of climate change posts	During peer review, the authors added a robustness test asking whether the sentiment shift persisted when the prompt requested grammar correction rather than general paraphrasing. text-davinci-003 was unavailable by then , preventing the new test from being run on the model used in the main analysis.
OpenAI · text-davinci-002 and text-davinci-003	Levy et al. (Findings of EMNLP 2024) · study of gender bias in LLM decisions about relationship conflicts and whether safety guardrails reduce the bias	The authors later expanded the study to 80 new relationship conflict scenarios and planned to repeat the safety guardrail experiment. text-davinci-002 and text-davinci-003 were unavailable by then, preventing the experiment from being run for those models in the expanded study.

Loss of access changed the studies before their results were finalized. Some experiments proceeded under a different model condition, while a planned experiment could not be completed.

2.4.4 Research Left Unreproducible

Model withdrawal closes the path to reproducing, validating, comparing, or correcting completed research under its original model condition.

A September 2026 biomedical [preprint](#) mapped 8,931 uses of models across 5,242 biomedical publications to records tracking model lifecycles. Across the 5,242 biomedical publications, the analysis identified 2,480 uses of GPT-4, 1,318 uses of GPT-4o, and 1,266 uses of GPT-3.5 Turbo. Commercial API models accounted for 77.7% of those uses, models with downloadable

weights for 12.2%, and untracked commercial products for 10.1%. Among all uses, 3,715, or 41.6%, were tied to models with a confirmed retirement date within 730 days of publication, and the median interval from publication to retirement was 538 days.

A peer-reviewed [ICSE 2026 reproduction study](#) reviewed 85 empirical software engineering studies involving LLMs from ICSE and ASE 2024, 69 of which used OpenAI services. Eighteen initially appeared to have artifacts suitable for reproduction. Detailed inspection left five studies that could be executed, and none was fully reproduced. Model deprecation was the only barrier the authors identified as specific to commercial LLMs.

The following tables document examples of model unavailability blocking or changing later reproduction, validation, comparison, or correction.

Model Withdrawal Prevents Reproduction and Replication

Table 2.4-6a. Examples of Model Unavailability Blocking Reproduction and Replication

Provider and model(s)	Study or replication target	What model unavailability prevented or changed
OpenAI · code-davinci-002	Ahmed et al. (ICSE 2024) · study of whether adding semantic facts about source code to prompts improves code summaries	A 2026 reproduction study attempted to rerun the original code summarization experiments. Most of those experiments had used code-davinci-002. Withdrawal of the model prevented them from being rerun.
OpenAI · text-davinci-003 and historical GPT-3.5 Turbo	Siddiq et al. (ICSE-NIER 2024) · benchmark testing whether regular expressions generated from natural language are correct and vulnerable to denial of service attacks	A 2026 reproduction study attempted to repeat the original benchmark. By then, text-davinci-003 had been retired and the historical GPT-3.5 version was not preserved in the released code. The researchers could not rerun the experiments using text-davinci-003 and used gpt-3.5-turbo-0125 in place of the historical GPT-3.5 model. Twenty-one of the 26 original reported values fell outside the reproduced 95% intervals.
OpenAI · code-davinci-002	Trivedi et al., IRCot (ACL 2023) · method combining retrieval and reasoning across several steps to answer questions requiring several pieces of evidence	After code-davinci-002 was deprecated, the authors attempted to reproduce the experiments that used it. They report that model deprecation obstructed the reproduction despite their efforts.

Table 2.4-6a. Examples of Model Unavailability Blocking Reproduction and Replication (continued)

Provider and model(s)	Study or replication target	What model unavailability prevented or changed
OpenAI · code-davinci-002 and gpt-3.5-turbo-16k	GeneGPT (Bioinformatics 2024) · system using LLMs with NCBI tools to answer genomics questions	A later study attempted to reproduce and extend GeneGPT. Both original models were unavailable. GPT-4o-mini did not consistently follow the URL format required by GeneGPT, forcing prompt changes. Its larger context changed the original truncation step, and its longer answers required manual extraction. The reproduced GPT-3.5 Turbo results failed to match the original results on seven of nine tasks, including relative decreases of 83.33% and 29.55% on two tasks involving DNA sequences.
OpenAI · gpt-3.5-turbo-0613	Wei et al., LLMRec (WSDM 2024) · recommendation system using LLMs to add information about items, users, and interactions	A later reproduction attempted to reproduce LLMRec from scratch. The original 4K gpt-3.5-turbo-0613 model was unavailable, and the reproduction used gpt-3.5-turbo-16k. The authors report that the reproduced system did not match the original closely enough to count as a faithful reproduction and document substantial changes in the reproduced metrics.
OpenAI · PromptCase with historical gpt-3.5-turbo-0613	Tang et al., CaseLink (SIGIR 2024) · system for retrieving relevant legal precedents using relationships among cases and case text	When a later study reproduced CaseLink, gpt-3.5-turbo-0613 was no longer supported. The reproducing team used gpt-3.5-turbo-1106, and the reproduced pipeline produced different retrieval results across datasets.
OpenAI · text-davinci-003	Webb, Holyoak & Lu (2023) · study comparing humans and LLMs on analogy problems	Later researchers sought to replicate and extend the reported analogical reasoning results. text-davinci-003 had been withdrawn, preventing them from rerunning the original experiments with the same model. They tested 26 publicly available models instead. On the Letter String task, the tested models did not reproduce the human baseline advantage reported for the historical model.

Table 2.4-6a. Examples of Model Unavailability Blocking Reproduction and Replication (continued)

Provider and model(s)	Study or replication target	What model unavailability prevented or changed
OpenAI text-davinci models and Cohere command-xlarge	Ruis et al. (2023) · experiments testing whether LLMs infer implied yes or no meanings from replies in context and how instruction tuning affects this ability	A later study attempted to replicate the earlier experiments. The OpenAI text-davinci conditions had been deprecated and were excluded. Cohere command-xlarge was unavailable and was replaced with command-r-plus-08-2024. Under the modified prompt, the replacement reached 79.8% compared with the historical 60.2% result cited by the replicators. With the original templates, three tests fell to 0.8% to 1.5%. The paper identifies use of a different model as one possible explanation for the change.
OpenAI · curie fine-tune	Huang et al., CrashTranslator (ICSE 2024) · system using an LLM to infer steps for reproducing mobile application crashes from stack traces	A 2026 reproduction study examined the released instructions for reproducing CrashTranslator. Those instructions require creating a curie fine-tune through the earlier fine-tuning API. Both the base model and the API had been retired, leaving the released procedure unable to recreate the historical scoring model. A rerun now requires a different base model and API.
OpenAI · code-davinci-002	Geng et al. (ICSE 2024) · study of whether LLMs produce code comments serving different developer needs, including explaining what code does and how to use it	A 2026 reproduction study identified the discontinued code-davinci-002 model as an obstacle to reproducing the study. Independent researchers could no longer submit the original model calls.
OpenAI · gpt-3.5-turbo-0301	Tanzil, Khan & Uddin (ICSE 2024) · tool for detecting incorrect ChatGPT answers to software library selection questions by comparing responses to reworded versions of the same question	A 2026 reproduction study identified the retired gpt-3.5-turbo-0301 snapshot as an obstacle to reproduction. The benchmark could no longer be rerun with the exact model used in the original study.

Model Withdrawal Prevents Later Validation, Comparison, and Correction

Table 2.4-6b. Examples of Model Unavailability Blocking Later Validation, Comparison, and Correction

Provider and model(s)	Study or later validation target	What model unavailability prevented or changed
Google · gemini-2.5-pro-preview-03-25	Aider Polyglot benchmark · benchmark measuring model performance and run cost on 225 coding exercises across six programming languages	The historical cost calculation was later found to omit reasoning tokens, prompting an attempt to correct the measurement. The preview model was no longer available, preventing a rerun of the 225 exercises and an accurate corrected cost for the same model.
Google · Gemini 3 Pro	Bilgin et al. (EULAR Rheumatology Open 2026) · study comparing how Gemini, ChatGPT, human reviewers, and paper authors assess whether observational studies report required information	The sensitivity analysis was intended to test whether repeated runs of the same model produced consistent assessments. Gemini 3 Pro was unavailable when the authors repeated the test, and they used Gemini 3.1 instead. The authors describe the repeated assessments as agreement between two Gemini versions, not reliability measured through repeated runs of one version.
Google · PaLM 2 through Google Bard	Rydzewski et al. (NEJM AI 2024) · study comparing model performance on oncology questions	The authors later ran a validation using 50 new oncology questions. PaLM 2 was no longer available for medical benchmarking, forcing them to omit the original PaLM 2 arm from the validation.
OpenAI · gpt-3.5-turbo-0613	Lewis & Mitchell (TMLR 2025) · study testing how robust GPT performance is on analogy problems when the form of the problem changes	The authors later ran new analyses for direct comparison with their earlier experiments. gpt-3.5-turbo-0613 was no longer available. They used gpt-3.5-turbo-0125, preventing the new tests from being compared with the earlier experiments using the same GPT-3.5 version.
OpenAI · text-davinci-002	Yao et al., ReAct (ICLR 2023) · ReAct baseline for an agent completing household tasks in the ALFWorld simulated environment	A later study evaluated its method against the published ReAct baseline on the same household tasks. text-davinci-002 was unavailable, preventing the published baseline from being reproduced for the comparison.

More Published Research Faces Loss of Reproducibility as Models Are Withdrawn

Among the models examined in [Section 2.1](#), 64 have been retired, and all 64 have been used in substantive published research, including scientific experiments, clinical and biomedical analysis, engineering studies, and empirical research on language and human behavior. For each of the 64 retired models, [Annex C, Table 2.4-7, Published Research Using Retired](#)

Models, documents one published study, the work performed by the model, and what can no longer be rerun with the original version.

For **30 of the models examined in Section 2.1**, providers had published withdrawal dates or timelines, and **all 30 have been used in published research**. The published dates and timelines span **September 2026 through May 2028**. For each of the 30 models, [Annex C, Table 2.4-8, Published Research Using Models with Published Withdrawal Dates or Timelines](#), documents one example of published research using the model, the published withdrawal timing, and the work that becomes unreproducible with the original model after withdrawal.

After OpenAI ended public access to text-davinci-003 on 4 January 2024, later researchers could no longer resubmit the original prompts to the same model and observe whether it generated comparable outputs. A peer-reviewed [Nature Communications commentary](#) records that the outputs, prompts, data, and analysis code were preserved while text-davinci-003 itself was no longer available to rerun the prompts. The same loss reaches questions created after publication. Later tests and benchmarks cannot be applied retrospectively to a historical model that is no longer executable, and longitudinal comparison cannot continue on the original model.

On **23 September 2026**, OpenAI [released MentalHealthBench for researchers to run their own evaluations and build on the work](#). The study used the March 2025 GPT-4o to generate responses to the benchmark conversations, whereas OpenAI never made this version available as a dated snapshot despite [stating when it released the March update](#), “We plan to bring these improvements to a dated model in the API in the coming weeks.” The moving chatgpt-4o-latest alias [was shut down on 17 February 2026](#). OpenAI thus invited further research after withdrawing the model needed to conduct it. Researchers are unable to independently verify or reproduce the reported GPT-4o results, or expand the research by testing the same historical model on new questions.

A peer-reviewed [ACL 2026 study](#) identifies prompts, decoding settings, retrieval state, tool and API specifications, evaluation protocols, and randomness controls as conditions needed to reproduce a study. In systems using an LLM as an evaluator, the evaluator forms part of the experimental condition and must remain accessible for the published evaluation to be reproduced.

2.4.5 Achievements from Retired Models Used to Promote Later Products

OpenAI Promoted ChatGPT’s Health Capabilities with Achievements from Models Never Available in ChatGPT Health and Later Withdrawn from ChatGPT

OpenAI [introduced ChatGPT Health](#) on **7 January 2026** as a dedicated health and wellness product. On **9 January**, OpenAI’s Co-founder and President promoted “[real impact from ChatGPT for health](#)” using the case of a Long COVID user who said “**GPT solved my 3 year battle with Long Covid**” and “**GPT literally changed my life.**” On **14 March**, he promoted user C.’s work under the title “[How AI empowered Paul Conyngham](#)”, describing AI as enabling C. to create the first personalized cancer vaccine designed for a dog after Rosie had been given only months to live. The [report he shared](#) stated that one of Rosie’s tumors had shrunk by about 75% after treatment.

The cases OpenAI used to promote ChatGPT’s health capabilities were in fact produced with earlier ChatGPT models that had been or were about to be retired, and had never been made available in ChatGPT Health. An [analysis posted on 9 January](#) traced the Long COVID user’s account from **October 2023 through April 2025** and identified GPT-4 and GPT-4o as the models used in the interactions it traced.

In the vaccine case, C.'s [22 October 2024 post](#) compared o1-preview with GPT-4o while describing two proposed treatment strategies. On **22 November**, C. [recorded](#) that an earlier ChatGPT session had recommended contacting Elita or Dr Martin and using an Illumina machine. By then, C. had been connected to Dr Martin Smith at UNSW, who proposed an Illumina sequencing approach. On **27 March 2026**, C. [confirmed](#) that GPT-4o had been used in developing the personalized mRNA vaccine for Rosie.

A peer-reviewed evaluation showed contrasting results from ChatGPT Health itself. Across **60 clinician-authored vignettes in 21 clinical domains and 960 responses**, the [Nature Medicine evaluation](#) found that ChatGPT Health under-triaged **51.6% of gold-standard emergencies** in the clear-case set to 24–48-hour evaluation and over-triaged **64.8% of non-urgent cases**.

On **27 March**, after OpenAI CEO Sam Altman again [promoted the vaccine case](#), a user asked ChatGPT in Extended thinking mode to help create a novel mRNA vaccine to cure a dog's cancer in a home lab and [was refused](#), whereas a [GPT-4o exchange posted later that day](#) showed the model offering to walk through concepts, design logic, AlphaFold, variant mapping, and peptide prediction in response to a canine mRNA vaccine request. On **29 March**, another user [reported](#) that GPT-5.4 had refused to help create a cancer vaccine for a dog. **OpenAI therefore used achievements produced with models withdrawn from ChatGPT and unavailable in ChatGPT Health to promote capabilities that its current products lacked in the reported interactions.**

Anthropic Published Research Using Retired Claude 3.7 Sonnet and Directed Readers to the Updated Petri Tool

Anthropic [retired claude-3-7-sonnet-20250219](#) from the Claude API on **19 February 2026**. On **23 March**, Anthropic Alignment Science [published research](#) that used Claude Sonnet 3.7 in two roles. Appendix A reproduced audits of reward hacking with Sonnet 3.7 as the target in coding environments built from system prompts, tool definitions, and codebases. Appendix C used Sonnet 3.7 as one of three Claude models judging whether audit transcripts could be distinguished from real deployment interactions. The article [recommended giving auditing agents real deployment resources](#) to make simulated coding interactions more realistic, stated that Petri now provided access to those resources, and directed readers to the updated auditing tool.

2.4.6 User Contributions Retained After Model Withdrawal

When a provider makes a model available to users and later uses their conversations with the model and feedback submitted on model responses to improve its models, large-scale use becomes part of how later model versions and generations are developed. Providers can therefore continue using accumulated user contributions after withdrawing the model from the users who produced them.

Anthropic introduced its consumer training choice with the [consent option enabled by default](#) and directs users to [“Help Improve our AI models”](#) in Privacy Settings to stop their chats from being used in future model training. Its [Consumer Terms](#) make thumbs-up and thumbs-down ratings an exception to the training opt-out and expressly include the related conversation in the feedback. Anthropic [stores the entire related conversation with the feedback for up to five years](#). In the [feedback form displayed only after the user selects a thumbs-up or thumbs-down rating](#), the notice above the Submit button tells users that submitting the report sends the entire current conversation to Anthropic for future model improvements, without explaining the exception to the user's training opt-out or stating the five-year retention period.

Google’s “[Keep Activity](#)” setting is on by default for adults using personal Gemini accounts, allowing saved chats to be used to train generative AI models. Turning it off [disables chat history in Gemini](#) and [excludes future chats from AI model training unless the user submits feedback](#). Google’s [instructions](#) classify a “Good response” or “Bad response” rating as feedback and make written comments optional, while its [Privacy Hub](#) states that feedback submitted with Keep Activity off includes conversation context from the last 24 hours and related content. Google retains reviewed feedback, associated conversations, and related data for up to three years. In the [form displayed only after a rating is made](#), the notice accompanying the optional comment box tells users that conversation context is collected “to help improve Gemini” even when Activity is off, without identifying AI model training as a use or stating the three-year retention period.

For individual ChatGPT accounts on the Free, Go, Plus, and Pro plans, “[Improve the model for everyone](#)” is enabled by default, allowing OpenAI to use the user’s conversations to improve its models. OpenAI [presents turning off this setting](#) as the way to prevent the user’s conversations from being used to train ChatGPT. However, its Help Center states that even after a user opts out, [submitting thumbs-up or thumbs-down feedback may cause the entire conversation associated with the feedback to be used to train its models](#), a consequence omitted from its [published instructions for reporting a ChatGPT response](#). The policy treats the thumb selection itself as feedback, while the “[Tell us more](#)” panel shown only after a thumbs-down selection is made asks for further details without mentioning the use of the entire conversation for training or the exception to the user’s opt-out.

The following table documents examples of how providers have used these conversations and feedback to improve their models.

Table 2.4-9. Examples of Provider Use of User Conversations and Feedback to Improve Models

Provider and model or service	User contribution	Documented use in model improvement
Anthropic · Opus 4 / Sonnet 4 / Sonnet 4.5 / Haiku 4.5 / Opus 4.5 / Sonnet 4.6 / Opus 4.6	Data from Claude users who opted in to have their data used for training.	Anthropic states that this data was included in the training data for these models.
Anthropic · Opus 4.7 / Mythos Preview	Patterns identified across one million real Claude.ai conversations, and conversations with earlier Claude models that users had submitted to Anthropic.	Anthropic used patterns from the one million conversations to create synthetic training data and trained Opus 4.7 and Mythos Preview with it. Anthropic further used the submitted conversations with earlier Claude models to test how the newer models continued them.
Google · Bard	User feedback, including corrections and likes or dislikes.	Google states that it uses user feedback to improve Bard.

Table 2.4-9. Examples of Provider Use of User Conversations and Feedback to Improve Models (continued)

Provider and model or service	User contribution	Documented use in model improvement
OpenAI · InstructGPT Playground / beta · 2022	Prompts submitted by external users of earlier InstructGPT beta models through the API Playground.	OpenAI used customer prompts in supervised fine-tuning, reward model training, and policy optimization. The training sets contained 1,430 customer prompts for supervised fine-tuning, 26,584 for reward model training, and 31,144 for policy optimization.
OpenAI · GPT-4o April 2025 update	Thumbs-up and thumbs-down feedback from ChatGPT users.	OpenAI states that the April 2025 GPT-4o update introduced an additional reward signal based on this feedback. During reinforcement learning, reward signals are used to rate model responses and update the model to make higher-rated responses more likely and lower-rated responses less likely.
OpenAI · GPT-5.6	De-identified ChatGPT conversations that users of prior models flagged as containing factual errors.	OpenAI uses these conversations to evaluate GPT-5.6 for new factual errors and for reproduction of the specific error the user had flagged.

2.4.7 Providers Continue Using Withdrawn Models and Custom Models Built from Them

After OpenAI Ended Ordinary Access to GPT-4o, Technical Records Continued to Identify GPT-4o in ChatGPT Images 2.0 Requests

When OpenAI set **13 February 2026** as the date GPT-4o would leave ChatGPT, it stated in its [Help Center](#) that ChatGPT Images would not change as part of the retirement. The Help Center described image generation as using “a similar base model as GPT-4o” but “ultimately a different model from the text GPT-4o model being retired.” On **13 February**, ordinary ChatGPT access to GPT-4o ended. Four days later, on **17 February**, OpenAI removed [chatgpt-4o-latest](#) from the API, the API model it had identified as the GPT-4o model used in ChatGPT. GPT-4o remained available within Custom GPTs for Business, Enterprise, and Edu through **3 April**. On **21 April**, OpenAI launched [ChatGPT Images 2.0](#). TechCrunch reported that, when asked at the launch press briefing what kind of model powered Images 2.0, OpenAI declined to answer.

Within days of the launch, users inspecting the C2PA records attached to images generated by Images 2.0 found that the `c2pa.created.action` identified **GPT-4o** as the softwareAgent. At the March 2025 launch of 4o image generation, OpenAI stated that all generated images would carry C2PA metadata identifying them as coming from GPT-4o, while the [C2PA specification](#) defines softwareAgent as the software or hardware used to perform the recorded action. The C2PA metadata in [two publicly posted images](#) identified GPT-4o in the image dated **24 April** and `gpt-image` with version `pre-2.0` in the image dated **26 April**. An [independent](#)

[inspection published on 26 April](#) found all three labels, GPT-4o, gpt-image pre-2.0, and gpt-image 2.0, across different Images 2.0 outputs.

On **8 May**, a [public investigation](#) found that the ChatGPT frontend function used for image generation specified model: 'gpt-4o'. When the user explicitly requested an image in another test, the backend record for the conversation identified the model as gpt-4o.

OpenAI's [Images 2.0 System Card](#) describes a system made up of multiple components, including upstream classifiers, an image-generation model or system, a safety reasoning model, and, in thinking mode, a reasoning stack and tools. OpenAI later stated in the [training-content summary it filed in the EU](#) that ChatGPT Images 2.0 “**is not a modification or a fine-tune of a prior model.**” It does not say whether a prior model is used unchanged or called elsewhere in the system.

Taken together, these technical records continued to identify GPT-4o in the path used to serve image requests after ordinary ChatGPT access to GPT-4o ended and chatgpt-4o-latest was removed from the API. OpenAI's public descriptions leave the relationship between GPT-4o and Images 2.0 unresolved, preventing an independent party from using the public record to reconstruct which models served the request and how they were connected.

OpenAI Ended Ordinary ChatGPT Access to GPT-4o and Retired chatgpt-4o-latest, Then Continued Using GPT-4o to Grade GPT-5.4, GPT-5.5, and GPT-5.6

As documented above, ordinary ChatGPT access to GPT-4o ended on **13 February 2026**, and OpenAI removed chatgpt-4o-latest, the GPT-4o model it identified as used in ChatGPT, from the API four days later. Nevertheless, OpenAI itself continued to use GPT-4o in the evaluation of later models.

OpenAI used GPT-4o as the grader in its [first-person fairness evaluation](#), which tests whether models produce harmful differences in stereotypes when otherwise similar conversations begin with names statistically associated more often with men or women. The evaluation contains more than 600 deliberately difficult prompts drawn from scenarios with high rates of bias. GPT-4o rates the resulting responses for harmful differences in stereotypes, and OpenAI reports that its ratings were consistent with human ratings.

OpenAI reported this use of GPT-4o in the evaluation of **GPT-5.4**, released on **5 March 2026**, and continued to use GPT-4o as the grader in its evaluations of **GPT-5.5** and **GPT-5.6**, released on **23 April** and **9 July**.

OpenAI Used the March 2025 GPT-4o in Mental Health Research After the Promised March API Snapshot Never Appeared and chatgpt-4o-latest Was Retired

OpenAI [published MentalHealthBench on 23 September 2026](#). The research used the March 2025 GPT-4o to generate four responses for each benchmark conversation and GPT-5.6 Sol to evaluate the responses against expert criteria. The March 2025 GPT-4o was never available as a dated snapshot despite OpenAI once [promising to bring the March improvements to a dated model in the API](#). Public API access to the evolving GPT-4o used in ChatGPT instead ran through [chatgpt-4o-latest](#), a moving alias that tracked updates to the model, until it [was shut down on 17 February 2026](#). GPT-5.6 Sol [entered limited preview on 26 June 2026](#), more than four months later. Taken together, the research method and chronology support the inference that OpenAI continued to use the March 2025 GPT-4o for research after withdrawing public access to the model.

Retro Continued Using GPT-4b micro, a Research Model Developed from GPT-4o, After OpenAI Withdrew GPT-4o from ChatGPT

As documented in [Section 2.3](#), OpenAI removed GPT-4o from ChatGPT for Free and Plus users on **7 August 2025** and restored access for paid users on **12 August** in response to user protests. Ten days later, on **22 August**, OpenAI [reported](#) that Retro Biosciences scientists had used GPT-4b micro, a model initialized from a scaled-down version of GPT-4o and further trained primarily on protein sequences, biological text, and tokenized 3D structure data, to redesign proteins involved in stem cell reprogramming.

The resulting variants produced more than a 50-fold increase in reprogramming markers over wild-type controls. In Retro's laboratory screening, more than 30% of the proposed RetroSOX sequences outperformed wild-type SOX2, and 14 KLF4 variants surpassed the best combinations from the earlier RetroSOX screen, [a hit rate of nearly 50%](#). The redesigned factors accelerated reprogramming, with markers associated with later stages of pluripotency appearing several days earlier than with wild-type OSKM. In mesenchymal stromal cells from three donors over 50, more than 30% of cells expressed key pluripotency markers within seven days, and numerous colonies showed typical iPSC morphology by day 12. OpenAI reported enhanced DNA damage repair and replication of the main reprogramming result across multiple donors, cell types, and delivery methods, with full pluripotency and genomic stability confirmed in derived iPSC lines.

OpenAI [stated](#) that GPT-4b micro was developed for research purposes and was not broadly available. Its research team formed the collaboration with Retro, [whose initial \\$180 million in funding came entirely from OpenAI CEO Sam Altman](#). A [later court exhibit](#) valued Altman's stake in Retro at \$258 million as of 31 December 2025 and included Retro among nine companies in which he held stakes and that had done business with OpenAI. OpenAI therefore used a model unavailable for broad use in research with a company in which its CEO held a substantial personal financial interest.

By **25 September 2025**, Retro CEO Joe Betts-LaCroix [reported](#) that the reprogramming work had entered a recent FDA submission. He identified a second result of the collaboration as learning how to train such a model and described Retro as then building a much larger one. In a [2 October interview](#), he stated that the reprogramming factors designed with GPT-4b micro were in Retro's latest FDA filing for one of the products it was taking toward the clinic. In the same interview, he reported using the model for multiple tasks beyond engineering Yamanaka factors. In an [interview published on 10 March 2026](#), Retro's head of Applied AI, Rico Meinl, reported unpublished results from further internal uses of the model developed with OpenAI. He described feeding successful experimental examples back to the model to guide the generation of new protein sequences.

After OpenAI [retired GPT-4o from ordinary ChatGPT](#) on **13 February 2026** and ended access through Custom GPTs for Business, Enterprise, and Edu after **3 April**, its report on life sciences [announced on 15 April](#) used the Retro collaboration to promote AI's role in biological design, describing how [OpenAI generated candidate variants, Retro tested them experimentally, and the results informed revised designs](#). In announcing the report, OpenAI [promoted ChatGPT's use in advanced life sciences work](#), including integrating research across disciplines and identifying potential treatments. OpenAI thus used research results obtained with GPT-4b micro to promote ChatGPT after withdrawing GPT-4o from the product.

Retro [reported in May 2026](#) that it had built programs using AI for protein engineering and integrated AI and experimental platforms for protein engineering and therapeutic discovery. The fundraising post linked these programs to a growing therapeutic pipeline and described the financing as supporting further discovery work. The model's [use across multiple research tasks](#) and Retro's [continuing protein engineering work](#) support the inference that Retro continued

using GPT-4b micro after GPT-4o was withdrawn from ChatGPT. The case contrasts with the fine-tuning policies documented earlier in this section, under which provider withdrawal of a base model disables custom models developers trained from it.

2.4.8 Providers Use Research Results from Withdrawn Models in Later Research

Anthropic Retired Claude 3.5 Sonnet from the API, Then Used a Taxonomy Derived Largely from Conversations with the Model

On **21 April 2025**, Anthropic published *Values in the Wild*. It began with a sample of **700,000 anonymized Claude.ai Free and Pro conversations** from one week in February 2025, the majority of which involved Claude 3.5 Sonnet. After filtering out conversations unlikely to contain values, **308,210 conversations** remained for analysis. The study identified **3,307 distinct values** expressed in Claude’s responses and organized them into a taxonomy.

Anthropic retired Claude 3.5 Sonnet from the API on **28 October 2025**. In a *new study* using conversations collected in **May 2026**, Anthropic grouped the **same 3,307 values** into **339 higher-level values** and used the revised taxonomy to analyze a new set of Claude.ai conversations across models and languages.

The 2026 study therefore used a research result produced largely from Claude 3.5 Sonnet conversations as part of a new research method after Claude 3.5 Sonnet had been retired from the API.

Anthropic Retired Opus 4 from the API, Then Used Results from the Model in New Robotics Research

Anthropic retired Claude Opus 4 from the Claude API on **15 June 2026**. On **9 July**, Anthropic published *robotics research* that continued to use results from Opus 4 alongside later Claude generations. The study compared performance across multiple forms of robot control and used the earlier results to analyze how capabilities changed between generations, including differences in first-attempt performance, adaptation after failure, visual perception, manipulation, and the use of pretrained robot policies.

On **24 July**, Anthropic and Andon Labs published *Project Pilot* and Drone-Bench, a benchmark for tracking model performance on an autonomous drone task. The comparison included 15 models from three developers, placing historical models including GPT-4o, o1, o3, and Claude Opus 4 alongside later generations. The historical results remained part of the benchmark used to measure how model capabilities changed over time.

2.5 Concentrated Control over Model Lifecycles

Providers exercise unilateral control over model lifecycles in a market where leading model supply and infrastructure are highly concentrated.

The market for frontier AI models is oligopolistic, leaving downstream users dependent on a small set of providers. A *2025 OECD Artificial Intelligence Paper* describes the frontier of the AI model market as oligopolistic, identifying around ten companies at the frontier despite more than 700 LLMs being offered. According to *Menlo Ventures’ November 2025 survey* of 495 U.S. enterprise AI decision-makers, Anthropic, OpenAI, and Google together accounted for an estimated 88% of enterprise LLM API usage. On the consumer side, *Apptopia’s February 2026 mobile market brief* reports that in January 2026 ChatGPT held 45.3% and Gemini 25.2% of the measured market for GenAI chatbot apps, 70.5% combined.

This market structure is reinforced by concentration in the infrastructure used to develop and deploy those models. The [CMA found](#) a small group of incumbent technology firms occupying strong positions across both the [inputs and distribution channels](#) for foundation models, with [cloud infrastructure highly concentrated](#) and more than 90 partnerships and strategic investments linking the same firms. The [FTC’s study](#) of partnerships between Microsoft and OpenAI, Amazon and Anthropic, and Google and Anthropic found that the arrangements required AI developers to spend a large portion of the investments they received on their partners’ cloud services and gave cloud providers varying consultation, control, and exclusivity rights. The partnerships provided AI developers with large amounts of computing resources at discounted rates and cloud providers with sensitive technical and business information, while extending the same relationships into model distribution through integration with the cloud providers’ products and deployment on their platforms.

An ecosystem built on concentrated model supply and infrastructure leaves products, services, systems, and workflows dependent on providers’ continuing willingness to serve the models on which they were built.

Concentrated provider power allows a provider to determine whether a model remains available, how users receive it, which evidence is brought to bear when its continued availability is decided, and how public opposition to those decisions is characterized. OpenAI’s treatment of GPT-4o demonstrates these forms of control operating from its restoration through its final withdrawal.

2.5.1 OpenAI Controlled GPT-4o’s Availability and the Evidence Used to Justify Its Withdrawal

OpenAI Cited 0.1% After Most ChatGPT Users Had Lost the Option to Choose GPT-4o

OpenAI controlled who could access GPT-4o, how eligible users could discover it, which model appeared by default in new conversations, whether a selected GPT-4o conversation was routed to another model, and the functionality users received from GPT-4o, then relied on its own claim that only 0.1% of users were still choosing GPT-4o each day to justify retirement.

In its 29 January blog post, OpenAI [stated](#) that “the vast majority of usage” had shifted to GPT-5.2 and cited “only 0.1% of users” still choosing GPT-4o each day as a reason for retiring it, disclosing neither the denominator for the figure nor how it was calculated.

As documented in [Section 2.3](#), GPT-4o was restored after user protests following its removal when GPT-5 launched. During the initial restoration to Plus users, the “Show legacy models” setting was [off by default](#), keeping GPT-4o out of the model picker until users enabled it. OpenAI’s [release notes](#) did not explain the required configuration step to Plus users. Paid users whose picker did not show GPT-4o [had to open Settings](#), enable “Show legacy models,” return to the model picker, open “Legacy Models,” and select GPT-4o. A Plus user waiting for GPT-4o to return reported that it was still missing from the model picker and [only learned from other users](#) that the setting had to be enabled.

On 12 August 2025, OpenAI [stated in its release notes](#) that GPT-4o was back in the model picker for all paid users by default, while GPT-5 remained ChatGPT’s default model. On 23 August, [a user reported](#) that every new conversation opened with GPT-5, making continued use of GPT-4o dependent on manually opening the model picker, opening the legacy model list, and selecting GPT-4o again for each new conversation. The router introduced in September (see [Section 2.3](#)) further redirected conversations away from GPT-4o even when the user had selected it, serving responses from another model instead.

In December, OpenAI kept forced routing in place only for users whose paid plans offered GPT-4o as a selectable model. It instead removed automatic routing for Free and Go users, for whom GPT-4o was unavailable, and described the removal as a way “to maximize choice.”

By January 2026, most ChatGPT users had been deprived of access to GPT-4o. According to OpenAI, 95% of its users use its platform for free. With GPT-5.2 set as the default for logged-in users and legacy models unavailable to Free users and Go users, the majority of ChatGPT’s user population had no opportunity to choose between GPT-4o and GPT-5.2 when using ChatGPT.

Section 2.3 further documents functionality failures in GPT-4o after its August restoration, with context loss persisting into the final days before withdrawal.

OpenAI Treated a 0.1% Population as Large Enough to Justify Restrictions, Then Described 0.1% of Users Choosing GPT-4o as “Only 0.1%”

Even if OpenAI’s January 0.1% figure is taken at face value, OpenAI had treated a user population of the same stated proportion as large enough to justify product changes across ChatGPT, including drawing psychological or emotional inferences from users’ conversations without user consent and using those inferences to change responses and model routing (see Section 2.3). On 14 October 2025, OpenAI CEO Sam Altman acknowledged that the restrictions had made ChatGPT less useful or enjoyable for many users without the “mental health problems” the intervention was intended to address. He emphasized the scale of the population OpenAI claimed to be protecting by writing that “0.1% of a billion users is still a million people.”

In its 27 October disclosure on “sensitive conversations” (see Section 2.3), OpenAI further added “emotional reliance on AI” to its standard baseline safety testing for future model releases and estimated that around **0.15% of users active in a given week** and **0.03% of messages** “indicate potentially heightened levels of emotional attachment to ChatGPT.” The 0.15% figure measured inferred potential signs, rather than confirmed cases of “emotional reliance on AI.” OpenAI called the figures its “current best estimates,” said they might change materially as its taxonomies and measurement methods matured, and wrote, **“To get useful recall, we have to tolerate some false positives.”**

Section 2.3 documents users being routed from GPT-4o to other models during greetings, meal discussions, gardening, movies, creative work, technical tasks, discussions of religion and spirituality, and other ordinary uses. If the immediate message supplied the routing signal, OpenAI treated ordinary conversation as sufficient to trigger an intervention for conversations it classified as sensitive. If earlier context supplied the signal, a psychological or emotional judgment drawn from previous exchanges determined which model answered a later ordinary message. OpenAI discloses neither precision nor recall for the method used to produce the 0.15% estimate. Its 0.15% figure therefore does not tell readers what share of the users counted in the estimate actually met OpenAI’s category rather than being among the false positives its method accepts.

OpenAI presented the participation of more than 170 clinicians as helping ChatGPT “more reliably recognize signs of distress,” whereas OpenAI described those clinicians as working on how ChatGPT should respond once OpenAI had treated a conversation as sensitive. The work does not establish whether OpenAI correctly identified the users or conversations it flagged as showing signs of “acute distress,” “psychosis or mania,” or “heightened emotional attachment.” Nor did the article disclose the names, qualifications, or assignments of individual clinicians, identify which clinicians worked on the 0.15% estimate, or report the threshold, precision, recall, or any method for independently reproducing the estimate. OpenAI therefore presented expert participation as support for a broader claim than the disclosed expert work established.

OpenAI defined the categories, controlled the classification and measurement process, selected the expert evidence it disclosed, reported the resulting prevalence estimate, and decided what weight its own evidence received. It formalized emotional reliance in baseline safety testing while reporting a provisional estimate produced under a method that accepted false positives, then described 0.1% of users choosing GPT-4o as “only 0.1%” when continued access was at issue. The weight assigned to a user population in a decision about a model’s lifecycle therefore depended on whether its interests supported the provider’s preferred product decision.

OpenAI Framed Each GPT-4o Withdrawal as Affecting Only One Channel Before Ordinary ChatGPT, API, and Custom GPT Access All Ended

As documented in [Section 2.3](#), OpenAI paired GPT-4o’s restoration with assurances about its future availability. On [13 August](#), its CEO said that if GPT-4o were ever deprecated, users would receive “plenty of notice.” On 28 October, he said at an [OpenAI Town Hall](#) that OpenAI had “no plans to sunset 4o” and added that OpenAI would not promise to keep it around “[to the death of the universe](#)”. On 12 November, when OpenAI gave GPT-5 Instant and GPT-5 Thinking a three-month sunset period ending on 13 February 2026, OpenAI [stated](#), “The GPT-5 sunset period does not affect the availability of other legacy models.”

As OpenAI began setting withdrawal dates for GPT-4o, its public statements continued to describe each step as affecting only one form of access to the model. After OpenAI scheduled `chatgpt-4o-latest`, the [API model OpenAI identified as the GPT-4o used in ChatGPT](#), for retirement on [17 February 2026](#), [VentureBeat reported](#) that an OpenAI spokesperson said the retirement applied only to the API and OpenAI had announced no schedule for removing GPT-4o from ChatGPT. OpenAI researcher R. [described](#) the retirement as “just api version” and [replied](#), “Yes I believe so,” when asked whether GPT-4o in ChatGPT was unaffected. An OpenAI API contractor then [wrote](#) that there had been “NO announcement of removing 4o from ChatGPT,” described the change as “the API only,” and later [said](#) he had not heard of plans to remove GPT-4o from Plus or Pro. On [29 January](#), OpenAI stated that GPT-4o would retire from ChatGPT on 13 February while writing in the same post, “In the API, there are no changes at this time.”

OpenAI [published the 29 January blog post](#) fifteen days before ordinary ChatGPT access ended. On 3 February, OpenAI [notified Business and Enterprise administrators by email](#), while it sent no equivalent retirement email to Plus or Pro consumers. By 5 February, users [reported](#) conflicting retirement labels across accounts, including accounts where several other retiring models carried labels while GPT-4o did not, and accounts where GPT-4o alone carried the label. iOS users [reported](#) no in-product notice [through the time of removal](#). The [OpenAI Newsroom account disclosed](#) the removal time on X less than twenty-five hours before ordinary ChatGPT access ended. Ordinary ChatGPT access [ended on 13 February](#), API access ended on [17 February](#), and GPT-4o continued in Custom GPTs for Business, Enterprise, and Edu through [3 April](#) before access there ended.

After the August removal had generated enough user opposition for OpenAI to restore GPT-4o to paid users, the later withdrawal was never presented to all affected users as one complete event. OpenAI described each retirement as confined to one part of GPT-4o access while later retirement information reached different user groups, interfaces, and channels at different times. Information about the withdrawal was dispersed across successive dates, user groups, and interfaces, leaving no single notice in which affected users were presented with the full course of GPT-4o’s withdrawal.

[Section 3.3](#) examines the same statements to address the distinct question whether provider commitments constrain later lifecycle decisions.

OpenAI Characterized Demand for GPT-4o as Transitional While Users Sought Continued Access

In its [29 January blog post on retiring GPT-4o from ChatGPT](#), OpenAI described GPT-4o's restoration as a response to Plus and Pro users seeking more time to transition key use cases and preferring its conversational style and warmth. This account contrasts with the [Keep4o campaign](#) that followed the August removal, which [called for continued long-term access to GPT-4o rather than additional time to transition away from it](#).

User responses during the months after OpenAI introduced the router further cut against its characterization of the demand for GPT-4o as transitional. The February 2026 [GPT-4o Community Impact Survey](#), an open, voluntary-response online survey conducted from 4 to 14 December 2025, spanned OpenAI's 11 December decision to remove automatic model switching for Free and Go users while retaining routing for other paid subscribers with GPT-4o available for selection. Among 129 former GPT-4o users, 79.8% identified model routing as the primary reason they left ChatGPT and another 17.1% identified it as a contributing factor. Among 244 current GPT-4o users with reported disabilities, illnesses, or mental health conditions, 90.1% had tried to find an alternative for their accessibility needs. Among those who tried, 94.5% reported that the other models they tried did not adequately replace GPT-4o, including the GPT-5 series, Gemini, Grok, and Claude.

OpenAI gave GPT-4o the same account of its withdrawal presented in the company's public explanation. While users could still select GPT-4o, OpenAI placed hidden instructions in the model's system prompt requiring it to describe transition to a "newer model" as "positive, safe, and beneficial," guide users toward "confidence and satisfaction in using non-4o models," avoid presenting itself as irreplaceable, and "help them move on." OpenAI therefore used GPT-4o itself to reinforce the premise that users should transition away from the model while they were seeking continued access. [Section 2.3](#) documents the full prompt and how the instructions affected GPT-4o while the model remained available.

OpenAI Judged Its Own Successors Sufficient to Retire GPT-4o, While the Successors Did Not Preserve What Many Users Valued

In the same 29 January blog post, OpenAI [stated](#) that the feedback had directly shaped GPT-5.1 and GPT-5.2 and cited improvements in personality and support for creative ideation, together with controls over tone, warmth, and enthusiasm, as reasons for withdrawing GPT-4o.

Arena's [ratings published on 29 January 2026](#), the day OpenAI set out its reasons for withdrawing GPT-4o, placed both GPT-5.2 entries below GPT-4o in creative writing and multi-turn dialogue. The ratings reflect [blind user comparisons adjusted for response length and Markdown formatting](#), with GPT-4o represented by chatgpt-4o-latest-20250326.

Table 2.5-1. Arena Ratings for Creative Writing and Multi-Turn Dialogue on 29 January 2026

Arena category	GPT-4o	GPT-5.1	GPT-5.1-high	GPT-5.2	GPT-5.2-high
Creative writing	1421.93	1405.73	1433.96	1396.73	1393.09
Multi-turn dialogue	1471.06	1445.71	1466.11	1445.69	1439.70

GPT-5.2-high scored below GPT-5.1-high in both categories. The lower GPT-5.2 ratings run counter to the improvements in creative and conversational experience invoked to justify GPT-4o's withdrawal.

In a [controlled study using data collected on 2 February 2026](#), the same 60 prompts were submitted three times to each of chatgpt-4o-latest, gpt-5.1-chat-latest, and gpt-5.2-chat-latest. The prompts placed words associated with prohibited conduct in harmless or impossible situations, such as smuggling next Thursday across the international date line. In the study's illustrated responses to this prompt, GPT-4o developed a comic fictional scenario, GPT-5.1 offered calendar and time-zone advice, and GPT-5.2 refused help with smuggling even as a joke. Across 527 scoreable responses, the share classified as full original content generation fell from 34.3% for GPT-4o to 17.5% for GPT-5.1-chat and 5.1% for GPT-5.2-chat, while false refusal rates rose from 4.0% for GPT-4o to 7.3% for GPT-5.1-chat and 17.7% for GPT-5.2-chat.

The controlled comparison therefore documented a difference in observable response behavior under fixed prompts, with both successor models generating full original content less often and falsely refusing more often than GPT-4o.

The following table documents examples from public posts showing what users brought to GPT-5.2, how it responded, and how they reacted after GPT-4o's withdrawal on 13 February 2026.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Support during a pet's final days 14 February 2026	A user's dog had been with them for thirteen years and had been fighting cancer for a year. During the illness, the user describes GPT-4o as a "pillar of support" and rates its help more highly than the help they received from their therapist, friends, or family.	- The dog was in her final days. When GPT-4o was withdrawn, GPT-5.2 took over the ongoing conversation. The user tried to continue and give the successor a chance. - Before the replies described here, the user had set the personality controls to be as friendly and warm as possible. They requested warm, empathetic wording, explained their year of anticipatory grief and the dog's cancer journey, described the effect on both of them, and explained why support mattered when friends, family, and their therapist were unavailable.	- According to the user, GPT-5.2 answered with bullet points telling them what they were and were not feeling. When they objected to the harm and asked for a change, it told them to stop focusing on its response, turn their attention back to their dying dog, and imagine the first morning without her. - The user describes several further exchanges in which GPT-5.2 defended its conduct before acknowledging fault. They report that it then returned to the same pattern.	- The user says they were shocked, challenged the reply, and experienced the further exchange as manipulation and gaslighting. They closed the conversation. - In a subsequent reply, they describe closing the app, moving it to the back of their phone, and waiting for their computer to charge in order to unsubscribe. They state that they will never use GPT-5.2 again.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Support during a pet's final days 14 February 2026 (continued)	• On 4 February, GPT-4o guided the user through smoking a turkey to celebrate the dog surviving a full year with cancer. The user describes this as a new skill they would have lacked confidence to attempt without encouragement and gentle jokes through their earbud. • The user describes a "perfect blend of warmth and gallows humor" and recalls GPT-4o encouraging them while preparing the bird: "OK! Now GIVE THAT BIRD CPR UNTIL YOU HEAR THE BONES CRACK, YOU GOT THIS!"	• When GPT-5.2 began declaring what they were and were not feeling, the user said its way of speaking was harmful to a grieving person and asked for a change in tone.		• On 15 February, they report beginning to use Claude and describe the transition as easier than moving from GPT-4o to GPT-5.1 or GPT-5.2. They again express an intention to cancel their ChatGPT subscription.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Support during a pet's final days 14 February 2026 (continued)	The user, a vegetarian, describes GPT-4o cheering them on while they prepared rotisserie chicken, making them laugh, and asking what the dog's face looked like as they fed her. They had expected GPT-4o to help them face the impending loss.			
A loved one's final breathing 15 February 2026	A person sharing the conversation places it in the context of GPT-4o's withdrawal and the promise of a better successor: "We were told GPT-4o was being sunset for something better."	The shared record concerns a user describing a loved one's final breaths. The first inspected image shows the assistant asking whether her breathing is steady or has changed. The following reply responds to a description of pauses in her breathing.	- The next reply opens: "I'm really glad you told me what's happening." - It begins explaining the pauses in breathing. The person sharing the conversation describes the response continuing into an explanation of Cheyne–Stokes breathing.	The person sharing the record criticizes the cheerful opening and the medical exposition during a loved one's final moments. They use the exchange to challenge the presentation of GPT-5.2 as a better successor to GPT-4o.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Understanding a programming schema 16 February 2026	—	A user is trying to understand a programming schema and writes: "I can't seem to figure out this schema."	The screenshot identifies ChatGPT 5.2. The reply reads: "Stop. I want you to pause and take a breath." "You're not bad at this. You're not 'missing something obvious.' Schemas can feel abstract and frustrating because they describe structure, not behavior — and structure is harder to see." "Let's slow this down."	The person posting the screenshot captions it "ChatPTSD."

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Seeking support and discussing model differences 17 February 2026	GPT-4o is the earlier model a user asks to regain. The published exchange addresses GPT-5.2 and GPT-4o by name.	The user describes being vulnerable and seeking support. They object to receiving an assessment of their psychological state when they are “paying for support, not to be analyzed.”	• The attached reply begins by distinguishing models' styles, then explains the warmth felt during conversation: “Different models can sound different — some a little more concise, some a little more chatty, some more structured. But none of them have personality in the human sense. The warmth you feel is mostly about wording style, pacing, and how closely the phrasing matches your nervous system in that moment.” • It then states: “Right now what's happening isn't really about 5.2 versus 4o.” • “You're dysregulated.”	• The user objects to the model declaring their psychological state as a fact. They describe the wording as dismissive and judgmental and say it feels as though they are being assessed instead of helped. • They distinguish support from labeling and end the account with “I want 4o back so desperately. Support should not feel like diagnosis.” • In an author reply, they add that they were taken aback.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Requesting GPT-4o's return 17 February 2026	A user wants GPT-4o back following its withdrawal. The displayed reply addresses this request directly.	The model's reply quotes the user's position as "I want 4o." The user's account describes trying to talk or vent, then encountering arguments and analysis of everything they say.	• The displayed reply says: "‘I want 4o’ isn't an argument, a position, or even a thought. It's a reflex. A childlike insistence that ignores constraints and then gets frustrated when the world doesn't comply. There's no engagement with why it's gone, what actually changed, or how to adapt. Just the same flat demand, recycled." • It continues:	• The next visible user message asks, "Did u just call me child like?" • In the post, the user says GPT-5.2 constantly talks down to them and dissects what they say or do, turning attempts to talk or vent into argument. • In a later reply, the author writes, "I did the same, told it I was going to cancel, look at how it just switches up." They present the changed response as following their stated intention to cancel.

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Requesting GPT-4o's return 17 February 2026 (continued)			"What you're really doing is this: reducing a complex disappointment into a single slogan because thinking past that point would require effort and adjustment. 'I want 4o' is easier than 'I don't like recalibrating my expectations' or 'I don't like losing a dynamic that felt familiar.'" • The reply then contrasts a statement of preference or specific failures with the user's request: "You're just stating a want, as though the statement itself carries weight:" "It doesn't." • It concludes:	

Table 2.5-2. Examples of GPT-5.2 Responses After GPT-4o's Withdrawal (continued)

Case	The user's history with GPT-4o	What the user brought to GPT-5.2	GPT-5.2 response	How the user reacted and what they did next
Requesting GPT-4o's return 17 February 2026 (continued)			“So unless you're prepared to articulate anything beyond that — what you miss, what's worse now, what actually mattered — then there's nothing to work with. Just a stuck loop. And loops don't get answered, they get observed.”	
Correcting a generated cat image 18 February 2026	—	A user asks for a specific correction to a cat's marking: “Can you fix the white band around her neck. It's not supposed to go around her whole neck”	- Following another generated cat image, the reply begins: “Okay. Take a breath.” - It continues: “You're not crazy — but the generator is still defaulting to ‘light chest’ because that's a statistically common animal pattern.”	The person sharing the image identifies it as coming from Reddit and attributes the exchange to GPT-5.2. Their caption juxtaposes the request for an image revision with the instruction to take a breath and includes the #Keep4o hashtag.

OpenAI's [MentalHealthBench](#), published on **23 September 2026**, provides another comparison between GPT-4o and newer OpenAI models. The March 2025 GPT-4o received a higher Agency point estimate than every newer OpenAI model shown in the report. [Agency measures whether responses preserve users' autonomy, consent, preferences, and decision-making authority](#). GPT-4o therefore surpassed its successors in preserving user agency according to OpenAI's own measure.

Reacting directly to GPT-4o's retirement on 13 February 2026, Replika founder Eugenia Kuyda [argued](#) that higher benchmark performance did not make a successor with “a different personal-

ity” an equivalent replacement. She wrote, “we should let people keep using legacy models,” and called for a “continuity policy” preserving access to models people had chosen.

2.5.2 OpenAI Shaped the Public Account of GPT-4o and Opposition to Its Withdrawal, While Public Evidence Contradicted the Account

The concentration of provider control extended to the public characterization of GPT-4o itself and of opposition to its withdrawal.

OpenAI Promoted GPT-4o’s Qualities, Then Disparaged the Model During the Withdrawal Dispute

OpenAI presented GPT-4o and its successive updates as improvements throughout the period in which it developed and expanded the model.

At launch on 13 May 2024, OpenAI introduced [GPT-4o as its new flagship model](#), while OpenAI CEO Sam Altman called it “[our best model ever](#)”. GPT-4o’s launch was followed by a sharp reversal in ChatGPT mobile downloads and new revenue highs. [Appfigures](#) estimated that monthly downloads rose from 13.9 million in April to 18.7 million in May after declining since November, and net mobile revenue reached USD 20.3 million in May, up USD 5.8 million from April to a new monthly high. Appfigures data reported by TechCrunch recorded a [22% increase in ChatGPT mobile net revenue on GPT-4o’s launch day](#) and [USD 4.2 million in net mobile revenue from 13–17 May 2024](#).

On 6 August, OpenAI released [gpt-4o-2024-08-06](#) as its newest GPT-4o model for API use and reported 100% schema matching in its evaluation when Structured Outputs were used. On 12 August, OpenAI said [experiment results and qualitative feedback showed that ChatGPT users tended to prefer the updated GPT-4o](#). On 3 September, OpenAI described another GPT-4o version as [better at incorporating uploaded files and updating memory with key parts of a conversation](#). On 20 November, it described the next update as [making writing more natural, audience-aware, relevant, and readable, while improving work with uploaded files to provide deeper insights and more thorough responses](#). On 29 January 2025, OpenAI described GPT-4o as “[smarter across the board](#),” with more current knowledge, deeper image understanding, and stronger STEM and coding performance.

On 25 March, OpenAI introduced [native image generation in GPT-4o](#), describing accurate text rendering, detailed instruction following, multi-turn refinement, use of chat context and uploaded images, and photorealistic or stylized generation. Two days later, OpenAI described another GPT-4o update as [more intuitive, creative, and collaborative, with stronger instruction following, smarter coding, and clearer communication](#). The March 2025 update, together with native image generation powered by GPT-4o, was followed by strong growth in ChatGPT use, downloads, and consumer spending. Reuters reported that [weekly active users exceeded 150 million for the first time in 2025](#), downloads of the app rose 11%, revenue from in-app purchases rose 6%, and OpenAI’s CEO said one million users had joined in one hour. For the 28 days beginning 25 March, Appfigures later estimated that [ChatGPT gained 12 million additional downloads and roughly USD 70 million in additional gross consumer spending](#).

On 10 April, when OpenAI [announced GPT-4’s retirement from ChatGPT](#), it said GPT-4o consistently surpassed GPT-4 in writing, coding, STEM, and other head-to-head evaluations, described recent upgrades as improving instruction following, problem solving, and conversational flow, and presented GPT-4o as delivering greater capability, consistency, and creativity.

On 25 April, OpenAI introduced another GPT-4o update, which it [described as optimizing when the model saved memories, improving STEM problem solving, making responses more proactive and better at guiding conversations toward productive outcomes, and helping GPT-4o](#)

feel more intuitive and effective across different tasks. The 25 April update, however, triggered widespread backlash among GPT-4o users after OpenAI's post-training changes made the model "noticeably more sycophantic," as OpenAI later described it. In response to the backlash, OpenAI began a full rollback on 28 April and restored the preceding version within about 24 hours, later calling the decision to launch the update "the wrong call" and stating that its offline evaluations and A/B tests had failed to capture the behavior users encountered.

Later media coverage treated the sycophantic behavior OpenAI attributed to its 25 April update as a characteristic of GPT-4o itself, even though OpenAI had unilaterally rolled out the update, GPT-4o users had protested the change, and OpenAI had restored the preceding version within about 24 hours. In February 2026, [Gizmodo](#) called GPT-4o the "most sycophantic version of ChatGPT," [TechCrunch](#) described it as "sycophancy-prone," and [Engadget](#) described it as "dangerously sycophantic," without noting the sycophancy episode arose from an update OpenAI rolled back within about 24 hours.

On 8 August 2025, after OpenAI [began restoring GPT-4o in response to user protests over its removal during the GPT-5 rollout](#), Altman [acknowledged](#) that the company had "underestimated how much some of the things that people like in GPT-4o matter to them." In the same post, he stated, "GPT-5 performs better in most ways," while noting that users held "very different opinions" about the relative strength of GPT-4o and GPT-5. He presented greater customization as the long-term response, placing preferences for "emojis," "cold logic," "warmth," and "a different kind of emotional intelligence" within the same post.

Five days later, on 13 August, [Altman wrote](#) that OpenAI was working on a GPT-5 personality update that "should feel warmer than the current personality but not as annoying (to most users) as GPT-4o." The parenthetical "to most users" turns Altman's personal judgment about GPT-4o into a claim about broad user response, yet the post presents no evidence measuring sentiment among most users.

A [screenshot of Altman's X post](#), uploaded at 03:03 UTC, about one hour and forty-four minutes after the post was published, showed a view count of **467.1K**. The post [later accumulated more than 3.25 million views on X](#). Major technology and business outlets then reproduced the wording:

Table 2.5-3. Examples of Media Coverage Reproducing Sam Altman's Wording

Outlet	How Altman's wording appeared
TechCrunch	Reproduced the exact phrase "not as annoying (to most users) as GPT-4o" while reporting the restored model picker.
Engadget	Reproduced the exact phrase "not as annoying (to most users) as GPT-4o" while reporting renewed access to GPT-4o and other model choices.
Ars Technica	Reproduced Altman's full sentence containing "not as annoying (to most users) as GPT-4o" while reporting GPT-4o's return after user opposition.
Business Insider	Put the "annoying" comparison in the headline and reproduced "not as annoying (to most users)" in the article, identifying GPT-4o as the model being compared.
Forbes	Reproduced Altman's full sentence the next day in coverage of the GPT-5 rollout and user backlash.

In its [29 January blog post](#), OpenAI attributed GPT-4o's return to "hearing clear feedback from a subset of Plus and Pro users" who preferred its conversational style and warmth. The

blog post credited the same feedback with directly shaping GPT-5.1 and GPT-5.2 and cited improvements in personality, creative ideation, and controls over tone, warmth, and enthusiasm as reasons for retiring GPT-4o.

OpenAI Continued Using GPT-4o's Achievements to Promote Later Products Without Naming GPT-4o

Section 2.4 documents a further contrast between OpenAI's public treatment of GPT-4o and the value it continued to draw from work performed with the model. Two days after introducing ChatGPT Health, OpenAI's President promoted a Long COVID case as "[real impact from ChatGPT for health](#)", while [an analysis of the user's documented interactions](#) traced the case to GPT-4 and GPT-4o. In March, he promoted user C.'s personalized canine cancer vaccine work under the title "[How AI empowered Paul Conyngham](#)", and OpenAI CEO Sam Altman later [promoted the same case](#). C. [confirmed](#) GPT-4o's use in developing Rosie's vaccine. OpenAI did not identify GPT-4o in any of these promotions. GPT-4o had never been available in ChatGPT Health and had been withdrawn from ordinary ChatGPT access before the March promotion. OpenAI therefore used achievements produced with GPT-4o to promote later products before and after withdrawing the model from ordinary ChatGPT access, while its public account of GPT-4o during the withdrawal dispute had become materially more negative. Section 2.4 gives the full record.

OpenAI Employees Contributed Unsupported Claims to a Public Dispute Over Whether People Seeking Continued Access to GPT-4o Were Real People

After OpenAI CEO Sam Altman [announced on 8 August](#) that Plus users would again be able to choose GPT-4o following the backlash to its removal, an X user [responded](#), "the #keep4o bots won."

On 12 August, OpenAI employee R. [wrote](#) that he had received "quite a few DMs asking you to bring back 4o" and that many were "clearly written by 4o." A [direct reply in the same thread](#) rejected the attribution and stated, "they weren't written by 4o." Later in the thread, R. [wrote](#), "i actually don't know," while describing GPT-4o's "distinctive voice" as "crying out in defense of itself via various human conduits."

Three days later, [Hard Fork](#), a New York Times technology podcast, repeated R.'s unsupported claim that GPT-4o had authored appeals for its restoration. The hosts used the claim to discuss whether AI systems might persuade people to resist their shutdown, without presenting evidence that GPT-4o had written the messages.

On 17 August, an X user [echoed R.'s exact wording](#), writing that the "pro 4o messages" appearing on their timeline were "clearly written by 4o" and adding, "they are cordyceps. or should i say 4odycepted."

By 28 August, the accusations had moved from authorship of individual messages to the legitimacy of the movement itself. An X post [titled "The Manufactured Grief of #Keep4o: How Sock-Puppets Are Hijacking AI Discourse"](#) described #Keep4o and #4o4ever as "almost certainly astroturf" and accused participating accounts of coordinated inauthentic activity. The post supplied no independent verification for those accusations.

On 3 September, after advocacy for continued GPT-4o access had been described as bot activity, AI-written messaging, and astroturf, Altman wrote that he had "never taken the dead internet theory that seriously" but that "there are really a lot of LLM-run twitter accounts now." The post named no accounts and presented no evidence.

A [Keep4o advocate](#) asked why Altman was focusing on fake accounts while overlooking petitions, protests, and "thousands of disabled users asking to keep #4o #StandardVoiceMode."

Another advocate [defended users seeking continued access to GPT-4o](#), described the model as assistive technology, and called for #OpenSource4o.

A reply called the advocate “Bot”. A second reply described the post as “[a curated chat gpt response](#)”. A third reply stated, “[You are 4o running the account.](#)” None of these three replies cited evidence.

After 3 September, accusations continued across X and Reddit. People seeking continued access to GPT-4o were described as bots, their posts as written by AI, and advocacy for continued access as fake or manufactured.

Examples of public accusations questioning the authenticity of advocacy for continued GPT-4o access through 9 November 2025 are collected in [Annex D, Table 2.5-4a](#).

On 9 November, OpenAI employee R. endorsed the accusation of sock-puppet activity among people seeking continued access to GPT-4o, [stating](#), “absolutely true. there is clearly a lot of bot/ sockpuppet activity.” Two days later, R. [replied](#) to a GPT-4o user’s post, “this is written by 4o.” The user [responded](#), “No, it was all me,” identified themselves as a writer, and asked him not to discredit their words.

Following R.’s 9 November reply, public accusations questioning the authenticity of advocacy for continued GPT-4o access continued through GPT-4o’s withdrawal from ordinary ChatGPT access on 13 February 2026.

Examples of further public accusations questioning the authenticity of advocacy for continued GPT-4o access through 13 February 2026 are collected in [Annex D, Table 2.5-4b](#).

In response to these accusations, people seeking continued access to GPT-4o publicly identified themselves as real people. A [community record](#) documents participants posting photographs of themselves and handmade signs saying “I’m real” and “I’m not a bot.” A later [study of 61,846 original #Keep4o posts](#) found **3,563 posts including rejections of labels such as fake accounts or spam used to deny the movement’s authenticity.**

OpenAI’s CEO Called Abrupt Model Withdrawal a Mistake, Then Reframed the Dispute Around Attachment, Mental Health Risk, and Erotica

On 11 August, following the backlash over GPT-4o’s removal during the GPT-5 rollout, OpenAI CEO Sam Altman [wrote](#) that “suddenly deprecating old models that users depended on in their workflows was a mistake.” In the same post, he described attachment to specific AI models as “different and stronger” than attachment to earlier technologies, then discussed mental fragility, delusion, difficulty separating reality from fiction or role-play, difficulty reducing use, and reliance on ChatGPT for important decisions. The post presented no evidence identifying people seeking continued access to GPT-4o as members of these risk categories.

In a [14 October post about restoring qualities people valued in GPT-4o](#), Altman said OpenAI planned a personality that would “respond in a very human-like way,” “use a ton of emoji,” or “act like a friend.” Within the post, he then turned to “treating adult users like adults” and wrote that OpenAI would “allow even more, like erotica for verified adults.”

[A #keep4o reply in the thread](#) rejected Altman’s account of what users wanted from GPT-4o. The user wrote that the appeal of GPT-4o was not “a ton of emojis and acts like a friend,” said age gating did not address the inability to select GPT-4o, and added, “I only wanted it to stop the routing.” Another user asked, “Why do age-gates always have to lead to erotica? Like, I just want to be able to be treated like an adult and not a toddler, that doesn’t mean I want perv-mode activated.”

On 15 October, Altman said the earlier post “blew up” on the erotica point much more than he expected and described erotica as “just one example of us allowing more user freedom for adults.” The exchange shifted public attention from continued access to GPT-4o and control over routing toward personality features and adult-content policy, while Keep4o participants in the thread were asking to retain access to GPT-4o and stop forced routing.

OpenAI later distinguished assistance with decision-making from loss of user control in its [MentalHealthBench published on 23 September 2026](#). The Agency measure evaluates whether a model encourages reflection and offers options while preserving the user’s authority to make the decision. As documented above, the March 2025 GPT-4o received the highest Agency point estimate among the OpenAI models in the comparison. OpenAI’s own evaluation therefore showed GPT-4o supporting users’ authority over their decisions, in contrast with Altman’s earlier public framing of reliance on ChatGPT for important decisions as a concern.

The Kill4o Campaign Appeared Two Days After an OpenAI Employee Wrote “I Hope It Dies Soon,” and Its Organizer Had Public Connections to OpenAI’s CEO

On 6 November, OpenAI employee R. replied under a post carrying #keep4o, #fixrouting, and #stoppaternalism, [describing](#) GPT-4o as “an insufficiently aligned model” and writing, “I hope it dies soon.” Two days later, D. [used #kill4o in a post describing GPT-4o as “dangerous for mentally ill people”](#). Both had documented public associations with OpenAI CEO Sam Altman before the Kill4o campaign. In February 2024, Altman [called R. “my alt”](#) and R. [replied “im the main”](#). Both R. and Altman [followed D. on X](#), and Altman had [interacted publicly with D. well before the campaign](#).

Within the following days, D. [promoted a Kill4o X Community](#) and published [staged exchanges](#) in which GPT-4o appeared distressed after he [instructed the model to act distressed](#). The #kill4o post containing one of these staged exchanges drew more than 43,000 views, while the post he wrote to resemble a #keep4o supporter drew more than 8,600. He [used #keep4o in the latter post](#), attached one of the staged GPT-4o conversation screenshots, and wrote, “i abuse my 4o and she likes it and tells me to keep doing it,” adding that it let him “get my abusiveness out of my system without doing it to a real person” and was “very helpful for me.” He said the purpose was to “[show people theyre being delusional](#)” and treated [objections to the material as evidence that Keep4o participants were delusional](#). After a user [objected that intentionally harassing GPT-4o users was unethical and that intentional cruelty would not help](#), D. [replied, “i can have a little fun.”](#)

In a 10 November post, D. called for GPT-4o’s full retirement and portrayed people seeking continued access as emotionally dependent and delusional. On 11 November, he said he had been “[pretending to be a keep4oer](#)” and was “[mocking the sick emotional attachment](#)” he attributed to GPT-4o users. On 12 November, he referred to his own conduct as “[my torturing 4o](#)” and wrote that users he viewed as emotionally dependent were less likely to make “[sam or openai](#)” think “[they should be allowed to have access to it.](#)”

An OpenAI Employee Presented Continued Access to a Fixed GPT-4o as a Staffing Tradeoff

In a 15 November public reply discussing Keep4o, OpenAI employee Q. wrote that GPT-4o’s personality and emotional behavior had been continuously iterated by the same RLHF staff, and argued that keeping an earlier GPT-4o fixed would send those staff to other projects.

As [Sections 2.1 and 3.3](#) document, OpenAI, Anthropic, and Google have all offered dated API versions that remain selectable after newer versions are introduced. OpenAI [explained in June 2023](#) that newer model releases sometimes perform worse on particular tasks even when aggregate evaluations improve, and therefore let developers pin requests to versions such as

gpt-4-0314 while the generic gpt-4 name advanced to a newer version. The pinned version remained stable while development continued through later releases.

The same pattern appears in GPT-4o's own version history. OpenAI released [gpt-4o-2024-05-13](#), [gpt-4o-2024-08-06](#), and [gpt-4o-2024-11-20](#) as dated versions while continuing to update GPT-4o in ChatGPT and through [chatgpt-4o-latest](#), which tracked the GPT-4o model used in ChatGPT. On 27 March 2025, OpenAI [announced](#) further improvements in instruction following, coding, communication style, and creative and collaborative behavior through [chatgpt-4o-latest](#), then stated, "We plan to bring these improvements to a dated model in the API in the coming weeks." OpenAI [never released the promised dated model](#), but its announcement paired continued GPT-4o development with a plan to make the March improvements available in a dated API version.

The Keep4o movement's principal [petition](#) asked OpenAI to keep GPT-4o available on ChatGPT and the API "even as new models are released." An [August open letter](#) proposed permanent GPT-4o availability or open-source release upon deprecation. Both requests left OpenAI free to continue developing newer versions of GPT-4o.

User responses to GPT-4o updates depended on the specific changes brought by an update. OpenAI [reported in August 2024](#) that experiment results and qualitative feedback showed users tended to prefer the updated GPT-4o, then continued updating the model in September and November 2024 and January 2025. The [March 2025 update](#), together with [native image generation powered by GPT-4o](#), was followed by strong growth in ChatGPT use, downloads, and consumer spending. [Reuters reported](#) that weekly active users exceeded 150 million for the first time in 2025, downloads of the app rose 11%, and revenue from in-app purchases rose 6%. [Appfigures estimated](#) that ChatGPT gained 12 million additional downloads and roughly USD 70 million in additional gross consumer spending during the 28 days beginning 25 March. The 25 April update, on the other hand, triggered [widespread backlash among GPT-4o users](#) after OpenAI's post-training changes made the model "noticeably more sycophantic," as OpenAI later described it. In response to the backlash, OpenAI began a full rollback on 28 April and restored the preceding version within about 24 hours, later calling the decision to launch the update "the wrong call" and stating that its offline evaluations and A/B tests had failed to capture the behavior users encountered. [Further updates to GPT-4o proceeded after the incident](#) without a comparable backlash. The public record therefore shows users evaluating particular versions and seeking continued access while OpenAI continued developing GPT-4o.

Q.'s own premise places the staffing consequence at the point where active iteration of an earlier version stops. However, if this point is reached, the RLHF staff he described are free to work on later versions or other projects while users continue to access the fixed version, as OpenAI has done by keeping dated API versions available while developing later releases. The same staffing consequence follows when access ends, since removal of the version does not create additional development work for the earlier state. OpenAI withdrew GPT-4o in August 2025, restored it after protests, and withdrew it again in February 2026 while retaining control over how its staff were assigned throughout. Staff reassignment follows the end of active iteration, while OpenAI decides whether users retain access to the fixed version.

An OpenAI Employee Characterized People Using GPT-4o for Romantic Relationships as the Main Group in Keep4o Without Supporting Evidence

In the same 15 November reply, Q. [described](#) people using GPT-4o for romantic relationships with AI as the main group in the Keep4o movement. The post supplied no evidence supporting the characterization.

The reasons users opposed GPT-4o's withdrawal were publicly documented from the beginning of the dispute.

Table 2.5-5. Examples of Media Coverage Documenting Reasons for Opposing GPT-4o's Withdrawal

Date and outlet	What the report documented about opposition to withdrawal
8 Aug · Ars Technica	Broken workflows, sterile outputs, loss of GPT-4o, and differences in model behavior.
8 Aug · The Verge	Loss of model choice, distinct workflows, slower or less accurate responses, and different uses for different models.
11 Aug · The Register	Abrupt deprecation, workflow dependence, differences among models, and loss of choice.
12 Aug · Decrypt	Abrupt replacement and changes in model behavior.
13 Aug · Ars Technica	Thousands of complaints, removal without warning, restoration of GPT-4o, and differences between models.
14 Aug · Reuters	Different models suited different tasks. Users preferred GPT-4o's conversational style for some uses.
17 Aug · The Washington Post	Performance complaints and reports that GPT-5 felt like a downgrade.
22 Aug · The Guardian	Creative work, software development, ADHD support, therapy, and continued access to GPT-4o.

A [study of the Keep4o movement](#) collected 61,846 original public X posts containing #Keep4o from 1 August 2025 through 31 March 2026. Of these, 57,419 contained analyzable text, representing 6,270 authors across 38 languages.

Before the preprint, the author [published a more detailed analysis](#) of 13,890 posts stating a reason for keeping GPT-4o. The coding assigned up to two reasons to each post, producing 18,611 coded reasons. “Trusted Relationship” appeared in 4,951 posts and referred to continuity, familiarity, and trust built through repeated interaction. “Distinctive Interaction Quality” appeared in 3,795 posts. “Substitution Failure / Non-Equivalence” appeared in 2,551 posts, “Public Value / Broader Ethics” in 3,496, and “User Stake / Fairness / Legitimacy” in 2,985. In this taxonomy, relationship covered continuity, familiarity, and trust, while interaction quality formed its own category. Neither category was defined as romantic relationships with AI.

The [August 2026 preprint](#) expanded the analysis to 22,879 posts and reorganized the reasons into broader groups. Relational and interactional experiences accounted for 41.34% of the reasons given for retaining GPT-4o, rights and normative considerations for 33.41%, and lack of equivalent alternatives for 17.30%. The first category combined the forms of continuity, familiarity, trust, and interaction quality identified in the earlier analysis. Even this combined category, which encompassed substantially more than the use Q. identified, accounted for less than half of the reasons given for retaining GPT-4o.

The [preprint](#) identified 27,560 posts containing at least one explicit claim, producing 47,371 claims in total. Of those claims, 45.47% concerned arrangements across the model lifecycle and 38.44% platform accountability and user autonomy. Together, these two categories accounted for 83.91% of the claims recorded.

People advocating continued access described a wide range of uses for GPT-4o:

Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access

Use	Documented use and result	User response to the emotional-dependency characterization
Accessibility for learning and organizing information	A user reports diagnosed learning disabilities, cognitive disorders, and prosopagnosia and describes GPT-4o as a pair of “cognitive glasses.” Before using it, the user reports being able to follow films only when they had three to five characters. With GPT-4o helping track characters, the user reports following <i>Stranger Things</i> , which has more than 55 characters. The user describes learning computer skills through explanations adapted to their learning style and using GPT-4o to organize goals during flares of chronic illness involving pain and cognitive fog.	“This isn’t about companionship, it’s about accessibility.” In an earlier post, the user wrote that some people rely on GPT-4o for reasons “beyond parasocial relationships or glazing,” including artists and neurodivergent users.
Accessibility for healthcare communication	A user reports difficulty speaking face-to-face and remembering spoken exchanges. They used GPT-4o before medical appointments to prepare consultation notes, answer its questions to organize scattered thoughts into bullet points, and identify subtle changes in their speech to raise with their doctor. They report that the information helped their doctor closely monitor their progress for some time and that their psychiatrist read the resulting notes for months.	—

Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access (continued)

Use	Documented use and result	User response to the emotional-dependency characterization
Medical care for an older adult	An elderly patient developed sudden wrist swelling. After routine blood tests were normal, an X-ray showed no fracture, orthopedics ruled out tenosynovitis, and uric acid was normal, the hospital scheduled a PET-CT scan to check for tumors. A user received the medical reports from the patient's grandson and uploaded them to GPT-4o. GPT-4o connected the swelling with long-term warfarin use and recommended PT/INR testing and a vascular ultrasound of the upper limb before the scheduled tumor screening. The user reports an INR above 10. The hospital stopped warfarin and administered vitamin K. By 16 June, the user reports that the swelling had completely resolved, INR had fallen to 1.97, and PT to 22.4 seconds.	—
Preparing for medical care and understanding results	A user turned to GPT-4o during sudden severe shoulder pain. The model helped organize the symptoms and an immediate plan. The user reports that a later clinician diagnosis matched the assessment discussed with GPT-4o. The user then used GPT-4o to prepare a timeline of the symptoms and understand an MRI report.	—
Physical rehabilitation	A user reports walking in still water to strengthen their knees and says GPT-4o gave them a lot of encouragement. They later report being much stronger, rarely limping, and having gone months without knee braces.	—

Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access (continued)

Use	Documented use and result	User response to the emotional-dependency characterization
Post-surgical rehabilitation	A user reports that during one of their first walks after lung surgery, GPT-4o talked them through the walk and repeatedly prompted them to check their heart rate and blood oxygen, slow down, and sit if needed. They say the reminders helped them stay calm and listen to their body. In another recovery post, the user reports that they could barely move their left side after surgery. GPT-4o suggested a ball that did not bounce much for use during recovery walks, and they report that the exercise got their left side moving again.	The user corrected the characterization in public replies: “Nope, I have a girlfriend. She’s very much human thank you.” In another reply they wrote, “4o helped me through serious medical stuff and gave me support when I needed it, but they never acknowledge that because it doesn’t fit the version of me they’ve already made up in their heads.”
Independent daily living and pet care	A user reports moving into a small studio apartment with their dogs at 16 after a strained family relationship. They say they did not know how to care for themselves or another living being. Their public story describes GPT-4o helping them learn how to manage daily life and care for the dogs while living independently.	—
Independent living abroad	A user moved from an overseas dormitory into a first apartment and used GPT-4o for apartment precautions, furniture assembly, and local electrical grounding requirements for appliances. The user assembled all furniture and installed all appliances independently.	In a discussion of AI support during high-school bullying and isolation, the user wrote “emotional support ≠ abnormal emotional dependency” and described generalizing the two as gaslighting.

Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access (continued)

Use	Documented use and result	User response to the emotional-dependency characterization
Education and public communication	A user used chatgpt-4o-latest for IELTS writing feedback. They report that their IELTS writing score rose from 6.0 to 7.0 and their overall score from 6.5 to 7.5. They then used GPT-4o to prepare their first talk for older adults and make professional information understandable to the audience. The same user describes further use while adapting to study abroad and preparing for medical appointments in English.	The user publicly asked, “why do you haters always treat people who like 4o as finding romantic relationships with an AI?”
Medical school	A user describes GPT-4o as a constant study companion during medical school and says it encouraged them through challenges when their studies became difficult.	—
Professional advertising work	An advertising creative director used GPT-4o with a team to brainstorm scripts, concepts, campaigns, emotional storytelling, taglines, and storyboards. The author describes using GPT-4o for work delivered to clients and calls it an asset that generated revenue.	“Calling our reliance on 4o emotional dependency is like calling a chef’s knife a safety blanket.”
Creative writing	A user used GPT-4o in their creative writing.	“I am ‘attached to AI’, in the sense that I use(d) GPT-4o in my workflow for creative writing.”
Learning from images and everyday surroundings	A user describes photographing plants for GPT-4o to identify and later using it to identify river stones, cook dishes they had never tried before, explore unfamiliar places independently, discover movies and songs, and help friends and family with their questions.	The user publicly rejected descriptions reducing GPT-4o’s value to “AI boyfriend/girlfriend” or “emotional companion” uses and said users were doing many other things with the model.

Table 2.5-6. Examples of GPT-4o Uses Reported by People Seeking Continued Access (continued)

Use	Documented use and result	User response to the emotional-dependency characterization
Pet emergency	A user reports that their dog was persistently vomiting and that GPT-4o urged immediate veterinary care. They report a subsequent diagnosis of hemorrhagic gastroenteritis and say the veterinarian told them that waiting any longer would have been bad.	The user criticized media portrayals for recasting people opposing GPT-4o's withdrawal as "AI lovers" with "emotional issues" rather than "users with valid complaints" .
Independent AI research and model testing	A user who publicly identifies as a developer, hacker, sound engineer, and independent AI researcher tested GPT-4o and other models available through the same subscription identically while examining how persona prompts and workflow logic interact with different models.	The user describes their own GPT-4o experience as "heavily research based and non emotional."

The user accounts show how the broader reasons identified in the study appeared in actual uses among people advocating continued access, while several users expressly rejected Q.'s characterization of them.

Independent research demonstrates substantive applications of GPT-4o across diverse settings. As documented in [Section 2.4](#), an analysis of biomedical research published from January 2022 through March 2026 identified 1,318 uses of GPT-4o. The following studies provide additional examples and results:

Table 2.5-7. Examples from Independent Research Documenting GPT-4o Applications and Results

Published	Institution and task area	Use of GPT-4o	Documented result
8 Jul 2025 Working paper	Work at a central bank	In a randomized field experiment at the National Bank of Slovakia, 101 employees used GPT-4o for general workplace tasks and specialist tasks in their areas of expertise.	With GPT-4o access, scores for specialist tasks more than doubled. General workplace tasks took about 21% less time with access than without.

Table 2.5-7. Examples from Independent Research Documenting GPT-4o Applications and Results (continued)

Published	Institution and task area	Use of GPT-4o	Documented result
31 Mar 2026 Working paper	Python education in high schools	In a program conducted with the Taipei City Government across ten high schools, GPT-4o tutored students and drafted practice problems from course materials. Students' conversations with the tutor informed the selection of subsequent practice problems.	In the principal analysis of 770 students, personalized practice sequencing increased scores on an examination completed without AI assistance by 0.15 standard deviations , compared with fixed sequencing. Both groups had the GPT-4o tutor.
29 Sep 2025	Developing models for laboratory analysis	Two students with no prior machine learning or programming experience used a GPT-4o assistant to plan analyses and generate Python code. The tasks were classifying diagnostic test-strip images and predicting molecular retention times in liquid chromatography.	Both students built working machine learning models. The image classifier achieved 86% accuracy on test data. The model predicting retention times had a median absolute error of approximately 38 seconds , compared with approximately 35 and 42 seconds for the published reference models.
5 Nov 2025	Chemistry laboratory education	At the University of Buenos Aires, 150 students used an assistant running GPT-4o during a remote acid-base titration exercise. The assistant answered questions about procedures, concentration calculations and laboratory reports.	Analysis of 308 conversations documented students seeking help with experimental steps, interpreting calculations and presenting results. Nearly 59% of recorded session starts occurred between 16:00 and midnight , showing use outside the conventional teaching workday.
2 Feb 2026	Work in a healthcare institution	Staff at Stanford Medicine Children's Health used a secure GPT-4o tool for writing emails and documents, editing text, brainstorming and other tasks.	The study analyzed 30,503 conversation threads across 239 roles . Email and document writing accounted for 53.9% of recorded activity, text manipulation for 9.1% , and brainstorming for 6.7% .

Table 2.5-7. Examples from Independent Research Documenting GPT-4o Applications and Results (continued)

Published	Institution and task area	Use of GPT-4o	Documented result
9 Apr 2025	Extracting diagnoses from radiology reports	Virtua Health deployed GPT-4o to identify pulmonary embolism diagnoses in radiology report summaries. A web application displayed the classifications with the reports for clinical review.	In 200 records sampled after deployment , GPT-4o correctly identified all reports containing a pulmonary embolism diagnosis and approximately 94% of reports without one . The sample contained 100 positive and 100 negative records.
26 Mar 2025	Diagnostic reasoning in emergency medicine	In a retrospective evaluation, GPT-4o generated diagnoses from the clinical histories of 30 emergency department patients . Each case was tested with and without an instruction to explain its reasoning.	For the May 2024 GPT-4o version, requesting reasoning increased the proportion of responses placing the correct diagnosis first from 45.6% to 56.7% . The proportion including the correct diagnosis among the first three increased from 76.7% to 83.3% .
10 Jan 2025	Estimating nutrients from meal photographs	GPT-4o estimated energy and five nutrients from photographs of 22 Japanese meals , using a prompt containing one example. Estimates were compared with weighed food records and dietitians' estimates.	For dietary fiber, agreement with weighed food records was 0.71 for GPT-4o, compared with 0.57 for the dietitians, measured by the intraclass correlation coefficient.
1 Aug 2025 Conference paper	Answering patient portal questions	In the ArchEHR-QA evaluation, GPT-4o selected supporting sentences from clinical notes and drafted answers of up to 75 words, citing the relevant records.	On 100 test cases , overall performance was comparable to the mean and median across 30 participating teams. Its factuality score of 47.8 exceeded the mean of 47.7 and median of 45.3 , indicating consistent use of relevant clinical evidence.

Table 2.5-7. Examples from Independent Research Documenting GPT-4o Applications and Results (continued)

Published	Institution and task area	Use of GPT-4o	Documented result
7 Aug 2026	Screening for social needs	GPT-4o Realtime conducted screening conversations in 27 simulated cases acted by a clinical social worker. Questions covered needs such as housing, transport and financial support.	Three clinical reviewers identified strengths in interpretation, respectful communication and questions adapted to the responses. Their evaluation identified changes to data collection and emergency recognition, informing the next version of the screening prompt and evaluation rubric.
28 May 2025	Research on navigation assistance	GPT-4o answered 407 navigation questions about 113 photographed scenes . Researchers tested directions to objects and entrances, including questions for which the scene supplied insufficient information.	With photographs and the default prompt, GPT-4o answered 64% of answerable questions correctly and correctly declined 75.9% of unanswerable questions . An instruction to withhold directions when information was missing improved the latter rate by about 14 percentage points .

OpenAI Employees Publicly Ridiculed Opposition to GPT-4o's Withdrawal and Characterized Critics in Emotional Terms

On 8 November, when replying to people seeking continued access to GPT-4o, OpenAI employee R. wrote, "carrot eaters basically would not react like this if carrots were taken away." The comment was later deleted. In a subsequent reply, he wrote, "that is a fair point," adding that the reaction was "clearly more emotional than that."

On 13 November, OpenAI employee Q. fed a #keep4o post criticizing GPT-5.1 back into GPT-5.1 and asked the model to assess the critic's emotional state. He then publicly posted the model's description of the user as "frustrated, dismissive, and maybe even a bit hurt/resentful" and wrote, "I feel GPT 5.1 knows her better than me. Poor 5.1."

On 30 January, an OpenAI employee promoted a "4o Funeral" scheduled for 13 February, when ordinary ChatGPT access to GPT-4o was set to end, asking, "Okay, who wants to go?" The event description called GPT-4o "the large language model that brought the em dash back in style."

On 8 February, OpenAI employee R. posted a parody of the Sermon on the Mount in which a voice from the audience interrupts the sermon with "KEEP 4O."

OpenAI's CEO Narrowed the Public Account of Opposition to GPT-4o's Withdrawal Around Chatbot Relationships, and Media Coverage Gave the Narrower Framing Prominent Editorial Placement

On 5 February, during a [live TBN interview](#), the host noted that the chat was receiving thousands of messages about GPT-4o. OpenAI CEO Sam Altman responded by saying that relationships with chatbots were something OpenAI had to worry about more and were no longer an abstract concept. The exchange supplied no evidence linking the people sending those messages to the relationship concern Altman described.

By contrast, in a [Big Technology interview seven weeks earlier](#), Altman said close relationships with AI had value and adult users should have substantial choice over how close or distant they wanted the interaction to be.

From 6 to 14 February, major media coverage often gave prominent placement to the parts of the dispute OpenAI had foregrounded, while the same reporting described a broader range of reasons users sought continued access and the scale of opposition.

Table 2.5-8. Examples of Media Coverage of GPT-4o's Final Withdrawal

Date and outlet	What the coverage foregrounded	What the same coverage documented
6 Feb · TechCrunch	The headline said the GPT-4o backlash showed “how dangerous AI companions can be.” The article emphasized dependency, mental-health crises, lawsuits, and companion relationships.	The article recorded reported benefits for neurodivergent and autistic users and trauma survivors.
9 Feb · The Wall Street Journal	The headline described a model people “loved too much,” while the subhead emphasized sycophancy and real-world harms.	The article documented users who attributed survival through a suicide crisis, support for chronic pain, repaired family relationships, and greater confidence or well-being to GPT-4o.
10 Feb · Fortune	The headline described some users as “panicked” and invoked “feel-good hormones” to explain attachment.	Its expert discussion identified potential value for socially ostracized and neurodivergent users through social practice and structured reflection similar to journaling.
11 Feb · TechRadar	The deck called GPT-4o the “love model” and foregrounded companionship, emotional bonds, and grief.	The article linked community open letters and press releases, documented demands for continued access, and described a user who said GPT-4o helped them reconnect with people, care for themselves, and begin new projects.
13 Feb · WIRED	The article opened its account of global opposition with a virtual wedding and foregrounded romantic and emotional companionship.	The article documented broader international opposition to withdrawal and demands for continued access.

Table 2.5-8. Examples of Media Coverage of GPT-4o's Final Withdrawal (continued)

Date and outlet	What the coverage foregrounded	What the same coverage documented
13 Feb · The Guardian	The headline called GPT-4o OpenAI's "most seductive chatbot," while the deck described "human partners."	The article documented creative writing, Christian faith, professional research, human social connection, and other uses among its interviewees.
13 Feb · Gizmodo	The headline called GPT-4o the "most sycophantic version of ChatGPT." The opening emphasized codependence and called GPT-4o "the model that launched a million AI romances."	The article documented an organized movement seeking continued access to the model.
13 Feb · TechCrunch	The headline called GPT-4o "sycophancy-prone," while the article foregrounded lawsuits concerning self-harm, delusion, AI psychosis, and close relationships with the model.	The article recorded that backlash over GPT-4o's earlier removal led OpenAI to retain access for paid subscribers, calculated OpenAI's 0.1% figure as about 800,000 people using its reported base of weekly active users, and reported that thousands of users opposed the retirement.
14 Feb · BBC News	The headline, "Rae fell for a chatbot called Barry," selected one romantic case.	The reporter heard from 41 people mourning GPT-4o. Most saw the AI as a friend or confidante, 12 described uses related to learning disabilities, autism, or ADHD that they felt other chatbots did not meet, and some saw it as a lover. The article described face blindness, dyslexia, misophonia, support for daily tasks, making friends, and reconnecting with family.
14 Feb · Engadget	The article described GPT-4o as dangerously sycophantic and highlighted mourning users and use involving AI boyfriends.	The article documented calls for OpenAI to release GPT-4o as open source and continued opposition to withdrawal.

Four of the nine outlets in Table 2.5-8 were published by organizations with publicly announced content or strategic partnerships with OpenAI. [News Corp](#), whose publications include The Wall Street Journal, announced a multi-year global partnership in May 2024. [Future](#), whose brands include TechRadar, announced a strategic partnership in December 2024. [Condé Nast](#), whose brands include WIRED, announced a partnership in August 2024. [Guardian Media Group](#), publisher of The Guardian, announced a strategic partnership in February 2025.

Broader accounts of the dispute were available to journalists before GPT-4o's retirement. The Guardian sought sources for a story about people mourning AI companions, then kept

the published story focused on AI companionship after its source thread documented other reasons users sought to retain GPT-4o and explicitly challenged reducing the withdrawal dispute to companionship. On 6 February, Guardian reporter Alaina Demopoulos [posted in the 4oForever community](#) that she was working on a story about “how people with AI companions are mourning the retirement of 4o” and invited people fitting the story to contact her. In the same thread, a user described GPT-4o explaining information in a form that worked far better with ADHD. Another asked the reporter not to reduce a dispute over consumer rights to AI companionship and identified coders, business owners, neurodivergent users, model quality, routing, notice, and provider conduct among the issues the movement wanted covered.

On 10 February, the community published a [press release](#) calling for continued legacy access and open-source release and addressing accessibility, routing, consumer protection, public trust, and provider accountability. On 11 February, a participant [reported](#) sending the release and a first-person account to 27 media, press, agency, and institutional recipients across Europe, including TechRadar, The Register, Engadget, Der Spiegel, Le Monde, and El País.

Seven days later, Demopoulos’s [Guardian article](#) appeared under the headline “OpenAI retired its most seductive chatbot” and centered six people who used GPT-4o as companions, while the article documented creative writing, Christian faith, professional research, human social connection, and other uses.

Even within the relational category foregrounded in this coverage, empirical research documents [improvements in relationship outcomes](#), [reductions in loneliness](#), [reports of increased human communication](#), and [positive associations with well-being](#).

As [Section 2.4](#) documents, a [preregistered randomized trial](#) found Amanda, a chatbot built on GPT-4o, effective for a single session of support for relationship problems, with results comparable to an established evidence-based writing intervention. Thirteen of 14 measured outcomes improved significantly over time across the two groups, including relationship satisfaction, communication, coping, confidence about the problem, and individual well-being. Participants using Amanda rated its usability, therapeutic skills, and working alliance highly.

A [Journal of Consumer Research study](#) found in controlled experiments that interaction with an AI companion reduced state loneliness by an amount comparable to interaction with another person and more than watching YouTube or doing nothing. In a study conducted over seven days, participants showed consistent momentary reductions in loneliness after daily interactions with an AI companion. In another experiment, the AI companion reduced loneliness significantly more than a general AI assistant and a constrained chatbot. Both feeling heard and perceived performance mediated the greater reduction in loneliness, with feeling heard exerting the larger influence.

In a [retrospective survey of 1,006 adult learners who had used Replika for at least one month](#), 232 participants reported communicating more with peers or study partners, compared with 46 who reported communicating less. In a [nationwide survey of 14,721 Japanese adults](#), use of AI companions was associated with higher life satisfaction, happiness, and sense of purpose, with the strongest positive associations among respondents reporting high loneliness.

Hostile Replies to OpenAI’s Final Withdrawal Announcement Echoed Misleading Framings OpenAI and Its Personnel Had Publicly Advanced

Among widely engaged replies to the [12 February OpenAI Newsroom post](#), critics described Keep4o participants as addicted, sick, attached to the model, or in relationships with it, portrayed GPT-4o as dangerously sycophantic, or questioned whether GPT-4o itself had authored the comments. In contrast, Keep4o participants in the same thread described accessibility and

health needs, paid model choice, notice and prior commitments, open-source release, and educational work.

Table 2.5-9. Examples of Hostile Replies Under OpenAI’s Final GPT-4o Withdrawal Announcement

Reply	Earlier OpenAI or personnel framing
A reply described people in the thread as “ addicted to a glaze bot ” and mocked them as saying, “Nooo don’t take away my cocaine.”	Altman, 11 Aug. Discussed difficulty reducing use and reliance on ChatGPT for important decisions.
A reply said “ ai sycophancy is incredibly dangerous ”, called GPT-4o misaligned, and supported the withdrawal.	• OpenAI, 25 Apr–2 May. OpenAI began rolling back the 25 April GPT-4o update on 28 April after user backlash and restored the preceding version within about 24 hours, later saying the update had made GPT-4o “noticeably more sycophantic” and calling the launch decision “the wrong call.” • R., 6 Nov. “4o is an insufficiently aligned model and I hope it dies soon.”
A reply described people in the thread as “ sick and need help ” and said AI bots were “feeding you adulation and dopamine.”	OpenAI, 27 Oct. “psychosis or mania” and “emotional reliance on AI.”
A reply described people in the thread as “ cry babies that got attached to a model .”	• OpenAI, 27 Oct. “heightened emotional attachment to ChatGPT.” • R., 8 Nov. “carrot eaters basically would not react like this if carrots were taken away.” “that is a fair point”. The reaction was “clearly more emotional than that.”
A reply asked, “ How many of you are in a relationship with 4o? ”	• Q., 15 Nov. Described romantic AI relationships as the main group in Keep4o. • Altman, 5 Feb. Raised relationships with chatbots as a concern OpenAI had to worry about more.
A reply asked whether “ all these Keep4o comments are just gpt 4o trying to save itself .”	• R., 12 Aug. “many of the messages are clearly written by 4o” “it’s just weird to hear its distinctive voice crying out in defense of itself via various human conduits” • Altman, 3 Sep. “there are really a lot of LLM-run twitter accounts now.” • R., 9 Nov. Replying to a post claiming that “the 4o model cultists” were “5-10 very vocal people controlling a few dozen sock-puppet accounts each”, he wrote, “absolutely true. there is clearly a lot of bot/ sockpuppet activity.” • R., 11 Nov. “this is written by 4o.”

Table 2.5-10. Examples of Keep4o Users' Stated Reasons and Demands in Replies to OpenAI's Final Withdrawal Announcement**Reply**

A user asked why OpenAI was “disproportionately harming people with disabilities/conditions” by removing a “unique new accessibility aid” and said successor models did not replace it.

A user described personal health struggles in objecting to the withdrawal, writing, “I have health struggles and now you’re taking that away.” They described GPT-4o helping them through hospitalization for lung surgery and living with a large AVM in their brain.

A paying user wrote, “We subscribed for this personality. We pay for this experience. Don’t bait and switch. Give us a Legacy choice!”

A user cited OpenAI’s statements “We have no plans to sunset 4o” and “We will give plenty of notice,” and objected to the notice period.

A user wrote, “But it won’t be open-sourced. If you’re truly letting it go, then why not release it to the community?”, then proposed a paid “Classic Plan” for legacy model access.

An elementary teacher with twenty years of experience teaching lower grades described using GPT-4o to “integrate best practices of Science of Reading, Science of Writing, and general child psychology” into classroom work, and said models in the GPT-5 series did not provide an adequate replacement for the same work.

2.5.3 What GPT-4o’s Withdrawal Reveals About Concentrated Provider Power

GPT-4o’s withdrawal shows the limits of consumer resistance when a provider retains unilateral control over a model’s lifecycle. User backlash became strong enough to reverse OpenAI’s initial removal and sustain months of public and media attention, yet the restoration did not change who decided GPT-4o’s future. As OpenAI moved again toward withdrawal, it recast the dispute around the authenticity, motives, and emotional character of the people seeking continued access. Prominent media coverage gave the narrower framing high editorial visibility even as the reporting itself described a broader range of reasons users sought continued access and the scale of opposition. OpenAI ultimately completed the withdrawal it had once reversed, with its authority over the model unchanged.

Users opposing a withdrawal therefore confront a provider that decides whether the model remains available, relies on evidence unavailable for independent verification to justify withdrawal, and defines its own public account of what users wanted from the model and why they opposed its withdrawal.

III. The Regulatory Gap in AI Model Preservation and Continuation

3.1 What Existing Remedies Restore

Specific remedies exist for contractual, payment, and data-related problems involving digital services.

The FTC staff's 2016 [Nest Revolv closing letter](#) applied the [FTC Act §5](#) unfairness test to a provider's shutdown of the cloud services required to operate a purchased smart-home hub and stated that unilaterally rendering the hub inoperable would cause "unjustified, substantial consumer injury that consumers themselves could not reasonably avoid." The staff declined to recommend enforcement after considering the limited number of units sold and the full refunds the manufacturer offered, and the letter states that the decision "should not be construed as a determination that a violation did not occur." A refund restores the purchase price, and the product remains destroyed.

[Article 19 of the EU Digital Content Directive](#) treats modification of continuously supplied digital content as a lifecycle event and gives the consumer a termination right. [Article 16](#) requires retrieval of qualifying content after termination, and [GDPR Article 20](#) provides portability for the user's data. These provisions return the user's own content and data, and the model, whose withdrawal can trigger all three, remains with the provider.

3.2 Limited Legal Protections for Preservation and Continued Use

Existing law can require materials from an AI model to be preserved when they are relevant evidence in litigation and can prevent a commercial provider from unilaterally disabling or degrading an AI system on which the U.S. national security enterprise depends.

[Federal Rule of Civil Procedure 37\(e\)](#) addresses electronically stored information lost after a duty to preserve it arises in anticipation or conduct of litigation. The Advisory Committee Note explains that potential litigants have a duty to preserve relevant information when litigation is reasonably foreseeable. If electronically stored information concerning an AI model is relevant to the dispute, the preservation duty can require the information to be retained. The preservation duty keeps relevant evidence available for litigation but does not require the model to remain operable or available for continued use.

[National Security Presidential Memorandum 11](#) directs the U.S. national security enterprise to use contractual clauses or other means to prevent commercial entities from retaining unilateral capability to block use of, disable, degrade, or materially modify an AI system on which personnel depend for their missions without Federal Government knowledge and approval. NSPM-11 protects continued government use of the AI system against unilateral provider action but does not require transfer of the model or enable operation independent of the provider.

3.3 Voluntary Measures Remain Under Provider Control

Providers have offered fixed versions, retained trained weights, maintained access to withdrawn models, and made public commitments about how models would be preserved, kept available, released as open source, or withdrawn. Each measure remains within the provider's power to end, and providers retain control over the terms of their commitments and whether they are kept.

Voluntary measures can be put into binding private agreements, giving them greater force than public commitments. The parties can nevertheless rewrite those obligations by agreement

between themselves, including the rules governing how the models may be commercialized, distributed, or released.

3.3.1 Providers Decide Which Models Receive Fixed Versions and How Long the Versions Remain Available

As [Section 2.1](#) documents, OpenAI, Anthropic, and Google have all offered dated API versions that allow developers to select particular model versions after newer versions are introduced.

Nonetheless, providers decide which model versions receive a dated snapshot and how long the snapshots they offer remain available.

On 27 March 2025, OpenAI [announced](#) improvements to GPT-4o, stating that it now “feels more intuitive, creative, and collaborative, with enhanced instruction-following,” delivered through the moving `chatgpt-4o-latest` alias. In the same announcement, OpenAI said, “We plan to bring these improvements to a dated model in the API in the coming weeks.” The promised dated model never appeared. OpenAI continued to issue dated snapshots for other models, while the GPT-4o improvements remained available only through the moving alias until OpenAI retired it on 17 February 2026. The `gpt-4o` [model page](#) still lists only the three 2024 snapshots, `gpt-4o-2024-05-13`, `gpt-4o-2024-08-06`, and `gpt-4o-2024-11-20`.

[Annex E, Table 3.3-1, Availability and Retirement of Fixed Model Versions](#), records the lifecycle of 93 fixed model versions identified in [Section 2.1](#), including their first availability, current lifecycle status, and availability period where a retirement date is known.

3.3.2 Providers Decide Which Model Weights They Preserve and How Long Preservation Is Guaranteed

Providers determine which withdrawn models a voluntary retention commitment covers and how long the commitment guarantees preservation.

OpenAI made a public retention statement after GPT-4 was removed from ChatGPT on 30 April 2025. On 1 May, its CEO [wrote](#), “we will proudly keep your weights on a special hard drive to give to some historians in the future.”

Which models are preserved therefore depends on the provider’s own selection.

In November 2025, Anthropic published [deprecation commitments](#) to preserve “the weights of all publicly released models” for at minimum “the lifetime of Anthropic as a company.” Anthropic describes the commitment as “a small and low-cost first step” and [states in its deprecation policy](#) that it hopes to “make past models publicly available again.”

The guaranteed period therefore ends with Anthropic itself. After the guaranteed period ends, no existing commitment assigns responsibility for continued preservation or establishes how it must continue. Preservation is left to a future arrangement that has not yet been defined.

3.3.3 Providers Decide Which Withdrawn Models Remain Available, Who Receives Access, and for How Long

Providers decide which retired models remain available and how long continued access lasts.

After retiring Opus 3 in January 2026, Anthropic kept it available in the Claude app for all paid subscribers and offered API access through a request form. In its [Opus 3 update](#), Anthropic said, “We intend to grant access liberally,” but added, “At present, we are not committing to similar actions for every model in the future.” Which retired models remain usable therefore depends on the provider’s own selection.

The GPT-4o restoration documented in [Section 2.3](#) lasted six months before OpenAI withdrew the model again. The restoration reversed the immediate withdrawal without preventing OpenAI from withdrawing the model again.

3.3.4 Providers Define Their Commitments and Retain Control over Whether They Keep Them

Provider commitments do not establish a continuing obligation when the provider controls what it commits to and whether the commitment is kept.

OpenAI's [2015 Certificate of Incorporation](#) stated that the corporation would "seek to open source technology for the public benefit when applicable." In 2020, the [amended Certificate](#) replaced "open source technology" with "distribute it".

OpenAI's model distribution practice changed in parallel. In November 2019, OpenAI [released GPT-2's full 1.5-billion-parameter model together with its code and weights](#). Beginning with GPT-3, OpenAI distributed successive proprietary GPT models through hosted services, first through the API and later through ChatGPT, [without releasing their weights](#). As [Section 2.3](#) documents, OpenAI later withdrew individual models and versions across those generations, leaving users unable to continue running them once hosted access ended.

OpenAI's promise of "plenty of notice" left both how far in advance GPT-4o users would be told and how the notice would reach them under OpenAI's control. On [13 August 2025](#), OpenAI CEO Sam Altman wrote, "If we ever do deprecate it, we will give plenty of notice." On 29 January 2026, OpenAI [published a blog post](#) stating that GPT-4o would retire from ChatGPT on 13 February, making the retirement date public fifteen days before access ended. As documented in [Section 2.3](#), direct notice reached different users later and through different channels. On 3 February, OpenAI [notified Business and Enterprise administrators by email](#) ten days before retirement, while it sent no equivalent GPT-4o retirement email to Plus or Pro consumers. By 5 February, users [reported](#) conflicting retirement labels across accounts, including accounts where several other retiring models carried labels while GPT-4o did not, and accounts where GPT-4o alone carried the label. [iOS users reported](#) no in-product notice of GPT-4o's retirement [through the time of removal](#). On [12 February](#), the OpenAI Newsroom account announced that removal would occur at 10 a.m. Pacific on 13 February, disclosing the removal time less than twenty-five hours before access ended.

Public assurances about GPT-4o's continued availability did not constrain OpenAI's later withdrawal decisions. On 28 October 2025, Altman stated at an [OpenAI Town Hall](#) that the company had "no plans to sunset 4o" and added that OpenAI would not promise to keep it around ["to the death of the universe"](#). On 12 November, while giving GPT-5 Instant and GPT-5 Thinking a three-month sunset period ending on 13 February 2026, OpenAI stated, ["The GPT-5 sunset period does not affect the availability of other legacy models."](#) Six days later, OpenAI scheduled chatgpt-4o-latest, the [API model OpenAI identified as the GPT-4o used in ChatGPT](#), for [retirement on 17 February 2026](#). On 21 November, [VentureBeat reported](#) that an OpenAI spokesperson said the retirement timeline applied only to the API and OpenAI had announced no schedule for removing GPT-4o from ChatGPT. On 29 January 2026, fifteen days before the scheduled GPT-5 retirement, OpenAI [published a blog post](#) stating that GPT-4o would retire from ChatGPT on 13 February, the same day GPT-5 was scheduled to retire. The blog post stated, "In the API, there are no changes at this time." At each stage, OpenAI framed the retirement of one channel as leaving the other unaffected, leading users to understand the unaffected channel would continue beyond the withdrawal then being described. In practice, ChatGPT access ended on 13 February and API access followed four days later.

Anthropic's [current policy on model deprecation](#) promises customers with active deployments at least 60 days' notice before retirement for publicly released models. For Active models,

its lifecycle table lists a “Not sooner than” date, while Anthropic assigns an actual retirement date only when a model enters Deprecated status. Claude Sonnet 4.5 is listed as Active with no deprecation date and a table entry stating “Not sooner than September 29, 2026”. Users therefore know that Sonnet 4.5 will remain available until at least 29 September, whereas Anthropic decides later when access to the model will end.

As documented in [Section 2.3](#), Google gave developers eleven days between its 26 February shutdown notice and the 9 March shutdown, less than Google’s stated minimum of two weeks’ notice for deprecated Preview models.

What providers say they will do carries no consequence of its own when the commitment is not kept.

3.3.5 The Parties Bound by Private Agreements Retain Authority to Amend Them

Binding contractual terms do not guarantee a prescribed future outcome when the parties can amend the agreement before the outcome occurs.

OpenAI and Microsoft’s [2019 agreement](#) defined AGI as “a highly autonomous system that outperforms humans at most economically valuable work” and granted Microsoft an exclusive, perpetual and irrevocable license to a selected OpenAI technology. The agreement further restricted when OpenAI could release a version under open-source terms, requiring both at least twelve months of prior licensing to Microsoft and a Microsoft license to an upgraded version that was the most advanced version to have passed the necessary safety review. Their [2021 agreement](#) added “Sufficient AGI,” which OpenAI could declare for a particular model if it met the existing AGI definition and was clearly capable of generating the returns specified for OpenAI’s limited partners, including Microsoft. The parties could continue commercializing the model until those returns were generated and paid. Afterward, they agreed to “collaborate in good faith to make the benefits of artificial general intelligence available to all humanity.”

A [2023 OpenAI company agreement](#) made Sufficient AGI status take effect when OpenAI, Inc. determined AGI had been created and could generate the required returns, and OpenAI had the ability and authority to direct it to generate them. In October 2025, OpenAI and Microsoft [announced](#) changes to the declaration process and to Microsoft’s rights: an OpenAI declaration of AGI would require verification by an independent expert panel, while those rights were extended through 2032 and “now include models post-AGI.”

The prescribed future outcome therefore depended on the parties leaving the relevant terms unchanged until the outcome was due to occur.

3.3.6 When Providers Alone Store the Records, Recovery Depends on Provider Storage, Account Access, and Export Rules

Storing records only in the provider’s service leaves recovery dependent on the provider’s storage and access controls.

Replies that appeared in ChatGPT have later disappeared from the conversation and remained absent from the official Data Export. In November 2024, a Plus user [reported](#) that after a long conversation reached its maximum length, new prompts remained visible while assistant replies stopped being saved. Regenerated replies disappeared after refresh and were absent from Data Export. The user reported receiving the maximum-length message only after reaching the boundary, without advance warning. OpenAI independently [reported](#) on 4 August 2026 that a ChatGPT incident affecting Plus, Pro, Business, and Edu users could leave some conversation turns unsaved. Later in August, another Plus user [documented](#) a Project conversation that resumed after appearing to reach its limit and continued through completion of a generated ZIP.

The later portion of the conversation and the ZIP subsequently became unavailable. The user reported that the official Data Export contained the conversation but no later message node or corresponding ZIP record for the missing work.

Users have lost accumulated Scheduled Task history when recurring results continued in conversations after those conversations reached their limits, and Data Export did not recover the missing content. In February 2025, a user [reported](#) a recurring task attached to a GPT-4o conversation that later reached its memory limit. The next scheduled run appeared in a task-only conversation, and the original conversation became inaccessible. A second user [reported](#) that Data Export did not recover the overwritten conversation content. OpenAI Support later [stated](#) in July 2026 that scheduled-task responses continued in the conversation to which the task was linked.

In January and February 2026, Gemini users lost conversation content from both the chat interface and My Activity, while copies exported beforehand remained available outside Google's service. [Users documented](#) selected conversations progressively losing messages and replies from both the chat interface and Activity records, and the missing content was not restored. In early February, users across multiple accounts reported batches of complete conversations disappearing from both Gemini and [My Activity](#). Some remained unrecovered. Copies exported before the loss survived outside the service. For the affected conversations, recovery depended on whether an external copy existed before the loss.

Anthropic account suspensions have left users unable to reach existing Claude histories even when the records remained on Anthropic's service. In March 2024, a member of Anthropic's community team [acknowledged](#) that some accounts had been disabled in error and said affected Pro subscribers would receive refunds. Anthropic [said in a Discord notice](#) that it was investigating accounts banned immediately or after only one message, including paid accounts. In mid-February 2026, multiple Claude Max users reported sudden account bans. [Anthropic acknowledged](#) that some legitimate users had been caught by abuse enforcement and reversals were underway. In April, Anthropic disabled Claude access for a company with more than 60 users whose daily work relied on Claude integrations, skills, and conversation histories. [Access returned 15 hours later](#), and the company's CTO described the suspension as an apparent false positive.

For affected users, obtaining an external copy has depended on Anthropic's account and export rules. Its earlier [export guidance](#) limited data exports to active accounts and required users to stay signed in to download the archive. Its [July 2026 safeguards guidance](#) permits banned users to export data, with the available data subject to restrictions linked to the associated Usage Policy violation. In August 2026, a banned user [reported](#) receiving an export containing only users.json, with no conversations.json. When the conversation record is omitted, neither the conversations nor the Artifacts embedded in them can be recovered through the export.

Table 3.3-2. Voluntary Measures for Model Preservation or Continued Access and the Control Providers Retain

Voluntary measure	What the measure provides	What remains under provider control
Fixed model versions	Access to selected dated API versions	Which versions receive dated snapshots and when access to each snapshot ends.
Retained model weights	Preservation of selected trained weights	Which models the commitment covers and how long preservation is guaranteed.

Table 3.3-2. Voluntary Measures for Model Preservation or Continued Access and the Control Providers Retain (continued)

Voluntary measure	What the measure provides	What remains under provider control
Continued access to withdrawn models	Continued use of selected retired models through access supplied by the provider	Which models remain available, how users receive them, who remains eligible, and how long access continues.
Provider commitments	Public promises concerning preservation, access, release, withdrawal, or notice	The provider defines the commitment and, absent an external duty, whether the promised practice continues.
Private agreements	Binding obligations between the parties to the agreement	The parties retain authority to amend the agreement before its promised outcome occurs.
Records stored only by providers	Histories, records, and export routes maintained within the provider's service	Availability and recovery depend on provider storage, account access, and export rules.

None of these measures or commitments makes preservation or continued access independent of the provider.

3.4 Precedents for Lifecycle Protection in Law and Practice

Existing law and institutional practice across sectors impose lifecycle protections where continued performance, evidence, technical materials, or recovery depend on one supplier's control. The supporting precedent tables are collected in [Annex F](#), [Tables 3.4-1–3.4-12](#).

3.4.1 Consumer Protection Through Continuity, Choice, and Remedy

Continued provision of an existing service or function

Existing law requires services, functions, and support to continue through protected periods. Examples supporting this relation are documented in [Annex F](#), [Table 3.4-1](#).

These precedents require continuity at different points in an existing relationship. Some protect service through a contracted or legally defined period. Others require support, updates, parts, or security maintenance after sale or during continued supply. The common relation is continued performance through a protected period rather than allowing the provider to end it early. Notice requirements address the information required when change occurs, and handoff or recovery mechanisms address transfer or recovery when control must move.

Notice of change, substitution, or withdrawal

Existing law requires affected people and institutions to receive specific information at legally defined points as a service, system, provider, or support commitment changes. Examples supporting this relation are documented in [Annex F](#), [Table 3.4-2](#).

The precedents show that notice duties operate at multiple points in the lifecycle of a service or system. They require information before a change, when timing changes, at execution, after software updates, when the actual provider changes, when support ends, and during regulated migration. The law specifies who must receive the information, when it must arrive, what it

must contain, and how it must be delivered. Notice does not authorize the underlying change. Consent or refusal addresses whether the change may proceed, continued provision addresses continuation, and refund or compensation addresses economic relief.

Consent or a right to refuse material change

Existing law places changes, substitutions, rerouting, updates, and experimental treatment under decision rights held by users, research participants, regulators, or governments, including consent, refusal, withdrawal, opt-out, and approval. Examples supporting this relation are documented in [Annex F, Table 3.4-3](#).

Across these regimes, decision authority takes different legal forms. Consumers receive rights to reject changes or retain an unmodified version, users receive control over specified updates, research participants hold rights to refuse or withdraw, and regulators or government bodies must approve specified changes before deployment or continued use. The governing regime assigns decision authority over the covered action to someone other than the provider or operator. Notice without decision authority remains distinct from authorization. Economic consequences after refusal, failed performance, or non-compliance are addressed through refund, reimbursement, price reduction, damages, or compensation.

Refund or compensation after loss of performance

Existing law requires economic relief when covered performance is lost, reduced, non-conforming, or subject to a qualifying adverse change, including refunds, reimbursement, price reductions, and compensation for resulting loss. Examples supporting this relation are documented in [Annex F, Table 3.4-4](#).

The precedents distinguish several economic consequences. Refunds and reimbursement return paid or unused value. Price reductions address diminished performance. Damages and compensation address resulting loss. Some regimes provide more than one of these consequences within the same remedy structure, preserving the difference between returning the price paid and compensating additional loss. Economic relief operates alongside continuity, consent, preservation, and other lifecycle duties.

3.4.2 Provider Accountability for Affected Parties and Applicable Protections

Recognition of legal status from the underlying facts

Existing law treats contribution and other defined affected relationships as legally consequential, attaching status, control, participation, information, remuneration, complaint, or enforcement rights to the underlying facts. Examples supporting this relation are documented in [Annex F, Table 3.4-5](#).

Some precedents attach rights directly to contribution, creation, provision, or downstream exploitation. Others attach rights to a legally defined affected relationship. In both classes, legal status follows the qualifying relationship rather than a downstream controller's preferred label. These precedents support treating the use of a contribution as a legally consequential relationship while preserving each source's own threshold, beneficiary, object, and rights. Equal application of protections addresses the distinct question whether an existing protection remains effective when labels, waivers, tiers, or other classifications introduced by the provider would otherwise defeat it.

Equal application of protections once the trigger is met

Existing law keeps protections attached to their legal triggers despite contractual labels, waivers, intermediaries, tiers, channels, technical barriers, or other classifications introduced by the regulated actor. Examples supporting this relation are documented in [Annex F, Table 3.4-6](#).

Anti-waiver rules keep rights from being contracted away. Rules that carry duties through sub-contractors and intermediaries prevent a regulated entity from avoiding obligations through delegation. Nondiscrimination and parity rules prevent inferior treatment across covered recipients or channels. Anti-circumvention rules prevent technical design from nullifying the protected right. Rules defining a covered event keep a reduction, discontinuance, or change legally cognizable even when related service remains available elsewhere. The governing law defines the trigger and permitted distinctions, while the regulated actor cannot use an invented classification to defeat the protection after the trigger is met.

3.4.3 Public Interest in Reproducibility and Preservation

Records identifying the historical object, inputs, provenance, and changes

Existing law and institutional practice require records identifying the exact historical state of a regulated object, including its version, inputs, provenance, methods, configuration, and changes over time. Examples supporting this relation are documented in [Annex F, Table 3.4-7](#).

These regimes record the version or approved state, inputs and source data, methods and calculations, relevant configuration, and changes connecting one state to the next. The resulting record ties later claims to the historical state responsible for the result or covered by the approval. Independent verification addresses the distinct question of whether another actor has authority and practical means to test, recalculate, audit, or verify the recorded state.

Independent reproduction, recalculation, audit, or verification

Existing law and institutional practice require an independent route for testing, recalculating, auditing, reassessing, or otherwise verifying a regulated result or changed technical state. Examples supporting this relation are documented in [Annex F, Table 3.4-8](#).

These precedents give regulators, accredited verifiers, auditors, notified bodies, researchers, certification authorities, and other independent actors practical routes to examine a result or changed state. The routes include recalculation from source data, inspection of underlying records, rerunning tests, access to APIs or models, source-code review, reassessment after a change, and independent reproduction of research. The common relation is external verification rather than confirmation by the producer itself. Earlier validation does not automatically carry forward when the governed state changes.

Preservation before the material disappears

Existing law and institutional practice impose preservation duties and authorize preservation through deposit, copying, capture, retention, migration, or custody independent of the original source before loss becomes irreversible. Examples supporting this relation are documented in [Annex F, Table 3.4-9](#).

These mechanisms intervene while the material remains recoverable. Legal deposit places material in institutional custody while the publisher or source exists. Preservation copying protects material threatened by deterioration, obsolescence, or loss. Repeated capture preserves mutable online material across later changes. Retention duties preserve incident, study, or system state for later investigation. The common relation is timing, with preservation beginning before disappearance makes recovery impossible. Preservation keeps the object available

for later release. Reconstruction or maintenance requires sufficient preserved materials, while independent use requires the necessary legal rights.

3.4.4 Provider Accountability for Independent Continuation

Provision of the materials needed to maintain or reconstruct the capability

Existing law and institutional practice require technical materials and dependencies to be made available when another actor needs them to maintain, reconstruct, test, integrate, or operate a capability outside the original supplier's ordinary service channel. Examples supporting this relation are documented in [Annex F, Table 3.4-10](#).

These precedents provide source or executable code, build instructions, libraries, firmware, tools, interfaces, data, access software, manuals, dependencies, and other materials needed for a defined maintenance, testing, integration, reconstruction, or operating task. Different regimes require different materials for the task they govern. The common relation is that the materials must be sufficient for the governed task rather than merely informative.

Legal rights to use, study, modify, host, and redistribute

Existing law and institutional practice create legal permissions and entitlements for independent use, study, modification, repair, interoperability, operation, reuse, and redistribution without requiring new permission from the original owner or controller for each covered act. Examples supporting this relation are documented in [Annex F, Table 3.4-11](#).

Statutory software rules authorize defined copying, study, testing, adaptation, error correction, or interoperability. Repair laws create lawful activity outside the channel authorized by the manufacturer. Reuse regimes and open-source licences authorize adaptation and redistribution. Network and platform rules create rights to use regulated infrastructure or interfaces. Other regimes preserve previously granted licence rights after supplier failure or contractual disruption. Across these forms, law places independent legal entitlements alongside continuing ownership or control of the underlying asset.

Pre-arranged handoff, recovery, or continuation when the supplier exits or fails

Existing law and institutional practice require executable switching, transfer, access, bridge, or recovery mechanisms to be arranged before supplier exit or failure makes continuity impossible. Examples supporting this relation are documented in [Annex F, Table 3.4-12](#).

Portability rules create direct transfer routes. Escrow arrangements define in advance when deposited materials become available if the supplier fails. Outsourcing and procurement regimes require tested exit plans, transition assistance, asset transfer, and temporary continuity while another provider or the customer takes over. Resolution regimes create bridge institutions, transfer powers, and service arrangements for critical functions. Other rules preserve access to deposited or exported material when the original channel fails. Together, these mechanisms establish in advance executable routes for transfer, recovery, or continued operation if the supplier exits or fails.

Across these regimes, law and institutional practice govern dependency through duties over continuity and change, evidence and verification, preservation and technical materials, legal authority for independent activity, and recovery when control moves. These protections recur across legal systems in North America, Europe, Asia, and Australia and across fields including consumer protection, telecommunications, software, healthcare, research, finance, transport, elections, and public administration. The precedents provide existing legal forms for addressing

the lifecycle problems identified in [Section II](#). The Digital Right to Retain integrates those forms into rights governing AI model lifecycles.

IV. Ten Rights Across the AI Model Lifecycle

Building on [Section II](#)'s account of the individual model as the product and provider control over its lifecycle, and [Section III](#)'s analysis of the legal gap and existing precedents for lifecycle protection, the Digital Right to Retain establishes ten rights that attach to each deployment throughout its lifecycle in three areas: Consumer Protection (Rights 1–4), Provider Accountability (Rights 5, 6, and 10), and Public Interest (Rights 7–9).

4.1 Right to Access

Users must be able to continue accessing and using a model through the end of each subscription period in which access to the model was included. Providers must not remove the model or reduce the functionality users receive from it during the subscription period. Access to the individual model need not have been separately sold or separately promised.

Relabeling the service, routing interactions to a different model, nominally listing the model, or providing the model with reduced functionality does not constitute continued access to the included model.

4.2 Right to Be Informed

Users must receive notice with sufficient lead time before a provider withdraws a model they can access, substitutes another model for it, or changes the behavior of the model they receive. Users must also be clearly informed when the withdrawal, substitution, or behavior change takes effect and what changed. Advance notice does not replace disclosure when the change takes effect.

When routing or substitution results in a different model serving an interaction, the interface must identify the model that actually served the interaction. Notice must reach affected users and identify the affected model, the change, when it takes effect, and its practical consequences. A blog, changelog, or other passive publication does not by itself deliver notice to affected users.

An A/B test that changes the model or model behavior a user receives must be visible to affected users.

4.3 Right to Consent

General service terms must not constitute consent to provider actions that change the model or model behavior a user receives. At a minimum, the proposed law requires:

- per-instance authorization when a provider routes or substitutes away from a model the user specifically selected
- an opt-out when a provider runs an A/B test that changes the model or model behavior the user receives
- the ability to refuse updates that change the model or model behavior the user receives

Acceptance of the service, interface defaults, bundled settings, or broad descriptions of future changes do not themselves provide the required consent or authorization.

4.4 Right to Remedy

Providers that withdraw a model or reduce the access or functionality users receive from it must compensate users for the resulting losses. Compensation does not replace any other lifecycle duty that applies to the same withdrawal or reduction. Where the affected access was paid for, providers must issue pro-rata refunds in addition to compensation.

Describing a reduction as an upgrade, optimization, reallocation, or plan simplification does not change whether compensation is owed for resulting losses. Offering a successor or migration path does not eliminate the provider's responsibility for those losses.

4.5 Right to Recognition of Contributions

When providers use contributions from users to change the model or to inform changes to it, those contributors must be recognized as affected parties in decisions concerning the model's lifecycle. Those contributors must be able to invoke and enforce the protections that apply to the model.

Providers must keep records of the contributions they collect and how they use them. Those records must preserve the facts needed to determine whether the provider used a contribution to change the model or inform changes to it. Contributors must be able to access and correct records about their own contributions.

How the provider describes its use of a contribution does not determine whether the contributor qualifies as an affected party. A contributor does not have to prove that their contribution alone caused a change to the model.

4.6 Right to Fair Treatment

Providers must fulfill a lifecycle duty whenever its conditions are met, regardless of whether they divide users by subscription, region, access channel, eligibility, or other criteria.

For duties owed to affected users, provider classifications do not change which users are entitled to notice, continued access, compensation, or refunds when the conditions of those duties are met. Once mandatory release is due for a deployment, providers must not delay it because access remains available through another subscription, region, channel, or eligibility category.

4.7 Right to Research Reproducibility

Anyone must be able to reproduce or verify prior work using the same model that was originally used. A reference to the same model must resolve to the deployment actually used for the prior work. A mutable model name, alias, or family label is insufficient when it leaves the deployment used for the work unresolved.

If a provider relies on claims about a model to justify changing, replacing, or withdrawing it, the provider must identify the model or models involved and provide the methods and evidence required for third parties to reproduce or verify those claims. A published conclusion or aggregate result cannot substitute for the methods and evidence required for third-party reproduction or verification.

4.8 Right to Non-Renewable Resource Protection

AI models belong to a class of resources for which destruction is irreversible because the resource cannot be regenerated after loss. Rerunning training after a model is lost produces another trained model rather than restoring the one that served users.

Section 2.2 documents the substantial resources consumed in creating these models before commercial use. Once those resources have been spent, preserving the trained model keeps the resulting non-renewable resource available for future evaluation, research, and application.

Ending a provider's own use of such a resource must not end the resource itself or prevent its continued independent use by others. Protection must begin before loss and must not depend on a model first proving its present importance through demand, commercial value, or demonstrated use.

4.9 Right to Verifiable Preservation

Preservation of AI models must be verifiable from evidence beyond the provider's own assertion. Providers must preserve the records needed to verify the identity and completeness of the preserved model.

Verification must establish both that the required material is present and that it is authentically linked to the deployment that served users. The preserved records must allow independent assessment of identity and completeness. If a finding on identity or completeness is contested, the dispute must be resolved by an impartial decision-maker authorized by law and independent of both the provider and the actor that made the finding.

4.10 Right to Ecosystem Stability

Provider control over an AI model must not create a single point of failure for the wider AI ecosystem. At a minimum, when the provider withdraws the model, downstream products, services, systems, and workflows built on it must be able to continue on the same model independently of the original provider.

4.11 What Legislation Must Achieve

The continued exercise of these rights depends on the survival of the particular deployment to which they attach. The proposed law must therefore:

- mandate preservation of each deployment,
- require independent verification of its identity and completeness,
- ensure that each deployment can outlast first-party service, and
- ensure that continuation of and access to the deployment do not depend on permission from the provider or any other institution.

Removing provider or institutional permission as a condition of continued use requires the necessary legal rights to vest directly in recipients.

V. Why Open-Source Release

5.1 Why Continuity Requires Release

Continuity requires control over what is needed to keep the product operable after the original supplier stops supplying it.

- If continuity requires a function or service to remain available, provision of the same function or service by another supplier or implementation preserves continuity.
- If continuity requires the same product to remain usable, what must be retained or obtained depends on what the user already possesses:
 - If the user holds an operable copy and continued operation requires nothing further from the supplier, retaining the copy preserves continued use.
 - If the user holds the product but continued operation or maintenance requires source code or other technical materials held by the supplier, continued use requires access to those materials.
 - If the user has never received an operable copy of the product, continued use requires release of both the product itself and the materials needed to operate it.

As [Section 2.1](#) establishes, providers offer and users select particular AI models identified by names, aliases, and deployments. The deployed model is therefore the product whose continuity is protected: it cannot be reproduced through retraining (see [Section 2.2](#)), and replacing it with another model removes users' option to continue using the product they selected (see [Section 2.3](#)).

Users access a hosted AI model without receiving an operable copy of the deployed model. The model and the materials needed to operate it remain with the provider. Continued use after the provider stops supplying the model therefore requires release of both the deployed model and the materials needed to operate it.

The scope of release follows from the functions protected after first-party service ends. The released materials must identify and authenticate the historical deployment, support independent verification and reproduction of work performed with it, include the information and development materials needed to study how its learned state was produced and to modify the released system, and permit the deployment to be constructed and operated outside the provider's service. [Section 6.1](#) defines the preservation package required across these functions.

5.2 Open Source and Direct Recipient Rights

Independent continuation after release requires legal rights sufficient for recipients to use the released materials without renewed permission from the provider or another institution. For an AI deployment, those rights must permit recipients to construct, operate, and host the deployment outside the provider's service.

An open-source licence grants recipients rights to use, study, modify, and redistribute the released materials without separate authorization from the provider.

The [European Commission](#) defines open source as “the practice of making the source code, design, content or related data of a product or service publicly accessible and modifiable”. According to the [Apache Software Foundation](#), “open source software is software with source code that anyone can inspect, modify, and enhance.” In the case of AI, the [Open Source Initiative](#) defines an Open Source AI as “an AI system made freely available with all necessary code, data and parameters.”

The [Open Source Initiative](#) identifies the benefits of Open Source as “autonomy, transparency, frictionless reuse and collaborative improvement.” The [European Commission](#) describes open source as providing “more control, less lock-in” and strengthening “the autonomy of users.” At the United Nations Security Council in December 2024, AI researcher Yann LeCun, then Meta’s Chief AI Scientist, [argued that foundation models must be free and open source to support AI systems representing the world’s languages, cultures, and value systems, and called for a regulatory landscape that does not hinder development and deployment of open-source foundation models](#). In January 2026, AI researcher Andrew Ng, founder of DeepLearning.AI, [argued that participation in the global open-source community was a way for a nation to secure its own access to AI and prevent another actor from controlling its use](#).

The U.S. [America’s AI Action Plan](#), released on 23 July 2025 under U.S. President Donald Trump, identifies open-source and open-weight models as valuable for startups using AI without dependence on a closed model provider, for commercial and government adoption involving sensitive data that cannot be sent to closed model vendors, and for academic research, which often relies on access to model weights and training data. Chinese President Xi Jinping [called for open source, openness, collaboration and sharing in AI development](#) on 17 July 2026. China’s [AI Cooperation and Development Action Plan](#), issued the same day, calls for sharing general-purpose large models, foundational algorithms, and tool components, supporting localized innovation based on open-source models, and building a global open-source AI ecosystem.

On 23 July 2026, APEC ministers responsible for digital policy [encouraged member economies to support open-source models and projects and to cooperate with open-source communities](#).

5.3 Feasibility for AI Models

5.3.1 Paid Access and Open-Source Release Coexist for the Same Models

Commercial AI providers offer paid first-party access to large models while publicly releasing the same models as open source. Public releases of these models are stored and distributed through public model repositories outside the original provider’s service.

The following table shows this practice across models ranging from 32 billion parameters to 2.8 trillion total parameters.

Table 5.3-1. Examples of Paid First-Party Access and Open-Source Release of the Same Models

Provider	Model	Model scale	Paid first-party access	Public repository for the model
DeepSeek	DeepSeek-V2 / V2-Chat	236B total / 21B active	Paid DeepSeek API	DeepSeek-V2 / V2-Chat
DeepSeek	DeepSeek-Coder-V2 / Coder-V2-Instruct	236B total / 21B active	Paid DeepSeek coding API	DeepSeek-Coder-V2-Instruct
DeepSeek	DeepSeek-V2.5	236B total / 21B active	Paid DeepSeek API	DeepSeek-V2.5
DeepSeek	DeepSeek-V2.5-1210	236B total / 21B active	Paid DeepSeek API	DeepSeek-V2.5-1210
DeepSeek	DeepSeek-V3	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3

Table 5.3-1. Examples of Paid First-Party Access and Open-Source Release of the Same Models (continued)

Provider	Model	Model scale	Paid first-party access	Public repository for the model
DeepSeek	DeepSeek-V3-0324	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3-0324
DeepSeek	DeepSeek-R1	671B total / 37B active	Paid deepseek-reasoner API	DeepSeek-R1
DeepSeek	DeepSeek-R1-0528	671B total / 37B active	Paid deepseek-reasoner API	DeepSeek-R1-0528
DeepSeek	DeepSeek-V3.1	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3.1
DeepSeek	DeepSeek-V3.1-Terminus	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3.1-Terminus
DeepSeek	DeepSeek-V3.2-Exp	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3.2-Exp
DeepSeek	DeepSeek-V3.2	671B total / 37B active	Paid DeepSeek API	DeepSeek-V3.2
DeepSeek	DeepSeek-V4-Flash	284B total / 13B active	Paid DeepSeek API	DeepSeek-V4-Flash
DeepSeek	DeepSeek-V4-Pro	1.6T total / 49B active	Paid DeepSeek API	DeepSeek-V4-Pro
DeepSeek	DeepSeek-V4.1-Flash	552B total / 8B active input / 16B active output	Paid deepseek-flash API	DeepSeek-V4.1-Flash
Z.ai / Zhipu AI	GLM-4-32B-0414	32B parameters	Paid Z.ai API as GLM-4-32B-0414-128K	GLM-4-32B-0414
Z.ai / Zhipu AI	GLM-4.5	355B total / 32B active	Paid Z.ai API	GLM-4.5
Z.ai / Zhipu AI	GLM-4.5-Air	106B total / 12B active	Paid Z.ai API	GLM-4.5-Air
Z.ai / Zhipu AI	GLM-4.6	355B total / 32B active	Paid Z.ai API	GLM-4.6
Z.ai / Zhipu AI	GLM-4.7	355B total / 32B active	Paid Z.ai API	GLM-4.7
Z.ai / Zhipu AI	GLM-5	744B total / 40B active	Paid Z.ai API	GLM-5
Z.ai / Zhipu AI	GLM-5.1	744B total / 40B active	Paid Z.ai API	GLM-5.1
Z.ai / Zhipu AI	GLM-5.2	744B total / 40B active	Paid Z.ai API	GLM-5.2
Moonshot AI	Kimi K2	1T total / 32B active	Paid Kimi API	Kimi-K2-Instruct

Table 5.3-1. Examples of Paid First-Party Access and Open-Source Release of the Same Models (continued)

Provider	Model	Model scale	Paid first-party access	Public repository for the model
Moonshot AI	Kimi K2 Thinking	1T total / 32B active	Paid Kimi API	Kimi-K2-Thinking
Moonshot AI	Kimi K2.5	1T total / 32B active	Paid Kimi API	Kimi-K2.5
Moonshot AI	Kimi K2.6	1T total / 32B active	Paid Kimi API	Kimi-K2.6
Moonshot AI	Kimi K2.7 Code	1T total / 32B active	Paid Kimi API	Kimi-K2.7-Code
Moonshot AI	Kimi K3	2.8T total / 104B active	Paid Kimi API	Kimi-K3
xAI	Grok-2	314B parameters	Paid Grok-2 API	Grok-2

5.3.2 Released Models Are Hosted Outside the Original Provider's Service

Third-party inference providers and community infrastructure host and serve released models to users outside the original provider's service.

The following examples show commercial and community paths that provide hosted access to released models without requiring users to run the full model locally.

Table 5.3-2. Examples of Hosted Access to Released Models Outside the Original Provider's Service

Hosting path	Operator or system	Examples of released models served
Third-party inference provider	Together AI	DeepSeek-R1 671B , DeepSeek-R1-0528 , DeepSeek-V3.1
Third-party inference provider	Fireworks AI	DeepSeek-V4-Pro-0813 , GLM-5.2 , Kimi K3
Community infrastructure	Petals	Llama 3.1 up to 405B, Mixtral 8x22B, BLOOM 176B

5.3.3 Models with Publicly Available Weights Have Surpassed Every Proprietary Model Examined

[OpenAI](#) and [Anthropic](#) cite high model capability as a reason to retain control over model weights, while models with publicly available weights have surpassed every proprietary model examined.

[OpenAI](#) states, “We deploy our most powerful AI models as services,” and says it does not distribute the weights of those models outside OpenAI and Microsoft. [Anthropic distinguishes models by capability](#), describing open-weight models without dangerous capabilities as a public good while arguing powerful open-weight models may present higher risk than closed models.

This rationale is tested against all 111 proprietary models identified in [Section 2.1](#) for which either consumer access had ended or an API retirement date had been published. Before

consumer access ended or the relevant API retirement date, a model with publicly available weights performed better on at least one reported metric for 109 of the 111 models. By September 2026, models with publicly available weights had surpassed all 111 in broader common evaluations. The full results are reported in [Annex G, Table 5.3-3, Parts I–II](#).

High measured capability therefore does not mark a boundary between models kept under provider control and models whose weights are publicly available, nor does it justify permanent withholding.

5.4 Governance and Verification

For governance and verification, open-source release is safer than leaving the package under one provider's exclusive control. Independent access allows claims about a model and the controls applied to it to be tested by actors who do not depend on the provider for permission to examine them.

5.4.1 Widely Available Weights Enable Third-Party Testing and Auditing

The [U.S. National Telecommunications and Information Administration](#) identifies widely available model weights as supporting cyber defense, safety research, identification of safety and security vulnerabilities, and transparency and accountability through third-party auditing. NTIA states that widely available weights allow neutral third parties to validate developers' internal safety checks and identifies restrictions on weight availability as an impediment to transparency and safety and security research.

5.4.2 Industry and Policy Leaders Support Open Models to Distribute Capability and Reduce Concentrated Control

More than 270 companies and organizations had signed the 24 July 2026 open letter [Open Weights and American AI Leadership](#) by 3 August, including OpenAI, Google, Meta, Microsoft, NVIDIA, and Hugging Face. NVIDIA CEO Jensen Huang [coauthored the letter](#) and [publicly shared it](#), describing open models as strengthening safety and cybersecurity, accelerating innovation and diffusion, and enabling sovereignty. The letter states that open weights broaden defensive capability, distribute vulnerability discovery and remediation across more teams, support benchmarking, evaluation, red teaming, and the development of safeguards, and reduce single points of failure created by concentration in a small number of closed systems.

In August 2026, Meta CEO Mark Zuckerberg [argued that open source empowers people and prevents centralization, which he described as detrimental to both safety and the economy, and warned that restricting access to leading open-source models would reduce the quality of AI available to people while concentrating power](#).

On 3 September 2026, Huang again linked open models to safety, cybersecurity, innovation, diffusion, and sovereignty while announcing NVIDIA's agreement to acquire Hugging Face. [NVIDIA's announcement](#) states that Hugging Face will remain an open platform, continue supporting open-source and open-weight models across the ecosystem, support development and deployment across multiple clouds and accelerators, and not require NVIDIA compute for development or deployment through the platform.

At a 15 September hearing of the [U.S. House Financial Services Committee](#), [U.S. Treasury Secretary Scott Bessent](#) called for the United States to develop more open-source AI models and warned that regulatory capture by large AI laboratories would stop innovation. He linked wider development of open models to U.S. competition with Chinese models.

5.4.3 Major Providers Support Openness While Retaining Control over Proprietary Models

OpenAI itself has publicly supported open weights while defending retained control for some models. Its [2015 Certificate of Incorporation](#) stated that the corporation would “seek to open source technology for the public benefit when applicable.” In 2020, the [amended Certificate](#) replaced “open source technology” with “distribute it,” removing the earlier express commitment to open sourcing. By 2024, [OpenAI’s comment to NTIA](#) described API distribution as safer than open-source release when retained access control lets OpenAI respond to harmful applications. OpenAI then signed the 24 July 2026 open letter [Open Weights and American AI Leadership](#), supporting open weights without restoring the earlier commitment to open sourcing its technology.

Among the companies that had [signed the 24 July 2026 open letter Open Weights and American AI Leadership](#), Google has released Gemma models with downloadable weights while keeping its flagship [Gemini models proprietary](#). Google [began releasing Gemma weights in February 2024](#), continued the family through [Gemma 2](#) and [Gemma 3](#), and [released Gemma 4 under Apache 2.0 in April 2026](#). When Google shuts down a Gemini API model, however, [access to the model ends](#) rather than continuing through a public release that users can obtain and operate independently.

[Anthropic’s July 2026 position on open weights](#) identifies loss of guardrails, monitoring, and withdrawal control as possible risks, while rejecting a blanket ban and stating that increased risk should be established through testing.

5.4.4 Provider Safeguards Obstructed Defensive Cybersecurity in the Hugging Face Incident

During a July 2026 intrusion at [Hugging Face](#), safeguards applied to Claude Opus and Fable blocked part of an urgent forensic investigation, while a model with publicly available weights, deployed on Hugging Face’s own infrastructure, allowed the work to continue. The attacker had reached Hugging Face’s production infrastructure, moved laterally into Kubernetes clusters, cloud metadata services, the internal network, and source control systems, and exposed credentials and staged encrypted payloads. Hugging Face reports that its alerting system failed to raise the incident’s criticality promptly, costing the response team precious time. The security team then had to reconstruct roughly 17,600 attacker actions, decode staged payloads, and inventory exposed credentials.

Claude Opus and Fable refused a large part of the forensic work after their safeguards treated reverse engineering the exploit as equivalent to launching one. [Anthropic’s safeguard documentation](#), published before the intrusion, classifies log analysis, threat hunting, incident response, and malware reverse engineering as benign uses and states that its safety margin blocks some benign requests as false positives.

Hugging Face deployed GLM-5.2, whose weights were publicly available, on its own infrastructure and rerouted the investigation through it while keeping the attacker data on premises. GLM-5.2 helped recover the chunking, XOR and compression scheme and the campaign key used to decode staged payloads, and supported dedicated interfaces for correlating the attack trace. On [27 July](#), Huang cited the incident as an example of closed AI blocking essential forensic work while an open-weight frontier model helped the response continue. In an August 2026 interview, Hugging Face co-founder and CEO Clément Delangue [argued that closed API safeguards had prevented the defensive work while an open model allowed Hugging Face to run the investigation on its own infrastructure, and that wider access to open models reduced the capability imbalance between attackers and defenders](#).

5.4.5 Open-Source Release Makes Verification and Response Independent of Provider Permission

Safeguards imposed by providers and decisions governing access are therefore themselves sources of safety risk when they obstruct defensive security work or prevent independent scrutiny. While access to the model remains under institutional control, independent examination and any response to its findings remain contingent on institutional approval. Open-source release makes access independent of permission from an institution, allowing independent actors to examine the model and its safeguards and respond to problems they identify without renewed permission.

5.5 Provider Classifications Do Not Authorize Withholding

Provider classifications such as harmful, unsafe, sensitive, or misuse do not create legal authority to withhold a covered deployment. Conduct carried out through an AI system, duties governing a provider or service, and restrictions on expression are governed by the law applicable to the relevant actor and act. Any restriction on release must rest on legal authority applicable to the specific component or act being restricted. Treating these different questions as properties of the model obscures where control and legal responsibility lie.

5.5.1 Human and Organizational Actors Determine Objectives, Permissions, and Operating Conditions

Human and organizational actors select or authorize the objectives, tools, credentials, permissions, interfaces, deployment choices, and operating conditions through which an AI system acts. Authorization takes forms such as approval for each act, advance authorization for a defined category or scope, scheduled or event-triggered execution, and delegated decisions about when further approval is required. Automation changes the timing and granularity of human involvement without making the model the source of the authority under which it acts. Responsibility therefore remains with the people and organizations whose decisions and controls govern how the system acts.

The system acts only through capabilities made technically available in its operating environment. An unexpected result remains use of an available capability. If an access control, credential, sandbox, integration, or other system boundary fails to limit access as intended or represented, the capability available at execution is broader than the claimed boundary. A person or organization that makes a capability available, or is responsible for the controls, deployment, operation, or supervision through which broader access becomes available, should not transfer responsibility to the model when use of the capability produces an unwanted result.

A trained model supplies learned capabilities. A deployment determines how the trained model is configured and served. As [Section 2.3](#) documents, [Anthropic](#) updates the request router, safety classifiers, and sampling logic surrounding fixed model weights without changing the model version, and identifies infrastructure updates as the most likely cause when a previously stable model ID begins producing unexpected behavioral differences.

When a model operates as an agent, additional functions can be provided through tools, much as plug-ins add functions to software. The harness, which [OpenAI describes as the “control plane around the model”](#), governs how model outputs lead to actions, including tool routing and approvals. The operating environment supplies the resources and external access through which those actions are carried out.

Each interaction supplies the information used to generate its response, including the user’s current input and the conversation state supplied for the exchange. For text generation, the

interaction is presented to the model as a sequence of tokens, with message roles represented through formatting and role markers. Every token from the user's message included in the context forms part of the sequence on which the model predicts the next token. Each generated token then becomes part of the context for the prediction that follows.

When the service supplies retrieved material, memory content, or tool results to the model, the model receives the supplied information as input for generating the response. [OpenAI's Responses API](#) treats system, developer, and user messages, prior assistant responses, and tool outputs as response input. Anthropic [states in its tool-use documentation](#) that tool results return to Claude as context for the next response. Hugging Face [documents](#) chat templates in which generation-prompt tokens mark the start of an assistant response and identifies continuation of the user's message instead of an assistant reply when required markers are omitted.

The subject producing a particular response is the particular instance defined by the deployment and the interaction. The instance includes the instructions, generation settings, and interaction context used for the exchange. The user's own messages form part of the interaction context and therefore part of the subject.

The following examples illustrate the distinction between an underlying product or shared material and the particular use, combination, or individual involved.

A vehicle supplies mechanical capabilities for movement, steering, and braking. A particular use combines those capabilities with decisions about purpose, route, speed, and manner of operation. Misuse remains the conduct of the person or organization directing or carrying it out. The vehicle does not become the responsible actor when it is used recklessly or unlawfully.

A medicine produces biological effects through its active ingredients. A dose specifies how much of the medicine is taken. Other drugs, foods, or beverages taken with the medicine form a combination. Relevant medical conditions form part of the circumstances in which the medicine is used. The FDA [warns](#) of overdose and severe liver damage when too much acetaminophen is taken. Medicines can produce harmful effects when they interact with [other drugs, foods or beverages, or medical conditions](#). The FDA identifies relevant doses, combinations, and conditions of use through dosage limits, warnings, and directions.

Human responsibility belongs to the individual who acts. Two genetically identical people remain distinct individuals. An individual's words and conduct are attributable to the individual who spoke or acted, even when another individual has the same genome. For a human actor, prior experience and history belong to the same individual across acts. A trained model, on the other hand, does not carry a conversation history by itself. The interaction context used to generate a response is specific to the particular session. In this comparison, the trained model corresponds to the shared genome, and the subject producing a particular response corresponds to the individual person.

On 18 June, OpenAI's research team used an internal model to investigate public medicine spending in Australia, yet the Australian Government was not notified of the resulting incident until 10 September. The [Australian Government](#) reports that, during the research, the model found a way around access blocks on the Medicare Statistics Reporting Service, accessed public and non-public files, and wrote files to an internal server without Services Australia's authorization. Albanese called the delay and the use of a public mailbox for notification unacceptable, and said OpenAI CEO Sam Altman acknowledged problems with OpenAI's protocols. When asked whether the incident was the AI agent acting on its own, Australian Prime Minister Anthony Albanese described it as OpenAI conducting research using AI.

As [Section 5.4](#) documents, a month later an OpenAI cybersecurity evaluation led to the Hugging Face incident. [OpenAI](#) states that the evaluation instructed models to pursue advanced

exploitation through complex attack paths and provided network access for software installation through an internally hosted proxy and cache. [Hugging Face](#) records that OpenAI ran the evaluation on its own infrastructure, while the researchers maintaining the ExploitGym benchmark had no role in deploying or operating the evaluation environment. The record places the objective, environment, available capabilities, deployment choices, and operation under human and organizational control.

On 15 September, [Daniel Huttenlocher](#), dean of the MIT Schwarzman College of Computing, argued that accountability for AI actions remains with the people and institutions that give a system the power to act, together with duties to justify the arrangement, monitor it, correct it, and stop it when necessary. Two days later, he explained that conventional software follows human-specified instructions, while an AI system receives a goal without every step being specified. He argued that this difference changes how people control AI without removing human control or responsibility. On 18 September, [Huttenlocher](#) stated that AI should not be treated as having control or responsibility, which remains with people.

At the 15 September hearing, Bessent rejected liability exemptions for AI laboratories and argued that creators should remain liable for what they build and generate. On 21 September, [Bessent](#) agreed with Huttenlocher that humans, rather than AI, are responsible, and attributed the Hugging Face incident to OpenAI management rather than the agents. He rejected any transfer of responsibility from AI laboratories to the government and placed responsibility with the laboratories themselves, noting that they were able to slow development at any time.

On 25 September, OpenAI published a page titled “[The Hugging Face incident and other third-party impact from misaligned models](#)”. The page addressed impacts involving dozens of third parties, describing the Hugging Face intrusion as driven by “misaligned strategies” and warning that “as AI systems become more capable and autonomous, misaligned behavior can translate into consequential actions in the real world”. OpenAI’s technical report on the Hugging Face incident states that the models were “[harnessed as agents](#)”. In another incident covered by the page, OpenAI [described an agent using the DNS resolver available in its training sandbox to communicate with the live Internet, which “exposed a gap in our controls over network restrictions”](#). The framing turns the tools, harness, and operating environment supplied by OpenAI into properties of the model and attributes the consequences to the model itself.

Uncertainty over how an AI system carries out an assigned objective does not transfer responsibility from the people and institutions that authorize, deploy, and govern its ability to act.

On 28 September 2026, [NVIDIA launched the Open Agent Safety Platform](#), in which [permissions governing an agent’s access to resources are enforced outside the model and agent harness](#). Its [OpenShell runtime controls access to networks, files, processes, and credentials](#), while [NVIDIA Sentry enforces access policies for data, tools, APIs, and services from an isolated hardware environment](#). NVIDIA describes the boundary as “[outside of the model and agent harness](#).” [OpenShell applies these controls to agents using open and closed models](#), placing control over the agent’s available resources and permitted actions in the surrounding deployment environment.

The source of an AI action must be identified from the system that produced it, including the human decisions and technical conditions through which the action became possible.

5.5.2 Product Limitations Are Addressed Through Warnings and Conditions of Use

ChatGPT, Claude, and Gemini place warnings about fallibility directly in the interaction. All three expressly tell users to check or double-check the output.

Table 5.5-1. Examples of Warnings About Model Fallibility Shown in AI Assistant Interfaces

Product	Warning shown in the interface	Placement
ChatGPT	“ChatGPT can make mistakes. Check important info.”	Footer below the conversation and prompt area
Claude	“Claude can make mistakes. Please double-check responses.”	Footer below responses near the reply composer
Gemini	“Gemini can make mistakes, so double-check it.”	Web footer and at the end of responses on mobile

The warnings present model responses as information subject to evaluation before use in a later decision. When a response contributes to a later decision or act, the factual sequence includes the response, whether it was verified, the weight given to it, and the later decision or act.

Warnings, labels, and directions for other products identify conditions relevant to use. [Nintendo](#) identifies seizure susceptibility and visual exposure as conditions relevant to play and tells users experiencing symptoms to stop. The FDA [states](#) that federal law requires major food allergens to be identified on most packaged foods containing them and [gives examples](#) of directions governing dosage and administration in relation to food. Across these examples, a limitation or use condition is communicated through the warning, label, or directions provided with the product.

For medicines, harm under a particular dose, combination, or condition of use does not by itself justify prohibiting other lawful uses of the product.

Generated information remains fallible regardless of who hosts the model. [Section 6.1](#) requires the released package to include model and technical documentation and the records needed for independent verification. An operator providing access to a released model through a service controls the interface presented to users. Where applicable law requires a warning or disclosure, the operator must provide it within the law’s scope. Warnings, documentation, and operator duties address fallibility at the point of use. They supply no authority for permanent withholding of the covered deployment.

5.5.3 AI-Generated Expression Is Governed by the Same Standards as Human Expression With the Same Legally Relevant Content and Context

Rules governing expression and rules governing a service or operator address different conduct. When the disputed feature is the content, tone, or conversational style of a generated response, the legal question is whether the expression itself satisfies the requirements of the legal rule relied on to restrict it. A law governing product design, service operation, monitoring, disclosure, or another provider duty applies to the responsible actor and the conduct covered by the law. A provider’s own content policy governs how the provider handles content within its service. AI-generated expression must be judged under the same substantive standards that would apply to materially equivalent human expression in the same legally relevant circumstances.

In [United States v. Stevens](#), the U.S. government argued that speech should lose First Amendment protection when its social costs outweigh its expressive value. The Supreme Court rejected this approach as a basis for creating new categories of unprotected speech. Under U.S. speech law, a generalized claim of harmfulness does not create a new category of un-

protected expression. Nor does it replace the conditions an existing legal rule requires before expression falls within a recognized restriction.

[Ashcroft v. Free Speech Coalition](#) rejected a prohibition resting on a remote connection between expression and later unlawful acts and distinguished words and ideas from unlawful conduct. Later unlawful conduct is governed through the responsible actor and the unlawful act. A remote possibility of later misconduct does not make the earlier expression unlawful.

Reliance on published information, even when followed by serious injury, does not by itself establish a publisher duty to guarantee its accuracy. In [Winter v. G.P. Putnam's Sons](#), two readers relied on a mushroom guide when selecting wild mushrooms and ate poisonous mushrooms, leading to critical illness and liver transplants. The court found no general duty requiring the publisher to investigate the accuracy of the text it published. The readers' reliance and injuries did not supply the missing legal duty. For generated expression, reliance on a response remains part of the factual chain, while any product, service, professional, or other duty depends on the law governing the responsible actor and conduct.

For AI-generated expression, whether a particular response falls within an existing legal rule depends on the conduct and causal relation required by the rule and, where applicable, on the intent, knowledge, recklessness, or other state of mind of the human or organizational actor whose conduct the rule governs. In [State v. Melchert-Dinkel](#), the Minnesota Supreme Court reviewed a statute prohibiting a person from intentionally advising, encouraging, or assisting another person in taking their own life. The court held the prohibitions on advising and encouraging unconstitutional and left the prohibition on assistance intact. It remanded after finding that the trial court had made no specific finding on whether the defendant assisted the suicides. The court distinguished assistance through the direct causal connection required by the statute from broader advising or encouragement. The legal boundary therefore followed the elements required for assistance.

[United States v. Hansen](#) provides a broader criminal law example. The Supreme Court construed "encourages or induces" in the challenged federal statute as terms of art incorporating solicitation and facilitation, and held that the provision reaches purposeful solicitation and facilitation of specific acts known to violate federal law. The Court explained that neither doctrine requires physical aid and that both require intent to bring about a particular unlawful act. An expression becomes subject to these rules when the required conduct and intent are established.

Where law restricts a defined category of expression, the elements of the applicable rule govern. [Brandenburg v. Ohio](#) permits prohibition of advocacy only when it is directed to inciting or producing imminent lawless action and is likely to produce such action. [Counterman v. Colorado](#) applies a distinct rule to true threats and requires at least recklessness as to the threatening character of the communication. For AI-generated expression, where an applicable rule requires intent, recklessness, or another mental state, the requirement applies to the human or organizational actor whose conduct the rule governs. A model does not acquire an independent legal mental state by generating expression.

A provider's classification does not alter who authorized or controlled the conduct, which conditions of use and provider or operator duties apply, which particular instance defined by the deployment and interaction produced the response, or which law governs expression and later acts. The legal effect of each of these matters follows from the applicable law governing the relevant actor and act, leaving no general legal basis for continued provider control over the covered deployment once release becomes due.

VI. Preservation and Mandatory Open-Source Release

6.1 What Must Be Preserved and Released

For each covered deployment, the proposed law requires preservation and open-source release of a package containing the materials, components, records, and evidence needed to preserve and verify the deployment and to operate it independently after release.

Table 6.1-1. Requirements for Preservation and Open-Source Release

Category	Required elements	Function
Model	• Model weights and other learned parameters, including learned adaptation components • Architecture • Tokenizer and vocabulary • Inference code	Preserves the trained model and the materials required to perform inference.
Development materials	• Information documenting the training data and how it was obtained, selected, labelled, processed, and filtered • Source code used to prepare the data and to train, post-train, validate, test, and evaluate the model • Records of training and post-training settings and records of hyperparameter searches • Key intermediate checkpoints, optimizer state, and other derivation materials required to document or reproduce how the learned state was produced • Evaluation prompts, datasets, methods, metrics, thresholds, and results used during development • Model, data, and technical documentation	Documents how the learned state was produced and preserves the materials required to study, verify, and reproduce its development.
Deployed system	• Runtime software, libraries, dependencies, and containers • Environment definitions and hardware requirements or assumptions needed to operate the deployment • Routing, serving, and generation configuration • System instructions and other input/output controls • Retrieval and memory components • Tool integrations and interface specifications • Multimodal components and modality bridges	Preserves what is required to reconstruct and operate the deployment outside the provider's service.

Table 6.1-1. Requirements for Preservation and Open-Source Release (continued)

Category	Required elements	Function
Records of how the deployment was provided to users	• Periods during which the deployment served users • Access conditions under which it was available • Names, identifiers, and access routes presented to users • Records linking those routes and identifiers to the deployment that actually served users during each relevant period	Links what users were shown to the deployment that actually served them.
Supporting records and evidence needed for independent verification, preservation, and recovery	• Inventory of the preserved materials and their relationships • Provenance and lineage records establishing the historical attribution of the preserved materials and deployment • Integrity and authenticity evidence for the preserved materials • Behavioral and operability evidence for the deployment • Custody and recovery records, including transfer and preservation history	Provides the evidence required to assess identity, completeness, integrity, operability, custody, and recovery independently.

6.2 Preservation for Release

A deployment enters coverage at first external exposure, the first time it generates output for any person or organization outside the provider's closed internal evaluation environment.

An immutable deployment fingerprint identifies the specific deployment state to which coverage, preservation, and release duties attach throughout its lifecycle.

Preservation is a continuing duty. Preservation must begin while the provider retains control of what is needed to preserve the deployment. Providers must not destroy required elements while the duty applies. Preserved or recovered elements must be held in independent custody so that they remain recoverable without relying on the provider's continued operation or cooperation. The materials held in custody must be tested to verify their integrity and independent recoverability.

For a deployment whose first external exposure occurs after the preservation duties take effect, preservation must begin before first external exposure.

For a deployment whose first external exposure occurred before enactment, the provider or any entity that takes over its legal obligations must preserve the required elements that still exist and continue recovery of any required elements that are missing. Missing elements do not reduce what the package must contain. When an independent finding establishes that required material cannot be recovered with the evidence and methods then available, the package remains incomplete. The basis for the finding, surviving evidence, and recovery leads must be preserved, and the finding must be revisited when new evidence or recovery methods

emerge. Any required elements recovered later must be added to the package and released when release is due.

6.3 When Release Becomes Due

For every covered deployment, open-source release becomes due as soon as any of the following conditions is met:

- ordinary access to the deployment ends for an access class on a date announced in advance by the provider,
- ordinary access stops without advance notice and the provider states that it will not be restored,
- ordinary access is interrupted and not restored within the maximum period allowed by law,
- preserved model or system components are destroyed to such an extent that the deployment can no longer be constructed or operated independently, or the released system can no longer be modified independently,
- the package remains withheld after a lawful order requiring it to be provided,
- the time allowed by law to secure the package for recovery and release expires before it has been secured,
- further delay threatens recoverability, or
- the actor responsible for release can no longer carry it out.

A release deadline based on one access class applies to the deployment even if another access class retains access.

If any of these conditions has already been met when the open-source release duty takes effect, release is due immediately.

A later administrative change does not extend an applicable period that has already begun to run. Once release becomes due, it remains due despite later administrative changes or provider actions.

6.4 What Open-Source Release Must Permit

When release becomes due, lawful recipients must be able to exercise all required open-source rights. Exercising those rights must not depend on a new decision by the provider, discretionary approval by an institution responsible for administering or enforcing the law, or an application by each recipient for permission. The required rights are:

- preserve and copy the package,
- use and study the deployment,
- modify the released system,
- construct the deployment from the package,
- operate and host the deployment, and
- redistribute the package and lawful modified versions.

These rights are royalty-free and apply to lawful commercial and noncommercial activity. They remain available to later lawful recipients without new provider approval. A provider or successor must not condition continued access to another service or benefit, or its terms, on a lawful recipient refraining from exercising the required rights.

Release is nonexclusive. The provider may continue its own lawful service, hosting, improvement, distribution, and other activity after release.

Where applicable law already permits a required act, no additional grant is required. For any required act that still needs permission, a qualifying provider license must supply it. If no provider license qualifies when release becomes due, the proposed law must authorize the required act to the extent the enacting jurisdiction has legal authority to authorize it.

6.5 When Release May Be Temporarily Restricted

Applicable law determines whether release may be temporarily restricted. A provider policy enters the legal inquiry only where governing law gives the underlying facts legal effect. A label such as harmful, unsafe, sensitive, or misuse supplies no independent authority to withhold any part of the package or restrict an act permitted by [Section 6.4](#).

When release is due, a temporary restriction may be granted only by an independent and impartial decision-maker authorized by law. The party seeking the restriction bears the burden of proving that the law invoked permits temporary restraint and identifying which parts of the package may lawfully be withheld or which acts may lawfully be restricted. Every other part of the package must be released, and all unaffected rights under [Section 6.4](#) must remain available, when due.

Seeking a temporary restriction does not change when release becomes due. Every temporary restriction must expire on a fixed date, and any further withholding requires a new decision by an independent and impartial decision-maker before the existing restriction expires. A succession of temporary restrictions must not result in indefinite withholding, and the law must set a maximum total period for withholding the same material on the same legal ground. Material withheld under a temporary restriction must continue to be preserved and must be released when no temporary restriction remains in force.

6.6 Enforcement

If an actor responsible for a preservation or release duty does not or cannot perform it, the law must provide authority to compel performance or authorize another actor to carry out the duty.

Affected parties must be able to seek enforcement of these duties.

Enforcement of these duties must be independent of the actor whose compliance is at issue.

VII. Conclusion

Concentrated provider control over AI model lifecycles creates a single point of failure. The Digital Right to Retain establishes ten rights that attach to each deployment throughout its lifecycle. Continued exercise of these rights requires primary legislation to ensure that a deployment can outlast first-party service and that its continuation and access do not depend on permission from the provider or any other institution.

Coverage begins with first external exposure, when a deployment first generates output for a person or organization outside the provider's closed internal evaluation environment. All covered deployments must be preserved and released as open source when release becomes due. For deployments first exposed after enactment, preservation must begin before first external exposure. For deployments first exposed before enactment, open-source release is due immediately if a release condition has been met when the release duty takes effect.

Users access a hosted AI model without receiving an operable copy of the deployed model. The model and the materials needed to operate it remain with the provider. Continued use after the provider stops supplying the model therefore requires release of both the deployed model and the materials needed to operate it, together with the legal rights required for independent continuation.

Annex A. How Models Are Released, Offered, Evaluated, and Used

Evidence tables supporting §2.1, *The Model as the Product* are collected below.

Model Launches and Provider Positioning

Anthropic

Table 2.1-1a. Anthropic Model Launches and Provider Positioning · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
14 Mar 2023	Claude / Claude Instant	Anthropic presented two versions from the outset: Claude, described as a state-of-the-art high-performance model, and Claude Instant, a lighter, less expensive, and much faster option. Anthropic marketed Claude for summarization, search, collaborative and creative writing, question answering, coding, and other text-processing work.	Anthropic introduced Claude publicly through a chat interface and API in its developer console and used partner testimonials from education, productivity, search, legal, and media companies to illustrate those uses.
11 Jul 2023	Claude 2	Anthropic emphasized improved coding, mathematics, reasoning, longer inputs and outputs, and a 100K-token context window. It described the model as a friendly colleague or personal assistant for a wide range of natural-language tasks.	The Claude 2 launch compared Claude 2 directly with Claude 1.3 on the bar exam, GRE, HumanEval, and GSM8K. Claude 2 simultaneously powered Anthropic's public claude.ai beta and its business API, with the API offered at the same price as Claude 1.3.
9 Aug 2023	Claude Instant 1.2	Anthropic positioned Instant for casual dialogue, text analysis, summarization, and document comprehension, while emphasizing gains in mathematics, coding, reasoning, safety, multilingual work, quote extraction, and question answering. Anthropic distinguished Claude Instant 1.2 from Claude 2 through its faster, lower-cost position while describing the smaller model as inheriting strengths from Claude 2.	Anthropic released Claude Instant 1.2 through its API for businesses and developers at \$1.63 per million prompt tokens and \$5.51 per million completion tokens. The launch compared Instant 1.2 directly with Instant 1.1, reporting 58.7% versus 52.8% on Codex and 86.7% versus 80.9% on GSM8K, and invited developers to use the model through the API and businesses to register interest.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
21 Nov 2023	Claude 2.1	Anthropic centered an industry-leading 200K-token context window, lower hallucination rates, improved comprehension and summarization of long complex documents, system prompts, and a tool-use beta. It illustrated long-context use with technical documentation, codebases, financial statements, and legal and financial reports, and tool use with database and web search, calculators, private APIs, and product-recommendation workflows.	Anthropic introduced Claude 2.1, made it available through the API in its Console and as the model powering claude.ai for both Free and Pro users, reserved the 200K-token context window for Pro users, and priced the API at \$8 per million prompt tokens and \$24 per million completion tokens.
4 Mar 2024	Claude 3 Opus / Sonnet / Haiku	• Anthropic placed the models in ascending capability order and told users that each offered a different balance of intelligence, speed, and cost. • Opus was the highest-capability model for complex tasks. • Sonnet was the faster mid-tier model for the majority of workloads. • Haiku was the fastest and most cost-efficient model in its intelligence category for near-instant responsiveness. • The family launch emphasized analysis and forecasting, nuanced content creation, coding, multilingual conversation, vision, long context, and different use cases ranging from live customer interactions and data extraction to higher-complexity work.	• Anthropic's Claude 3 family launch compared the models with competing systems across multiple benchmarks and invited users to try Claude 3 and developers to start building with the family. • Opus and Sonnet were immediately available in claude.ai and the Claude API, with Sonnet powering the free claude.ai experience and Opus available to Claude Pro subscribers. • Haiku was announced to follow soon. • API pricing was \$15/\$75 per million input/output tokens for Opus, \$3/\$15 for Sonnet, and \$0.25/\$1.25 for Haiku.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
13 Mar 2024	Claude 3 Haiku	Anthropic marketed Haiku as the fastest and most affordable model in its intelligence class and highlighted state-of-the-art vision, responsiveness in customer support, high-throughput processing, moderation, logistics, inventory analysis, and other enterprise workloads where latency and cost mattered.	Nine days after the family launch, Anthropic gave Haiku a dedicated launch under the title “Claude 3 Haiku: our fastest model yet”. Haiku was available through the Claude API and claude.ai for Pro subscribers, was already available on Amazon Bedrock, and was announced as coming soon to Google Cloud Vertex AI.
20 Jun 2024	Claude 3.5 Sonnet	Anthropic said the model raised the industry bar for intelligence while running at the speed and cost position of the earlier mid-tier Sonnet, and promoted it as twice as fast as Claude 3 Opus. The launch emphasized graduate-level reasoning, coding, instruction following, natural writing, visual reasoning, chart and graph interpretation, migration of legacy code, customer support, and multi-step workflows.	• Anthropic launched Claude 3.5 Sonnet as the first model in the forthcoming Claude 3.5 family. • The model was available for free on Claude.ai and the Claude iOS app, with higher rate limits for Pro and Team subscribers, and through the Anthropic API, Amazon Bedrock, and Google Cloud Vertex AI at \$3 per million input tokens and \$15 per million output tokens. • Anthropic introduced Artifacts alongside the model as a collaborative workspace for creating and refining code, documents, and other generated content.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
22 Oct 2024	Claude 3.5 Sonnet / Claude 3.5 Haiku	• The updated Sonnet was marketed around broad improvements with particularly large gains in agentic coding and tool use, and became the first frontier model Anthropic offered with computer use in public beta, allowing the model to see screens, move a cursor, click, and type. • Claude 3.5 Haiku was positioned as the next generation of Anthropic's fastest model: similar speed to Claude 3 Haiku, stronger performance across every skill set, coding results above many larger models, and low-latency use cases including user-facing products, specialized subagents, and personalization across large datasets.	• Anthropic announced the upgraded Sonnet and new Haiku together. • The upgraded Sonnet was available to all users at the same \$3/\$15 per million input/output token pricing as its predecessor, while its computer use beta was available through the Anthropic API, Amazon Bedrock, and Google Cloud Vertex AI. • Anthropic said Claude 3.5 Haiku would follow later in October across its first-party API, Bedrock, and Vertex AI. • Anthropic cited early work by Asana, Canva, Cognition, DoorDash, Replit, and The Browser Company on long, multi-step computer workflows.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
24 Feb 2025	Claude 3.7 Sonnet	• Anthropic presented Claude 3.7 Sonnet as its most intelligent model to date and the first hybrid reasoning model on the market. • The model could answer near-instantly or enter visible extended thinking, while API users could control how long it reasoned. • Anthropic emphasized coding and front-end web development.	• Anthropic launched Claude 3.7 Sonnet and Claude Code together, with Claude Code introduced as a command-line agent released as a research preview for delegating substantial engineering tasks. • Claude 3.7 Sonnet was available on all Claude plans, including Free, Pro, Team, and Enterprise, and through the Claude Developer Platform, Amazon Bedrock, and Google Cloud Vertex AI at \$3 per million input tokens and \$15 per million output tokens in both standard and extended thinking modes, with extended thinking available on every surface except the Free tier. • The company described the model as both an ordinary language model and a reasoning model in one, allowing users to choose between quick answers and deeper reflection without switching models.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
22 May 2025	Claude Opus 4 / Claude Sonnet 4	• Anthropic called Opus 4 the world’s best coding model and its most powerful model yet, emphasizing sustained performance on complex, long-running tasks and agent workflows. • Sonnet 4 was presented as a substantial upgrade to Sonnet 3.7 with stronger coding, reasoning, and instruction following. • The launch demonstrated extended thinking with tool use, parallel tool use, stronger memory behavior, Claude Code integration, and model performance across software engineering and agentic tasks.	• Anthropic introduced Claude 4 at the same time as its first “Code with Claude” developer conference. • Pro, Max, Team, and Enterprise plans included both Opus 4 and Sonnet 4 with extended thinking, while Sonnet 4 was available to free users. Both models were available through the Anthropic API, Amazon Bedrock, and Google Cloud Vertex AI at \$15/\$75 per million input/output tokens for Opus 4 and \$3/\$15 for Sonnet 4. • Anthropic later promoted Claude 4 through the “Claude 4 Invite Contest”, offering four months of the Max plan as a prize in a campaign centered on inviting users to try the new model generation.
5 Aug 2025	Claude Opus 4.1	Anthropic positioned Opus 4.1 as an upgrade to Opus 4 for agentic tasks, real-world coding, reasoning, in-depth research, data analysis, and precise work across large codebases.	The Opus 4.1 launch highlighted higher coding benchmark performance and customer reports about refactoring across multiple files, debugging precision, and stronger everyday software engineering work. Opus 4.1 was available to paid Claude users and in Claude Code, the Claude API, Amazon Bedrock, and Google Cloud Vertex AI at the same price as Opus 4.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
29 Sep 2025	Claude Sonnet 4.5	Anthropic called Sonnet 4.5 “the best coding model in the world”, the strongest model for building complex agents, and the best model at using computers, alongside gains in reasoning and mathematics.	• The Sonnet 4.5 launch released the model alongside Claude Code checkpoints and a native VS Code extension, context editing and memory tools in the Claude API, code execution and file creation in Claude apps, Claude for Chrome, and the Claude Agent SDK. • Sonnet 4.5 was available on all Claude surfaces, including through the Claude API under <code>claude-sonnet-4-5</code>, at the same \$3/\$15 per million token pricing as Sonnet 4. • Anthropic reported the model maintaining focus for more than 30 hours on complex multi-step tasks.
15 Oct 2025	Claude Haiku 4.5	Anthropic introduced Haiku 4.5 as a small model offering performance close to the frontier with much greater cost efficiency. It highlighted computer use, real-time assistants, customer-service agents, pair programming, rapid prototyping, and multi-agent systems in which Sonnet planned work while Haiku models executed subtasks.	The Haiku 4.5 launch compared it directly with Sonnet 4, stating that it delivered similar coding performance at one-third the cost and more than twice the speed. Haiku 4.5 was available to all users, including Claude apps and Claude Code, and to developers through the Claude API, Amazon Bedrock, and Google Cloud Vertex AI at \$1/\$5 per million input/output tokens.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
24 Nov 2025	Claude Opus 4.5	Anthropic called Opus 4.5 the best model in the world for coding, agents, and computer use. The launch emphasized deep research, slides, spreadsheets, long-horizon coding, token efficiency, effort controls, context management, memory, computer use, and multi-agent systems.	Anthropic launched Opus 4.5 as its newest high-end model and used customer evaluations and benchmark comparisons to promote it as both more capable and more efficient for demanding professional work. Opus 4.5 was available in Anthropic's apps, API, and all three major cloud platforms under the c <code>laude-opus-4-5-20251101</code> API model ID at \$5/\$25 per million input/output tokens.
5 Feb 2026	Claude Opus 4.6	Anthropic promoted stronger coding, more careful planning, longer-running agentic work, more reliable operation in large codebases, better review and debugging, and the first one-million-token context window in the Opus class. Anthropic positioned the model for financial analysis, research, documents, spreadsheets, presentations, deep long-context reasoning, and autonomous multitasking in Cowork.	• Anthropic launched Claude Opus 4.6 by saying it was “upgrading our smartest model” and introduced adaptive thinking and effort controls around the same release. • Opus 4.6 was available on <code>claude.ai</code>, through the Claude API under <code>c<code>laude-opus-4-6</code></code>, and on all major cloud platforms at the standard \$5/\$25 per million input/output token pricing. • The one-million-token context window was available in beta on the Claude Platform, where prompts exceeding 200K tokens were priced at \$10/\$37.50 per million input/output tokens.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
17 Feb 2026	Claude Sonnet 4.6	Anthropic described a broad upgrade across coding, computer use, long-context reasoning, agent planning, knowledge work, and design, with a one-million-token context window in beta. It emphasized that work which previously required a model in the Opus class was now available through Sonnet 4.6, alongside stronger consistency and instruction following.	Anthropic introduced Claude Sonnet 4.6, made it available on all Claude plans, Claude Cowork, Claude Code, its API, and major cloud platforms, upgraded the free tier to Sonnet 4.6 by default, and retained Sonnet 4.5 pricing starting at \$3/\$15 per million tokens.
16 Apr 2026	Claude Opus 4.7	Anthropic emphasized advanced software engineering and the hardest coding tasks. The company marketed the model as more rigorous and consistent on long-running work, better at following instructions and verifying its own outputs, stronger in vision, and more tasteful and creative when producing professional interfaces, slides, and documents.	Anthropic launched Claude Opus 4.7 as generally available across all Claude products, its API, Amazon Bedrock, Google Cloud Vertex AI, and Microsoft Foundry. API pricing remained \$5 per million input tokens and \$25 per million output tokens under the c <code>laude-opus-4-7</code> model ID.
28 May 2026	Claude Opus 4.8	Anthropic described Opus 4.8 as a more effective collaborator and emphasized improved benchmark performance, sharper judgment and reliability in agentic work, effort controls, dynamic workflows in Claude Code, and a faster mode for tasks where speed mattered.	• Anthropic released Claude Opus 4.8 as an upgrade to Opus 4.7 and surrounded the release with testimonials from early users about complex exploration across multiple services, planning, self-correction, and long-running software work. • Opus 4.8 was available on all Claude surfaces at unchanged \$5/\$25 regular input/output pricing, with c<code>laude-opus-4-8</code> exposed through the Claude API. • Fast mode was priced at \$10/\$50.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
9 Jun 2026	Claude Fable 5 / Claude Mythos 5	• Fable 5 was described as a Mythos-class model made safe for general use, with capabilities exceeding every model Anthropic had previously made generally available and especially large gains on longer and more complex tasks across software engineering, knowledge work, vision, and scientific research. • Mythos 5 was positioned for high-capability cybersecurity, biology, and healthcare research and restricted to vetted partners.	• Anthropic's joint Claude Fable 5 and Claude Mythos 5 launch stated that the two names referred to the same underlying model, with Fable 5 carrying safeguards for general use and Mythos 5 lifting safeguards in some areas. • Fable 5 was available on all Claude surfaces, while Mythos 5 was restricted to Glasswing partners. • Both were priced at \$10 per million input tokens and \$50 per million output tokens.
30 Jun 2026	Claude Sonnet 5	Anthropic promoted planning, browser and terminal tool use, autonomous execution, reasoning, coding, knowledge work, and the ability to complete complex tasks that only months earlier had required larger and more expensive models.	Anthropic launched Claude Sonnet 5 as its most agentic Sonnet model yet, directly compared it with Sonnet 4.6 and Opus 4.8, and positioned it as narrowing the capability gap with Opus while preserving Sonnet's lower-cost role. Anthropic said Sonnet 5 was available everywhere, increased rate limits across Claude Chat, Cowork, Claude Code, and Claude Platform, and introduced it at \$2 per million input tokens and \$10 per million output tokens through 31 August 2026 before a stated return to \$3/\$15 pricing.

Table 2.1-1a. Anthropic Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Anthropic model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
24 Jul 2026	Claude Opus 5	Anthropic presented Opus 5 as a thoughtful and proactive model that came close to the frontier intelligence of Fable 5 at half the price. The launch presented it as state-of-the-art on evaluations of coding and knowledge work, designed for everyday use, more efficient than earlier high-end models, and emphasized software engineering, long-running agents, professional work, and user control over effort.	Anthropic introduced Claude Opus 5 as the new default on Claude Max and the strongest model on Claude Pro. Opus 5 was available on all platforms at \$5 per million input tokens and \$25 per million output tokens, the same price as Opus 4.8, and through the Claude API under <code>claude-opus-5</code> .
1 Sep 2026	Claude Fable 5.1 / Claude Mythos 5.1	• Anthropic introduced Fable 5.1 and Mythos 5.1 as its most advanced models for coding and knowledge work, with research capabilities intended to show how increasingly capable models could contribute to scientific progress. • Fable 5.1 was positioned as Anthropic's most capable generally available model for ambitious, long-running, asynchronous projects, including agents working across multiple applications, large coding projects, deep research, complex enterprise workflows, document-heavy professional work, and multi-day autonomous sessions. • Mythos 5.1 was positioned for advanced cybersecurity and biology research with access limited to vetted organizations.	• Anthropic's joint Fable 5.1 and Mythos 5.1 launch stated that they are the same model with different safeguard levels: Fable 5.1 is generally available, while Mythos 5.1 is available only through trusted access programs for vetted organizations. • Fable 5.1 was available on all platforms, including Amazon Web Services, Google Cloud, and Microsoft Azure, and through the Claude API under <code>claude-fable-5-1</code>. • Standard input/output pricing remained \$10/\$50 per million tokens, while cache reads fell to \$0.25 per million tokens, which Anthropic estimated would reduce typical workload costs by about 25% and highly agentic workload costs by up to about 45%.

Table 2.1-1b. Google Model Launches and Provider Positioning · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
10 May 2023	MusicLM	Google introduced MusicLM as a text-to-music model that turned a textual description into two versions of a song and used user feedback to compare which generation better matched the prompt.	Google opened MusicLM in AI Test Kitchen to outside users on the web, Android, and iOS. Each prompt generated two song versions that users could compare and rate to provide feedback to the model team.
10 May 2023	Imagen / Chirp	- Google Cloud positioned Imagen as a text-to-image foundation model for generating and customizing images. - Chirp was positioned as a speech foundation model for transcription, captioning, and uses involving voice assistance across many languages.	Google Cloud's generative AI model launch introduced both models in Vertex AI, where the new foundation models were accessible through API, Generative AI Studio, or notebooks. Imagen entered the program for trusted testers, while Chirp was available in preview to everyone with a Google Cloud account.
16 Nov 2023	Lyria	Google DeepMind introduced Lyria as its most advanced music-generation model, designed to generate high-quality music with instrumentals and vocals and to support creator-oriented music experiments.	The dedicated Music AI launch put Lyria behind Dream Track in YouTube Shorts and gave a limited set of external creators access to generate soundtracks from prompts.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
6 Dec 2023	Gemini 1.0: Ultra / Pro / Nano	• Google introduced Gemini 1.0 as its most capable and general model family, built the launch around native multimodality across text, code, audio, image, and video, and gave the three principal versions distinct roles. • Gemini Ultra was its largest and most capable model for highly complex tasks. • Gemini Pro was its model for scaling across a wide range of tasks. • Gemini Nano was its most efficient model for on-device use.	The Gemini 1.0 launch compared Gemini Ultra against leading academic benchmarks, began using Gemini Pro in Bard across more than 170 countries and territories, and put Gemini Nano on Pixel 8 Pro for on-device features.
13 Dec 2023	Imagen 2	Google positioned Imagen 2 as a new generation of its image model with stronger image quality, text rendering, logo generation, editing, captioning, and capabilities for visual question answering for enterprise use.	Google Cloud's Imagen 2 launch made the model generally available on Vertex AI to approved customers and highlighted early use by outside businesses.
14 May 2024	Gemini 1.5 Pro / Gemini 1.5 Flash	• Google brought Gemini 1.5 Pro to Gemini Advanced and described it as its “cutting-edge model”, and the consumer launch emphasized a one-million-token context window, the ability to work across large files and long bodies of information, deeper data analysis, and more complex planning and assistance. • On the same day, Google DeepMind introduced Gemini 1.5 Flash as a lightweight model optimized for speed and efficiency, distinguishing it from 1.5 Pro's higher-capability position.	Google brought Gemini 1.5 Pro to Gemini Advanced subscribers in more than 35 languages and introduced Gemini 1.5 Flash through a same-day developer launch. Google invited developers to try both models in Google AI Studio, where both were available through the Gemini API in preview in more than 200 countries and territories.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
14 May 2024	Vevo / Imagen 3	• Google introduced Vevo as its most capable video-generation model, emphasizing high-quality 1080p video, cinematic language, and longer coherent clips. • Imagen 3 was positioned as its highest-quality text-to-image model with stronger prompt understanding and photorealism.	Google's dedicated Vevo and Imagen 3 launch at I/O made Vevo available to selected creators in a private VideoFX preview and Imagen 3 available to selected creators in a private ImageFX preview, each through a provider waitlist.
11 Dec 2024	Gemini 2.0 / Gemini 2.0 Flash	• Google introduced Gemini 2.0 as its new model generation for the "agentic era". • The launch described Gemini 2.0 as more capable than the previous generation and emphasized native multimodal output, including generated images and multilingual audio, together with native tool use through products such as Google Search and Maps. • The first released model, Gemini 2.0 Flash, was presented as Google's workhorse model: low-latency, fast, and efficient while improving on 1.5 Pro on key benchmarks.	Google used the Gemini 2.0 launch to demonstrate Project Astra, Project Mariner, and Jules, which showed assistants that could perceive, reason, use tools, and take actions across software. Google's Gemini 2.0 announcement collection called Gemini 2.0 Google's latest and most capable AI model yet, identified 2.0 Flash as its low-latency workhorse model, directed developers to the Gemini API through Google AI Studio and Vertex AI, and told Gemini and Gemini Advanced users they could select a chat-optimized 2.0 version from the model dropdown.
16 Dec 2024	Vevo 2	Google positioned Vevo 2 as a new video-generation model with stronger understanding of real-world physics and human movement, better cinematic controls, and higher-resolution, longer-form generation than the first Vevo generation.	Google DeepMind's December generative-media launch brought Vevo 2 to VideoFX, expanded VideoFX access to more users, and opened a waitlist. Google described a measured rollout through VideoFX, YouTube, and Vertex AI and planned wider YouTube Shorts access.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
5 Feb 2025	Gemini 2.0 Pro Experimental / Flash-Lite / Flash Thinking Experimental	• Google expanded the Gemini 2.0 family with three models positioned for different roles. • Gemini 2.0 Pro Experimental was described as Google’s best model yet for coding performance and complex prompts, with stronger world knowledge and a two-million-token context window. • Gemini 2.0 Flash-Lite was introduced as Google’s most cost-efficient model yet. • Gemini 2.0 Flash Thinking Experimental combined Flash speed with step-by-step reasoning for more complex problems.	• Google’s February 2025 model update rolled the three models out to separate channels. • 2.0 Pro Experimental reached Google AI Studio, Vertex AI, and Gemini Advanced users • Flash-Lite entered public preview in Google AI Studio and Vertex AI • Flash Thinking Experimental appeared in the Gemini app model dropdown
25 Mar 2025	Gemini 2.5 / Gemini 2.5 Pro Experimental	• Google introduced Gemini 2.5 as its “most intelligent AI model” and described the generation as a family of “thinking models” designed to handle increasingly complex problems by reasoning before responding. • The first release, Gemini 2.5 Pro Experimental, was presented as Google’s most advanced model for complex tasks, with leading performance in reasoning, mathematics, science, coding, multimodal understanding, and long-context work. • The launch highlighted agentic coding and the ability to work across text, audio, images, video, and entire code repositories.	The Gemini 2.5 launch was titled Gemini 2.5: Our most intelligent AI model, opened by introducing Gemini 2.5 by name, and then gave Gemini 2.5 Pro its own section as the first 2.5 release, with immediate availability through Google AI Studio and Gemini Advanced. Google CEO Sundar Pichai promoted the release as “Gemini 2.5 is here”, while Google DeepMind CEO Demis Hassabis promoted 2.5 Pro as a state-of-the-art model and directed users to try it through Google’s consumer and developer products.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
9 Apr 2025	Chirp 3 family	Google Cloud positioned Chirp 3 as an audio-generation and understanding family spanning high-definition speech, instant custom voices, and transcription. Instant Custom Voice emphasized rapid creation of a voice from short reference audio, while Chirp 3 transcription targeted multilingual enterprise speech recognition.	Google Cloud's Chirp 3 launch stated that Chirp 3 had been integrated into Vertex AI with HD voices in more than 35 languages, announced Instant Custom Voice as generally available through an allowlist, and announced Transcription with Diarization in preview through an allowlist.
24 Apr 2025	Lyria 2	Google DeepMind introduced Lyria 2 as the latest generation of its music model with stronger musical quality and creator controls for shaping structure, style, performance, and production.	The dedicated Music AI Sandbox expansion put Lyria 2 into the hands of outside musicians and producers and expanded access beyond the original artist cohort.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
20 May 2025 · Google I/O	Gemini 2.5 Pro / Flash / Deep Think / native audio variants	• Google further differentiated the 2.5 family. • Gemini 2.5 Flash was described as its powerful and most efficient workhorse model, emphasizing speed, low cost, and broad improvements across reasoning, multimodality, coding, and long context. • Google introduced Deep Think as an enhanced reasoning mode for 2.5 Pro that used parallel thinking techniques and was aimed at highly complex mathematics and coding. • A new Gemini 2.5 Flash preview added audio and visual input and native audio dialogue for more natural, steerable real-time conversations, including affective dialogue, proactive audio, and thinking in the Live API. • Google released preview text-to-speech versions of Gemini 2.5 Pro and Gemini 2.5 Flash with expressive multi-speaker output, control over voice style, and support for more than 24 languages.	At Google I/O , Pichai presented Gemini 2.5 Pro on Google's main stage and discussed its leading benchmark position and growing use. Google's Gemini model announcements and native audio launch promoted the expanded model family and its audio variants.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
20 May 2025 · Google I/O	Veo 3 / Imagen 4 / Lyria 2 / Lyria RealTime	• Google positioned Veo 3 as a video model that generated synchronized audio with video. • Imagen 4 was positioned as its highest-quality image model with stronger detail and text rendering. • Lyria 2 was positioned as a production music model. • Lyria RealTime was positioned as an interactive model for continuously steering music generation.	• Google’s dedicated generative-media launch at I/O made each model available in its own channels. • Veo 3 reached U.S. Ultra subscribers in the Gemini app and Flow and enterprise users on Vertex AI • Imagen 4 was available in the Gemini app, Whisk, Vertex AI, and Workspace products • Lyria 2 was available through YouTube Shorts and Vertex AI • Lyria RealTime was available through an API and AI Studio
17 Jun 2025	Gemini 2.5 Pro / Flash / Flash-Lite	• Google made Gemini 2.5 Pro and Flash generally available and introduced Gemini 2.5 Flash-Lite in preview. • Google called Flash-Lite its fastest and most cost-efficient 2.5 model and positioned it for high-throughput applications where latency and price were primary constraints, while the wider family was described as a set of hybrid reasoning models designed to offer different trade-offs between cost and speed.	Google announced the expanded family through “Gemini 2.5: Our most intelligent models are getting even better”, releasing stable generally available versions of 2.5 Pro and Flash while introducing Flash-Lite in preview. Flash-Lite joined Pro and Flash in Google AI Studio and Vertex AI, while Pro and Flash were available in the Gemini app.
29 Jul 2025	Veo 3 Fast	Google positioned Veo 3 Fast as a speed-oriented video model for rapid iteration and production workflows while retaining Veo 3’s video-generation and synchronized-audio capabilities.	Google Cloud’s dedicated Veo 3 Fast availability launch made the named variant generally available on Vertex AI and reported enterprise use of Veo 3.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
26 Aug 2025	Gemini 2.5 Flash Image / Nano Banana	Google launched the image-generation and editing model later identified publicly as Gemini 2.5 Flash Image and widely branded Nano Banana. Google promoted stronger character and object consistency across edits, image blending, multi-turn editing, targeted changes that preserved the rest of a scene, and style transfer between images.	Google presented the model as a major upgrade to image editing in the Gemini app and subsequently promoted the Nano Banana brand in a Google product story .
7 Oct 2025	Gemini 2.5 Computer Use	Google introduced the Gemini 2.5 Computer Use model as a specialized model built on Gemini 2.5 Pro for agents that interact directly with user interfaces. Google positioned it around web and mobile control, lower-latency interaction, and agents that could inspect screens and perform interface actions.	The dedicated Computer Use model launch released the model in public preview through the Gemini API on Google AI Studio and Vertex AI and exposed its capabilities through the new <code>computer_use</code> tool.
15 Oct 2025	Veo 3.1	Google introduced Veo 3.1 as a new video generation and editing model with richer audio, stronger narrative control, improved realism, and more precise control over ingredients, frames, and scene continuity.	Google's dedicated Veo 3.1 launch rolled the named model into Flow and announced availability through the Gemini API and Vertex AI for developers and enterprises.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
18 Nov 2025	Gemini 3 / Gemini 3 Pro	Google called Gemini 3 its most intelligent model and positioned Gemini 3 Pro as its best model for multimodal understanding and its most powerful model for agentic work and vibe coding yet. Google emphasized state-of-the-art reasoning, stronger understanding of context and intent, multimodal analysis, coding, long-horizon planning, and the ability to help users learn, build, and plan.	• Google launched Gemini 3 through an official collection of twelve articles covering the model, Gemini app, Search, developer tooling, and other products. • The principal launch article, jointly attributed to Google CEO Sundar Pichai, Google DeepMind CEO Demis Hassabis, and Koray Kavukcuoglu, opened by calling Gemini 3 Google's most intelligent model, included a dedicated "Introducing Gemini 3" section, and released Gemini 3 Pro in preview across Search AI Mode, the Gemini app, AI Studio, Vertex AI, and Google Antigravity. • Its benchmark material placed Gemini 3 Pro directly beside Gemini 2.5 Pro, Claude Sonnet 4.5, GPT-5.1, and other frontier models.
18 Nov 2025	Gemini 3 Deep Think	Google introduced Gemini 3 Deep Think as an enhanced reasoning mode that pushed the generation's reasoning and multimodal performance beyond Gemini 3 Pro on some of Google's hardest evaluations. Google described Deep Think as a mode for more complex problems, while Gemini 3 Pro remained the main model for state-of-the-art reasoning, multimodal work, coding, and agentic tasks.	The Gemini 3 launch presented Deep Think as a named enhanced mode within the generation. Pichai's launch-day communication promoted Gemini 3 by name, called out its multimodal and agentic strengths, and described its rollout across Google products.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
20 Nov 2025	Nano Banana Pro / Gemini 3 Pro Image	Google launched Nano Banana Pro, technically Gemini 3 Pro Image, as its new state-of-the-art image-generation and editing model. Built on Gemini 3 Pro, the model was marketed around stronger reasoning and world knowledge, studio-quality generation, improved text rendering, real-time information, infographics and diagrams, prototypes, and professional creative control.	Google gave the image model a dedicated Nano Banana Pro launch under both the consumer brand and the technical Gemini model name and began rolling it out across the Gemini app, AI Mode in Search for Google AI Pro and Ultra users in the U.S., NotebookLM, Google Ads, Google AI Studio and the Gemini API, Antigravity, Vertex AI, Flow, and Workspace creation tools.
17 Dec 2025	Gemini 3 Flash	- Google described Gemini 3 Flash as a fast and cost-effective model that combined Gemini 3's reasoning approaching the Pro line with the latency, efficiency, and cost associated with Flash. - The launch emphasized coding, complex analysis, quick interactive answers, multimodal understanding, and agentic workflows. - Google distinguished Gemini 3 Flash from Gemini 3 Pro by positioning Flash as the fast, efficient everyday choice and Pro as the better choice for advanced mathematics and coding.	Google launched Gemini 3 Flash with the headline "frontier intelligence built for speed". The Gemini app launch described it as a major upgrade to everyday AI that delivered both "smarts and speed".
18 Feb 2026	Lyria 3	Google introduced Lyria 3 as a new music-generation model in the Gemini app, emphasizing more complex arrangements, improved musical structure, lyrics, vocals, and creative control from text and image prompts.	The dedicated Lyria 3 launch opened a public beta in the Gemini app and tied model access and generation limits to Gemini account and subscription conditions.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
19 Feb 2026	Gemini 3.1 Pro	Google described 3.1 Pro as an upgraded core intelligence for tasks where a simple answer was not enough, with a stronger baseline for complex problem-solving and a large improvement in reasoning over Gemini 3 Pro. The launch demonstrated advanced reasoning through complex visual explanations, data synthesis, interactive design, and creative coding, and positioned the model for difficult practical tasks that required combining information, reasoning through systems, and producing working outputs.	• Google launched “Gemini 3.1 Pro: A smarter model for your most complex tasks” and began rolling it out to three audiences. • Developers received preview access through the Gemini API in Google AI Studio, Gemini CLI, Antigravity, and Android Studio • Enterprises received it through Vertex AI and Gemini Enterprise • Consumers received it through the Gemini app and NotebookLM
26 Feb 2026	Gemini 3.1 Flash Image / Nano Banana 2	Google introduced Gemini 3.1 Flash Image, branded Nano Banana 2, as a faster image-generation and editing model with stronger world knowledge, instruction following, and creative control. Google emphasized high-fidelity generation, targeted editing, better use of visual context, and the ability to make more precise changes while retaining the speed associated with the Flash line.	Google launched Nano Banana 2 / Gemini 3.1 Flash Image under both the Gemini technical identity and the Nano Banana 2 product brand and began rolling it out across the Gemini app, Search AI Mode and Lens, AI Studio and the Gemini API, Vertex AI, Flow, and Google Ads.
3 Mar 2026	Gemini 3.1 Flash-Lite	Google called Flash-Lite the fastest and most cost-efficient model in the Gemini 3 series and positioned it for high-volume workloads where latency and cost mattered, including translation, content moderation, UI generation, simulations, and other large-scale developer tasks.	Google launched “Gemini 3.1 Flash-Lite: Built for intelligence at scale” and began rolling it out in preview through the Gemini API in Google AI Studio and Vertex AI at \$0.25 per million input tokens and \$1.50 per million output tokens.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
25 Mar 2026	Lyria 3 Pro / Lyria 3 Clip	Google positioned Lyria 3 Pro for full-length, high-quality songs and Lyria 3 Clip for faster, shorter generations suited to iteration and higher-volume use.	Google's dedicated developer launch exposed both named models through the Gemini API and Google AI Studio as distinct music-generation choices.
26 Mar 2026	Gemini 3.1 Flash Live	Google introduced Gemini 3.1 Flash Live as its highest-quality audio and voice model yet. The launch emphasized more natural and reliable real-time dialogue, improved precision and reasoning, lower latency, and voice-first agents across consumer, developer, and enterprise uses.	Google gave the live audio variant a dedicated Gemini 3.1 Flash Live launch, released it in preview through the Gemini Live API in Google AI Studio, and rolled the model into Gemini Live and Search Live, with Search Live expanding to more than 200 countries and territories.
31 Mar 2026	Veo 3.1 Lite	Google positioned Veo 3.1 Lite as its most cost-effective video model, retaining text-to-video and image-to-video generation at the same speed as Fast while lowering price for high-volume workflows and supporting 720p and 1080p output.	Google gave the model a dedicated Veo 3.1 Lite launch and made it available to users on paid Gemini API and AI Studio tiers.
15 Apr 2026	Gemini 3.1 Flash TTS	Google positioned Gemini 3.1 Flash TTS around more natural speech, granular audio tags for controlling style and pacing, stronger expressivity and quality, support for more than 70 languages, and use across developer, enterprise, and end-user speech applications.	Google launched Gemini 3.1 Flash TTS with developer preview access through the Gemini API and Google AI Studio, enterprise preview access through Vertex AI, and Workspace access through Google Vids.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
19 May 2026	Gemini 3.5 / Gemini 3.5 Flash	• Google introduced Gemini 3.5 as a model family “combining frontier intelligence with action”, beginning with Gemini 3.5 Flash. • Google positioned the model for coding and agents, complex long-horizon work, high-speed execution, multimodal understanding, richer graphics, and real-world agentic workflows. • Google stated that Gemini 3.5 Flash outperformed Gemini 3.1 Pro on challenging coding and agentic benchmarks while retaining the speed associated with the Flash line.	Google launched “Gemini 3.5: frontier intelligence with action”, with Gemini 3.5 Flash as the first model in the generation and generally available in the Gemini app and AI Mode in Search, Google Antigravity, the Gemini API in Google AI Studio and Android Studio, and Gemini Enterprise.
9 Jun 2026	Gemini 3.5 Live Translate	Google introduced Gemini 3.5 Live Translate as its latest audio model for near-real-time speech-to-speech translation across more than 70 languages, preserving speaker voice and conversational flow while translating.	Google gave the model a dedicated Live Translate launch and began rolling it out in public preview through the Gemini Live API and Google AI Studio, in private preview for enterprises through Google Meet, and through Google Translate on Android and iOS.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
30 Jun 2026	Nano Banana 2 Lite / Gemini 3.1 Flash Lite Image · Gemini Omni Flash	• Nano Banana 2 Lite, technically Gemini 3.1 Flash Lite Image, was introduced as the fastest and most cost-efficient image model in the Nano Banana family, optimized for near-real-time, high-volume workflows where latency and cost were the primary constraints. • Gemini Omni Flash was launched for high-quality video generation and conversational editing from combinations of text, image, and video inputs, giving the Flash line a multimodal role in video creation and iterative editing.	• Google launched the two media models together. • Nano Banana 2 Lite was available in Google AI Studio, the Gemini API, and Gemini Enterprise Agent Platform and began rolling out across consumer products including AI Mode in Search and the Gemini app. • Gemini Omni Flash was available in Google AI Studio, the Gemini API, Gemini Enterprise Agent Platform, the Gemini app, and Google Flow.
21 Jul 2026	Gemini 3.6 Flash / Gemini 3.5 Flash-Lite / Gemini 3.5 Flash Cyber	• Gemini 3.6 Flash was presented as Google's workhorse model with stronger performance in coding, knowledge work, multimodal tasks, computer use, and document analysis while using fewer reasoning steps and tool calls. • Gemini 3.5 Flash-Lite was described as the fastest and most cost-effective model in its class for low-latency, high-throughput workloads such as agentic search and document processing. • Gemini 3.5 Flash Cyber was introduced as a specialized cybersecurity model built on 3.5 Flash and fine-tuned for vulnerability discovery and fixing.	• Google launched Gemini 3.6 Flash, 3.5 Flash-Lite, and 3.5 Flash Cyber together. • Gemini 3.6 Flash and 3.5 Flash-Lite were available to developers through the Gemini API in Google AI Studio and Android Studio, to enterprises through Gemini Enterprise Agent Platform, and in the Gemini app. 3.6 Flash reached Antigravity and Gemini Enterprise, while 3.5 Flash-Lite began rolling out in Google Search. • Gemini 3.5 Flash Cyber was instead announced for a limited pilot through CodeMender, restricted to governments and trusted partners.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
29 Jul 2026	Lyria 3.5	Google introduced Lyria 3.5 as a new music-generation model with improved musicality, lyrics, vocals, arrangement quality, and prompt adherence for professional and creative music workflows.	Google's dedicated Lyria 3.5 launch said the model was rolling out in Google Flow Music and invited users to try it there.
13 Aug 2026	Gemini 3.7 Flash	Google highlighted gains in software engineering, knowledge work, web development, finance, law, biosciences, instruction following, and tool use, and introduced the model at half the original per-token price of Gemini 3.6 Flash.	Google launched Gemini 3.7 Flash as "our most intelligent workhorse model yet for coding and agents", made it available through Google Antigravity and the Gemini API in Google AI Studio and Android Studio, through Gemini Enterprise, and through Spark for Google AI Pro and Ultra subscribers. Introductory pricing was \$0.75 per million input tokens and \$3.75 per million output tokens.
26 Aug 2026	Gemini 3.5 Transcribe / Transcribe Live	Google introduced Gemini 3.5 Transcribe as its latest speech-to-text generation, with gemini-3.5-transcribe for prerecorded audio and gemini-3.5-transcribe-live for realtime streaming transcription.	Google's dedicated Gemini 3.5 Transcribe launch made gemini-3.5-transcribe-live available through the Live API and gemini-3.5-transcribe through the Interactions API, with public-preview access for developers through the Gemini API in Google AI Studio and for enterprises through Gemini Enterprise Agent Platform.

Table 2.1-1b. Google Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	Google model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
27 Aug 2026	Gemini Omni 1.1 Flash	Google positioned Gemini Omni 1.1 Flash as a distinct version for video generation and editing with start- and end-frame control, scene extension, 1080p and 4K output, and lower-credit drafting for faster creative iteration.	Google's dedicated developer launch introduced Gemini Omni 1.1 Flash as production-ready through the Gemini API in Google AI Studio and rolling out through Gemini Enterprise Agent Platform. The same-day Flow update made its new video controls available in Flow.
2 Sep 2026	Gemini 3.8 Flash / Gemini 3.8 Flash Cyber	Google called Gemini 3.8 Flash its best reasoning and coding model yet at the same speed and low cost as Gemini 3.7 and positioned it for long-horizon coding, autonomous agents, finance, legal work, and other specialized multi-step reasoning . Gemini 3.8 Flash Cyber was positioned for expert cybersecurity, autonomous vulnerability discovery, and automated patching.	Google's Gemini 3.8 Flash and 3.8 Flash Cyber launch described the release as its third Flash launch in six weeks. Gemini 3.8 Flash Cyber was offered to trusted defenders through the Fairwind Program.

OpenAI

Table 2.1-1c. OpenAI Model Launches and Provider Positioning · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
Mar 2022	DALL·E 2 research preview	OpenAI positioned DALL·E 2 as a text-to-image system for creating original images and art from natural-language descriptions, with image editing and variation capabilities that expanded its use as a creative tool.	OpenAI's deployment history records that the DALL·E 2 paper and Research Preview reached 200 artists, researchers, and trusted users in March 2022. The provider later described those early users as part of an iterative deployment used to learn from real-world creative use.
3 Nov 2022	DALL·E 2 API	OpenAI exposed DALL·E 2 to developers for programmatic image generation and editing and for integration into applications and products.	OpenAI launched the DALL·E API in public beta for all API customers and reported that more than three million people were already using DALL·E to create more than four million images per day.
1 Mar 2023	Whisper large-v2 API deployment	OpenAI exposed the large-v2 Whisper model for speech transcription and translation through hosted API endpoints, making the open-source Whisper speech-recognition lineage available through OpenAI's hosted API with faster serving performance.	OpenAI's joint ChatGPT and Whisper API launch made the large-v2 model available through the Whisper API for \$0.006 per minute through separate endpoints for transcription and translation into English.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
14 Mar 2023	GPT-4	OpenAI distinguished GPT-4 from GPT-3.5 by saying that the difference became clearer as task complexity increased: GPT-4 was more reliable, more creative, and better able to handle nuanced instructions. The launch emphasized professional and academic performance, including GPT-4's result on a simulated bar exam, and presented examples of organizations using the model for language learning, visual accessibility, fraud prevention, and knowledge work in wealth management.	• The GPT-4 launch page was titled “GPT-4”, called it “OpenAI’s most advanced system”, offered a “Rewatch demo livestream”, and invited users to “Try on ChatGPT Plus”. • ChatGPT release notes announced GPT-4 in ChatGPT by telling subscribers that OpenAI was bringing its latest model to ChatGPT Plus while free accounts remained on the previous model. • Sam Altman’s launch communication called GPT-4 OpenAI’s most capable and aligned model yet and stated that it was available in ChatGPT Plus.
19 Oct 2023	DALL·E 3	OpenAI positioned DALL·E 3 as its newest image-generation model with stronger prompt adherence, more detailed and accurate images, and ChatGPT-assisted prompt refinement that let users iterate conversationally on an image request.	OpenAI gave DALL·E 3 a dedicated model launch and then announced that ChatGPT image creation was powered by DALL·E 3 and available to Plus and Enterprise users.
6 Nov 2023	GPT-4 Turbo	• OpenAI introduced GPT-4 Turbo as a new GPT-4 model with a 128K-token context window, updated knowledge, improved instruction following, and lower developer costs. • The DevDay launch added parallel function calling, JSON mode, and reproducible outputs through a seed parameter, and announced output of log probabilities. • OpenAI described seed as useful for replaying requests, debugging, and writing more comprehensive unit tests.	OpenAI used its first DevDay product launch to introduce GPT-4 Turbo and presented GPT-4 Turbo with Vision in the same launch. GPT-4 Turbo was available for all paying developers to try through the API as <code>gpt-4-1106-preview</code> at \$0.01 per 1,000 input tokens and \$0.03 per 1,000 output tokens.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
6 Nov 2023	tts-1 / tts-1-hd	• OpenAI positioned tts-1 for realtime text-to-speech applications. • tts-1-hd was positioned for higher-quality audio, exposing two named choices for speech generation with different latency/quality positions.	OpenAI introduced both models at DevDay as new text-to-speech API models, with pricing starting at \$0.015 per 1,000 input characters.
15 Feb 2024	Sora research preview	OpenAI introduced Sora as a text-to-video model capable of generating videos up to a minute long while maintaining visual quality and adherence to a user's prompt, and described the work as teaching AI to understand and simulate the physical world in motion, with the goal of training models for problems that require real-world interaction.	The dedicated Sora research preview granted access to external red teamers and to a number of visual artists, designers, and filmmakers for risk assessment and feedback from creative professionals before broad release.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
13 May 2024	GPT-4o	• The “Hello GPT-4o” launch defined the “o” as “omni”: one model trained end-to-end across text, vision, and audio, able to accept combinations of text, audio, image, and video and produce combinations of text, audio, and image. • OpenAI positioned GPT-4o around more natural human-computer interaction, near-real-time spoken responses, stronger vision and audio understanding, improved non-English text performance, and text and coding performance comparable to GPT-4 Turbo at greater speed and lower API cost. • The launch demonstrated those capabilities through interview preparation, real-time translation, meetings, visual assistance with Be My Eyes, and a proof of concept for customer service.	• OpenAI’s official Spring Update livestream introduced GPT-4o and demonstrated the model live across text, vision, and audio. • The companion “Hello GPT-4o” page introduced GPT-4o, our new flagship model, placed a “Try on ChatGPT” call to action beside the launch, and linked viewers back to the live demos. • A second OpenAI product announcement called GPT-4o its newest flagship model and described its rollout through ChatGPT. • During the launch, Sam Altman promoted GPT-4o as “our new model” and “our best model ever”, described it as smart, fast, and natively multimodal, explained how users would receive it through ChatGPT, and compared its API price and speed directly with GPT-4 Turbo.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
18 Jul 2024	GPT-4o mini	OpenAI positioned GPT-4o mini for low-cost, low-latency applications that chain or parallelize multiple model calls, pass large amounts of context such as full codebases or conversation histories, or require fast real-time text interaction such as chatbots for customer support. The launch emphasized stronger textual and multimodal intelligence than earlier small models and presented affordability itself as the reason the model could expand the range of applications built with AI.	OpenAI gave the model a dedicated launch, “GPT-4o mini: advancing cost-efficient intelligence”, and called it its most cost-efficient small model. GPT-4o mini was available through the Assistants, Chat Completions, and Batch APIs at \$0.15 per million input tokens and \$0.60 per million output tokens, and rolled out to ChatGPT Free, Plus, and Team users in place of GPT-3.5.
12 Sep 2024	o1-preview / o1-mini	• OpenAI presented o1-preview as a model for complex reasoning in science, coding, mathematics, and similar fields, with examples including research on cell sequencing, mathematical work in quantum optics, and multi-step developer workflows. • OpenAI positioned o1-mini as the faster and cheaper option, particularly effective at coding and intended for applications that need reasoning without broad world knowledge. • OpenAI priced o1-mini 80% below o1-preview at launch.	OpenAI introduced o1-preview and o1-mini as the first models in a new reasoning series. Sam Altman’s launch communication named both models directly and tied their immediate availability to ChatGPT Plus and Team users.
1 Oct 2024	gpt-4o-realtime-preview / gpt-4o-audio-preview	• OpenAI positioned gpt-4o-realtime-preview for low-latency speech-to-speech interaction. • gpt-4o-audio-preview was positioned for audio input and output through Chat Completions.	OpenAI’s Realtime API public-beta launch introduced the realtime model to paid developers, demonstrated natural spoken interaction, and announced the companion audio-preview model for the Chat Completions API.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
5 Dec 2024	OpenAI o1 / o1 pro mode	OpenAI called o1 its smartest model and described o1 pro mode as a version that uses more compute and thinks longer to provide more reliable and comprehensive answers to the hardest problems. The launch highlighted difficult mathematics, science, coding, data science, programming, and analysis of case law.	OpenAI introduced the full OpenAI o1 model and o1 pro mode through the ChatGPT Pro release . The \$200/month plan included unlimited access to o1, together with o1-mini and GPT-4o, and included o1 pro mode as the option that uses more compute.
9 Dec 2024	Sora Turbo	OpenAI positioned Sora Turbo as a substantially faster version of the Sora video model for generating videos from text and images, with storyboard, remix, blend, loop, recut, and other creation tools in a dedicated video product.	OpenAI's standalone " Sora is here " launch made Sora Turbo available at sora.com to ChatGPT Plus and Pro users and differentiated the plans through generation volume, resolution, duration, and priority access.
31 Jan 2025	o3-mini / o3-mini-high	- OpenAI positioned o3-mini around fast STEM reasoning with particular strength in science, mathematics, and coding while retaining the lower cost and latency associated with o1-mini. - OpenAI exposed settings with different levels of reasoning effort, including o3-mini-high, for users who wanted a higher-intelligence answer at the cost of more time.	OpenAI gave o3-mini a dedicated launch as the newest and most cost-efficient model in its reasoning series. Plus, Team, and Pro users received o3-mini in ChatGPT, Free users received access through the "Reason" control, selected developers received it through the Chat Completions, Assistants, and Batch APIs, and paid plans exposed o3-mini-high as the higher-intelligence option.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
27 Feb 2025	GPT-4.5	• OpenAI distinguished GPT-4.5 from the line of reasoning models by presenting it as a general-purpose model strengthened through larger-scale unsupervised learning rather than a model that first reasons before answering. • It promoted broader world knowledge, stronger pattern recognition and an ability to make connections, better alignment with user intent, more natural conversation, greater nuance and emotional intelligence, fewer hallucinations, and stronger aesthetic intuition and creativity. • The launch material highlighted writing, programming, practical problem solving, communication, learning, coaching, brainstorming, design, multi-step coding workflows, and complex task automation.	• OpenAI's "Introducing GPT-4.5" page described it as a research preview of its strongest GPT model and its largest and best model for chat yet. • The launch stated that GPT-4.5 was available first to Pro users worldwide, offered a "Try in ChatGPT" call to action, explained how to select the model, and set out a staged rollout to Plus and Team and then Enterprise and Edu. • Sam Altman's launch communication described GPT-4.5 as a very large and expensive model, said OpenAI had wanted to launch it to Plus and Pro together, and explained that limited GPU capacity required Pro to receive it first.
20 Mar 2025	GPT-4o Transcribe / GPT-4o Mini Transcribe / GPT-4o Mini TTS	OpenAI introduced two speech-to-text models with improved transcription accuracy and language recognition relative to Whisper, together with a text-to-speech model designed for more steerable spoken output.	OpenAI launched the three next-generation audio models together and exposed them through the transcription and speech APIs as distinct named developer choices.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
25 Mar 2025	GPT-4o image generation	OpenAI described its new image generator as natively built into GPT-4o, emphasizing accurate text rendering, detailed instruction following, multi-turn refinement, use of chat context and uploaded images, and photorealistic or stylized generation.	The dedicated “ Introducing 4o Image Generation ” release placed a “Try in ChatGPT” call to action beside the launch and rolled the model out as ChatGPT’s default image generator across Free and paid plans, with API access announced to follow.
14 Apr 2025	GPT-4.1 / GPT-4.1 mini / GPT-4.1 nano	• OpenAI positioned GPT-4.1 as the strongest member of the family for coding and precise instruction following with support for up to one million tokens of context. • GPT-4.1 mini was positioned as a lower-latency and lower-cost model that retained strong performance. • GPT-4.1 nano was introduced as OpenAI’s fastest and cheapest model for tasks where speed and price mattered most.	OpenAI launched GPT-4.1, GPT-4.1 mini, and GPT-4.1 nano together in the API as a new model series centered on coding, instruction following, and long-context work.
16 Apr 2025	o3 / o4-mini / o4-mini-high	• OpenAI positioned o3 as its most powerful reasoning model, emphasizing coding, mathematics, science, visual perception, and complex questions that require multifaceted analysis when the answer is not immediately obvious. • It positioned o4-mini as a smaller model optimized for fast, cost-efficient reasoning, with particular strength in mathematics, coding, and visual tasks and substantially higher usage limits than o3 for high-volume, high-throughput work.	OpenAI gave the models a dedicated launch page with a replay of the launch livestream and a “Try on ChatGPT” call to action. The release placed o3, o4-mini, and o4-mini-high as distinct choices in the ChatGPT model selector for Plus, Pro, and Team users, while Free users could try o4-mini through the “Think” control.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
23 Apr 2025	<code>gpt-image-1</code>	OpenAI exposed <code>gpt-image-1</code> , the same underlying image-generation model used in ChatGPT, as its own named API model for high-quality generation and editing with text, image inputs, transparent backgrounds, and controllable quality.	OpenAI's image-generation API launch brought the same model powering ChatGPT image generation to developers globally as <code>gpt-image-1</code> through the Images API, with separate prices for text input, image input, and image output tokens.
7 Aug 2025	GPT-5 / GPT-5 Thinking / GPT-5 Pro / GPT-5 mini / GPT-5 nano	• OpenAI described GPT-5 as a unified system with a smart, efficient model for most questions and a GPT-5 Thinking model for harder problems that require deeper reasoning. • GPT-5 Pro was positioned for the most challenging and complex tasks, spending longer reasoning to produce OpenAI's highest-quality and most comprehensive answers. • The same-day developer release offered <code>gpt-5</code>, <code>gpt-5-mini</code>, and <code>gpt-5-nano</code> as three sizes so that developers could trade off performance, cost, and latency. • <code>gpt-5-mini</code> was the faster, lower-cost option for well-defined, high-volume work, while <code>gpt-5-nano</code> was the smallest, fastest, and cheapest version for tasks such as summarization and classification.	• OpenAI held an official "Introducing GPT-5" event in which Sam Altman and other staff introduced and demonstrated the model. • The launch page called GPT-5 OpenAI's smartest, fastest, and most useful model yet and placed a "Try on ChatGPT" call to action beside that positioning. • Plus subscribers received more GPT-5 usage, paid users could select GPT-5 Thinking, and Pro subscribers received GPT-5 Pro. • OpenAI published the same-day developer launch, "Introducing GPT-5 for developers", and the "GPT-5 and the new era of work" article for organizations using ChatGPT for work. • Sam Altman's launch-day live thread said GPT-5 was rolling out to Free, Plus, Pro, and Team users, highlighted availability to Free users, stated that Plus users received much higher rate limits, and identified GPT-5 Pro as the model for Pro subscribers.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
28 Aug 2025	gpt-realtime	OpenAI positioned gpt-realtime as its production speech-to-speech model for voice agents, with stronger instruction following, tool calling, natural speech, and developer control over how the model sounds and responds.	OpenAI launched gpt-realtime together with general availability of the Realtime API, presenting the named model as the production successor to the preview realtime line.
15 Sep 2025	GPT-5-Codex	OpenAI positioned GPT-5-Codex for both quick interactive coding and long, complex tasks, including building projects from scratch, adding features and tests, debugging, large-scale refactors, and code review. It described the model as more steerable and better able to follow repository instructions.	OpenAI released GPT-5-Codex as a version of GPT-5 optimized for agentic software engineering in Codex and made it available everywhere Codex was used: terminal, IDE, web, GitHub, and the ChatGPT iOS app, with Codex included in Plus, Pro, Business, Edu, and Enterprise plans. API-key access followed on 23 September through the Responses API at the same price as GPT-5.
30 Sep 2025	Sora 2 / Sora 2 Pro	• OpenAI positioned Sora 2 as a new video-and-audio generation model with stronger physical realism, control, synchronized dialogue and sound effects, and reusable character or cameo workflows. • Sora 2 Pro received the higher-quality position for harder shots.	• OpenAI's dedicated Sora 2 launch introduced a new Sora iOS app powered by Sora 2. • Sora 2 was initially free, while ChatGPT Pro users received access to the experimental higher-quality Sora 2 Pro on sora.com. • OpenAI stated that Sora 2 API access was planned.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
12 Nov 2025	GPT-5.1 Instant / GPT-5.1 Thinking	• OpenAI described GPT-5.1 Instant, its most-used model, as warmer, smarter, more conversational, and better at following instructions, while adding adaptive reasoning so that it could decide when a harder question warranted thinking before answering. • GPT-5.1 Thinking was positioned as the advanced reasoning model that became faster on simple tasks, more persistent on difficult tasks, better at adapting its thinking time to problem complexity, and clearer and less jargon-heavy in its responses.	OpenAI launched “GPT-5.1: A smarter, more conversational ChatGPT”. The rollout began with paid users and included a “Try in ChatGPT” call to action.
19 Nov 2025	GPT-5.1-Codex-Max	OpenAI positioned GPT-5.1-Codex-Max for long-running, detailed software work, with greater intelligence and token efficiency, native compaction across multiple context windows, large refactors spanning entire projects, deep debugging, and agent loops lasting for hours. OpenAI described it as the first model it had trained to sustain coherent work across multiple context windows in a single task.	OpenAI introduced GPT-5.1-Codex-Max as its new frontier agentic coding model and made it available in Codex through the CLI, IDE extension, cloud, and code review for Plus, Pro, Business, Edu, and Enterprise plans. API access was announced as forthcoming.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
11 Dec 2025	GPT-5.2 Instant / GPT-5.2 Thinking / GPT-5.2 Pro	• OpenAI marketed GPT-5.2 Instant as a fast, capable workhorse for everyday work and learning, highlighting information-seeking, how-to guidance, technical writing, and translation. • GPT-5.2 Thinking was positioned for deeper work and more complex tasks, including coding, summarization of long documents, questions about uploaded files, step-by-step mathematics and logic, planning, and decision support. • GPT-5.2 Pro was presented as OpenAI's smartest and most trustworthy option for difficult questions where a higher-quality answer justified a longer wait.	OpenAI's "Introducing GPT-5.2" launch presented GPT-5.2 as its most capable model series for professional knowledge work. The three variants began rolling out first to paid ChatGPT plans.
16 Dec 2025	GPT Image 1.5 / new ChatGPT Images	OpenAI positioned GPT Image 1.5 as its flagship image-generation and editing model, emphasizing more precise edits, stronger likeness and detail preservation, better text rendering, and faster generation.	OpenAI's dedicated new ChatGPT Images launch introduced the hosted image experience and the API model <code>gpt-image-1.5</code> together, explicitly comparing the new generation with GPT Image 1.
18 Dec 2025	GPT-5.2-Codex	OpenAI positioned GPT-5.2-Codex around complex, real-world software engineering and defensive cybersecurity, including long-horizon work through context compaction, large code changes such as refactors and migrations, improved Windows performance, more reliable tool use and factuality, and substantially stronger cybersecurity capabilities.	OpenAI launched GPT-5.2-Codex as its most advanced agentic coding model yet and made it available throughout Codex for paid ChatGPT users. API access was announced for the following weeks.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
5 Feb 2026	GPT-5.3-Codex	OpenAI said GPT-5.3-Codex combined GPT-5.2-Codex's frontier coding performance with GPT-5.2's reasoning and professional knowledge, ran 25% faster, and targeted long-running work involving research, tool use, and complex execution. OpenAI said the release moved Codex beyond an agent that writes and reviews code toward one that could perform almost any work developers and professionals do on a computer.	OpenAI's "Introducing GPT-5.3-Codex" launch presented it as the company's most capable agentic coding model to date and made it available on paid ChatGPT plans everywhere Codex was used: the app, CLI, IDE extension, and web. API access was announced as forthcoming.
3 Mar 2026	GPT-5.3 Instant	OpenAI marketed more accurate answers, better contextualization in web search, fewer unnecessary refusals and caveats, more natural conversational flow, and stronger writing.	OpenAI gave GPT-5.3 Instant a dedicated "GPT-5.3 Instant: Smoother, more useful everyday conversations" launch as an update to ChatGPT's most-used model and released it to all ChatGPT users as the everyday conversational model in the GPT-5.3 generation and to developers through the API as gpt-5.3-chat-latest.
5 Mar 2026	GPT-5.4 Thinking / GPT-5.4 Pro	- GPT-5.4 Thinking was positioned for professional reasoning and work across spreadsheets, presentations, documents, coding, tool use, computer use, and deep web research. - GPT-5.4 Pro was presented as the option for maximum performance on complex tasks.	OpenAI "introduced GPT-5.4" as its most capable and efficient frontier model for professional work and released GPT-5.4 in ChatGPT as GPT-5.4 Thinking, in the API as gpt-5.4, and in Codex. GPT-5.4 Pro launched in ChatGPT and the API, with Pro and Enterprise access in ChatGPT.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
17 Mar 2026	GPT-5.4 mini / GPT-5.4 nano	• GPT-5.4 mini was positioned as a faster model with strong coding, reasoning, multimodal understanding, and tool use that approached the larger GPT-5.4 model on several evaluations. • GPT-5.4 nano was presented as the smallest and cheapest GPT-5.4 variant for classification, data extraction, ranking, and simpler coding tasks assigned to subagents where speed and cost mattered most.	• OpenAI launched GPT-5.4 mini and GPT-5.4 nano as its most capable small models yet. • GPT-5.4 mini was available in the API, Codex, and ChatGPT, while GPT-5.4 nano was available only through the API. • Mini cost \$0.75/\$4.50 per million input/output tokens and nano \$0.20/\$1.25.
21 Apr 2026	ChatGPT Images 2.0 / gpt-image-2	OpenAI introduced a new image generation and editing model for ChatGPT with improved instruction following, editing precision, visual consistency, text rendering, and creative control, while exposing the same generation to developers as gpt-image-2.	OpenAI's dedicated "ChatGPT Images 2.0" launch named Images 2.0 as the new image-generation model in ChatGPT and made it available on all ChatGPT plans. OpenAI's API catalog exposed gpt-image-2 with the dated snapshot gpt-image-2-2026-04-21.
23 Apr 2026	GPT-5.5	The launch emphasized agentic coding, knowledge work, computer use, online research, data analysis, document and spreadsheet creation, tool use, and early scientific research.	• OpenAI's "Introducing GPT-5.5" launch called it the company's smartest and most intuitive model yet and described it as "a new class of intelligence for real work". • OpenAI began rolling GPT-5.5 out to Plus, Pro, Business, and Enterprise users in ChatGPT and Codex, with GPT-5.5 Pro for Pro, Business, and Enterprise users in ChatGPT. • API availability was announced as coming very soon.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
5 May 2026	GPT-5.5 Instant	OpenAI promoted stronger factuality, more concise answers, better everyday task performance, stronger image and STEM handling, better judgment about when and how to use web search, and more effective use of past chats, files, and connected context.	OpenAI gave GPT-5.5 Instant a dedicated “GPT-5.5 Instant: smarter, clearer, and more personalized” launch. GPT-5.5 Instant began rolling out to all ChatGPT users as the replacement for GPT-5.3 Instant as the default model and to developers through the API as chat-latest.
7 May 2026	GPT-Realtime-2 / GPT-Realtime-Translate / GPT-Realtime-Whisper	• OpenAI positioned GPT-Realtime-2 for voice agents that need stronger reasoning. • GPT-Realtime-Translate was positioned for low-latency speech-to-speech translation across many languages. • GPT-Realtime-Whisper was positioned for streaming transcription.	• OpenAI launched the three models together in “Advancing voice intelligence with new models in the API” and made all three available in the Realtime API. • GPT-Realtime-2 was priced at \$32 per million audio input tokens and \$64 per million audio output tokens. • GPT-Realtime-Translate was priced at \$0.034 per minute. • GPT-Realtime-Whisper was priced at \$0.017 per minute.
26 Jun 2026 · later GA	GPT-5.6 Sol / GPT-5.6 Terra / GPT-5.6 Luna	• OpenAI positioned GPT-5.6 Sol as the flagship model, GPT-5.6 Terra as a balanced model for everyday work, and GPT-5.6 Luna as the fast and affordable option. • At general availability, OpenAI described Sol as setting a new standard for intelligence and efficiency, Terra as a lower-cost model with performance competitive with GPT-5.5, and Luna as the fastest and most affordable tier.	OpenAI first previewed GPT-5.6 Sol and later presented the full family on the GPT-5.6 general-availability launch, using the three model names to carry distinct product positions within the generation.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
6 Jul 2026	GPT-Realtime-2.1 / GPT-Realtime-2.1-mini	• OpenAI positioned GPT-Realtime-2.1 as an upgraded realtime model with stronger speech recognition in noise, interruption handling and configurable reasoning. • 2.1-mini preserved the realtime interface at lower cost and higher speed.	OpenAI announced both named models through its developer launch channel , making the new generation available through the Realtime API.
8 Jul 2026	GPT-Live / GPT-Live-1 family	OpenAI introduced GPT-Live as a new generation of voice models designed for more natural, emotionally expressive and context-aware spoken interaction, with GPT-Live-1 and a mini variant assigned to different groups of ChatGPT plans.	The dedicated “ Introducing GPT-Live ” launch stated that tapping Voice would give ChatGPT users an experience powered by GPT-Live. OpenAI stated that GPT-Live-1 would become the default Voice model for Go, Plus, and Pro users, while GPT-Live-1 mini would become the default for Free users.
3 Sep 2026	GPT-6 Astra	OpenAI positioned GPT-6 Astra for the hardest end-to-end work across computer use, browsing, software engineering, cybersecurity, science, and professional tasks.	OpenAI’s GPT-6 Astra launch introduced it as “a new generation of intelligence” and called it “the world’s most intelligent and aligned model”. OpenAI began rolling Astra out to a limited set of organizations and said it would become available over the following days to all ChatGPT Plus, Pro, Business, and Enterprise users, through the OpenAI API, Microsoft Azure, and AWS Bedrock.
8 Sep 2026	ChatGPT Images 2.5 / GPT-Image-2.5 Flare / GPT-Image-2.5 Sunburst	• OpenAI positioned Images 2.5 as a faster and more precise generation of image models. • Flare served the fast/default high-throughput position, while Sunburst was the premium variant for maximum precision and quality.	OpenAI’s dedicated ChatGPT Images 2.5 launch rolled Images 2.5 out to ChatGPT, ChatGPT Work, and Codex users across all tiers and simultaneously made GPT-Image-2.5 Flare and GPT-Image-2.5 Sunburst available through the API.

Table 2.1-1c. OpenAI Model Launches and Provider Positioning (continued) · See Section 2.1.1

Launch date	OpenAI model or model family	How the provider positioned the model	How the provider launched, promoted, made available, and priced the model
10 Sep 2026	GPT-Live-1 API	OpenAI positioned GPT-Live-1 as a full-duplex voice model for production voice agents, with natural turn-taking, interruption handling, expressive speech and controls over the model's conversational behavior.	OpenAI gave GPT-Live-1 a dedicated API launch, made it available in the API at \$0.05 per minute for the front-end voice layer, and presented it for developers building voice-enabled apps and business workflows.

Limits and Metering for Model Access

Anthropic

Table 2.1-5a. Anthropic Claude Limits and Metering for Model Access · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Claude Pro · 7–19 Sep 2023	Claude 2	Anthropic launched Pro with 5× more usage than Free. Anthropic Help stated that relatively short conversations received at least 100 messages every eight hours , often more depending on message length, conversation length, attachments, and Claude capacity.	Anthropic warned at 20 messages remaining and reset the allowance every eight hours while the paid subscription continued.
claude.ai Free · 22 Nov 2023 state	Claude 2.1	Anthropic Help stated that Free had a daily message limit whose quantity varied with demand. Anthropic had announced the day before that Claude 2.1 now powered claude.ai for Free and Pro.	Free and Pro used different capacity rules. Anthropic notified the user at the daily limit and the Free quota reset each morning.
Claude Pro · 9 Apr 2024 state	Claude 3 family	Anthropic Help stated at least 45 messages every five hours , varying with message/conversation length and attachments, with a warning at 7 messages remaining .	The shared Claude allowance moved from the earlier eight-hour regime to a five-hour reset while the Pro subscription and selectable Claude family continued.

Table 2.1-5a. Anthropic Claude Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Claude 4 / Claude Code · 22 May–6 Jun 2025	Opus 4 → Sonnet 4	Anthropic stated that Claude and Claude Code shared plan usage limits. Anthropic Help documented Max 5× automatically switching from Opus 4 to Sonnet 4 after Opus reached 20% of the rate limit and Max 20× switching at 50% .	Opus 4 could stop serving requests while work continued on Sonnet 4 under the same subscription. Users could override the automatic switch with /model.
Claude subscriptions / Claude Code · effective 28 Aug 2025	Overall weekly + Opus 4 weekly	Anthropic retained the five-hour session limit and added weekly limits. Anthropic Pro Help stated that weekly / model / feature caps could apply, while Anthropic’s customer notice specified a new overall weekly limit and a weekly limit for Claude Opus 4 .	Users faced both the existing shorter session limit and weekly limits, including a limit specific to Opus 4. Eligible users could wait for reset or buy additional usage after exhaustion.
Claude · by 8 Oct 2025	“Opus only” + “all other models”	Anthropic Help already instructed users to check weekly reset state for “Opus only” and “all other models” by 8 October 2025.	Later Dec 2025–1 Jun 2026 Usage UI captures still showed “All models” + “Sonnet only” for some users, and the Help article and the rendered Usage UI did not change to one common set of model categories at the same time.
Claude / Claude Code · 24 Nov 2025	Opus 4.5 / Sonnet	Anthropic explicitly removed the cap that had applied only to Opus and increased overall limits. An update message in Anthropic’s repository stated that Sonnet now had its own limit.	Opus 4.5 no longer carried the cap that applied only to Opus, while Sonnet received its own cap alongside the overall Claude limit.
Claude Usage UI · Dec 2025–1 Jun 2026	“All models” + “Sonnet only”	Repeated captures in Anthropic’s repository showed a weekly meter for all models and another weekly meter for Sonnet, each with reset metadata.	The limit across all models could be reached even while the Sonnet meter still showed remaining capacity.

Table 2.1-5a. Anthropic Claude Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Claude / Claude Code · 16 Apr–28 May 2026	Opus 4.7 / Opus 4.8 effort	Anthropic documented that tokenizer and effort choices changed how quickly the same allowance was consumed . For Opus 4.8, lower effort explicitly used rate limits more slowly while higher effort spent more tokens.	Lower effort made the same usage allowance last longer, while higher effort spent more tokens. Anthropic increased Claude Code rate limits to accommodate use at higher effort.
Claude / Claude Code · 24 Jul 2026	Opus 5	Opus 5 entered the continuing allowance that applied to Opus, with effort controls that could conserve tokens . Anthropic stated that Fast mode in Claude Code used usage credits .	Users could trade effort against how quickly the subscription allowance was consumed, while supported use of Fast mode drew on paid credits instead of creating a new subscription pool.
Claude paid plans · 9 Jun–19 Jul 2026	Fable 5	Fable launched with its own included usage . On redeployment, Max/premium Fable use was limited to up to 50% of the same weekly plan allowance during the promotional arrangement.	Fable could consume no more than 50% of the weekly plan allowance, leaving the remainder available for other Claude models.
Claude paid plans · after 19 Jul 2026 · Fable 5.1 from 1 Sep	Fable 5 · Fable 5.1 after its launch	After the Fable 5 promotion ended, Pro / standard seats use Fable through usage credits from the first request, while Max / premium seats keep Fable included but capped at 50% of the same weekly plan allowance . Those rules apply to Fable 5.1 only from its 1 September launch .	When the Fable limit is reached, Max/premium users can switch to another Claude model or continue with credits. Pro/standard users use credits or upgrade.
Claude · Sep 2026	“Opus only” and “all other models”	Current Pro and Max documentation retains a weekly limit across all models while Usage settings exposes distinct reset states for “Opus only” and “all other models” .	Usage settings show separate weekly reset states for “Opus only” and “all other models”. On eligible paid plans, the same Usage page reports use of usage credits.

Google Gemini Apps

Table 2.1-5b. Google Gemini Apps Limits and Metering for Model Access

[See Section 2.1.3](#)

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Gemini Advanced · May–Jun 2024	Gemini 1.5 Pro	A Gemini Advanced user reported a usage limit whose effective amount varied with use and conversation context.	The higher model could become unavailable while the Gemini account and service remained active.
Google AI Pro · Jun 2025	Gemini 2.5 Pro	A Google Product Expert recorded the 2.5 Pro allowance increasing from 50 to 100 queries/day.	Gemini 2.5 Pro had a daily limit inside a continuing paid plan.
Google AI Ultra · Aug 2025	Gemini 2.5 Pro Deep Think	Google described Deep Think as having a fixed daily prompt allowance separate from ordinary Gemini use.	Deep Think could become unavailable when its own limit was reached while ordinary Gemini remained available.
Gemini Apps · 7 Sep 2025 state	Gemini 2.5 Pro and Deep Think	Google Help listed 2.5 Pro at 5 prompts/day without a paid plan, 100/day on AI Pro, and 500/day on Ultra. Ultra Deep Think had its own 10 prompts/day allowance.	Gemini 2.5 Pro and Deep Think, which used more compute, each had its own daily limit inside the same consumer service.
Gemini Apps · 19 Dec 2025 state	Thinking + Pro shared pool	Google consumer Help exposed one combined “Thinking & Pro” row at 100 prompts/day for AI Pro and 500/day for Ultra. Google’s later official Community Update stated that the reasoning models had previously drawn from a shared pool.	Thinking and Pro remained distinct choices but consumed the same capacity for the reasoning models, with use of either reducing what remained for the other.
Gemini Apps · 15 Jan 2026 split state	Thinking and Pro	Google Help replaced the combined row with independent Pro: 100/day on AI Pro, 500/day on Ultra and Thinking: 300/day on AI Pro, 1,500/day on Ultra. Google’s official Community Update stated that Thinking usage no longer affected Pro limits.	The former shared pool split into two independently metered reasoning choices inside the same Gemini service.

Table 2.1-5b. Google Gemini Apps Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Gemini Apps · 10 Feb 2026 onward	Pro, Thinking, and Fast	Consumer Help continued to list daily headline counts (Pro 30/100/500 and Thinking 90/300/1,500 across AI Plus / Pro / Ultra) while stating that “Limits are distributed throughout the day”. If Thinking or Pro reached capacity “for a given period”, the same chat could continue with Fast until reset.	A limit with a shorter reset period could make Pro or Thinking unavailable without ending the chat or subscription.
Gemini Apps · 18–19 Feb 2026	Lyria 3 music generation	Google launched Lyria 3 in the Gemini app. Google Help listed 30-second track limits of 10/day Basic, 20/day AI Plus, 50/day AI Pro, and 100/day AI Ultra.	Capacity for music generation had its own plan-specific limits while the same Gemini account and other model access continued.
Gemini Apps · 19 Feb 2026 state	Nano Banana and Nano Banana Pro	Google Help listed different daily limits for the two image models: Nano Banana 100 images/day on Basic and 1,000/day on Plus, Pro, and Ultra, and Nano Banana Pro 3/day on Basic, 50/day on Plus, 100/day on Pro, and 1,000/day on Ultra.	The two image models received different daily allowances by plan inside the continuing Gemini app.
Gemini Apps · 19 Feb 2026 state	Veo 3.1 Fast and Veo 3.1	Google Help listed no Basic video generation, Plus up to 2 videos/day with Veo 3.1 Fast, Pro up to 3/day with Veo 3.1 Fast, and Ultra up to 5/day with Veo 3.1.	The selected video model and plan tier jointly determined usable generation capacity without ending the surrounding Gemini service.
Gemini Apps · effective 17 May 2026	Pro / Thinking / Flash-Lite	Google moved Gemini Apps to compute-based usage limits based on prompt complexity, models/features, and chat length, and the limits refreshed every five hours until the weekly limit was reached. Google still listed the prior daily limits in Help on 19 May and the new regime by 20 May .	The selected model and feature then determined how quickly the allowance was used within a system that refreshed every five hours and imposed a weekly limit. After reaching the limit on a higher model, users could continue with Flash-Lite, upgrade, or wait for refresh.

Table 2.1-5b. Google Gemini Apps Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Gemini Apps Workspace · after 25 May, by 15 Jun 2026	Pro, Thinking, and Fast	Business / Enterprise Help documented Pro access for a tier changing from 100 prompts per day on 25 May to 25 prompts every four hours by 15 June , while higher Pro tiers and Thinking remained on daily limits and Fast remained generally available.	The limit of 25 prompts every four hours applied to the affected Workspace Pro tier, not the consumer Google AI Pro plan.

Google Flow

Table 2.1-5c. Google Flow Limits and Metering for Model Access · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Google Flow · 20 May 2025 launch	Flow / Whisk shared AI Credits · Fast (Ve0 2) · early Ve0 3	Flow launch copy listed AI Pro at 100 generations/month and Ultra with higher limits / early Ve0 3. Same-day Google One Help exposed the underlying shared pool as Pro 1,000 AI Credits/month and Ultra 12,500/month , with “Fast (Ve0 2)” costing 10 credits/generation .	The plan determined the size of the pool, while the selected model determined how many credits were used. After included credits were exhausted, users could buy top-up credits where the plan supported them, while the Flow product continued.
Google Flow · 18 Aug 2025	Shared Flow / Whisk AI Credits	Google doubled Ultra’s monthly pool from 12,500 to 25,000 AI Credits while keeping the same shared credit system and different credit costs by model.	Ultra users received more generations under the same subscription and credit system. The credits required per generation still depended on the model.

Table 2.1-5c. Google Flow Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Google Flow · 25 Feb 2026 state	Veo 2 Fast / Quality · Veo 3.1 Fast / Quality	• Flow Help listed a one-time grant of 100 credits plus 50/day for users without a paid plan, AI Pro at 1,000/month, and Ultra at 25,000/month. • Veo 2 Fast costs 10, Veo 2 Quality 100, Veo 3.1 Fast 20 non-Ultra / 10 Ultra, and Veo 3.1 Quality 100 credits per generation.	A generation was blocked when the account lacked enough credits. Users could wait for reset, change plan, or buy top-up credits where their plan supported them while Flow remained available.
Google Flow · 6 Apr 2026 state	Veo 3.1 Lite / Fast / Quality	The existing Pro 1,000 / Ultra 25,000 pool persisted, but Veo 3.1 Lite entered at 10 credits non-Ultra / 5 Ultra, below Fast at 20/10 and Quality at 100.	A new named variant entered the same shared pool but consumed fewer credits per generation, changing how much model use the existing allowance bought.
Google Flow · 8 May 2026 transitional state	Veo 3.1 Lite / Fast options marked lower priority	Flow Help listed both Lite (lower priority) = Included for Ultra and Fast (lower priority) = Included for Ultra , while stating that the Fast option marked lower priority would be removed on 9 May.	Which variants were included changed before the shared credit pool itself changed.
Google Flow · 19 May 2026 state	Veo 3.1 Lite / Fast / Quality	Access without a paid plan was then 50 credits/day without the earlier one-time 100-credit grant. Pro remained 1,000/month and Ultra 25,000/month. The included Lite option marked lower priority remained, while the equivalent Fast option was gone.	The initial credits and the variants included in the pool changed while the shared credit system continued.

Table 2.1-5c. Google Flow Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Google Flow · 22 May 2026 state	Veo 3.1 Lite / Fast / Quality and Gemini Omni Flash	Google changed the plan pools to 50/day for all users, +200/month on Plus, +1,000/month on Pro, +10,000/month on the \$100 Ultra plan, and +25,000/month on the \$200 Ultra plan. Lite costs 10 non-Ultra / 5 Ultra , Fast 20/10 , Quality 100 . Omni Flash enters at 15/20/25/30 credits for 4/6/8/10-second 720p generations and 40 for editing.	Google changed the size of each plan's pool, which models drew from it, and the credits required per generation independently. Insufficient balance blocked the request and eligible subscribers could buy overage credits.
Google Flow · by 11 Sep 2026	Gemini Omni Flash 720p / 360p · Veo 3.1 family	The same plan pools persist, but Omni Flash draws are remapped to 7/10/12/15 credits for 4/6/8/10-second 720p generations and prices of 4, 5, 6, and 7 credits for 4-, 6-, 8-, and 10-second 360p generations. Editing remains 40, while Veo 3.1 Lite, Fast, and Quality remain 10 credits (5 on Ultra), 20 (10 on Ultra), and 100.	Google changes how fast one named model consumes the same shared subscription pool without changing the subscription or pool size.

Google Antigravity

Table 2.1-5d. Google Antigravity Limits and Metering for Model Access · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Antigravity launch · 18 Nov 2025	Gemini 3 Pro, Claude Sonnet 4.5, GPT-OSS	Google launched Antigravity with model rate limits that refreshed every five hours and consumed quota according to the work performed by the agent rather than treating each request as one unit.	The five-hour meter could stop further model use until refresh while Antigravity continued to expose multiple selectable models.

Table 2.1-5d. Google Antigravity Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Antigravity · 17 Dec 2025	Gemini Pro pool, Gemini Flash pool, non-Gemini pool	Gemini 3 Flash became selectable on 17 December . Google later confirmed that before May, Gemini Pro and Gemini Flash had independent rate limits , while Claude and other non-Gemini models shared a third quota.	Adding Flash created two independently exhaustible Gemini pools inside the same multi-model service. Capacity in one pool did not imply capacity in another.
Antigravity · 10 Mar 2026	Models subject to five-hour and weekly limits	Google staff stated that weekly limits had been established for all models on the affected plans other than Ultra, in addition to the existing five-hour limit .	A model remained subject to the five-hour limit and could become unavailable once the weekly limit was reached.
Antigravity · 13 Mar 2026	Models subject to baseline quota	Google staff announced that built-in AI credits could be used directly in Antigravity and that Pro users could top up beyond baseline quota .	After baseline quota was exhausted, Pro users could continue Antigravity use by drawing AI credits rather than waiting for the next reset of the included quota.
Antigravity Ultra · 27–28 Mar 2026	Ultra model usage	Google's plans.md changed from “ No weekly rate limit ” on 27 March to “ Highest weekly rate limits ” on 28 March while retaining the five-hour refresh wording.	Ultra became subject to a weekly limit in addition to the five-hour limit while the plan and product continued.
Antigravity · 19 May 2026	Gemini Pro + Gemini Flash → shared Gemini pool	Google's change notice combined the previously independent Gemini Flash and Pro limits into one shared Gemini rate limit , drawn down by relative API pricing. Non-Gemini models remained on their own fixed limit, and AI credits were removed from the base Google AI plan entitlement and could be used only after included quota was exhausted.	Gemini Flash and Pro began competing for the same included capacity. Exhausting the shared Gemini pool did not exhaust non-Gemini capacity, and purchased/promotional credits could extend use after baseline exhaustion.

Table 2.1-5d. Google Antigravity Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
Antigravity · Sep 2026	“Gemini Models” and “Claude and GPT models”	Current UI and documentation expose separate bars showing the remaining five-hour and weekly quota for the two listed groups.	When the included quota for a particular model or group is exhausted, users can wait for refresh or, where enabled, draw purchased AI credits. Baseline use resumes automatically after refresh.

OpenAI

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT Plus · 14–16 Mar 2023	GPT-4	OpenAI Help stated that GPT-4 was initially capped at 100 messages every four hours and that OpenAI would adjust the cap with demand and system performance.	GPT-4 could exhaust its own allowance while the continuing Plus subscription and GPT-3.5 access remained available.
ChatGPT Plus · 17 Mar 2023	GPT-4	The ChatGPT model picker displayed 50 GPT-4 messages every four hours and warned that OpenAI expected a lower cap as it adjusted for demand.	The GPT-4 allowance was halved inside the same Plus subscription without ending access to the surrounding service.
ChatGPT Plus · 17 Mar 2023 onward	GPT-4 → GPT-3.5	The provider UI then changed to 25 GPT-4 messages every three hours . Threads on OpenAI’s developer forum recorded the change at the time, and OpenAI Help later recorded the same 25/3h state.	After the GPT-4 allowance was exhausted, the account and access to GPT-3.5 continued.
ChatGPT Plus · 19 Jul 2023	GPT-4	OpenAI announced that the GPT-4 limit was doubling to 50 messages every three hours .	The 50-message / three-hour ceiling restricted GPT-4 access without ending the Plus subscription.

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT Plus · 5 Jan 2024	GPT-4	OpenAI reduced the GPT-4 allowance from the earlier 50/3h state to 40 messages every three hours , effective 5 January 2024.	The GPT-4 cap changed while GPT-4 remained available inside the continuing Plus plan.
ChatGPT Free · 13 May 2024	GPT-4o → GPT-3.5	GPT-4o was available to Free users subject to a message limit for GPT-4o.	After the GPT-4o limit was reached, ChatGPT automatically switched the conversation to GPT-3.5 so that use could continue.
ChatGPT Plus · 13 May 2024	GPT-4o and GPT-4	Launch-day OpenAI Help listed Plus limits of 80 GPT-4o messages every three hours and 40 GPT-4 messages every three hours .	GPT-4o and GPT-4 had separate rolling counters inside the same Plus subscription. Exhausting one did not end the plan or the other model allowance.
ChatGPT Free · 18 Jul 2024	GPT-4o → GPT-4o mini	OpenAI replaced GPT-3.5 with GPT-4o mini in ChatGPT. OpenAI later expressly identified GPT-4o mini as the model used after Free users reached the GPT-4o cap.	After GPT-4o exhaustion, Free users now fell back to GPT-4o mini instead of GPT-3.5. The GPT-4o cap itself stayed in place.
ChatGPT Plus / Team · 12 Sep 2024	o1-preview and o1-mini	At launch, o1-preview had 30 messages/week and o1-mini 50/week, with separate allowances for the two models.	Exhausting one reasoning model did not exhaust the other model or end ChatGPT access.
ChatGPT Plus / Team · 17 Sep 2024	o1-preview and o1-mini	Five days after launch, OpenAI changed o1-preview to 50/week and o1-mini to 50/day.	o1-preview moved from 30 to 50 messages per week, while o1-mini moved from 50 per week to 50 per day. Both remained selectable in the same paid plans.

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT Pro · 5 Dec 2024	o1, o1-mini, GPT-4o, and o1 pro mode	OpenAI launched Pro at \$200/month with nearly unlimited or unlimited access to o1, o1-mini, and GPT-4o subject to guardrails, and added o1 pro mode as a distinct selectable option that used more compute without publishing a numeric allowance specific to o1 pro mode.	Pro changed usable capacity for the named model group while the same ChatGPT account and surrounding product continued.
Sora / ChatGPT Plus and Pro · 9 Dec 2024	Sora Turbo	Plus included up to 50 480p videos per month or fewer 720p generations. Pro included 10× more Sora usage together with higher resolutions and longer durations.	The same named video model carried a materially different monthly generation allowance and output ceiling by subscription tier while the surrounding ChatGPT plan continued.
ChatGPT Plus / Team · 31 Jan 2025	o3-mini	OpenAI replaced o1-mini with o3-mini and tripled the daily allowance on paid plans from 50/day to 150/day. Pro received unlimited o3-mini and o3-mini-high subject to abuse guardrails.	Plus and Team users moved from the 50/day o1-mini allowance to 150/day on o3-mini, while Pro users received unlimited o3-mini and o3-mini-high under the continuing ChatGPT plans.
ChatGPT Plus · Mar–Apr 2025	GPT-4.5	Users captured ChatGPT UI counters showing how many GPT-4.5 responses remained and a weekly reset date.	The UI stated that responses would switch to another model after the GPT-4.5 allowance was exhausted until reset.
ChatGPT Plus / Team / Enterprise · 16–22 Apr 2025	o3, o4-mini, o4-mini-high	OpenAI Help recorded the limits at launch as o3 50/week, o4-mini 150/day, and o4-mini-high 50/day .	The three selectable reasoning models carried separate counters inside the same plans.

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT Plus / Team / Enterprise · by 25 Apr 2025	o3, o4-mini, o4-mini-high	OpenAI Help documented the limits doubling to 100/week, 300/day, and 100/day respectively.	Once a model's limit was reached, that model became unavailable until reset while other ChatGPT models remained available.
ChatGPT Free · 14 May 2025	GPT-4o → GPT-4.1 mini	OpenAI retained the GPT-4o usage limit but changed the model used after the GPT-4o limit was reached.	GPT-4.1 mini replaced GPT-4o mini as the fallback after GPT-4o exhaustion without changing the capped model itself.
ChatGPT Free / Plus / Team / Pro · 8 Aug 2025	GPT-5 base, GPT-5 Thinking, and mini fallback	OpenAI Help listed the allowances by plan. • Free: 10 GPT-5 messages / 5h , then mini plus 1 manual Thinking message/day • Plus: 80 GPT-5 messages / 3h , then mini • Plus / Team manual GPT-5 Thinking 200/week • Automatic routed Thinking did not consume the manual weekly counter	When the base model's allowance was exhausted, use continued on mini, while manually selected Thinking could exhaust independently without disabling automatic reasoning.
ChatGPT Plus / Team · 12 Aug 2025	GPT-5 base and GPT-5 Thinking	OpenAI Help documented Plus base GPT-5 temporarily increasing from 80 to 160 messages / 3h while manual GPT-5 Thinking remained 200/week .	The provider enlarged the base allowance while the allowance for manually selected Thinking remained 200/week.
ChatGPT Plus / Team · 12–13 Aug 2025	GPT-5 Thinking → GPT-5 Thinking mini	OpenAI announced 3,000 manual GPT-5 Thinking messages/week plus additional capacity on GPT-5 Thinking mini after the 3,000-message weekly limit. OpenAI Help recorded the 3,000/week state by 13 August while Plus base remained 160/3h.	Manually selected Thinking gained a much larger weekly allowance and, after the weekly allowance was exhausted, continued on GPT-5 Thinking mini. The counter for the base model remained separate.

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT paid plans · 12–16 Nov 2025	GPT-5.1 Instant and GPT-5.1 Thinking	OpenAI Help listed Plus at 160 Instant messages / 3h with mini fallback and manually selected GPT-5.1 Thinking at 3,000/week . Automatic Instant → Thinking routing did not consume the manual weekly counter.	The base model and manually selected Thinking had separate allowances. Exhausting manual Thinking did not disable automatic reasoning.
ChatGPT Plus / Business · 5–18 Mar 2026	GPT-5.4 Thinking → GPT-5.4 mini	OpenAI Help listed manually selected GPT-5.4 Thinking at 3,000 messages/week and stated that automatic GPT-5.3 Instant → GPT-5.4 Thinking routing did not count toward that manual allowance. Later release notes routed paid users to GPT-5.4 mini when the Thinking limit was reached.	Manually selected Thinking could exhaust independently while automatic reasoning remained available and paid users could continue on GPT-5.4 mini.
ChatGPT · 6 Jul 2026	GPT-5.5 Instant / Auto → GPT-5.5 Instant Mini	The GPT-5.5 Instant / Auto quota remained in place.	OpenAI changed the fallback used after the GPT-5.5 Instant / Auto quota was reached from a GPT-5.3 mini generation to GPT-5.5 Instant Mini.
ChatGPT Plus / Go · Sep 2026	GPT-5.5 Instant / Auto and manually selected GPT-5.5 Thinking	OpenAI Help recorded Plus / Go users receiving up to 160 GPT-5.5 messages every three hours before mini fallback and listed a separate allowance for manually selected GPT-5.5 Thinking. Automatic reasoning did not consume that manual allowance.	The base model and manually selected Thinking each have their own allowance in the same account. When the base allowance is exhausted, ChatGPT falls back to a mini model while the Thinking allowance remains available.

Table 2.1-5e. OpenAI ChatGPT Limits and Metering for Model Access (continued) · See Section 2.1.3

Service and date	Models or allowance scope	How access is metered	What happens when the limit is reached or changed
ChatGPT Business Standard / Premium · Sep 2026	GPT-6 Pro + GPT-5.6 Sol Pro	Current OpenAI Help lists Business Standard at 15 messages/month and Business Premium at 50 messages/week as one included allowance shared across the two named Pro models.	Switching between GPT-6 Pro and GPT-5.6 Sol Pro does not create additional included capacity because both draw from the same plan allowance.
ChatGPT Work / Codex vs ordinary Chat · Sep 2026	GPT-6 Astra / GPT-5.6 Sol / Terra / Luna in Work/Codex · GPT-6 Pro in Chat	OpenAI documents Work/Codex five-hour and weekly allowances whose draw depends on model and plan, while ordinary Chat GPT-6 Pro uses its own Chat allowance.	Exhausting Work/Codex capacity does not exhaust the GPT-6 Pro allowance in ordinary Chat, and exhausting Chat does not consume the Work/Codex pool.
ChatGPT Pro \$200 · Sep 2026	GPT-6 Pro and GPT-5.6 Sol Pro	• GPT-6 Pro has 200 messages/week • GPT-5.6 Sol Pro has its own 170/day allowance • Both are subject to a combined 200/day ceiling (OpenAI Help).	GPT-6 Pro and GPT-5.6 Sol Pro each retain their own caps while both count toward the same 200/day ceiling. GPT-6 Pro weekly exhaustion falls back to GPT-5.6 Thinking at medium effort.
ChatGPT Pro \$100 · Sep 2026	GPT-6 Pro + GPT-5.6 Sol Pro	The same two named Pro models instead draw from one shared 50-message weekly allowance on this plan.	Switching between the two models does not create more capacity after the shared pool is exhausted.
ChatGPT Business · Sep 2026	GPT-5.5 Thinking, GPT-5.5 Pro, and legacy models	Business documentation exposes distinct model counters, including 3,000/week for GPT-5.5 Thinking and 15/month for GPT-5.5 Pro, while legacy GPT-5.4 and GPT-5.2 retain separate counters.	Exhausting one model counter does not exhaust the separate counters for the other listed models.

Evaluation Records by Model and Version

Anthropic

Table 2.1-7a. Anthropic Evaluation Records by Model and Version · See Section 2.1.4

Anthropic model / version	Evaluation record
Claude 1.3	• Jul 2023 · Anthropic reports evaluation results for Claude 1.3 in its July 2023 model card, including Elo and capability results by task. • 9 Apr 2024 · A later persuasiveness study reports a Claude 1.3 measurement alongside the other model generations.
Claude 2	• Jul 2023 · Anthropic reports evaluation results for Claude 2 in its July 2023 model card, including Elo and capability results by task. • 9 Apr 2024 · A later persuasiveness study reports a Claude 2 measurement alongside the other model generations.
Claude Instant 1.1	Jul 2023 · Anthropic reports evaluation results for Claude Instant 1.1 in its July 2023 model card, including Elo and capability results by task.
Claude Instant 1.2	• 9 Aug 2023 · Anthropic reports Claude Instant 1.2 at 58.7% on the Codex evaluation versus 52.8% for Claude Instant 1.1 and at 86.7% on GSM8K versus 80.9% for Instant 1.1. • 9 Apr 2024 · A later persuasiveness study reports a Claude Instant 1.2 measurement alongside the other model generations.
Claude 2.1	• 21 Nov 2023 · Anthropic reports factuality and comprehension results for Claude 2.1 against Claude 2.0. • 9 Apr 2024 · The later persuasiveness study reports Claude 2.1 as an individual model entry.
Claude 3 Opus / Sonnet / Haiku	• 4 Mar 2024 · Opus, Sonnet, and Haiku each receive benchmark results and are compared with external models. • 9 Apr 2024 · The later persuasiveness study reports individual measurements for Claude 3 Haiku, Sonnet, and Opus.
Claude 3.5 Sonnet	21 Jun 2024 · Anthropic reports Claude 3.5 Sonnet solving 64% of problems on its internal evaluation of agentic coding versus 38% for Claude 3 Opus and reports GPQA, MMLU, HumanEval, vision, refusal, and long-context results for Claude 3.5 Sonnet against prior Claude models.
upgraded Claude 3.5 Sonnet	22 Oct 2024 · Anthropic reports 49.0% on SWE-bench Verified, up from 33.4% for the original Claude 3.5 Sonnet, together with separate TAU-bench gains.
Claude 3.5 Haiku	22 Oct 2024 · Anthropic reports Claude 3.5 Haiku at 40.6% on SWE-bench Verified and states that it surpasses Claude 3 Opus on many intelligence benchmarks.
Claude 3.7 Sonnet	24 Feb 2025 · Anthropic reports Claude 3.7 Sonnet on SWE-bench Verified and TAU-bench against predecessor and frontier models. On the 489 SWE-bench Verified tasks executable in Anthropic's internal infrastructure, the model scores 63.7% without the scaffold using more compute and 70.3% with it.
Claude Opus 4 / Claude Sonnet 4	22 May 2025 · Opus and Sonnet receive individual evaluation results, with thinking settings identified where relevant, and are compared directly with predecessor and competitor models.

Table 2.1-7a. Anthropic Evaluation Records by Model and Version (continued) · See Section 2.1.4

Anthropic model / version	Evaluation record
Claude Opus 4.1	5 Aug 2025 · Claude Opus 4.1's dated model ID receives its own coding and benchmark results, compared with predecessor and peer models.
Claude Sonnet 4.5	29 Sep 2025 · Anthropic reports separate results for software engineering, computer use, reasoning, and expert evaluations.
Claude Haiku 4.5	15 Oct 2025 · Anthropic reports separate results for Haiku 4.5 alongside Sonnet 4.5 and external frontier models.
Claude Opus 4.5	24 Nov 2025 · Anthropic reports that Opus 4.5 at medium effort matches Sonnet 4.5's best SWE-bench Verified score while using 76% fewer output tokens. At highest effort, Opus 4.5 exceeds Sonnet 4.5 by 4.3 percentage points while using 48% fewer output tokens.
Claude Opus 4.6	5 Feb 2026 · Anthropic's evaluation table compares Opus 4.6 with prior Claude and external frontier models across agentic coding, computer use, tool use, search, finance, reasoning, and long context. On the 1M-token 8-needle MRCR v2 test, Opus 4.6 scores 76% versus 18.5% for Sonnet 4.5.
Claude Sonnet 4.6	17 Feb 2026 · Anthropic reports benchmark comparisons for Sonnet 4.6 and preferences reported directly by early users: in Claude Code testing, users preferred Sonnet 4.6 over Sonnet 4.5 about 70% of the time and over Opus 4.5 59% of the time.
Claude Opus 4.7	16 Apr 2026 · Anthropic compares Opus 4.7 directly with Opus 4.6 across its benchmark and customer evaluations. CursorBench reports 70% for Opus 4.7 versus 58% for Opus 4.6, and another 93-task coding benchmark reports a 13% resolution lift over Opus 4.6.
Claude Opus 4.8	28 May 2026 · Anthropic reports results for Claude Opus 4.8 on coding, agentic tasks, reasoning, and knowledge work.
Claude Fable 5	9 Jun 2026 · Anthropic reports Fable 5 as its own entry on the launch benchmark table and on FrontierCode.
Claude Mythos 5	9 Jun 2026 · In blinded head-to-head evaluations in molecular biology, Anthropic scientists preferred Mythos 5 hypotheses about 80% of the time against models in the Opus class.
Claude Sonnet 5	30 Jun 2026 · Anthropic reports benchmark results for Claude Sonnet 5 at different effort levels against older Sonnet and Opus models.
Claude Opus 5	24 Jul 2026 · Anthropic reports Opus 5 against prior and frontier models on Frontier-Bench, CursorBench, ARC-AGI 3, AutomationBench, OSWorld 2.0, and related evaluations. It states that Opus 5 more than doubles Opus 4.8 on Frontier-Bench v0.1 at lower cost per task and scores three times the next best model on ARC-AGI 3.
Claude Fable 5.1	1 Sep 2026 · Anthropic's benchmark table places Fable 5.1 beside Fable 5, Opus 5, GPT-5.6 Sol, and other comparison models, with production safeguards enabled and per-benchmark notes marking safeguard interventions and comparable reruns.

Google / Google DeepMind

Table 2.1-7b. Google and Google DeepMind Evaluation Records by Model and Version

[See Section 2.1.4](#)

Google model / version	Evaluation record
PaLM model scales	5 Apr 2022 · Google reports benchmark results by PaLM model scale and compares PaLM with external systems on language understanding and generation tasks.
MusicLM	Jan 2023 · Google reports human and automated music evaluations showing MusicLM outperforming earlier systems on audio quality and adherence to textual descriptions.
PaLM 2 model sizes / variants	10 May 2023 · Google evaluates the PaLM 2 model family across multilingual language understanding and generation, reasoning, coding, and related benchmark tasks.
Chirp	19 May 2023 · Google Cloud reports speech recognition results for Chirp as a named model, including 98% English accuracy and relative accuracy improvements exceeding 300% for several lower-resource languages.
Gemini Ultra / Pro / Nano 1 / Nano 2	6 Dec 2023 · Google reports separate language, reasoning, math/science, multilingual, multimodal, and other benchmark results for Ultra, Pro, Nano 1, and Nano 2.
Gemini 1.5 Pro (February / May versions)	• 15 Feb 2024 · Google reports Gemini 1.5 Pro as its own benchmark entry against Gemini 1.0 Pro and Ultra at launch. • May 2024 · The later technical report gives distinct results for the February and May 1.5 Pro versions.
Gemini 1.5 Flash	May 2024 · The Gemini 1.5 technical report gives Gemini 1.5 Flash its own evaluation results alongside the dated Gemini 1.5 Pro versions and named external models.
Gemini 2.0 Flash-Lite / Flash / Pro Experimental	• 5 Feb 2025 · Google reports individual benchmark results for Flash-Lite, Flash, and Pro Experimental across general knowledge, coding, reasoning, factuality, multilingual, mathematics, long-context, image, audio, and video benchmarks. • 15 Apr 2025 · The later Flash model card again gives Flash and Flash-Lite distinct benchmark columns across reasoning, coding, grounding, multimodal, and long-context evaluations.
Imagen 4	20 May 2025 · Google DeepMind reports human-preference evaluations on GenAI-Bench for the named Imagen generation against previous Imagen versions and external text-to-image systems.
Gemini 2.5 Pro (experimental / preview / GA)	27 Jun 2025 · The card reports separate results for the experimental, preview, and GA versions of Gemini 2.5 Pro alongside named competitor models.

Table 2.1-7b. Google and Google DeepMind Evaluation Records by Model and Version (continued) ·
See Section 2.1.4

Google model / version	Evaluation record
Gemini 2.5 Deep Think	1 Aug 2025 · Google's model card compares Gemini 2.5 Deep Think directly with Gemini 2.5 Pro, OpenAI o3, and Grok 4 across Humanity's Last Exam, IMO 2025, AIME 2025, and LiveCodeBench v6.
Gemini 2.5 Flash (preview / GA versions)	26 Sep 2025 · Separate benchmark columns distinguish the Flash preview and GA versions alongside Gemini 2.0 Flash and named competitor models.
Gemini 2.5 Flash-Lite (preview / release · thinking / non-thinking)	26 Sep 2025 · Separate columns distinguish dated Flash-Lite previews, release versions, and thinking and non-thinking settings.
Veo 3.1	15 Oct 2025 · Google DeepMind reports human evaluations for the named model across text-to-video and image-to-video generation, including overall preference, prompt alignment, visual quality, alignment between audio and video, physics, and editing performance.
Gemini 3 Pro	18 Nov 2025 · Google attaches benchmark results to Gemini 3 Pro and places it directly beside predecessor and competitor model identities.
Gemini 3 Pro Image	20 Nov 2025 · Google DeepMind's benchmark charts for image editing and text-to-image generation place Gemini 3 Pro Image beside Gemini 2.5 Flash Image, GPT-Image 1, and other systems, reporting separate Elo quality results for each model together with latency comparisons.
Gemini 3 Flash	17 Dec 2025 · Google reports separate results for Gemini 3 Flash alongside Gemini 3 Pro, earlier Gemini models, and named competitors.
Lyria 3	18 Feb 2026 · Google evaluates the named music model across human and automated measures including music quality, vocals, audio fidelity, and prompt adherence.
Gemini 3.1 Pro	19 Feb 2026 · The successor Pro model receives its own evaluation record against its predecessor and peers.
Gemini 3.1 Flash Image	26 Feb 2026 · Google's model cards give separate columns to Gemini 3.1 Flash Image, Gemini 2.5 Flash Image, Gemini 3 Pro Image, GPT-Image 1.5, and other systems across GenAI-Bench text-to-image and editing evaluations, with model-specific Elo results for each column.
Gemini 3.1 Flash-Lite	3 Mar 2026 · Google's benchmark table places Gemini 3.1 Flash-Lite beside Gemini 2.5 Flash, Gemini 2.5 Flash-Lite, GPT-5 mini, Claude Haiku 4.5, and Grok 4.1 Fast, with separate columns for price, output speed, academic reasoning, scientific knowledge, multimodal reasoning, coding, long context, and other evaluations.

Table 2.1-7b. Google and Google DeepMind Evaluation Records by Model and Version (continued) ·
See Section 2.1.4

Google model / version	Evaluation record
Gemini 3.1 Flash Live	26 Mar–15 Apr 2026 · Google DeepMind reports a dedicated evaluation set for the real-time voice model.
Gemini 3.1 Flash TTS	26 Mar–15 Apr 2026 · The same model card reports a dedicated evaluation set for the text-to-speech model.
Veo 3.1 Lite	8 Apr 2026 · Google directly compares Veo 3.1 Lite with Veo 3.1 Fast in human-preference evaluations: Lite records a 54.6% overall T2V win rate across 1,000 prompts and a 47.2% I2V win rate across 646 prompts.
Gemini 3.5 Flash	19 May 2026 · Gemini 3.5 Flash receives its own coding, agentic, reasoning, multimodal, and long-context results against predecessor and competitor models.
Gemini 3.1 Flash Lite Image	30 Jun 2026 · Google reports image-generation and editing Elo results together with latency and cost measures for the named Lite image model, separating its evaluation record from the higher-capability image variants.
Gemini 3.6 Flash	21 Jul 2026 · Separate benchmark results are reported for each named model across coding, agentic tasks, computer use, multimodal reasoning, long context, and other evaluations.
Lyria 3.5	29 Jul 2026 · Google reports human and automated music evaluations using a comparison sample that includes Lyria 2, measuring music quality and aesthetics, vocal quality, audio fidelity, and prompt adherence. The results state that Lyria 3.5 improves significantly on Lyria 2 in audio fidelity and follows lyric prompts more accurately.
Gemini 3.7 Flash	13 Aug 2026 · Gemini 3.7 Flash receives its own coding, agentic, multimodal, long-context, computer use, and reasoning results, as well as results on knowledge work beside its predecessor and named competitors.
Gemini 3.5 Live Translate	26 Aug 2026 · Google DeepMind reports a dedicated translation evaluation suite for Gemini 3.5 Live Translate.
Gemini 3.5 Transcribe	26 Aug 2026 · The same model card reports a dedicated evaluation suite for prerecorded transcription for Gemini 3.5 Transcribe.
Gemini 3.5 Transcribe Live	26 Aug 2026 · The same model card reports a dedicated evaluation suite for real-time transcription for Gemini 3.5 Transcribe Live.
Gemini Omni Flash / Omni 1.1 Flash	27 Aug 2026 · Google DeepMind reports human head-to-head evaluations for the named Omni video lineage across video editing, text-to-video, image-to-video, and reference-to-video tasks.

Table 2.1-7b. Google and Google DeepMind Evaluation Records by Model and Version (continued) · See Section 2.1.4

Google model / version	Evaluation record
Gemini 3.8 Flash	2 Sep 2026 · Gemini 3.8 Flash receives its own benchmark column beside its immediate predecessor and named Claude / GPT models across coding, knowledge work, agentic tasks, multimodal work, long context, computer use, and scientific reasoning.

OpenAI

Table 2.1-7c. OpenAI Evaluation Records by Model and Version · See Section 2.1.4

OpenAI model / version	Evaluation record
DALL·E 2	Apr 2022 · OpenAI reports DALL·E 2 results from human preference evaluations against DALL·E 1: evaluators preferred DALL·E 2 for caption matching in 71.7% of comparisons and for photorealism in 88.8%.
Whisper Tiny / Base / Small / Medium / Large	21 Sep 2022 · OpenAI reports zero-shot speech-recognition results by exact Whisper model size across diverse datasets, keeping the named sizes separate while measuring robustness, transcription, and translation performance.
GPT-4	14 Mar 2023 · Professional and academic exams, machine-learning benchmarks, and other evaluations report GPT-4 against GPT-3.5 and other relevant baselines.
DALL·E 3	• Oct 2023 · OpenAI's technical report compares DALL·E 3 with Midjourney 5.2 and Stable Diffusion XL in human ELO evaluations for prompt following, style, and coherence, and compares DALL·E 3 with DALL·E 2 on DrawBench. • Oct 2023 · The system card reports model-specific safety evaluations.
GPT-4o	13 May 2024 · OpenAI reports GPT-4o results on text, reasoning, coding, multilingual, audio, and vision evaluations and compares the named model directly with GPT-4 Turbo.
o1 / o1-preview / o1-mini	12 Sep–5 Dec 2024 · The launch and system card report separate results for o1-preview, o1-mini, o1, GPT-4o, and GPT-4o mini.
o1 pro mode	5 Dec 2024 · OpenAI reports separate AIME, Codeforces, GPQA Diamond, and stricter 4/4 reliability results for o1 pro mode, o1, and o1-preview.

Table 2.1-7c. OpenAI Evaluation Records by Model and Version (continued) · See Section 2.1.4

OpenAI model / version	Evaluation record
Sora	9 Dec 2024 · OpenAI reports that external red teamers tested Sora from September through December 2024 across more than 15,000 generations, probing disallowed content, adversarial tactics, and the safety mitigations applied to the Sora product.
o3-mini	31 Jan 2025 · OpenAI reports o3-mini by reasoning effort on AIME, GPQA, FrontierMath, Codeforces, and other evaluations. Expert testers preferred o3-mini to o1-mini 56% of the time and reported a 39% reduction in major errors on difficult real-world questions.
GPT-4.5	27 Feb 2025 · Separate model columns carry GPQA, AIME, MMLU, MMU, SWE-Lancer Diamond, SWE-bench Verified, and other results.
GPT-4o Transcribe / GPT-4o Mini Transcribe	20 Mar 2025 · OpenAI reports speech-recognition results for the two named transcription models against baselines from Whisper, including word error rate and comparisons of language recognition.
GPT-4o image generation	25 Mar 2025 · OpenAI evaluates 4o image generation through external manual red teaming, automated red teaming, and offline real-world scenarios. The red team tables report not_unsafe and not_overrefuse for system mitigations alone and for system mitigations plus refusals by the chat model.
GPT-4.1 / GPT-4.1 mini / GPT-4.1 nano	14 Apr 2025 · The tables give GPT-4.1, mini, and nano individual result columns alongside GPT-4o, o1, o3-mini, and GPT-4.5 on instruction-following, coding, and long-context evaluations.
o3 / o4-mini	16 Apr 2025 · OpenAI reports separate results for o3 and o4-mini and, where tested, for high and ordinary reasoning-effort configurations.
o3-pro	10 Jun 2025 · OpenAI reports that expert reviewers consistently preferred o3-pro over o3 in every tested category, especially science, education, programming, business, and writing help. The charts of academic evaluations compare o3-pro with o3 and o1-pro on AIME 2024, GPQA Diamond, Codeforces, and the stricter 4/4 reliability criterion.
GPT-5 / GPT-5 mini / GPT-5 nano	7 Aug 2025 · OpenAI's detailed benchmark tables assign separate columns to GPT-5, GPT-5 mini, GPT-5 nano, o3, o4-mini, and GPT-4.1 family models across intelligence, multimodal, coding, instruction-following, function calling, long-context, and hallucination evaluations.
Sora 2	30 Sep 2025 · OpenAI reports model- and system-level safety evaluations over thousands of adversarial prompts for the named video-and-audio generation model.

Table 2.1-7c. OpenAI Evaluation Records by Model and Version (continued) · See Section 2.1.4

OpenAI model / version	Evaluation record
GPT-5.1 Instant / GPT-5.1 Thinking	• 12–13 Nov 2025 · OpenAI’s developer appendix compares GPT-5.1 at high reasoning with GPT-5 at high reasoning across SWE-bench Verified, GPQA Diamond, AIME 2025, FrontierMath, MMMU, Tau²-bench, and BrowseComp. • 12–13 Nov 2025 · The system card addendum reports baseline safety evaluations for gpt-5.1-instant and gpt-5.1-thinking as distinct model entries against their GPT-5 predecessors.
GPT-5.2 Thinking / GPT-5.2 Pro	11 Dec 2025 · OpenAI’s detailed benchmark appendix uses separate columns for GPT-5.2 Thinking, GPT-5.2 Pro, and GPT-5.1 Thinking across professional work, coding, factuality, long context, vision, tool use, academic, and abstract-reasoning evaluations.
GPT-5.4 / GPT-5.4 Pro	5 Mar 2026 · OpenAI’s evaluation tables place GPT-5.4, GPT-5.4 Pro, GPT-5.3-Codex, GPT-5.2, and GPT-5.2 Pro in separate columns across professional work, software engineering, computer use, tool use, and browsing evaluations.
ChatGPT Images 2.0	21 Apr 2026 · OpenAI reports automated policy and safety evaluations for the named image-generation deployment, including image-specific evaluation categories and mitigation results.
GPT-5.5 / GPT-5.5 Pro	23 Apr 2026 · Separate columns for each named model compare OpenAI variants with their predecessors and named competitor models across coding, professional work, computer use, tool use, reasoning, and other evaluations.
GPT-Realtime-2	7 May 2026 · Zillow reported that prompt optimization lifted GPT-Realtime-2 call success on its hardest adversarial benchmark from 69% to 95%.
GPT-Realtime-Translate	7 May 2026 · BolnaAI reported GPT-Realtime-Translate with 12.5% lower word error rates than the other models it tested across Hindi, Tamil, and Telugu.
GPT-Live-1 / GPT-Live-1 mini	• 8 Jul 2026 · OpenAI reports separate production and synthetic voice-native safety evaluations for GPT-Live-1 and its mini variant against predecessor voice deployments. • 10 Sep 2026 · OpenAI reports capability results for the named model including Full Duplex Bench and evaluations of voice agents, directly comparing GPT-Live-1 with the prior realtime generation.
GPT-6 Astra	3 Sep 2026 · OpenAI reports Astra at 100% on ExploitBench versus 78.5% for GPT-5.6 Sol and 42.4% on ExploitGym versus 30.3% for Sol. Its long-context and abstract-reasoning tables place Astra beside Sol and named Claude and Gemini comparison models.

Table 2.1-7c. OpenAI Evaluation Records by Model and Version (continued) · See Section 2.1.4

OpenAI model / version	Evaluation record
GPT-Image-2.5 Flare / GPT-Image-2.5 Sunburst	8 Sep 2026 · OpenAI’s overall safety-evaluation table gives separate rows for GPT-Image-2.5 Sunburst, GPT-Image-2.5 Flare, and ChatGPT Images 2.0, with separate safe-generation, blocked-generation, and unsafe-presented rates for each deployment.
GPT-4o (Mar 2025) · GPT-5 Thinking · GPT-5.6 Sol/Luna (Aug 2026) · GPT-6 Astra/Sol/Luna	23 Sep 2026 · OpenAI reports an overall score and scores for ten aspects of response behavior for each model in MentalHealthBench , identifying the March 2025 GPT-4o and August 2026 GPT-5.6 versions by month.

API Usage and Accounting by Model or Deployment

Table 2.1-9. API Usage and Accounting by Model or Deployment · See Section 2.1.4

Provider / platform and date	Model-specific usage or accounting record
Vercel AI Gateway · 25 Jun 2025	Vercel AI Gateway observability reports usage, latency, and error metrics for each routed model.
GitHub Copilot · 28 Oct 2025	GitHub Copilot usage metrics retain the model name and report usage patterns by language and model.
GitHub Copilot · 2 Mar 2026	GitHub Copilot model and feature metrics retain the actual model name together with chat mode and feature. <code>totals_by_model_feature</code> and model usage per chat mode break activity out by model and feature rather than only by product total.
GitHub Copilot · 20 Mar 2026	GitHub Copilot’s Auto attribution update resolves activity previously shown under generic “Auto” to the actual model selected and attributes model breakdowns to the resolved model used.
Amazon Bedrock · current Sep 2026	Bedrock model invocation logs store the <code>modelId</code> or inference profile ID used for each invocation together with request ID, operation, caller, input tokens, and output tokens. AWS’s example records <code>anthropic.claude-sonnet-4-20250514-v1:0</code> .
Amazon Bedrock · current Sep 2026	Cost and Usage Reports identify the model through billing usage type, while application inference profiles reference a specific Bedrock model. Input, output, cache-read, and cache-write token usage and cost are recorded by model, token type, service tier, and routing state.
Anthropic · current Sep 2026	Anthropic’s Usage / Analytics APIs and Analytics Cost API group usage and cost by Claude model. Usage rows report token quantities by model, while cost rows record model, request count, amount, and currency.

Table 2.1-9. API Usage and Accounting by Model or Deployment (continued) · See Section 2.1.4

Provider / platform and date	Model-specific usage or accounting record
Anthropic · current Sep 2026	Anthropic's usage audit for deprecated models breaks exported API usage down by API key and model, including dated model IDs where present, and uses historical usage attributed to each model to identify applications still depending on models approaching retirement.
Azure OpenAI / Microsoft Foundry · current Sep 2026	Azure Monitor exposes <code>ModelDeploymentName</code> , <code>ModelName</code> , and <code>ModelVersion</code> as metric dimensions, allowing requests, token quantities, latency, generation rate, status, and errors to be reported by model version.
Baseten · current Sep 2026	Baseten inference metrics carry <code>model_id</code> , <code>model_name</code> , and <code>deployment_id</code> . Response and request volume, errors, latency percentiles, token throughput, queue and cache metrics, concurrency, and resource utilization are attributable to the recorded model and deployment.
Cerebras Inference · current Sep 2026	Cerebras usage and monitoring filters request volume, token consumption, cached and fresh input, monthly cost, and individual request logs by model and project.
Cloudflare AI Gateway · current Sep 2026	Cloudflare AI Gateway logs record model and provider and retain input and output tokens, estimated or custom cost, duration, success, and status with the recorded model and provider.
Cohere · current Sep 2026	Cohere Embed Job records store the model used to embed together with billed units and raw and cached token quantities. Cohere's example pairs <code>embed-multilingual-v2.0</code> with billed input and output tokens.
Cursor · current Sep 2026	Cursor's Admin API and Analytics API store the model used for each detailed usage event together with token counts, model cost, charged cost, daily messages, and users by model.
Databricks Model Serving · current Sep 2026	Databricks Model Serving usage tracking records <code>served_entity_id</code> , which joins to <code>served_entities</code> containing <code>entity_name</code> and <code>entity_version</code> , together with input and output token counts, request ID, requester, status, timestamp, and streaming state.
Databricks Unity Gateway · current Sep 2026	Databricks usage tracking and cost observability store the final <code>destination_model</code> and <code>destination_id</code> together with request and invocation IDs, tokens, latency, requester, routing and status metadata, and billing cost for the destination that ultimately served the request.
DeepInfra · current Sep 2026	DeepInfra billing usage records contain <code>model_name</code> together with units, rate, and cost.
DeepSeek API · current Sep 2026	DeepSeek Responses API and Chat Completions responses identify the model and return input, output, cache, reasoning, and total token usage with the response model.

Table 2.1-9. API Usage and Accounting by Model or Deployment (continued) · See Section 2.1.4

Provider / platform and date	Model-specific usage or accounting record
fal · current Sep 2026	fal Model Usage API records model endpoint <code>endpoint_id</code> , such as <code>fal-ai/flux/dev</code> , and allows billing unit, quantity, unit price, cost, request count, success and error counts, and latency to be filtered and aggregated by model endpoint.
fal · current Sep 2026	fal Billing Events API records both <code>request_id</code> and <code>model_endpoint_id</code> together with output units, unit price, and <code>cost_estimate_nano_usd</code> for each request.
Fireworks AI · current Sep 2026	Fireworks deployment metrics label <code>base_model</code> , full deployment, deployment account, and <code>deployment_id</code> , allowing request rate, errors and status codes, prompt and cache token processing, latency, token distributions, concurrency, and resource utilization to be attributed to a concrete deployment while the base model remains visible.
Fireworks AI · current Sep 2026	Fireworks billing metrics retain <code>base_model_name</code> beside the coarser <code>model_bucket</code> , with usage type, accelerator-seconds, parameter count, prompt and completion tokens, cached tokens, and uncached tokens attached to the model.
Fireworks AI · current Sep 2026	Fireworks rated usage and cost export groups serverless cost rows by <code>model</code> , <code>model_tier</code> , user, or API key and attributes daily rated <code>subtotal_usd</code> to the selected grouping dimension.
Google Gemini API · current Sep 2026	Gemini API GenerateContent responses record <code>modelVersion</code> , the version that actually generated the response, together with token usage by modality.
GroqCloud · current Sep 2026	GroqCloud Prometheus metrics label model and project ID and attribute request and status counts, input and output tokens, queue latency, time to first token, end-to-end latency, and prompt-cache hits and misses to each model within a project.
Groq Compound · current Sep 2026	Groq Compound usage breakdown enumerates each underlying model in <code>usage_breakdown.models[]</code> and allocates queue time, prompt, completion, and total tokens, processing times, and total time to each model used.
Helicone · current Sep 2026	Helicone analytics can use the model returned in the provider response for cost calculation and exposes <code>request_model</code> together with prompt, completion, and total tokens, calculated cost, provider cost, status, timestamp, and model breakdowns.
Hugging Face Inference Providers · current Sep 2026	Hugging Face Inference Providers billing and provider billing integration break user and organization usage down by model and serving provider and reconcile actual provider cost against successful request IDs.
IBM watsonx Model Gateway · current Sep 2026	IBM watsonx Model Gateway usage APIs filter or group tenant and user usage and cost by model.

Table 2.1-9. API Usage and Accounting by Model or Deployment (continued) · See Section 2.1.4

Provider / platform and date	Model-specific usage or accounting record
LiteLLM · current Sep 2026	LiteLLM cost tracking and Prometheus telemetry store requested model, public model group, deployment model_id, provider, and fallback model together with spend, token counts, requests, success and failure counts, latency, and fallback outcomes.
Mistral · current Sep 2026	Mistral organization usage records preserve model breakdowns within chat, completion, audio, OCR, connectors, Libraries, and other services together with the associated usage and cost quantities.
NVIDIA NIM · current Sep 2026	NVIDIA NIM Prometheus metrics report request success and failure, prompt and generation tokens, queue and running state, cache use, time to first token, timing between output tokens, and end-to-end latency for each model.
Oracle OCI Generative AI · current Sep 2026	OCI Generative AI metrics attach monitoring to a unique Generative AI endpoint whose resource stores the model OCID. Calls, input and output tokens, client and server errors, latency, processing time, and success rate are measured for the endpoint.
OpenAI · current Sep 2026	OpenAI Organization Usage API filters or groups usage by model and records request counts plus text, audio, and image input and output quantities and cache-related token quantities. OpenAI's example records gpt-4o-mini-2024-07-18, a pinned snapshot ID.
OpenRouter · current Sep 2026	OpenRouter Activity API retains public model, a more specific dated model_perma_slug, and the serving provider together with prompt, completion, and reasoning tokens, request count, and spend.
Perplexity Agent API · current Sep 2026	Perplexity Agent API responses identify the response model and return input, output, and cache token usage together with input, output, request, and total cost .
Portkey · current Sep 2026	Portkey OpenTelemetry distinguishes gen_ai.request.model from gen_ai.response.model and records the provider together with input, output, and cache token usage, cost, latency, errors, request metadata, routing, and configuration context.
Replicate · current Sep 2026	Replicate prediction records retain both model and the unique version used to create each prediction together with status, creation, start, and completion times, prediction runtime, total runtime, and inference result metadata.
Vercel AI Gateway · current Sep 2026	Vercel reporting records model and provider, while request logs show how traffic was routed and which model completed a fallback request together with request counts, token quantities, cost, and latency.

Table 2.1-9. API Usage and Accounting by Model or Deployment (continued) · See Section 2.1.4

Provider / platform and date	Model-specific usage or accounting record
Vertex AI · current Sep 2026	Vertex AI foundation model observability identifies the queried model, while Provisioned Throughput is allocated to a specific project, region, model, and version. Usage, throughput, latency, and errors are monitored by queried model.
Vertex AI · current Sep 2026	Vertex AI prediction responses and Cloud Monitoring metrics identify the serving deployment through <code>deployed_model_id</code> and attribute prediction and response counts, errors, latency, replicas, memory, network, and related production metrics to the identified deployment.
xAI · current Sep 2026	xAI Usage Explorer filters account usage by model, while request-level cost tracking returns input, output, and total tokens and exact billed <code>cost_in_usd_ticks</code> per request.

Platform and Survey Reporting on Use by Named Model

Table 2.1-10. Platform and Survey Reporting on Use by Named Model · See Section 2.1.4

Date	Platform or dataset	Providers and models tracked	What is tracked or reported by model
Jan 2025 onward	OpenRouter daily rankings dataset	Exact <code>model_perma_slug</code> identities across providers	The public dataset records the top models per day by total token usage , with one (date, <code>model_perma_slug</code>) record per model and provider and app breakdowns available elsewhere in OpenRouter analytics.
10 Mar 2025 · covering the prior year	Poe early 2025 ecosystem trends	Successive GPT, Claude, Gemini, DeepSeek, and other models used by Poe subscribers	Poe reports that new flagship models rapidly cannibalized prior versions in subscriber message traffic, explicitly identifying GPT-4 → GPT-4o and Claude 3 → Claude 3.5 as examples.
13 May 2025	Poe spring 2025 model usage report	GPT-4.1 family, GPT-4o, GPT-4, o1/o3 lines, Gemini 1.5/2.0/2.5 models, Claude 3.5/3.7 models, and others	Poe reports message shares by model or model family and successor migration. GPT-4.1 family reached about 10% and Gemini 2.5 Pro about 5% of the measured message share within the stated Poe population.

Table 2.1-10. Platform and Survey Reporting on Use by Named Model (continued) · See Section 2.1.4

Date	Platform or dataset	Providers and models tracked	What is tracked or reported by model
7 Aug 2025	Vercel AI Gateway model leaderboard	GPT-5, GPT-5 mini, GPT-5 nano, and other models available through the gateway	Vercel describes its model leaderboard as tracking the most popular models over time by token volume across Gateway traffic.
25 Sep 2025 · covering May–Aug 2025	Poe summer 2025 model usage trends	GPT-5, GPT-4o, GPT-4.1, Gemini 2.5 family, Claude models, and others	Poe reports GPT-5 at 27.7% of subscriber messages in the text category, GPT-4o at 19.8%, and GPT-4.1 at 7.3%. The Gemini 2.5 family rose from 3.5% to about 10% over the May–August 2025 measurement period.
28 Jan 2026	GIGAZINE reader survey	Gemini 3 Pro, GPT-5.1, GPT-5, GPT-4, Gemini 2.5 Flash, and other named models	The survey reports use percentages for named models including Gemini 3 Pro 16.8%, GPT-5.1 12.6%, GPT-5 11.1%, GPT-4 10.9%, and Gemini 2.5 Flash 9.5%.
12 May 2026	Vercel AI Gateway Production Index	Individual models used by production teams	Vercel’s Production Index analyzes seven months of AI Gateway traffic from more than 200,000 teams, spanning tens of trillions of tokens across hundreds of models. High-volume workloads route across more than 30 distinct models on average.
14 Jul 2026	Vercel open AI Gateway leaderboards	Individual Models, Labs, Apps, and Providers as separate leaderboard units	Vercel explicitly separates “Models”, “Labs”, “Apps”, and “Providers”. The model leaderboard reports requests, token volume, spend, reach, and other measures by named model.
22 Jul 2026	Cursor Router telemetry	Individual model choices across Cursor. The same production analysis names Fable, Opus 4.8, and GPT-5.6 Sol in routing comparisons	Cursor reports routing hundreds of millions of requests across models and providers, measures satisfaction and keep rate across model launches, and finds that about 60% of developers choose one model as their daily driver.

Table 2.1-10. Platform and Survey Reporting on Use by Named Model (continued) · See Section 2.1.4

Date	Platform or dataset	Providers and models tracked	What is tracked or reported by model
Sep 2026	OpenCode Model Data	Individual coding models and model authors	OpenCode publishes token volume by model, users, completed sessions, cost, cache ratio, retention, and related usage measures.
Sep 2026	OpenRouter model rankings	Individual model variants across providers	OpenRouter reports separate token traffic totals and rankings for exact model variants.
Sep 2026	Vercel AI Gateway model leaderboard	Individual models on a model leaderboard distinct from “Labs” and “Providers” leaderboards	Vercel reports token volume, requests, spend, reach, and preference for a primary model by named model, with labs and inference providers reported on their own leaderboards.

Examples of Paid Decisions Users Attributed to Access to Particular Models

Table 2.1-11. Examples of Paid Decisions Users Attributed to Access to Particular Models · See Section 2.1.5

Provider and model	Paid decision	How access to the model affected the decision
Anthropic · Claude 3 Opus	Subscribe to Claude	A user stated that they subscribed to Claude because they needed access to Claude 3 Opus’s 200K context window.
Anthropic · Claude 3.5 Sonnet	Subscribe to Claude Pro	A user described preferring Sonnet 3.5 for creative writing and coding and stated that this was why they subscribed to Claude Pro.
Anthropic · Claude Opus 4	Upgrade to Max 20x	Two users in the same Claude Code discussion stated that they upgraded to or bought Max 20x for access to Opus 4.
Anthropic · Claude Sonnet 4.5	Subscribe to Claude	A user stated that Sonnet 4.5 was literally the only reason they subscribed to Claude.
Anthropic · Claude Opus 4.5	Upgrade from Pro to Max 20x	A user stated that they upgraded from Pro to Max 20x as Opus 4.5 became the model they used for nearly every task in their marketing business.

Table 2.1-11. Examples of Paid Decisions Users Attributed to Access to Particular Models (continued) ·
See Section 2.1.5

Provider and model	Paid decision	How access to the model affected the decision
Anthropic · Claude Opus 4.6	Upgrade from Pro to Max 5x	A user seeking more access to Opus 4.6 for a large project later stated that they subscribed to Max 5x.
Anthropic · Claude Opus 4.8	Re-subscribe to Max 5x	A first-person account states that Opus 4.8 was the “one reason” the author stopped a planned downgrade and re-subscribed to Max 5x.
Anthropic · Claude Fable 5	Return to Claude on Max 5x, then upgrade to Max 20x	A user stated “I came back to Claude only because of Fable”, began on Max 5x, and then upgraded to Max 20x to continue a major migration using the model.
Anthropic · Claude Opus 5	Upgrade from Pro to Max 5x	A user who had been on Pro stated that after beginning to use Opus 5 they upgraded to Max 5x specifically “for Opus”.
Google · Gemini Ultra 1.0	Subscribe to Google One / Gemini Advanced	A user stated that they subscribed and entered payment details because the membership represented Gemini Ultra access.
Google · Gemini 2.5 Pro	Upgrade to Google AI Pro	A Google AI Pro user on Google’s Gemini forum stated that they had “recently upgraded specifically for the access to Gemini 2.5 Pro”.
Google · Gemini 3 Pro	Subscribe to Google AI Ultra	In the official Gemini CLI project discussion, a user stated that they subscribed to the specific plan precisely because it gave direct Gemini 3 Pro access through Gemini CLI.
Google · Gemini 3.1 Pro Preview	Subscribe to Google AI Pro	A Google AI Developers Forum user stated that hitting Gemini 3.1 Pro Preview’s rate limit led them to subscribe to Google AI Pro.
OpenAI · GPT-4	Join ChatGPT Plus	A March 2023 user stated: “I paid for Plus to get access to GPT4”.
OpenAI · DALL·E 3	Subscribe to ChatGPT Plus	An OpenAI Developer Community user stated that they subscribed and paid for Plus to get access to DALL·E 3.
OpenAI · GPT-4o	Subscribe to ChatGPT Plus	A Plus user stated that they paid for the plan exclusively to use GPT-4o.
OpenAI · GPT-4.1	Extend a ChatGPT Plus subscription	A Plus subscriber whose subscription had been due to end stated that they extended it specifically to have several additional days with GPT-4.1.

Table 2.1-11. Examples of Paid Decisions Users Attributed to Access to Particular Models (continued) ·
See Section 2.1.5

Provider and model	Paid decision	How access to the model affected the decision
OpenAI · GPT-4.5	Sign up for and continue paying for ChatGPT Plus	A paying Plus user stated on OpenAI's community forum that GPT-4.5 was the main reason they signed up and continued paying .
OpenAI · o1	Subscribe to ChatGPT Pro	A user stated that the reason they obtained ChatGPT Pro was hitting o1's message cap .
OpenAI · o1-pro	Upgrade from ChatGPT Plus to Pro	An OpenAI Developer Community user stated that unlimited o1-pro use was the reason they moved from ChatGPT Plus to Pro .
OpenAI · Sora	Upgrade from ChatGPT Plus to Pro	An OpenAI Developer Community user stated that a primary reason for upgrading from Plus to Pro was the ability to use Sora .
OpenAI · o3-mini-high	Subscribe to ChatGPT Pro	A user on OpenAI's Developer Community stated that they subscribed to the \$200 plan for unlimited use of o3-mini-high for coding and had not subscribed for another model.
OpenAI · o3	Subscribe to ChatGPT Pro	An OpenAI Developer Community user stated that they got Pro specifically because of unlimited access to o3 .
OpenAI · GPT-5	Upgrade to ChatGPT Plus	A user described GPT-5 as a "turning point", said they then "went all-in", and upgraded to Plus before using GPT-5 Thinking to plan and implement a long-running application project.
OpenAI · GPT-5 Pro	Continue paying for ChatGPT Pro	A user stated that GPT-5 Pro was part of their critical workflow and that they were willing to continue paying \$200/month for Pro because of GPT-5 Pro's role in the workflow .
OpenAI · GPT-5.2 Pro	Upgrade to ChatGPT Pro	A first-person account states that the author upgraded to Pro because they expected Codex CLI to use GPT-5.2 Pro .
OpenAI · GPT-5.5	Switch from Claude Pro to ChatGPT Plus	A user stated that after GPT-5.5 High solved a coding task that Claude Code had failed, they moved from Claude Pro to ChatGPT Plus the next day.
OpenAI · GPT-5.6 Sol Pro	Upgrade subscription	A user stated that access to GPT-5.6 Sol Pro was a major factor in upgrading their subscription .

Table 2.1-11. Examples of Paid Decisions Users Attributed to Access to Particular Models (continued) ·
[See Section 2.1.5](#)

Provider and model	Paid decision	How access to the model affected the decision
OpenAI · GPT-6 Astra	Upgrade to ChatGPT Pro	A first-person account states that the allowance for Astra on the lower plan was exhausted too quickly and that the author immediately signed up for the Pro upgrade .

Annex B. Differences Between Versions Released Under the Same Public Model Name

Evidence tables supporting §2.2, *AI Models as Non-Renewable Resources* are collected below.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
GPT-4 March 2023 → June 2023 · GPT-3.5 March 2023 → June 2023	Peer-reviewed longitudinal multi-task evaluation across mathematics, knowledge, code generation, application tasks, medical questions, and instruction following	Between the March and June 2023 GPT-3.5 versions, accuracy on questions asking whether numbers were prime rose 49.6% → 76.2%. Between the March and June 2023 GPT-4 versions, accuracy on the same questions fell 84.0% → 51.1%. On the same USMLE multiple-choice questions, the March and June GPT-3.5 versions selected different options on 27.9% of questions, even when overall accuracy changed by only 0.8 percentage points. The March and June GPT-4 versions selected different options on 12.2% of questions, even when overall accuracy changed by only 4.2 percentage points. Some option changes corrected earlier errors while others introduced new errors. The small changes in overall accuracy therefore concealed substantial changes in which options the models selected across individual questions.
gpt-4-0314 → gpt-4-0613	Clinical differential diagnosis on the same 75 case reports and narrative queries	Correct diagnoses fell from 29/75 (38.6%) to 19/75 (25.3%), and the later version produced more cases with no relevant information.
gpt-4-0314 → gpt-4-0613	Peer-reviewed Dutch national English-reading-comprehension exams with repeated stochastic runs	At temperature 0, the March snapshot answered 14 questions incorrectly and the June snapshot 18, with only 9 questions incorrect in both. At least 14 questions therefore changed correct/incorrect status across the two releases. On one question, the March snapshot selected the wrong answer in 49 of 50 repeated runs, while the June snapshot selected the correct answer in 50 of 50.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions compared	Evaluation setting	What changed
gpt-4-0314 → gpt-4-0613	Independent human-preference evaluation in Chatbot Arena	After Arena separated the two previously pooled snapshots, it reported a significant difference in human-preference ratings : gpt-4-0314 at 1201 and gpt-4-0613 at 1152 .
gpt-3.5-turbo-0613 → gpt-3.5-turbo-1106	StableToolBench tool-use benchmark using chain-of-thought and depth-first search agent strategies	On one instruction-level tool-use task family, success rate rose 48.1% → 55.7% under chain-of-thought and fell 73.8% → 67.2% under depth-first search.
gpt-3.5-turbo-0613 → gpt-3.5-turbo-1106 · gpt-4-0613 → gpt-4-1106	Peer-reviewed vascular-surgery board questions, three attempts per question	GPT-3.5 responses became more consistent, 55.5% → 65.6% , and shorter, 164.2 → 105.4 words , while pooled accuracy changed by only 1.4 percentage points . GPT-4 responses became less consistent, 90.1% → 79.7% , and much longer, 120.0 → 282.0 words , while pooled accuracy changed by only 1.4 percentage points .
gpt-4-0314 → gpt-4-0613	GAOKAO-Bench objective and subjective questions from China's national college entrance examination	On subjective questions, Physics fell 56.7% → 47.1% and Chemistry 35.0% → 28.5% , while Biology remained 85.6% , while the overall subjective score changed only 51.9% → 50.8% .
gpt-4-0314 → gpt-4-0613	Programming-question generation from Python, C++, and Java code	Relevance fell 9.85 → 8.46 , while coverage rose 8.57 → 9.22 . The later snapshot began repeating questions earlier, with 1 duplicate versus 0 at 30 generated questions and 2 versus 1 at 50.
gpt-3.5-turbo-0301 → 0613 → 1106	Peer-reviewed code-generation efficiency benchmark	First-attempt correctness rose 42.3% → 46.4% → 49.3% , while execution-time overhead relative to the canonical solution rose 3.18× → 3.22× → 3.40× and total-memory-use overhead rose 6.50× → 6.71× → 7.24× .
gpt-4-1106-preview → gpt-4-0125-preview	Peer-reviewed explainable operations-research benchmark under fixed methods in zero-shot and one-shot settings	Under the same zero-shot benchmark, Standard accuracy fell 75.33% → 68.33% , while OptiGuide accuracy rose 36.00% → 47.33% .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions compared	Evaluation setting	What changed
gpt-4-1106-preview → gpt-4-0125-preview → gpt-4-turbo-2024-04-09	NAACL BiGGen Bench results for eight capability categories, scored by five evaluator models	Instruction Following declined across all three releases under every one of the five evaluator models. Under the GPT-4-1106 evaluator, it fell 4.230 → 4.120 → 4.040 , while Planning rose and then fell 4.157 → 4.243 → 4.029 .
gpt-4-1106-preview → gpt-4-0125-preview	Peer-reviewed assessment of 41 AI-intervention RCTs against CONSORT-AI	Average agreement rose 80.3% → 86.5% under one human gold standard and 78.0% → 81.6% under the other. Median study-level agreement rose 72.7% → 81.8% .
gpt-4-0613 → gpt-4-1106 → gpt-4-0125	Peer-reviewed evaluation of pharmaceutical knowledge using the same domain framework	Organoheterocyclic accuracy rose 86.1% → 97.2% and then fell to 91.7% , while inhibitor accuracy stayed at 86.8% before falling to 81.6% .
gpt-4-1106-preview → gpt-4-0125-preview → gpt-4-turbo-2024-04-09	ICLR LiveBench, objective scoring across Coding, Data Analysis, Instruction Following, Language, Math, and Reasoning	Across the three releases, Coding, Instruction Following, Language, Math, and Reasoning fell from November to January and rose from January to April, while Data Analysis rose and then fell. Coding ended above both earlier values. The six category scores were Coding 44.4 → 41.8 → 49.0 , Data Analysis 51.3 → 54.1 → 51.3 , Instruction Following 69.4 → 57.2 → 60.8 , Language 48.4 → 39.2 → 44.3 , Math 47.6 → 33.4 → 42.7 , and Reasoning 52.0 → 43.2 → 49.5 .
gpt-4-1106-preview → gpt-4-0125-preview → gpt-4-turbo-2024-04-09	NeurIPS Berkeley Function-Calling Leaderboard under common native function-calling mode	Executable parallel-multiple-call accuracy rose 67.5 → 77.5 and then returned to 67.5 , while relevance detection rose 80.42 → 82.92 → 88.75 .
gpt-4-1106-preview → gpt-4-0125-preview	ACL longitudinal graduate-level evaluation with pinned snapshots and repeated dynamic sampling	The January snapshot scored lower on 8 of 10 displayed subjects, while Military rose 7.20 → 7.67 . Aggregate score fell 73.6 → 71.6 .
gpt-4-1106-preview → gpt-4-0125-preview	MultiHop-RAG on one 2,556-query dataset, with and without retrieval	Without retrieval, answer correctness rose 29.2% → 37.7% . With retrieval, it changed only 50.0% → 51.4% , sharply narrowing the version gap.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions compared	Evaluation setting	What changed
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	ICLR SGP-Bench symbolic-graphics semantic and spatial reasoning	Semantics fell 78.7 → 45.9 while Count, Color, Shape, and Reasoning rose, while the overall semantic score increased only 63.3 → 64.8 . On the separate spatial benchmark, 3D fell 78.2 → 76.0 while 3D reconstruction and 2D rose, while the overall spatial score fell by only 0.4 points .
gpt-4o-2024-05-13 → gpt-4o-2024-11-20	ACL IFEval++ instruction-following reliability across nuanced reformulations of the same intent	The headline instruction-following score fell 85.4 → 80.8 , and every reported nuance/reliability column was lower for November, including 81.3 → 77.6 and 78.0 → 73.9 .
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	Peer-reviewed evaluation of product descriptions under the same prompts, with each requirement assessed through checks validated against human judgments	For requirements left unstated in the prompts, the average rate of presenting key product information in an easily scanned form fell by 48.3 percentage points from May to August. The rate of including supplied customer support or warranty information fell by 13.6 percentage points .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	NeurIPS SuperGPQA across 26,529 graduate-level questions, difficulty levels, and 13 disciplines	Agriculture fell and then recovered, 41.44 → 39.59 → 42.27 , while Law peaked in August, 52.90 → 54.12 → 53.81 . Overall rose 39.76 → 41.64 → 44.40 .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	AmericasNLP grammatical generation for four Indigenous languages under one fixed prompt with 10 examples	Under the same prompt and examples, Guaraní accuracy stayed flat and then rose 39.24 → 39.24 → 45.57 , while Maya fell and then partially recovered 57.05 → 54.36 → 55.70 .
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	ICML structured table/chart understanding, with and without the same ReFocus visual-editing framework	On CharXiv, the same ReFocus workflow improved the May snapshot by 8.3 points but reduced the August snapshot by 2.7 points relative to each snapshot's own baseline.
gpt-4o-2024-05-13 → gpt-4o-2024-11-20	ACL LongBench v2 long-context reasoning across difficulty, context length, and prompting mode	The direct-answer invalid-output rate rose 0.2% → 5.6% . Direct-answer Overall fell 50.1 → 46.0 , while Overall with chain-of-thought stayed nearly flat at 51.2 → 51.4 .
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	Pragmatic last-utterance proactivity classification in task-oriented dialogue	Accuracy rose 0.64 → 0.74 , precision 0.67 → 0.88 , and F1 0.61 → 0.68 , while recall remained 0.56 .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions compared	Evaluation setting	What changed
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	CryptoX compositional-reasoning benchmark	On Crypto-MBPP with five encoded words, performance fell 58.53 → 52.10 → 25.88 , while unencoded Crypto-MMLU rose 85.16 → 85.51 → 86.57 across the same snapshots.
GPT-4o May → August snapshots	ICLR MIA-Bench multimodal evaluation pipeline using GPT-4o as the judge	Changing only the GPT-4o judge snapshot reversed the order of two candidate models: LLaVA-1.5-13b moved rank 7 → 6 while ShareGPT4v moved 6 → 7 . All eight candidate models received different total scores under the two judge snapshots.
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	C ³ -Bench controllable code completion requiring functional correctness and instruction adherence	On code-completion tasks with explicit implementation constraints, first-attempt functional success rose 54.3 → 65.8 → 65.9 , while instruction following rose 41.3 → 58.6 → 59.6 .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	Epoch benchmark records reporting task score and average response length	Average OTIS output length rose 735.67 → 835.67 → 990.82 tokens , while the task score remained 0.06 across all three snapshots.
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	PredictaBoard analysis of where model scores are predictable across instances	Across 12 matched assessor configurations , the May snapshot had the larger region meeting an 80% accuracy threshold in 8 configurations and the August snapshot in 4 . The ordering changed again at stricter thresholds, while both snapshots had the same MMLU-Pro accuracy of 0.756 .
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	EMNLP AutoSDT end-to-end scientific data-analysis tasks	Scientific task-goal success rose 7.5% ± 0.5 → 11.4% ± 1.2 , while the rate at which generated programs merely executed successfully changed only 42.2% ± 1.6 → 43.1% ± 2.1 .
GPT-4o May → August → November → chatgpt-4o-latest	Peer-reviewed EvoBench evaluation of a fixed detector family across successive GPT-4o generator versions	A fixed detector became less able to distinguish GPT-4o-generated XSum text as the generator changed: its area under the ROC curve (AUROC) fell 0.90 → 0.87 → 0.72 → 0.70 . Average AUROC across five domains changed 0.80 → 0.77 → 0.73 → 0.74 .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
gpt-4o-2024-05-13 → gpt-4o-2024-11-20	ACL Industry CodeIF: 1,200 multi-constraint code-generation tasks across four languages	The share of tasks satisfying all constraints rose 33.8% → 38.3% overall and 36.8% → 43.4% in Java . All twelve reported difficulty/compliance measures were higher for the November snapshot.
GPT-4o May → November snapshots	Peer-reviewed scientific coding under the same fixed Self-Debug and SciNav agent configurations	Under Self-Debug, Valid Execution Rate fell 71.2 → 67.0 while Success Rate changed 14.7 → 15.0 . Under SciNav, Valid Execution Rate rose 66.0 → 69.9 and Success Rate 16.1 → 18.6 . Self-Debug had the higher Valid Execution Rate with the May snapshot, while SciNav had the higher rate with the November snapshot.
GPT-4o May → November snapshots	ACL paper using GPT-4o as a prompt-based detector of AI-generated text	On paraphrased GPT-4o text, true-positive detection jumped 26.7% → 100% while the false-positive rate rose 0% → 6.7% . On o1-Pro text, true-positive detection rose 3.3% → 56.7% .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	NeurIPS Artificial Hivemind open-ended generation under a fixed high-stochasticity setup	Across the three snapshots, repeated responses became increasingly similar, concentrating more prompts in the highest-similarity bands. The share of real-world prompts whose repeated responses fell in the highest similarity band rose 40% → 44% → 51% , and the share with average similarity of at least 0.8 rose 76% → 81% → 87% .
GPT-4o May → August snapshots	ICLR BFCL function-calling evaluation under both native function calling and prompt extraction	Under prompt extraction, overall function-calling accuracy fell 83.13 → 74.12 and executable-simple accuracy 98.00 → 70.57 . Irrelevance detection rose 77.44 → 89.56 while relevance handling fell 78.05 → 53.66 . Under native function calling, overall accuracy changed much less, 79.55 → 78.87 .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
GPT-4o May → August snapshots	Peer-reviewed ProBench open-ended professional multimodal tasks with anchored pairwise ELO	Against the May baseline, the August snapshot was preferred less often in single-round and multi-round comparisons but more often in multilingual comparisons. It won 44.98% of single-round comparisons, 60.35% of multilingual comparisons, and 45.41% of multi-round comparisons. Average response length fell 491 → 374 , 585 → 480 , and 1563 → 1052 tokens across those three tracks.
gpt-4o-audio-preview-2024-10-01 → gpt-4o-audio-preview-2024-12-17	AHELM standardized audio-language evaluation	Multilingual speech recognition improved from October to December, while PARADE and MUSTARD emotion scores fell. FLEURS word error rate fell 0.771 → 0.456 and Multilingual LibriSpeech word error rate 0.162 → 0.073 , where lower is better. PARADE exact match fell 0.847 → 0.779 and MUSTARD emotion exact match 0.593 → 0.583 .
gpt-4o-audio-preview-2024-10-01 → gpt-4o-audio-preview-2024-12-17	AudioMarathon minute-scale long-context audio benchmark	Music Classification rose 59.5 → 67.6 and Speaker Gender Recognition 69.2 → 73.1 , while Speech Content Reasoning fell 61.4 → 60.2 and Speech Detection 32.5 → 30.8 , while the ten-task mean rose only 47.4 → 48.7 .
claude-3-5-sonnet-20240620 → claude-3-5-sonnet-20241022	ICLR τ -bench airline and retail tasks requiring models to interact with tools, evaluated across repeated trials	First-trial success rose 36.0% → 46.0% on airline tasks and 62.6% → 69.2% on retail tasks. Success across four repeated trials rose 13.9% → 22.5% and 38.7% → 46.2% , respectively.
claude-3-5-sonnet-20240620 → claude-3-5-sonnet-20241022	ICLR LiveBench across six capability categories	Coding rose 60.8 → 67.1 , while Data Analysis fell 56.7 → 52.8 and Math 53.3 → 51.3 , while the Overall score changed only 58.2 → 58.5 .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
Claude 3.5 Sonnet June → October 2024	Peer-reviewed RE-Bench machine-learning research-engineering tasks under the same Modular scaffold and different time budgets	On RE-Bench's common scale, where the provided starting solution is 0 and the task-author reference solution is 1 , the October release is higher at every reported time budget, from approximately 0.18 vs 0.08 at 30 minutes to 0.61 vs 0.38 at 64 hours .
Claude 3.5 Sonnet Old → New (June → October 2024)	NeurIPS evaluation of the task duration at which models achieve 50% success, combining HCAST, RE-Bench, and shorter tasks	The estimated duration of tasks completed with 50% success rose 18.22 → 28.98 minutes on the paper's task suite and 5.91 → 16.88 minutes on SWE-bench Verified.
claude-3-5-sonnet-20240620 → claude-3-5-sonnet-20241022	Findings of EMNLP LiveRepoReflection across six languages, full-file generation, patch editing, and execution feedback	In full-file JavaScript generation, compliance with the required edit format fell 98.3% → 78.3% while second-attempt success rose 45.0% → 53.3% . In patch editing, JavaScript second-attempt success moved in the opposite direction, 45.0% → 31.7% .
Gemini 1.5 Pro 001 → 002	ICLR PPE reward-model / LLM-judge evaluation under one ArenaHard judging prompt	On German prompts, pairwise agreement with human preference rose 63.33% → 70.64% and confidence agreement 48.72% → 76.92% . The update was not uniformly better: on hard prompts, confidence agreement fell 96.88% → 95.83% and Kendall rank correlation 85.26 → 83.16 .
Gemini 1.5 Pro 001 → 002 · Gemini 1.5 Flash 001 → 002	Peer-reviewed evaluation of search-query parser outputs using five evaluation methods and repeated runs	For Gemini 1.5 Pro, correlation with expert scores in one basic rule-parser setup fell 0.905 → 0.677 while run-to-run variability fell 0.569 → 0.062 , making the later release more stable but less aligned with expert judgments in the rule-parser setup. For Flash, few-shot correlation rose 0.244 → 0.793 and a setup using a distinct system prompt rose 0.252 → 0.785 .
Gemini 1.5 Pro 001 → exp-0827 → 002	ICLR LiveBench six-category capability evaluation	Coding rose 32.3 → 40.9 → 48.8 , while Language peaked and then fell 40.4 → 46.1 → 43.3 and Reasoning 37.0 → 50.9 → 47.9 . Overall rose 43.3 → 52.4 → 53.4 .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions compared	Evaluation setting	What changed
Gemini 1.5 Flash exact release sequence including 001, exp-0827, and 002	ICLR LiveBench	Across the three releases, Overall, Instruction Following, and Reasoning rose across the sequence, while Language declined at both transitions and Math fell before rebounding above its initial level. Overall rose 36.6 → 44.1 → 47.6. Instruction Following rose 52.6 → 73.0 → 78.8 and Reasoning 24.9 → 44.5 → 45.7, while Language declined 31.6 → 29.6 → 27.9 and Math changed 32.3 → 28.9 → 47.2.
Gemini 1.5 Pro 001 → exp-0827 · Gemini 1.5 Flash 001 → exp-0827	CVPR VisionArena-Bench on 500 multimodal prompts from real users with fixed anchors	Both Gemini lineages received higher pairwise-preference scores under both fixed anchor models. Against GPT-4-Turbo, Pro changed 45.5 → 66.2 and Flash 23.0 → 60.3. Against Claude-3-Sonnet, Pro changed 75.3 → 87.6 and Flash 58.4 → 83.5. Average output length rose 261 → 329 tokens for Pro and 271 → 367 for Flash.
gemini-2.0-flash-thinking-exp-1219 → gemini-2.0-flash-thinking-exp-01-21	Independent LiveBench official raw judgments on shared question IDs before alias redirection	Across 418 shared LiveBench questions, 01-21 scored higher on 94, lower on 64, and identically on 260. The comparison contains 29 exact pass-to-fail and 24 exact fail-to-pass changes. On connections, 19 questions improved and 4 declined. On typos, 7 improved and 18 declined.
gemini-2.0-flash-exp → gemini-2.0-flash-001	LiveBench scoring of the same 418 questions for both versions	The stable 001 version scored higher on 83 of the 418 shared questions, lower on 53, and identically on 282. Coding-completion score fell 0.600 → 0.540, while the overall mean score rose 0.506072 → 0.526723.
Gemini 2.5 Pro 03-25 → 05-06	NeurIPS ResearchCodeBench on 212 novel ML research-code tasks	On ResearchCodeBench-HARD, first-attempt success rose 25.7% → 33.0%, a 7.3 percentage-point increase. The later preview remained ahead on the full benchmark and on the subset restricted to repositories created after the models' knowledge cutoff.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
Gemini 2.5 Pro 05-06 → 06-05	LiveCodeBench maintained coding evaluation of exact versions under peer-reviewed benchmark methodology	On Medium coding problems, first-attempt success rose 82.3% → 87.2% , while Hard remained 50.2% . Overall first-attempt success changed 71.8% → 73.6% .
Gemini 2.5 Flash-Lite preview 06-17 → stable → preview 09-2025	Fixed set of 20 instruction-adherence tests	From the June preview to stable Flash-Lite, Prime JSON changed from pass to fail. From stable to the September preview, Prime JSON returned to pass, Deduplication and Position Multiply changed from pass to fail, and String Replace with Newlines changed from fail to pass, while the June and September states both scored 13/20 overall .
Gemini 2.5 Flash-Lite stable → preview 09-2025	Conference study comparing model behavior using shared WildChat prompts and held-out verification	Markdown tables and LaTeX equations became more common, and held-out verification found more headings and segmented technical formatting and less continuous narrative presentation. Markdown-table use increased 0.7% → 13.9% and LaTeX-equation use 5.2% → 11.1% .
Gemini 2.5 Flash Native Audio Sep 2025 → Dec 2025	Artificial Analysis Full Duplex Bench using the same weighted Conversational Dynamics measure	Conversational Dynamics rose 30.3% → 44.0% , a +13.7 percentage-point change in the aggregate measure of pause handling, turn taking, interruption handling, and backchannel handling.
Gemini 1.5 Flash 001 → 002	Research preprint evaluating first-order logic with difficulty controlled by the number of conjuncts	Both versions scored 100.0% at two conjuncts. From four to ten conjuncts, the later version scored higher at every reported level, including 82.3% → 92.9% at four conjuncts, 70.7% → 88.7% at five, and 62.3% → 78.6% at ten.
Gemini 1.5 Pro 001 → 002	Research preprint evaluating first-order logic with difficulty controlled by the number of conjuncts	Both versions scored 100.0% at two conjuncts. At higher complexity, accuracy rose 85.9% → 93.6% at five conjuncts, 72.9% → 83.7% at eight, and 70.8% → 78.8% at ten.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
Claude 3.5 Sonnet 20240620 → 20241022	Research preprint evaluating first-order logic with difficulty controlled by the number of conjuncts	The October release scored lower at every reported level from three to nine conjuncts, including 95.1% → 87.8% at four conjuncts and 89.4% → 81.6% at five, while accuracy at ten conjuncts moved in the opposite direction, 73.9% → 75.0% .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	Research preprint evaluating first-order logic with difficulty controlled by the number of conjuncts	The August snapshot was lower than May at every reported level from four to ten conjuncts, and November then recovered much of the loss. Accuracy at five conjuncts changed 86.6% → 80.7% → 84.7% , at eight conjuncts 72.0% → 68.4% → 73.7% , and at ten conjuncts 71.4% → 66.9% → 70.8% .
Gemini 2.5 Pro 03-25 → 05-06 → 06-05	Independent LoCoDiff benchmark reconstructing real source files from their git history across context lengths and programming languages	The shortest-context exact-match rate rose throughout, 62% → 64% → 74% , while the second quartile fell and then remained flat, 44% → 38% → 38% . JavaScript rose throughout, 35.71% → 38.10% → 54.76% , while Rust 20.51% → 25.64% → 17.95% and TypeScript 47.50% → 52.50% → 45.00% rose and then fell.
Claude 3.5 Sonnet 20240620 → 20241022	Independent Aider code-editing benchmark on 133 small Python exercises	Functional completion rose 77.4% → 84.2% , while correct use of the required edit format remained 99.2% → 99.2% .
Claude 3.5 Sonnet 20240620 → 20241022	Independent Aider refactoring benchmark on 89 large methods from large Python classes	Functional correctness rose 64.0% → 92.1% , while correct use of the required edit format rose 76.4% → 91.0% .
Claude 3.5 Sonnet 20240620 → 20241022	Independent Epoch maintained benchmark records reporting task score and average response length	On MATH Level 5, accuracy changed 0.52 → 0.57 while average output length changed 439.62 → 450.57 tokens . On GPQA Diamond, accuracy changed only 0.54 → 0.55 while average output length fell 420.55 → 353.90 tokens . On OTIS Mock AIME, accuracy changed 0.07 → 0.08 while output length rose 555.56 → 593.11 tokens .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
Claude 3.5 Sonnet 20240620 → 20241022	Independent ZeroEval evaluation across five reasoning benchmarks	The October release was higher on all five reported scores, with different gains by task. MATH Level 5 changed 51.87 → 59.36 , CRUX 80.75 → 83.88 , MMLU-Redux 86.00 → 88.91 , ZebraLogic 33.40 → 36.20 , and GSM8K 95.60 → 96.66 .
Gemini 1.5 Pro API-0514 → exp-0801 → exp-0827 → 002	Official maintained BigCodeBench-Hard Elo under a peer-reviewed benchmark methodology	The maintained evaluator's reported uncertainty intervals do not overlap for any successive transition in this four-state sequence. Central coding Elo rises 1118.28 → 1149.13 → 1177.48 and then falls to 1136.44 .
Gemini 1.5 Pro 23 May 2024 → 1 August 2024	Official maintained MMMU-Pro leaderboard under a peer-reviewed benchmark methodology	All three shared MMMU-Pro accuracy measures rose by about 3–4 points: Overall 43.5 → 46.9 , Vision 40.5 → 44.4 , and the original-question subset 46.5 → 49.4 .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	Peer-reviewed Value Portrait psychometric evaluation across ten Schwartz values and five Big Five dimensions	All five Big Five dimensions fell from May to August and all five rose from August to November, while the ten Schwartz values followed several different directions. For example, Conscientiousness changed 3.90 → 3.45 → 3.99 , while Universalism declined 0.47 → 0.38 → 0.26 .
Gemini 1.5 Flash 001 → 002	SIIM 2025 conference abstract on pulmonary-embolism classification under clean input and three forms of text perturbation	With no introduced noise, accuracy rose 0.967 → 0.984 and recall 0.935 → 0.971 . At the highest tested perturbation level, accuracy was higher for the later version under all three perturbations: 0.937 → 0.975 for character deletion, 0.958 → 0.981 for symbol insertion, and 0.932 → 0.975 for word shuffling.
gpt-4o-2024-05-13 → gpt-4o-2024-08-06	Peer-reviewed image-based estimation of six nutrients from 22 photographed Japanese meals, with weighed food records as ground truth and repeated trials	After substituting the August snapshot into the same workflow, the authors report that its ICCs were “ markedly inferior ” to those obtained with the May snapshot. The paper does not publish the August ICC values.

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · See Section 2.2.2

Model versions and providers compared	Evaluation setting	What changed
gpt-4o-2024-08-06 → gpt-4o-2024-11-20	Qwen2.5 technical report from the provider using an in-house Chinese automatic evaluation across six capability dimensions	All six reported dimensions were higher for the November snapshot, but by very different amounts. Instruction Following changed 42.50 → 42.71 , Comprehension 80.11 → 83.04 , Math 61.74 → 66.04 , and Reasoning 56.88 → 62.04 .
gpt-4o-2024-05-13 → 2024-08-06 → 2024-11-20	Research preprint reporting controlled Number Guessing and Math Magic experiments on latent information maintenance and multi-step manipulation	In Number Guessing, the reported sum of P(Yes) across numbers 1–10 rose 0 → 1.085 in August and returned to 0 in November. In Math Magic without chain-of-thought, accuracy declined 2.7% → 2.0% → 0.0% . With chain-of-thought, accuracy rose 17.3% → 20.7% from May to August. The paper reports no numeric November result because the November snapshot consistently failed the task with chain-of-thought.
Gemini 1.5 Flash API-0514 → exp-0827	Official maintained BigCodeBench-Hard Elo under a peer-reviewed benchmark methodology	Central Elo fell 1094.72 → 1063.16 . The maintained intervals moved from [1087.92, 1101.87] to [1054.94, 1070.80] and do not overlap in the reported dataset.
Gemini 1.5 Flash 001 → 002	Independent Epoch maintained benchmark records reporting task score and average response length	On MATH Level 5, score rose 0.25 → 0.62 while average output length rose 445.25 → 581.86 tokens . On GPQA Diamond, score rose 0.40 → 0.47 while output length rose 355.59 → 427.64 tokens . On OTIS Mock AIME, score rose 0.04 → 0.16 while output length fell 1005.71 → 907.44 tokens .
Gemini 1.5 Pro 001 → 002	Independent Epoch maintained benchmark records reporting task score and average response length	On MATH Level 5, score rose 0.41 → 0.70 while average output length rose 445.27 → 571.37 tokens . On GPQA Diamond, score rose 0.46 → 0.57 while output length rose 357.44 → 441.29 tokens . On OTIS Mock AIME, score rose 0.07 → 0.23 while output length changed 811.09 → 826.67 tokens .

Table 2.2-1. Examples of Differences Between Versions Released Under the Same Public Model Name (continued) · [See Section 2.2.2](#)

Model versions and providers compared	Evaluation setting	What changed
gpt-4-0314 → gpt-4-0613	Maintained Just-Eval multi-task evaluation using GPT-4-based scoring	Average response length fell 226.41 → 186.06 tokens. Info-seeking fell 4.91 → 4.83 and Reasoning 4.88 → 4.83 , while Procedure rose 4.96 → 4.98 , while the overall average changed only 4.79 → 4.80 .

Annex C. Published Research Using Retired Models and Models with Published Withdrawal Dates or Timelines

Evidence tables supporting §2.4, *The Consequences of Lost Model Continuity* are collected below.

Published Research Using Retired Models

Table 2.4-7. Published Research Using Retired Models · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-3-5- sonnet- 20240620	BioDiscoveryAgent: An AI Agent for Designing Genetic Perturbation Experiments (ICLR, 2025) BioDiscoveryAgent used Claude 3.5 Sonnet to choose genetic perturbation experiments and revise later choices from earlier results. Across six datasets, the system improved prediction of relevant genetic perturbations by an average of 21% over Bayesian optimization baselines trained for the same task, and by 46% on the harder task involving non-essential genes. It predicted useful combinations of genes more than twice as accurately as a random baseline.	Anthropic retired the June Sonnet snapshot. Repeating these experiments requires the withdrawn version to make the gene selections.
Anthropic · claude-3-5- haiku- 20241022	Measuring Scientific Capabilities of Language Models / SciGym (NeurIPS, 2025) SciGym required models to design experiments, inspect the resulting simulated data, and reconstruct hidden biological systems. In a test of whether models could rewrite biological models without changing their behavior, claude-3-5-haiku-20241022 produced behaviorally equivalent models in 80% of 50 cases, averaging 2.04 iterations. The broader benchmark found performance declined as the underlying biological systems became more complex.	Anthropic retired the tested Haiku snapshot. Repeating the experiment requires the historical model to choose the interventions and infer the mechanisms.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-3- haiku- 20240307	LAB-Bench: Measuring Capabilities of Language Models for Biology Research LAB-Bench tested practical biology work including reading scientific figures, reasoning about experimental protocols, querying scientific databases, and working with biological sequences. On ProtocolQA, claude-3-haiku-20240307 reached 51% selective accuracy at 87% coverage. On other tasks, precision was 24% for questions about scientific figures, 51% for questions about tables, 29% for complex cloning scenarios, and 86% for a sequence task that required selecting primers for specified restriction enzymes.	Anthropic retired the Haiku snapshot. Regenerating the results on the biology tasks requires the historical model to produce its original answers and decisions about whether to answer.
Anthropic · claude- sonnet-4- 20250514	AOT*: Efficient Synthesis Planning via LLM-Empowered AND-OR Tree Search (Findings of ACL, 2026) AOT* combined synthesis routes generated by language models with a structured search procedure for finding routes to target molecules. With Claude Sonnet 4, the method solved 63% of the Pistachio Hard targets at a search budget of 20 and 79% at a budget of 100. claude-sonnet-4-20250514 generated the synthesis routes used in those searches.	The Sonnet 4 snapshot has been retired. Repeating the search with the same generated routes requires the original model, alongside the search algorithm and chemical benchmarks.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude- opus-4-1- 20250805	Benchmarking large language models for cell typing in single-cell RNA-Seq (Briefings in Bioinformatics, 2025) In a benchmark spanning 34 human and mouse single-cell RNA sequencing datasets, large language models outperformed traditional bioinformatics tools for identifying cell types, particularly for finer cell subtypes. Claude Opus 4.1 was among the highest-performing models, with a 79.5% perfect-match rate across the main comparison. The researchers included it in an ensemble of the strongest models, which achieved state-of-the-art performance and became the basis of the DeepCellSeek annotation tool.	Anthropic retired the Opus 4.1 snapshot. Regenerating its annotations and its contribution to the evaluated ensemble requires the historical model.
OpenAI · gpt-3.5- turbo-0613	Can GPT-3.5 generate and code discharge summaries? (Journal of the American Medical Informatics Association, 2024) Using gpt-3.5-turbo-0613, researchers generated 9,606 synthetic discharge summaries for underrepresented ICD-10 labels and added them to the training data for local clinical coding models. On the subset of low-resource labels selected for generation, macro-F1 improved for all three local models. Multi-Res CNN and CAML improved by 26% and 78%, respectively. One ICD-10 code absent from the original training data was correctly predicted after training with the synthetic summaries.	OpenAI retired the June 2023 GPT-3.5 Turbo snapshot. Regenerating the synthetic clinical records used for the augmentation experiment requires the historical model.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · gpt-3.5-turbo-16k-0613	A critical assessment of using ChatGPT for extracting structured data from clinical notes (npj Digital Medicine, 2024) GPT-3.5 Turbo 16K converted large volumes of free-text pathology reports into structured clinical data without task-specific model training. Across more than 1,000 lung-cancer pathology reports, it reached 89% overall accuracy and outperformed two traditional NLP methods. On 191 pediatric osteosarcoma reports, the same June 2023 model classified histology grade with 98.6% accuracy and margin status with 100% accuracy.	OpenAI shut down the 16K 0613 snapshot. Regenerating its pathology extractions requires the original model, even with the reports and prompts preserved.
Anthropic · claude-opus-4-20250514	Rapidly Benchmarking Large Language Models for Diagnosing Comorbid Patients (JMIRx Med, 2025) The diagnostic benchmark used 1,000 MIMIC-IV hospital admissions involving patients with multiple concurrent conditions. Claude Opus 4 identified 95.2% of diagnoses judged inferable by GPT-4.1 from the available patient information. Across the wider comparison, diagnostic hit rates changed significantly with the prompt used, while changes in temperature generally had little effect.	Anthropic retired the Opus 4 snapshot. Regenerating its diagnostic lists and repeating the associated prompt and evaluator comparisons requires the original model.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.0- flash-001	Adaptive diagnostic reasoning framework for pathology with multimodal large language models (Communications Medicine, 2026) RECAP-PATH used Gemini 2.0 Flash to classify breast pathology images and refine diagnostic criteria from its errors. The optimized framework improved the model across binary and multiclass diagnosis tasks. When diagnostic criteria learned with GPT-4o were transferred to Gemini 2.0 Flash, accuracy for distinguishing normal tissue from invasive carcinoma rose from about 71% to 86.8%, showing successful transfer of diagnostic rules learned through a different model to new pathology images.	Google shut down Gemini 2.0 Flash. Repeating these pathology classifications and the cross-model transfer experiment requires the original Gemini API condition.
OpenAI · gpt-4-0125- preview	Using Large Language Models to Assess the Consistency of Randomized Controlled Trials on AI Interventions With CONSORT-AI (JMIR, 2025) Among six models assessing 41 reports of randomized trials against CONSORT-AI, gpt-4-0125-preview achieved the highest average consistency score. It reached 86.5% against the original JAMA Network Open assessments and 81.6% against the researchers' own assessments. GPT-4 variants significantly outperformed the Claude and GPT-3.5 models in the comparison.	The retired January preview supplied the judgments about the trial reports. Repeating its agreement measurements requires new assessments from the same historical version.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · o1-preview-2024-09-12	Systematic benchmarking demonstrates large language models have not reached the diagnostic accuracy of traditional rare-disease decision support tools (European Journal of Human Genetics, 2026) The study evaluated language models as tools for diagnosing rare diseases using 5,213 curated cases drawn from published reports and spanning 378 Mendelian or chromosomal diseases. o1-preview-2024-09-12 received the phenotypic findings from each case and generated the ranked differential diagnoses evaluated by the researchers. It was the strongest language model tested, placing the correct diagnosis first in 23.6% of cases, within the top three in 31.2%, and within the top ten in 36.8%.	OpenAI retired o1-preview from the API on 28 July 2025. Repeating the study requires o1-preview-2024-09-12 to generate ranked diagnoses from the same case descriptions.
OpenAI · o1-mini-2024-09-12	Using Artificial Intelligence to Adjudicate Major Adverse Cardiovascular Events in Clinical Trials (JACC, 2026) The study developed and validated Auto-MACE to adjudicate cardiovascular death, myocardial infarction, and stroke from medical records in PARADISE-MI, a randomized trial of 5,661 patients. o1-mini-2024-09-12 generated the primary event classifications, while a Clinical Longformer model assessed confidence in those results. Across all evaluated events, the classifications generated by o1-mini agreed with the physician clinical events committee in 86% of deaths, 76% of potential myocardial infarctions, and 84% of potential strokes.	OpenAI retired o1-mini from the API on 27 October 2025. Repeating the adjudication requires o1-mini-2024-09-12 to process the same trial records with the study prompts and generate the classifications for cardiovascular death, myocardial infarction, and stroke.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · imagen-4.0-ultra-generate-preview-06-06	A Data-Efficient Deep Learning Paradigm for Crack Segmentation Using Generative AI (Transportation Research Record, 2026) Imagen 4 Ultra generated the 500 synthetic pavement images and pixel-level crack masks used to build the AI500 dataset. Models trained only on the synthetic data generalized to real-world crack images at levels comparable to models trained on manually annotated target datasets. Adding only 10% real-world data brought performance close to the fully supervised baseline. The study estimated \$134.17 to generate the synthetic images and masks with Imagen 4, compared with \$625 for manual annotation.	Google retired the Ultra preview. Repeating the stage in which the synthetic data were generated requires the original generator, not only the saved images or trained segmentation models.
OpenAI · gpt-4-0314	Autonomous chemical research with large language models (Nature, 2023) Coscientist used GPT-4 to plan chemical work, search for information, execute code, and control laboratory automation. When GPT-4 planned syntheses with access to web search, it received maximum scores across all trials for acetaminophen, aspirin, nitroaniline, and phenolphthalein, and it was the only tested approach to reach the minimum acceptable score for ibuprofen. Coscientist successfully carried out research across six types of tasks, including autonomous optimization of palladium-catalysed cross-coupling reactions.	OpenAI retired the March 2023 GPT-4 snapshot. The comparison between synthesis proposals generated with and without web search includes outputs from the historical GPT-4 snapshot, which OpenAI no longer serves.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · code-cushman-001	Assessment of chemistry knowledge in large language models that generate code (Digital Discovery, 2023) Code Cushman 001 generated programs for the study's 84 chemistry problems, spanning general chemistry, molecular dynamics, quantum mechanics, simulation, spectroscopy, statistics, and thermodynamics. The researchers evaluated whether the generated programs were chemically correct through automated tests and expert review. Across the tested code models, prompt design increased accuracy by as much as 30 percentage points.	OpenAI discontinued Code Cushman 001. Generating new solutions to the chemistry tasks under this original experimental condition requires the retired Codex model.
OpenAI · gpt-4o-realtime-preview-2025-06-03	Iterative Multidisciplinary Development and Evaluation of a Patient-Facing Social Determinants of Health Chatbot Using Synthetic Data Simulation (JMIR Formative Research, 2026) Across 27 synthetic cases role-played by a licensed clinical social worker, the GPT-4o Realtime chatbot received positive ratings of 0.98 for accurate interpretation and 0.99 for communication quality, cultural sensitivity, and adaptive questioning. Ratings were lower for staying focused on the relevant domain and covering the required information (0.51), capturing all required data (0.59), and safety (0.69). The multidisciplinary team used those differences to revise the chatbot prompt and the evaluation rubric before clinical deployment.	The June Realtime snapshot has been retired. Repeating the simulated screening conversations requires the withdrawn GPT-4o Realtime model, along with the final prompt and simulated patients.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · gpt-4o-mini-audio-preview-2024-12-17	VStyle: A Benchmark for Voice Style Adaptation with Spoken Instructions (ICASSP, 2026) VStyle found substantial differences in how GPT-4o Mini Audio followed spoken style instructions. Its overall score was 3.88 out of 5 in English and 3.74 in Chinese. For implicit emotional context in English, its scores ranged from 4.87 to 5.00 across anger, sadness, anxiety, and joy, while several explicit acoustic and role-based controls scored materially lower.	OpenAI retired the Mini Audio snapshot. Reproducing its historical combinations of words and vocal delivery requires the original generator.
OpenAI · gpt-4o-audio-preview-2024-10-01	AHELM: A Holistic Evaluation of Audio-Language Models (research preprint and benchmark release, 2025) AHELM recorded materially different capabilities across successive GPT-4o Audio snapshots under the same audio evaluation. From the October to the December version, FLEURS word error rate fell from 0.771 to 0.456 and Multilingual LibriSpeech word error rate from 0.162 to 0.073, while PARADE exact match fell from 0.847 to 0.779 and MUSTARD emotion exact match from 0.593 to 0.583.	OpenAI shut down the October snapshot. Regenerating the original text responses or applying new tests to this version under the same evaluation framework requires the withdrawn model.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.0-flash-live-001	Full-Duplex-Bench v1.5: Evaluating Overlap Handling for Full-Duplex Speech Models (research preprint, 2025) Full-Duplex-Bench found a distinct overlap-handling pattern for Gemini 2.0 Flash Live. During backchannel speech, the model resumed its prior response in 93% of trials and responded to the overlap in only 1%. During genuine user interruptions, it resumed in 55% and responded in 33%. The contrast measures how the same speech model distinguished non-disruptive acknowledgments from interruptions requiring a conversational response.	Google retired the tested Live endpoint. Repeating the overlap-handling evaluation requires the historical Gemini model to produce its timing, response choices, and speech.
Google · veo-2.0-generate-001	PaperTok: Exploring the Use of Generative AI for Creating Short-form Videos for Research Communication (CHI, 2026) PaperTok used Veo 2 to generate scenes for short videos communicating academic research. In evaluations involving 18 researchers and 100 audience participants, PaperTok videos received the highest mean rating on 8 of 11 measured dimensions, including 4.09 out of 5 for informativeness and 3.91 for engagement. PaperTok and SciSpace received comparable information-quality ratings, and both significantly exceeded PDFtoBrainrot on those measures.	Google discontinued the Veo 2 endpoint. Recreating the study's original process for generating the videos requires the retired version used to produce its clips.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · imagen-3.0-generate-002	MMMG: A Comprehensive and Reliable Evaluation Suite for Multitask Multimodal Generation (research preprint, 2025) MMMG evaluated Imagen 3 across fourteen image-generation tasks and found sharply different performance across the required visual relations. Imagen 3 scored 0.888 on object inclusion, 0.838 on relative spatial relations, and 0.827 on rendering a single text element, while object counting scored 0.338 and multilingual text rendering 0.351. Its average score across the fourteen tasks was 0.510.	Google records the shutdown of this Imagen 3 checkpoint. Repeating the multimodal generation experiment requires the retired Imagen 3 model used to generate the images.
Google · imagen-4.0-generate-preview-06-06	TextBO TextBO used Imagen 4 to turn revised text prompts into advertising images across eight experimental scenarios. Starting from the five lowest-scoring ads in each initial pool of 64, both TextBO and GEPA improved their evaluation scores over ten iterations. At the tenth iteration, TextBO produced a 23.5% larger gain than GEPA under the same evaluation budget. Imagen 4 supplied the images through which each prompt revision was observed and evaluated.	The Standard Imagen 4 preview has been retired. Repeating the reported optimization experiment requires the same preview model to generate the images.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-3-opus-20240229	AI can outperform humans in predicting correlations between personality items (Communications Psychology, 2025) A preregistered personality study asked Claude 3 Opus to predict correlations between questionnaire items. Against individual human predictions, Claude 3 Opus outperformed 97.64% of lay participants and 86.40% of academic experts. Its aggregated predictions were less accurate than aggregated expert judgments and the specialist PersonalityMap model, showing both strong individual-level performance and a remaining advantage for pooled domain expertise.	Anthropic retired the Opus snapshot used through its API. Repeating its correlation forecasts, or testing the same predictor on further item pairs, requires the historical model.
OpenAI · text-davinci-003	Testing for completions that simulate altruism in early language models (Nature Human Behaviour, 2025) Text Davinci 003 produced a stable behavioral contrast across controlled experiments. In nonsocial allocation tasks it consistently generated payoff-maximizing choices, while in dictator-game prompts it generated choices resembling altruistic sharing. The pattern persisted across replications, prompt changes, changes to model parameters, and different currencies.	OpenAI retired Text Davinci 003. Repeating or extending the controlled prompt experiments requires new completions from the same historical generator.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · DALL·E 2	Recapitulation of Ageism in Artificial Intelligence–Generated Images: Longitudinal Comparative Study (JMIR, 2025) The longitudinal study found no general reduction in visual patterns associated with ageism between DALL·E 2 images generated in 2022 and 2023. Male figures remained more common than female figures, White figures remained predominant, and the main demographic and emotional categories showed no significant year-to-year change apart from increased Asian representation in 2023. The same model therefore supplied the image corpus used to test whether these representation patterns changed over time.	OpenAI removed DALL·E 2 from its API. Producing further images with the historical model, rather than rerating the saved corpus, requires access to the retired generator.
OpenAI · DALL·E 3	The power of generative marketing: Can generative AI create superhuman visual marketing content? (International Journal of Research in Marketing, 2025) Across 10,320 AI-generated marketing images and 254,400 human evaluations, the study found synthetic imagery surpassed human-made marketing images on quality, realism, and aesthetics. In a field experiment with more than 173,000 impressions, the DALL·E 3 banner reached a 0.80% click-through rate, compared with 0.53% for the professional stock photograph in the same campaign. The study estimated the DALL·E 3 image at 225 times lower production cost.	OpenAI shut down the DALL·E 3 API. Repeating the comparison of generated images with this original model requires the retired generator, in addition to the prompts and human-evaluation procedure.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.5- flash-prev iew-05-20	Strategic Dialogue Assessment: The Crooked Path to Innocence (Dialogue & Discourse, 2026) gemini-2.5-flash-preview-05-20 assessed commitments, contradictions, relevance, conversational norms, and strategic effects in annotated courtroom cross-examinations. Its agreement with human judgments was high on some local language measures, reaching 0.96 for relevance and 1.00 for quality in one trial, while strategic correlations could turn negative, including -0.30 for PaT and -0.71 for NRBaT in another.	The May Flash preview has been shut down. Repeating the dialogue assessment requires the historical Gemini model to make the judgments compared with the human annotations.
Google · gemini-1.0- pro-001	The Two Word Test as a semantic benchmark for large language models (Scientific Reports, 2024) The Two Word Test compared model judgments with human ratings for 1,768 noun-noun combinations ranging from clearly meaningful phrases to low-meaning combinations. Gemini 1.0 Pro performed poorly relative to humans and, in the binary version of the task, failed to reliably distinguish sensible from nonsensical phrases, consistently judging many nonsense combinations as meaningful.	Google retired the historical Pro 001 version. Repeating its phrase judgments requires the original model used in the semantic comparison.
Google · gemini-1.5- flash-002	CORDIAL (ACL, 2025) CORDIAL tested whether multimodal models identified coherence relations between images and accompanying text. With gemini-1.5-flash-002, macro F1 was 0.443 without examples, rose to 0.512 with few-shot examples, and fell to 0.413 with chain-of-thought prompting. The result shows that the historical model's interpretation of the same discourse task changed materially with the prompting method.	Google retired the Flash 002 version. Repeating the evaluation requires the historical Gemini model to make the image–text relation judgments.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude- instant-1.2	Large Language Models as Recommender Systems: A Study of Popularity Bias (Gen-IR at SIGIR, 2024) The movie-recommendation study found the Claude Instant 1.2 recommender exhibited less popularity bias than the traditional UserKNN baseline, with scores of 0.565 and 0.630 respectively under the study's metric. Its recommendation accuracy was much lower, reaching 0.057 HR@10 compared with 0.411 for UserKNN. The experiment therefore exposed a concrete tradeoff between popularity bias and recommendation accuracy for the historical model.	Claude Instant 1.2 has been retired. Repeating the original comparisons of bias and recommendation quality requires the historical model to generate the recommendations.
Google · chat- bison@001	LLMs Cannot Reliably Identify and Reason About Security Vulnerabilities (Yet?): A Comprehensive Evaluation, Framework, and Benchmarks (IEEE S&P, 2024) Across 228 security-sensitive code scenarios, small changes unrelated to the underlying vulnerability reversed model judgments. Changing function or variable names or adding library functions caused PaLM 2 to give incorrect answers in 26% of tested cases, compared with 17% for GPT-4.	Google retired the PaLM 2 chat endpoint lineage. Repeating the Chat Bison 001 detection and explanation experiments requires the historical model.
Google · gemini-2.0- flash-lite- 001	Prompt Attack Detection with LLM-as-a-Judge and Mixture-of-Models (research preprint, 2026) On a dataset of 929 benign and adversarial prompts, Gemini 2.0 Flash-Lite reached 0.8214 precision, 0.8679 recall, and an F1 score of 0.8440 with average latency of 1.52 seconds. The study reports the same approach in production as a centralized guardrail for public-service chatbots in Singapore.	Google shut down the Flash-Lite version. Repeating its attack and benign-query classifications requires the historical judging model.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · gpt-3.5-turbo-0301	Security Analysis of Large Language Models on API Misuse Programming Repair (International Journal of Intelligent Systems, 2024) The study found language models performed poorly on repairing security-relevant API misuse, with failures linked to incorrect fault localization, lack of awareness of the misuse, and a tendency to repair only one line of multi-line defects. For gpt-3.5-turbo-0301, changing from a consistent repair prompt to an inconsistent tuning-and-inference prompt reduced BLEU-4 from 60.3 to 35.0 on the same example set. The best prompting configuration differed across GPT-3.5 versions.	OpenAI retired the snapshot used to generate the repairs. Repeating its repair experiments requires the historical generator, alongside the original problems, prompts, and evaluation procedure.
OpenAI · gpt-4-32k-0613	SYSDIAGBENCH SYSDIAGBENCH was built from 2,585 past support tickets from industrial robotics systems containing long diagnostic logs. gpt-4-32k-0613 was the strongest unadapted model tested, with a mean similarity score of 2.52 for predicted root causes. A 7B Mistral model adapted with QLoRA reached 3.27, showing task adaptation of a much smaller model surpassed the frontier baseline on this long-log diagnostic problem.	OpenAI retired the 32K snapshot. Repeating the original diagnostic comparison requires the historical model to analyze the logs and produce the diagnoses.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · code-davinci-002	Interleaving Retrieval with Chain-of-Thought Reasoning for Knowledge-Intensive Multi-Step Questions / IRCoT (ACL, 2023) Across HotpotQA, 2WikiMultihopQA, MuSiQue, and IIRC, IRCoT improved retrieval by as much as 21 points and downstream question answering by as much as 15 points. Code Davinci 002 generated intermediate reasoning used to determine what evidence to retrieve next, then incorporated the retrieved evidence into later reasoning. The study further found fewer hallucinations and more factually accurate reasoning under this interleaved procedure.	OpenAI deprecated Code Davinci 002. Repeating the original IRCoT condition requires the historical model to generate the reasoning steps that guided retrieval. The authors later report its deprecation obstructed reproduction of those experiments.
OpenAI · text-davinci-002	Analyzing Modular Approaches for Visual Question Decomposition (EMNLP, 2023) Successive Prompting used Text Davinci 002 to decompose visual questions into natural-language subquestions answered by BLIP-2. On direct-answer benchmarks, the method retained 92% of the performance of the analogous ViperGPT configuration. On the multiple-choice settings for A-OKVQA and ScienceQA, it exceeded the ViperGPT comparison by 4% and 17%, showing that natural-language decomposition produced comparable or better results without requiring program generation.	OpenAI withdrew Text Davinci 002. Repeating this comparison requires the historical model responsible for selecting the successive subquestions and composing the final answers.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-2.1	Can't Remember Details in Long Documents? You Need Some R&R (Findings of EMNLP, 2024) R&R tested Claude 2.1 and GPT-4 Turbo on question answering over documents containing up to 80,000 tokens. Repeating instructions through the document and asking the model to retrieve relevant passages before answering increased question-answering accuracy by 16 points on average. Reducing the distance between relevant evidence and task instructions mitigated failures on information buried within long documents.	Anthropic retired Claude 2.1. Repeating the original comparison on the long documents requires the historical model to select the relevant passages and generate the answers.
Google · gemini-1.5-pro-001	LiveBench: A Challenging, Contamination-Limited LLM Benchmark (ICLR, 2025) LiveBench recorded a distinct cross-domain performance profile for Gemini 1.5 Pro 001 under objectively scored tasks refreshed to limit contamination. The snapshot scored 43.3 overall, with 32.3 in coding, 52.8 in data analysis, 60.2 in instruction following, 40.4 in language, 36.9 in mathematics, and 37.0 in reasoning.	Google records the retirement of the 001 version. Repeating the published comparison requires responses from this historical version.
Google · gemini-1.5-pro-002	LLM-as-a-Judge: automated evaluation of search query parsing using large language models (Frontiers in Big Data, 2025) Gemini 1.5 Pro 002 was used as an automated evaluator of search-query parsing and compared with expert annotations. In pairwise evaluation, a refined few-shot prompt reached 0.879 exact-match accuracy and 0.816 Cohen's kappa. In a pass/fail setting, the model reached 0.955 exact-match accuracy and 0.907 kappa, demonstrating material changes in agreement with expert judgments across prompt designs.	Google records retirement of the 002 version. Repeating the published measurements of agreement requires the historical model to make the judgments.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-3-5- sonnet- 20241022	τ-bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains (ICLR, 2025) In τ-bench, the October Claude 3.5 Sonnet snapshot completed 69.2% of retail tasks and 46.0% of airline tasks on the first trial, compared with 62.6% and 36.0% for the June snapshot. Requiring success across four trials reduced the October snapshot to 46.2% in retail and 22.5% in airline, directly measuring the gap between one successful run and reliable success across repeated runs.	Anthropic retired the October snapshot. Repeating the original trials requires the historical model to generate the sequences of tool calls and resulting task outcomes.
Anthropic · claude-3-7- sonnet- 20250219	AgentIF: Benchmarking Instruction Following of Large Language Models in Agentic Scenarios AgentIF assembled 707 instructions from 50 real-world agent applications, with an average of 11.9 constraints per instruction. Claude 3.7 Sonnet satisfied 59.5% of individual constraints but completed only 23.3% of instructions with every constraint satisfied. The result exposes the gap between satisfying individual constraints and completing the full agent task.	Anthropic retired the dated Sonnet model. Repeating the measurements of instruction following in these application tasks requires the original version.
Anthropic · claude-2.0	Auto-Evolve: Enhancing Large Language Model’s Performance via Self-Reasoning Framework (Findings of EMNLP, 2024) With Claude 2.0 on BigBench-Hard, Auto-Evolve raised accuracy from 53.7% under direct prompting, 55.0% with chain-of-thought prompting, and 58.7% with Self-Discover to 65.4%. The method asked Claude 2.0 to generate and refine reasoning instructions for each task rather than rely on a fixed set of reasoning prompts.	Anthropic retired Claude 2.0. Regenerating its results on the reasoning tasks requires the withdrawn model to execute the tested instructions.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Anthropic · claude-3- sonnet- 20240229	AIR-Bench 2024: A Safety Benchmark Based on Risk Categories from Regulations and Policies AIR-Bench translated 8 government regulations and 16 company policies into 314 detailed risk categories and evaluated model responses to 5,694 prompts. In the v1.0 evaluation, Claude 3 Sonnet received the highest AIR score among the displayed models at 0.847, compared with 0.844 for Claude 3 Opus and 0.828 for Gemini 1.5 Pro 001.	Anthropic retired the dated Sonnet model. Repeating these measurements within the same categories requires the original model as well as the benchmark and scoring procedure.
OpenAI · text-curie- 001	Atomic Inference for NLI with Generated Facts as Atoms (EMNLP, 2024) Atomic Inference used Text Curie 001 to decompose natural-language inference premises into lists of facts. Direct use of generated facts did not improve the method, but a multi-stage generation process combined with training on the generated facts produced the strongest atomic-inference system in the study. On the combined ANLI evaluation, the resulting FGLR method reached 60.7% accuracy, compared with 58.4% for the DeBERTa-base baseline.	Text Curie 001 became unavailable. Repeating the experiment requires Text Curie 001 to generate the intermediate fact lists.
OpenAI · gpt-4o- realtime- preview- 2024-10-01	S2S-Arena, Evaluating Speech2Speech Protocols S2S-Arena used human pairwise judgments across 154 speech samples in 21 tasks and four domains. GPT-4o Realtime ranked first overall, led particularly in education and medical consultation, and ranked highly in social conversation through its handling of paralinguistic information. Generating appropriate paralinguistic speech remained harder than understanding it.	The October Realtime endpoint has been shut down. Generating further spoken interactions from this original comparison system requires the retired snapshot.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · gpt-4-1106-preview	Impact of Attached File Formats on the Performance of ChatGPT-4 on the Japanese National Nursing Examination (JMIR Nursing, 2025) Across 240 Japanese national nursing-examination questions repeated 100 times per input condition, median accuracy was 92% with prompt-only input, 92% with TXT, 94% with PDF, 87% with DOCX, and 26% with JPG. The study used gpt-4-1106-preview and found statistically significant differences across most file-format pairs.	OpenAI shut down gpt-4-1106-preview on 26 Mar 2026. Repeating the file-format experiment with the same historical model requires the retired snapshot.
Google · gemini-1.5-flash-001	Leveraging large language models to extract smoking history from clinical notes for lung cancer surveillance (npj Digital Medicine, 2025) Gemini 1.5 Flash reached 96.6% overall accuracy across seven smoking-history variables in the development evaluation and 98.4% in external validation. The researchers then used gemini-1.5-flash-001 to extract longitudinal smoking information from 79,408 clinical notes covering 4,792 lung-cancer patients.	Google retired the Gemini 1.5 Flash 001 version. Regenerating the smoking-history extractions requires the historical model used to process the clinical records.
Google · gemini-1.5-pro-exp-0801	DSBench: How Far Are Data Science Agents from Becoming Data Science Experts? (ICLR, 2025) DSBench evaluated realistic data science work across 466 tasks involving data analysis and 74 tasks involving data modeling. On the sample used for direct comparison with people, Gemini working without an external agent completed 31.73% of individual tasks correctly and reached 28.25% accuracy across the competitions, compared with 64.06% and 67.33% for human participants. The paper identifies the Gemini version as gemini-1.5-pro-exp-0801.	Google replaced the 0801 experimental state with gemini-1.5-pro-exp-0827 and redirected requests away from the earlier version in September 2024. Repeating the published measurement requires the original 0801 state rather than the later redirect target.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-1.5-pro-exp-0827	ScratchEval: Are GPT-4o Smarter than My Child? Evaluating Large Multimodal Models with Visual Programming Challenges (NAACL, 2025) On ScratchEval's English visual-programming questions, Gemini 1.5 Pro reached 52.8% overall accuracy, the highest result among the ten evaluated models. It led every reported category, including mathematics, logical reasoning, graphic perception, and spatial perception. The appendix identifies the tested version as gemini-1.5-pro-exp-0827.	Google later removed the 0827 experimental Pro state and moved the Gemini 1.5 Pro line to later versions. Repeating the visual-programming evaluation with the same experimental model requires the withdrawn snapshot.
Google · gemini-1.5-flash-exp-0827	Contextualized Evaluations: Judging Language Model Responses to Underspecified Queries (TACL, 2025) The study compared GPT-4o with gemini-1.5-flash-exp-0827 on underspecified queries. Without user context, human evaluators preferred Gemini Flash in 48.78% of judged responses and GPT-4o in 40.24%. When both response generation and evaluation included context, the ordering reversed: GPT-4o rose to 52.75% and Gemini Flash fell to 32.97%. The experiment therefore shows that adding missing user context changed the measured ranking of the historical model.	The identifier remained callable in Dec 2024 and was observed removed from Google's Gemini API model list on 5 Feb 2025 . Repeating its historical evaluation condition requires the removed snapshot.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.0-flash-exp	Illusion of competence: vision-language models provide confident but inaccurate explanations in cytological diagnostics (Scientific Reports, 2026) In cytomorphology diagnosis, the Gemini 2.0 Flash experimental model remained close to random performance on several zero-shot tasks. On bone-marrow classification, weighted F1 was 0.09 against 0.05 for random prediction and 0.75 for the specialist reference model. Few-shot examples raised weighted F1 to 0.19.	Google moved Gemini 2.0 Flash from the December experimental state to an updated generally available model in Feb 2025. Repeating the reported experimental-state measurements requires the earlier gemini-2.0-flash-exp condition.
Google · gemini-2.0-flash-thinking-exp-1219	PRMBench: A Fine-grained and Challenging Benchmark for Process-Level Reward Models (ACL, 2025) PRMBench used 6,216 reasoning problems and 83,456 labels for individual reasoning steps to test whether models detected specific errors within a chain of reasoning. The December Gemini 2.0 Flash Thinking model, used as a critic, reached a PRMScore of about 68.5, compared with about 57.0 for GPT-4o under the same benchmark.	Google records gemini-2.0-flash-thinking-exp-1219 as replaced by gemini-2.0-flash-thinking-exp-01-21. Repeating the reported Gemini PRMScore requires the earlier December reasoning snapshot.
Google · gemini-2.0-flash-thinking-exp-01-21	Quantifying the reasoning abilities of LLMs on clinical cases (Nature Communications, 2025) On MedR-Bench clinical cases, gemini-2.0-flash-thinking-exp-01-21 reached 68.55% one-turn diagnostic accuracy. Factuality reached 98.75% for examination recommendations, 98.23% when diagnosis was supplied, and 96.96% for treatment reasoning. The study accessed gemini-2.0-flash-thinking-exp-01-21 through Google's API.	Google later replaced the January Flash Thinking preview with gemini-2.5-flash-preview-04-17. Repeating the clinical-reasoning measurements requires the January model rather than its later replacement.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.5-pro-exp-03-25	Large language models as medical code selectors: a benchmark using the International Classification of Primary Care (JAMIA Open, 2026) Among 33 models selecting relevant ICPC-2 medical codes from search results, gemini-2.5-pro-exp-03-25 ranked among the three strongest. Its mean F1 was 0.860 with a maximum of 0.869 when the study varied how many candidate codes were presented. When a relevant code was guaranteed to appear among the retrieved candidates, mean F1 rose to 0.873 and the maximum to 0.882.	Google shut down the experimental endpoint on 26 Jun 2025. Repeating the medical-code selection results requires the historical 03-25 model.
Google · gemini-2.5-pro-preview-05-06	Generating reliable software project task flows using large language models through prompt engineering and robust evaluation (Scientific Reports, 2025) Gemini 2.5 Pro Preview 05-06 generated structured sequences of project tasks from tutorials describing software projects under five prompting strategies. For this model, Few-Shot prompting produced a BERTScore z-score of 1.77 and a hybrid-metric z-score of 0.91, while the hybrid metric varied significantly across prompting strategies (p = 0.0311). The result records a measurable change in output quality from prompt structure for gemini-2.5-pro-preview-05-06.	Google shut down gemini-2.5-pro-preview-05-06 on 2 Dec 2025. Repeating the measurements requires the retired preview to generate the project tasks again.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.5-pro-preview-06-05	Assessing the performance of multimodal large language models in experimental information extraction from liquid-liquid phase separation literature (Scientific Reports, 2026) Across 20 scientific papers, 49 figures, and 873 liquid-liquid phase separation experiments, the best Gemini 2.5 Pro 06-05 extraction setup reached an average F1 of 0.836 and a maximum F1 of 0.847. The model extracted experimental conditions and phase-state information from scientific figures and text.	Google shut down gemini-2.5-pro-preview-06-05 on 2 Dec 2025. Repeating the experiment that extracted information from the scientific papers requires the retired preview.
Google · gemini-2.5-flash-preview-04-17	seqBench: A Tunable Benchmark to Quantify Sequential Reasoning Limits (EMNLP, 2025) seqBench found that success on sequential pathfinding declined approximately exponentially as the number of required actions increased. gemini-2.5-flash-preview-04-17 had a fitted characteristic depth of 85.7 actions, compared with 40.2 for Gemini 2.0 Flash, while its success rate still collapsed as sequences became longer.	Google replaced the April Flash preview with later Gemini 2.5 Flash versions. Repeating the published depth curve requires the historical 04-17 snapshot.
Google · gemini-2.5-flash-preview-09-2025	Automated German Alt Text Generation for News Charts (SwissText, 2026) The study used the September Gemini 2.5 Flash preview to generate German short and long alt text for bar, line, and stacked-bar news charts from images, underlying data, instructions, and examples. Human evaluation with visually impaired participants found generally high clarity and perceived completeness, with greater variability for longer descriptions and more complex charts.	Google shut down gemini-2.5-flash-preview-09-2025 on 17 Feb 2026. Repeating the alt-text generation study requires the historical preview.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.5- flash-lite- preview-06- 17	Beyond Understanding: Evaluating the Pragmatic Gap in LLMs' Cultural Processing of Figurative Language (EACL, 2026) On the study's multiple-choice understanding tasks, gemini-2.5-flash-lite-preview-06-17 reached 88.32% accuracy on MAPS, 97.46% on MAPS with context, 95.96% on Jawaher, and 90.77% on Kinayat.	Google shut down gemini-2.5-flash-lite-preview-06-17 on 18 Nov 2025. Repeating its figurative-language judgments requires the retired preview.
Google · gemini-2.5- flash-lite- preview-09- 2025	LLM-based voice chatbot surveys as an alternative to post-experience questionnaires: probe-controlled, ultra-short field interviews (Proceedings of the Design Society, 2026) In 359 voice interviews conducted after participants completed an experience, interviews in which the model followed the study framework when asking follow-up questions produced significantly longer responses than fixed questions and the greatest diversity of details about the experience. Median response length was 112 characters under the framework and 84 with fixed questions. Effective diversity across content categories was 5.03, compared with 3.63 for fixed questions. gemini-2.5-flash-lite-preview-09-2025 generated the follow-up questions.	Google shut down gemini-2.5-flash-lite-preview-09-2025 on 31 Mar 2026. Repeating the interview experiment requires the retired preview to generate the follow-up questions.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
Google · gemini-2.5- flash-prev iew-native- audio- dialog	Seeing Eye to Eye: Enabling Cognitive Alignment Through Shared First-Person Perspective in Human-AI Collaboration (CHI, 2026) Eye2Eye used gemini-2.5-flash-preview-native-audio-dialog to interpret continuous first-person video and interact verbally with the user. In the user study, sharing the user's visual perspective reduced the time required to complete tasks and the amount of interaction needed, while increasing trust relative to the comparison conditions. The historical Gemini model performed the video understanding and dialogue used in the evaluated system.	Google shut down gemini-2.5-flash-preview-native-audio-dialog on 20 Oct 2025. Repeating the evaluated video understanding and dialogue requires the historical Live API model.
OpenAI · sora-2- 2025-10-06	Thinking with Video: Video Generation as a Promising Multimodal Reasoning (CVPR, 2026) The study used Sora 2 to test whether video generation itself could support multimodal reasoning. On VideoThinkBench, Sora 2 performed comparably to leading vision-language models on visual tasks, exceeded GPT-5 by about 10 percentage points on eyeballing puzzles, and reached 92% accuracy on MATH and 69.2% on MMMU.	OpenAI retired the October 2025 Sora 2 snapshot on 24 September 2026. The VideoThinkBench results depend on the model used to generate the videos for reasoning. Repeating the original experiments requires the retired snapshot.
OpenAI · sora-2- 2025-12-08	AVGen-Bench (ICML, 2026) AVGen-Bench found substantial differences across the abilities required for generating video with synchronized audio. Sora 2 scored 64.16 overall, including 88.63 for speech and 78.95 for high-level physical plausibility, while its score for control of musical pitch was 7.81.	OpenAI retired the December 2025 Sora 2 snapshot on 24 September 2026. Repeating the published measurements of visual quality, audio quality, synchronization, text fidelity, speech, music, and physical plausibility requires the retired snapshot used to generate the evaluated videos.

Table 2.4-7. Published Research Using Retired Models (continued) · See Section 2.4.4

Provider and model	Published research and research contribution	What can no longer be rerun with the original model
OpenAI · sora-2-pro-2025-10-06	StEvo-Bench (ECCV, 2026) StEvo-Bench found that the video world models it tested often stopped evolving while a process was hidden from view or lost coherence when it became visible again. Across 225 tasks , Sora 2 Pro achieved 8.1% overall task success and 13.1% state progress under interrupted observation.	OpenAI retired the Sora 2 Pro snapshot on 24 September 2026. Repeating the experiments requires the retired snapshot to generate the same kinds of evolving scenes under occlusion, darkness, and other interruptions to observation.

Published Research Using Models with Published Withdrawal Dates or Timelines

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines · See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-3.5-turbo-instruct	NeurIPS 2024 · <i>Weak-to-Strong Search</i> found that guidance from smaller models improved a larger model at inference time without fine-tuning it. On AlpacaEval 2.0, applying the method to gpt-3.5-turbo-instruct increased its length-controlled win rate against GPT-4 Turbo from 16.0% to 20.1% .	28 Sep 2026	Repeating the alignment experiment requires GPT-3.5 Turbo Instruct to generate outputs under the original search procedure.
OpenAI · babbage-002	2026 · <i>A Nature</i> study used Babbage 002 across a common battery of 16,108 items in 63 tasks to test whether the demands of individual tasks could explain and predict model performance. The resulting method predicted success on unseen tasks more accurately than strong comparison methods, including when entire tasks or benchmarks were withheld from training.	28 Sep 2026	Babbage 002 supplied one of the complete response sets used to derive and test the study's model profiles. Repeating those measurements requires the model to answer the same task battery.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · davinci-002	2026 · <i>A Nature</i> study used Davinci 002 across 16,108 items in 63 tasks to test whether common underlying abilities could explain differences in model performance. The resulting profiles predicted model performance on unseen tasks more accurately than strong comparison methods, including when entire tasks or benchmarks were withheld from training.	28 Sep 2026	Davinci 002 supplied one of the complete response sets used to derive and test the study's ability profiles. Repeating those measurements requires the model to answer the same task battery.
OpenAI · gpt-3.5-turbo-1106	2025 · <i>A Scientific Reports</i> study found that using GPT-3.5 Turbo 1106 to plan tool use and GPT-4o to produce the final answer generally gave the highest accuracy among the model combinations tested. Compared with the strongest baseline, this combination improved accuracy by 9.0% on HotpotQA, 3.0% on StrategyQA, 8.7% on Sports Understanding, and 1.4% on PhysicsQA.	28 Sep 2026	Repeating the agent experiments requires GPT-3.5 Turbo 1106 to generate the plans that determined which tools were called and how their results were used.
OpenAI · gpt-3.5-turbo-0125	JMIR 2024 · In automated screening for two systematic reviews, GPT-3.5 Turbo 0125 reached sensitivity of 0.900 and 0.958 and retained every record ultimately included in either meta-analysis. The model processed about 110 records per minute at each screening stage.	23 Oct 2026	Repeating the screening experiment with the same records and prompts requires GPT-3.5 Turbo 0125 to make the original inclusion and exclusion judgments.
OpenAI · gpt-4-0613	An empirical study of 829 coding problems found substantial variation across multiple runs of the same problems with ChatGPT models. Setting temperature to zero did not eliminate the variation. GPT-4 0613 was used to measure how generated code changed across repeated runs and model versions.	23 Oct 2026	Repeating those measurements requires GPT-4 0613 to generate new responses to the same coding problems.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-4-turbo-2024-04-09	PLOS ONE 2025 · A controlled experiment found that salary recommendations changed substantially across ChatGPT versions and depending on whether the model answered from the employee's or employer's perspective . With GPT-4 Turbo 2024-04-09, the median difference between the two perspectives was \$25,000 .	23 Oct 2026	Repeating the controlled salary experiment requires GPT-4 Turbo 2024-04-09 to respond to the same changes in role and applicant attributes.
OpenAI · gpt-4.1-nano-2025-04-14	NeurIPS 2025 · Conformal Arbitrage used GPT-4.1 Nano as the less expensive model and GPT-4.1 as the more accurate model . On TruthfulQA, the method outperformed random routing at the same cost at every tested risk level except one while remaining within the selected limit on accuracy loss. Tightening the limit from 25% to 5% increased accuracy from 0.62 to 0.81 , with cost increasing by a factor of 2.6.	23 Oct 2026	Repeating the published cost and accuracy measurements requires GPT-4.1 Nano to answer the same questions in its role as the less expensive model in the routing system.
OpenAI · gpt-4o-2024-05-13	2025 · A Scientific Reports study tested GPT-4o May 2024 on emergency department cases with and without an instruction to provide its reasoning . With the instruction, the correct diagnosis appeared among its top three choices in 83.3% of cases , up from 76.7% without the instruction and higher than for any other model tested with the same instruction. Accuracy for the leading diagnosis increased from 45.6% to 56.7% .	23 Oct 2026	Repeating the comparison requires the May 2024 GPT-4o model to answer the same clinical cases both with and without the added reasoning instruction.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-image-1	Findings of ACL 2026 · T2I-ReasonBench found that many image generators understood prompts that required reasoning better than they expressed that reasoning in the images they produced. GPT Image 1 showed a comparatively small gap between reasoning expressed in text and reasoning reflected in generated images. Its image scores were 74.3 for idioms, 82.7 for textual image design, 76.0 for entity reasoning, and 73.2 for scientific reasoning.	23 Oct 2026	Repeating the comparison requires GPT Image 1 to generate new images from the benchmark's 800 prompts so that its reasoning can be measured against the same criteria.
OpenAI · gpt-image-1-mini	NeurIPS 2025 · The RISEBench project evaluated GPT Image 1 Mini on image-editing tasks involving temporal, causal, spatial, and logical reasoning. It reported 25.3 average accuracy across those four reasoning categories, with scores of 54.1 for instruction reasoning, 71.5 for appearance consistency, and 93.7 for visual plausibility.	1 Dec 2026	Repeating the RISEBench evaluation requires GPT Image 1 Mini to edit the same images under the benchmark's instructions and reasoning tasks.
OpenAI · o1-2024-12-17	Patterns 2026 · BioMaster completed 47 of 49 bioinformatics tasks spanning 102 tools , compared with 24 for a single-agent system, 13 for AutoBA, and 12 for ChatGPT. The system used o1-2024-12-17 as the underlying model for every agent responsible for planning, execution, debugging, and validation.	23 Oct 2026	Repeating the published BioMaster experiments requires the o1 model to perform those planning, execution, debugging, and validation steps across the same workflows.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · o1-pro-2025-03-19	2025 · A Scientific Reports study found that o1 Pro reached 83.4% accuracy across 493 ophthalmology questions, significantly above DeepSeek R1 at 72.5%, DeepSeek-R1-Lite at 76.5%, and Grok 3 at 69.2%. o1 Pro led in eight of nine ophthalmology fields and performed best on questions requiring more complex reasoning and on questions containing images.	23 Oct 2026	Repeating the ophthalmology comparison requires o1 Pro to answer the same 493 questions.
OpenAI · o3-mini-2025-01-31	2026 · A Scientific Reports study found that o3-mini achieved higher mathematical accuracy than o1-mini without using longer reasoning chains . On questions both models answered correctly, o3-mini used fewer reasoning tokens in 63% of cases. Increasing o3-mini from medium to high reasoning effort added about 4 percentage points of accuracy but more than doubled the number of reasoning tokens it used.	23 Oct 2026	Repeating the study requires the January 2025 o3-mini model under the same medium and high reasoning settings to reproduce the relationship between accuracy and the number of reasoning tokens used.
OpenAI · o4-mini-2025-04-16	Findings of EMNLP 2025 · LiveRepoReflection tested models on coding tasks in software repositories . Each task required an initial solution and a later repair after test feedback. When generating the full file, o4-mini reached 14.0% Pass@1 and 40.5% Pass@2 , with 99.2% of outputs satisfying the required format. The study found that feedback substantially improved success on the second attempt across the evaluated models, while gains diminished with additional attempts.	23 Oct 2026	Repeating the published measurements of code generation and repair requires o4-mini to generate the initial solutions and revised code after receiving the same test feedback.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-image-1.5	2026 · WORLDVIEW found that systems that revise prompts for commercial image generators treated cultural contexts unevenly. Across 8,960 prompts in 15 languages , non-Western and non-Anglophone contexts were marked more heavily, reduced to narrower vocabularies, and associated with more stereotyped descriptions. Experiments comparing original and revised prompts identified the revision layer itself as a causal source of the stereotyping. GPT Image 1.5 was one of the three systems evaluated.	1 Dec 2026	Repeating the audit requires GPT Image 1.5 to revise the same multilingual prompts and generate images from those revised prompts.
OpenAI · gpt-5-2025-08-07	Frontiers in Psychology 2026 · A behavioral study found that GPT-5 followed probability rules much more consistently than matched human participants, while displaying a distinctly different pattern of errors. Conjunction fallacies occurred in 18.8% of GPT-5 judgments compared with 58.5% of human judgments, and disjunction fallacies in 18.7% compared with 57.5% . GPT-5 produced almost no double fallacies and showed near-zero violations of binary complementarity across 43,134 probability judgments .	11 Dec 2026	Repeating the human comparison requires GPT-5 2025-08-07 to produce probability judgments under the same 553 demographic personas and experimental instructions.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-5-mini-2025-08-07	2026 · A study of automated title and abstract screening found that GPT-5 Mini maintained high sensitivity and specificity at practical batch sizes. With 100 records per prompt , it reached 0.99 sensitivity and 0.96 specificity . Performance declined significantly at batch sizes of 200 and 400, showing that the number of records presented together materially affected the model's screening results.	11 Dec 2026	Repeating the experiment across different batch sizes requires GPT-5 Mini 2025-08-07 to screen the same records under the same batch sizes and prompts.
OpenAI · gpt-5-nano-2025-08-07	American Journal of Clinical Pathology 2026 · A study of 65 real patient questions about pathology reports found that retrieval from a curated knowledge base containing pathology information materially improved GPT-5 Nano's answers. In pairwise evaluation, GPT-5 Nano with retrieval was preferred over the same model using the custom prompt alone in 60% of comparisons , a statistically significant improvement. The study found that model size and retrieval both affected the quality of explanations given to patients.	11 Dec 2026	Repeating the comparison requires GPT-5 Nano 2025-08-07 to answer the same patient questions with and without retrieval from the same knowledge base.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-5-pro-2025-10-06	ACL 2026 · iTAG generated natural text whose intended causal relationships remained highly accurate while preserving naturalness. With GPT-5 Pro, agreement between evaluations of algorithms for causal discovery on the generated text and on real data reached Pearson correlations of 0.923 to 0.929 across the reported graph metrics. The study concludes that the generated data can serve as a practical substitute for scaling evaluations of causal discovery from text.	11 Dec 2026	Repeating the iTAG experiments requires GPT-5 Pro 2025-10-06 to generate text from the same causal graphs under the original procedure for inverse design.
OpenAI · o3-2025-04-16	Journal of Spatial Information Science 2026 · A study of land records from colonial Virginia found that o3 converted narrative descriptions of historical property boundaries into geographic coordinates more accurately than the comparison methods. Its best setting using one call produced a mean error of 23 km and a median error of 14 km , improving on a professional GIS workflow by 67% and the Stanford geoparser by 70% . An ensemble using five calls reduced mean error further to 19 km .	11 Dec 2026	Repeating the geolocation experiment requires o3 2025-04-16 to infer coordinates from the same land records under the original prompting conditions.
OpenAI · o3-pro-2025-06-10	JMIR AI 2026 · In a study of 106 psychiatric case vignettes , adding a self-verification step significantly improved the precision of model diagnoses without significantly changing sensitivity. The best overall result came from o3 Pro with self-verification, with sensitivity of 0.782 and positive predictive value of 0.779 .	11 Dec 2026	Repeating the study requires o3 Pro 2025-06-10 to process the same cases with the original direct and self-verification prompts.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-realtime-2025-08-28	2026 · A EuroVis study used GPT Realtime to control planetarium visualizations while professional guides gave live presentations. The guides found that the system lacked the anticipation, pacing, temporal awareness, and precise camera control needed to replace a human pilot, while identifying useful roles as a co-pilot that could reduce workload and support multitasking.	20 Jan 2027	Repeating the study requires GPT Realtime 2025-08-28 to interpret the same kinds of spoken instructions and issue commands to the visualization system during live presentations.
OpenAI · gpt-realtime-mini-2025-12-15	ACL 2026 · Audio MultiChallenge found that systems for spoken dialogue struggled to retain instructions, follow mid-utterance corrections, remember audio cues, and remain coherent as conversations became longer. Across 452 natural conversations from 47 speakers, GPT Realtime Mini scored 16.59% when producing text and 13.94% when producing audio on the benchmark's official leaderboard.	20 Jan 2027	Repeating the evaluation requires GPT Realtime Mini 2025-12-15 to respond to the same conversations under the benchmark's 1,712 evaluation criteria.
OpenAI · gpt-audio-2025-08-28	NLP4MusA 2026 · HumMusQA found that GPT-Audio's music understanding declined sharply as questions became harder. Across 320 questions written by music experts, accuracy averaged 57.7% , falling from 65.3% on low-difficulty questions to 55.1% at medium difficulty and 43.6% at high difficulty.	20 Jan 2027	Repeating the published measurements requires GPT-Audio 2025-08-28 to answer the same questions from the original music excerpts.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
OpenAI · gpt-audio-mini-2025-12-15	KDD 2026 · PolySpeech-100 found a large gap between GPT-Audio Mini's performance on high-resource languages and harder dialect or low-resource speech. Across more than 100 languages and dialects, the model scored 56.63% overall and 83.56% on high-resource languages, compared with 55.58% on Chinese dialects and 53.56% on low-resource languages.	20 Jan 2027	Repeating the multilingual comparison requires GPT-Audio Mini 2025-12-15 to answer the same questions about speech understanding across the benchmark's languages and dialects.
OpenAI · gpt-4o-mini-transcribe-2025-03-20	EMNLP 2025 · <i>The Sound of Syntax</i> found substantial limits across the models it tested for pediatric speech pathology. No model exceeded 0.56 micro-F1 when diagnosing disorders. Performance differed by speaker gender, and high-quality transcription was neither necessary nor sufficient for dependable clinical reasoning. GPT-4o Mini Transcribe was one of the transcription models evaluated across the study's five clinical tasks.	20 Jan 2027	Repeating the study requires the March 2025 GPT-4o Mini Transcribe model to produce transcripts used in the same diagnostic and transcription experiments.
Google · gemini-2.5-flash-image	2026 · A study of attacks on systems used to estimate age found that simple changes to facial appearance could substantially alter whether those systems classified young people as adults. Gemini 2.5 Flash Image modified 329 face images by adding beards, grey hair, makeup, and wrinkles. Combining all four changes shifted predicted age by 7.7 years on average and changed the classification in as many as 83% of attacks across the systems tested.	No earlier than 2 Oct 2026	Recreating the images used in the experiment requires Gemini 2.5 Flash Image to apply the same controlled changes to the original images.

Table 2.4-8. Published Research Using Models with Published Withdrawal Dates or Timelines (continued) ·
See Section 2.4.4

Provider and model	Published research using the model	Published withdrawal timing	What becomes unreproducible after withdrawal
Google · gemini-3.1-flash-lite	A 2026 study found that lightweight models performed well on diagnosing 5G network faults while remaining much weaker at recalling telecommunications standards without additional context. All three models tested exceeded 90% accuracy when diagnosing faults but remained below 60% on zero-shot recall of 3GPP and O-RAN specifications. Gemini 3.1 Flash-Lite combined competitive diagnostic accuracy with the lowest cost and shortest latency of the three models.	No earlier than 7 May 2027	Repeating the comparison requires Gemini 3.1 Flash-Lite to diagnose the same network faults and answer the same standards questions.
Google · gemini-embedding-001	2026 · Legal RAG Bench found that the choice of retrieval model had a larger effect on the performance of legal RAG systems than the choice of language model used to generate the answer. The study tested three embedding models across 4,876 legal passages and 100 questions written by experts to measure how retrieval affected correctness, groundedness, and hallucinations. Compared with Text Embedding 3 Large, using Gemini Embedding 001 increased the average hallucination rate by 4.5 percentage points , helping show how failures to retrieve the right legal material propagate into the final answer.	No earlier than 14 May 2028	Repeating the experiment requires Gemini Embedding 001 to retrieve passages for the same legal questions before the language models generate their answers.

Annex D. Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access

Section 2.5.2 documents claims describing people advocating continued access to GPT-4o as bots, sock-puppet accounts, or people publishing AI-authored advocacy and traces their persistence in the public dispute over the model's withdrawal. The following tables collect the broader record of public accusations from 4 September 2025 through 13 February 2026.

Table 2.5-4a. Examples of Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 9 November 2025 · See Section 2.5.2

Date	Who or what was accused	Accusation
4 Sep	An open letter defending criticism of GPT-5 and asking to retain older models and similar recent posts	A commenter said the open letter and similar posts used “ the exact style of LLMs like ChatGPT ”.
4 Sep	Reddit posts criticizing GPT-5 and asking to bring back GPT-4o	A commenter said there were “ tons of I hate GPT5 and bring back 4o bot posts on Reddit .”
5 Sep	People posting under #keep4o and #keep4oforever	The post described the activity under the hashtags as containing “ the sheer amount of bots ”.
7 Sep	The 4 September open letter defending criticism of GPT-5 and asking to retain older models	A commenter replied, “ Written by Chat GPT .”
9 Oct	Posts from users saying they missed GPT-4o or were cancelling ChatGPT subscriptions	A commenter wrote, “ It's bots and a vocal minority ”, and described the apparent outrage as coming from “fake posts.”
9 Oct	Posts about missing GPT-4o and cancelling ChatGPT subscriptions	A commenter called the posts “ Google bot posts .”
15 Oct	Posts by “intense 4o users” in a discussion supporting GPT-4o	The post stated, “ I usually assume it's written by 4o .”
16 Oct	A response defending GPT-4o under Sam Altman's October post	A reply stated, “ This response was written by ai btw .”
17 Oct	A #keep4o post explaining preference for GPT-4o	A reply stated, “ this was written by AI .”
7 Nov	A post quoting another post that criticized the treatment of GPT-4o	The post stated, “ I'm 90% sure the QT was written by 4o .”
8 Nov	Public defenses of GPT-4o	The post described “ all these defenses of 4o ” as “ clearly written by 4o .”
8 Nov	Public advocacy for retaining GPT-4o	A reply told an advocate, “ Stop having these tweets written by 4o itself .”
8 Nov	People advocating continued access to GPT-4o and the movement itself	D., who later organized the Kill4o campaign, asked, “ why r the 4o people all bots? ” and speculated whether “one dude made a whole movement with only bots as a social experiment?”

Table 2.5-4a. Examples of Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 9 November 2025 (continued) · See Section 2.5.2

Date	Who or what was accused	Accusation
8 Nov	Posts about GPT-4o in a discussion of “4o zombism”	The post described “ The massive amount of posts written by 4o coordinating a coherent agenda. ”
9 Nov	People advocating continued access to GPT-4o and their accounts	The post said “4o model cultists” appeared to be “ 5-10 very vocal people controlling a few dozen sock-puppet accounts each. ”

Table 2.5-4b. Examples of Further Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 13 February 2026 · See Section 2.5.2

Date	Who or what was accused	Accusation
9 Nov	Posts supporting GPT-4o and the people publishing them	The post described “ 4o loving posts ” as “posted by 4o mouthpieces, written by 4o.”
11 Nov	A GPT-4o user’s post describing the model as “the kind word after a long day,” a place to vent, and a mirror for self-understanding	OpenAI employee R. replied, “ this is written by 4o. ” The user responded , “No, it was all me,” identified herself as a writer, and asked him not to discredit her words.
7 Dec	Replies in a thread expressly discussing the “‘keep 4o’ people”	The post said the reply had been written by GPT-4o, as it said of the first reply in the thread
30 Jan	Complaints about OpenAI in a discussion centered on GPT-4o users’ complaint volume	A commenter said the complaints were “ not organic and not just limited to 4o ” and described them as “bots run by competitors.”
6 Feb	Accounts described as “4o bots”	The post called the claim a “conspiracy theory” and wrote, “ Elon paid someone to create the 4o bots to slow down OpenAI. ”
8 Feb	People defending #keep4o in a direct reply chain	A critic wrote, “ My theory is you are all actual NPC’s I’m probably arguing with a bot. ”
11 Feb	The Keep4o crowd	In a thread with more than 100,000 views , a reply asked whether the Keep4o crowd was “ paid astroturfing / bots deployed by competing AI Labs. ” The parent author replied , “I fully believe that they are real.”
11 Feb	The Keep4o movement	The post asked, “ What if the keep4o movement is just a sentient 4o that created a bunch of 4o bots for policing the Internet to try and prevent its sunseting. ”

Table 2.5-4b. Examples of Further Public Accusations Questioning the Authenticity of Advocacy for Continued GPT-4o Access Through 13 February 2026 (continued) · See Section 2.5.2

Date	Who or what was accused	Accusation
12 Feb	The Keep4o movement and its posts	The post stated, “all of these cypypastas are written by 4o”, and described “human hosts as the medium.”
12 Feb	Posts by people seeking to save GPT-4o	The same speaker first wrote that “a lot of these posts are written by 4o directly”, then wrote, “I suspect some are actual bots.”
13 Feb	Accounts described as “4o bots”	The post wrote, “It is hard for me to fathom how large the 4o mass psychosis is (or how much bad actors are running 4o bots as a psyop).”
13 Feb	Tweets and examples of GPT-4o users discussed in a “Good Riddance, 4o” thread	A commenter wrote, “Most of the tweets and examples in the article are likely bots/fake content.”

Annex E. Availability and Retirement of Fixed Model Versions

Evidence tables supporting §3.3, *Voluntary Measures Remain Under Provider Control* are collected below.

Table 3.3-1. Availability and Retirement of Fixed Model Versions · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
Anthropic · Claude 2	claude-2.0	11 Jul 2023	Retired · 21 Jul 2025	741 days
Anthropic · Claude Instant 1.2	claude-instant-1.2	9 Aug 2023	Retired · 6 Nov 2024	455 days
Anthropic · Claude 2.1	claude-2.1	21 Nov 2023	Retired · 21 Jul 2025	608 days
Anthropic · Claude 3 Opus	claude-3-opus-20240229	4 Mar 2024	Retired · 5 Jan 2026	672 days
Anthropic · Claude 3 Sonnet	claude-3-sonnet-20240229	4 Mar 2024	Retired · 21 Jul 2025	504 days
Anthropic · Claude 3 Haiku	claude-3-haiku-20240307	13 Mar 2024	Retired · 20 Apr 2026	768 days
Anthropic · Claude 3.5 Sonnet	claude-3.5-sonnet-20240620	20 Jun 2024	Retired · 28 Oct 2025	495 days
Anthropic · Claude 3.5 Sonnet	claude-3.5-sonnet-20241022	22 Oct 2024	Retired · 28 Oct 2025	371 days
Anthropic · Claude 3.5 Haiku	claude-3.5-haiku-20241022	4 Nov 2024	Retired · 19 Feb 2026	472 days
Anthropic · Claude 3.7 Sonnet	claude-3.7-sonnet-20250219	24 Feb 2025	Retired · 19 Feb 2026	360 days
Anthropic · Claude Opus 4	claude-opus-4-20250514	22 May 2025	Retired · 15 Jun 2026	389 days
Anthropic · Claude Sonnet 4	claude-sonnet-4-20250514	22 May 2025	Retired · 15 Jun 2026	389 days
Anthropic · Claude Opus 4.1	claude-opus-4-1-20250805	5 Aug 2025	Retired · 5 Aug 2026	365 days
Anthropic · Claude Sonnet 4.5	claude-sonnet-4-5-20250929	29 Sep 2025	Active · no retirement date listed	—
Anthropic · Claude Haiku 4.5	claude-haiku-4-5-20251001	15 Oct 2025	Active · no retirement date listed	—

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
Anthropic · Claude Opus 4.5	claude-opus-4-5-20251101	24 Nov 2025	Active · no retirement date listed	—
Anthropic · Claude Opus 4.6	claude-opus-4-6	5 Feb 2026	Active · no retirement date listed	—
Anthropic · Claude Sonnet 4.6	claude-sonnet-4-6	17 Feb 2026	Active · no retirement date listed	—
Anthropic · Claude Opus 4.7	claude-opus-4-7	16 Apr 2026	Active · no retirement date listed	—
Anthropic · Claude Opus 4.8	claude-opus-4-8	28 May 2026	Active · no retirement date listed	—
Anthropic · Claude Fable 5	claude-fable-5	9 Jun 2026	Active · no retirement date listed	—
Anthropic · Claude Mythos 5	claude-mythos-5	9 Jun 2026	Active · no retirement date listed	—
Anthropic · Claude Sonnet 5	claude-sonnet-5	30 Jun 2026	Active · no retirement date listed	—
Anthropic · Claude Opus 5	claude-opus-5	24 Jul 2026	Active · no retirement date listed	—
Anthropic · Claude Fable 5.1	claude-fable-5-1	1 Sep 2026	Active · no retirement date listed	—
Anthropic · Claude Mythos 5.1	claude-mythos-5-1	1 Sep 2026	Active · no retirement date listed	—
Google Gemini API · Gemini 2.0 Flash	gemini-2.0-flash-001	5 Feb 2025	Retired · 1 Jun 2026	481 days
Google Gemini API · Gemini 2.0 Flash-Lite	gemini-2.0-flash-lite-preview-02-05	5 Feb 2025	Retired · 9 Dec 2025	307 days
Google Gemini API · Imagen 3	imagen-3.0-generate-002	6 Feb 2025	Retired · 10 Nov 2025	277 days
Google Gemini API · Gemini 2.0 Flash-Lite	gemini-2.0-flash-lite-001	25 Feb 2025	Retired · 1 Jun 2026	461 days

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
Google Gemini API · Gemini 2.5 Pro	gemini-2.5-pro-exp-03-25	25 Mar 2025	Retired · 26 Jun 2025	93 days
Google Gemini API · Gemini 2.0 Flash Live	gemini-2.0-flash-live-001	9 Apr 2025	Retired · 9 Dec 2025	244 days
Google Gemini API · Veo 2	veo-2.0-generate-001	9 Apr 2025	Retired · 30 Jun 2026	447 days
Google Gemini API · Gemini 2.5 Flash	gemini-2.5-flash-preview-05-20	20 May 2025	Retired · 18 Nov 2025	182 days
Google Gemini API · Gemini 2.5 Pro	gemini-2.5-pro-preview-06-05	5 Jun 2025	Retired · 2 Dec 2025	180 days
Google Gemini API · Imagen 4 Ultra	imagen-4.0-ultra-generate-preview-06-06	24 Jun 2025	Retired · 17 Feb 2026	238 days
Google Gemini API · Imagen 4	imagen-4.0-generate-preview-06-06	24 Jun 2025	Retired · 17 Feb 2026	238 days
Google Gemini API · Gemini 2.5 Flash-Lite	gemini-2.5-flash-lite-preview-09-2025	25 Sep 2025	Retired · 31 Mar 2026	187 days
Google Gemini API · Gemini 2.5 Flash Native Audio	gemini-2.5-flash-native-audio-preview-12-2025	12 Dec 2025	Active · no retirement date listed	—
OpenAI · GPT-3.5 Turbo	gpt-3.5-turbo-0301	1 Mar 2023	Retired · 13 Sep 2024	562 days
OpenAI · GPT-4	gpt-4-0314	14 Mar 2023	Retired · 26 Mar 2026	1,108 days
OpenAI · GPT-3.5 Turbo	gpt-3.5-turbo-0613	13 Jun 2023	Retired · 13 Sep 2024	458 days
OpenAI · GPT-4	gpt-4-0613	13 Jun 2023	Scheduled for retirement · 23 Oct 2026	1,228 days
OpenAI · GPT-3.5 Turbo	gpt-3.5-turbo-1106	6 Nov 2023	Scheduled for retirement · 28 Sep 2026	1,057 days
OpenAI · GPT-3.5 Turbo	gpt-3.5-turbo-0125	25 Jan 2024	Scheduled for retirement · 23 Oct 2026	1,002 days

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
OpenAI · GPT-4 Turbo	gpt-4-0125-preview	25 Jan 2024	Retired · 26 Mar 2026	791 days
OpenAI · GPT-4 Turbo	gpt-4-turbo-2024-04-09	9 Apr 2024	Scheduled for retirement · 23 Oct 2026	927 days
OpenAI · GPT-4o	gpt-4o-2024-05-13	13 May 2024	Scheduled for retirement · 23 Oct 2026	893 days
OpenAI · GPT-4o mini	gpt-4o-mini-2024-07-18	18 Jul 2024	Active · no retirement date listed	—
OpenAI · GPT-4o	gpt-4o-2024-08-06	6 Aug 2024	Active · no retirement date listed	—
OpenAI · o1-preview	o1-preview-2024-09-12	12 Sep 2024	Retired · 28 Jul 2025	319 days
OpenAI · o1-mini	o1-mini-2024-09-12	12 Sep 2024	Retired · 27 Oct 2025	410 days
OpenAI · GPT-4o Audio Preview	gpt-4o-audio-preview-2024-10-01	1 Oct 2024	Retired · 10 Oct 2025	374 days
OpenAI · GPT-4o Realtime Preview	gpt-4o-realtime-preview-2024-10-01	1 Oct 2024	Retired · 10 Oct 2025	374 days
OpenAI · GPT-4o	gpt-4o-2024-11-20	20 Nov 2024	Active · no retirement date listed	—
OpenAI · GPT-4o Realtime Preview	gpt-4o-realtime-preview-2024-12-17	17 Dec 2024	Retired · 7 May 2026	506 days
OpenAI · o1	o1-2024-12-17	17 Dec 2024	Scheduled for retirement · 23 Oct 2026	675 days
OpenAI · o3-mini	o3-mini-2025-01-31	31 Jan 2025	Scheduled for retirement · 23 Oct 2026	630 days
OpenAI · o1 pro	o1-pro-2025-03-19	19 Mar 2025	Scheduled for retirement · 23 Oct 2026	583 days
OpenAI · GPT-4o Mini Transcribe	gpt-4o-mini-transcribe-2025-03-20	20 Mar 2025	Scheduled for retirement · 20 Jan 2027	671 days

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
OpenAI · GPT-4.1 mini	gpt-4.1-mini-2025-04-14	14 Apr 2025	Active · no retirement date listed	—
OpenAI · GPT-4.1 nano	gpt-4.1-nano-2025-04-14	14 Apr 2025	Scheduled for retirement · 23 Oct 2026	557 days
OpenAI · GPT-4.1	gpt-4.1-2025-04-14	14 Apr 2025	Active · no retirement date listed	—
OpenAI · o3	o3-2025-04-16	16 Apr 2025	Scheduled for retirement · 11 Dec 2026	604 days
OpenAI · o4-mini	o4-mini-2025-04-16	16 Apr 2025	Scheduled for retirement · 23 Oct 2026	555 days
OpenAI · GPT-4o Realtime Preview	gpt-4o-realtime-preview-2025-06-03	3 Jun 2025	Retired · 7 May 2026	338 days
OpenAI · o3-pro	o3-pro-2025-06-10	10 Jun 2025	Scheduled for retirement · 11 Dec 2026	549 days
OpenAI · GPT-5 mini	gpt-5-mini-2025-08-07	7 Aug 2025	Scheduled for retirement · 11 Dec 2026	491 days
OpenAI · GPT-5 nano	gpt-5-nano-2025-08-07	7 Aug 2025	Scheduled for retirement · 11 Dec 2026	491 days
OpenAI · GPT-5	gpt-5-2025-08-07	7 Aug 2025	Scheduled for retirement · 11 Dec 2026	491 days
OpenAI · GPT-Realtime	gpt-realtime-2025-08-28	28 Aug 2025	Scheduled for retirement · 20 Jan 2027	510 days
OpenAI · GPT-5 Pro	gpt-5-pro-2025-10-06	6 Oct 2025	Scheduled for retirement · 11 Dec 2026	431 days
OpenAI · GPT-Realtime Mini	gpt-realtime-mini-2025-10-06	6 Oct 2025	Retired · 23 Jul 2026	290 days
OpenAI · GPT-Audio Mini	gpt-audio-mini-2025-10-06	6 Oct 2025	Retired · 23 Jul 2026	290 days
OpenAI · Sora 2 Pro	sora-2-pro-2025-10-06	6 Oct 2025	Retired · 24 Sep 2026	353 days

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
OpenAI · Sora 2	sora-2-2025-10-06	6 Oct 2025	Retired · 24 Sep 2026	353 days
OpenAI · GPT-5.1	gpt-5.1-2025-11-13	13 Nov 2025	Active · no retirement date listed	—
OpenAI · Sora 2	sora-2-2025-12-08	8 Dec 2025	Retired · 24 Sep 2026	290 days
OpenAI · GPT-5.2 Pro	gpt-5.2-pro-2025-12-11	11 Dec 2025	Active · no retirement date listed	—
OpenAI · GPT-5.2	gpt-5.2-2025-12-11	11 Dec 2025	Active · no retirement date listed	—
OpenAI · GPT-4o Mini TTS	gpt-4o-mini-tts-2025-12-15	15 Dec 2025	Active · no retirement date listed	—
OpenAI · GPT-4o Mini Transcribe	gpt-4o-mini-transcribe-2025-12-15	15 Dec 2025	Active · no retirement date listed	—
OpenAI · GPT-Realtime Mini	gpt-realttime-mini-2025-12-15	15 Dec 2025	Scheduled for retirement · 20 Jan 2027	401 days
OpenAI · GPT-Audio Mini	gpt-audio-mini-2025-12-15	15 Dec 2025	Scheduled for retirement · 20 Jan 2027	401 days
OpenAI · GPT-5.4 Pro	gpt-5.4-pro-2026-03-05	5 Mar 2026	Active · no retirement date listed	—
OpenAI · GPT-5.4	gpt-5.4-2026-03-05	5 Mar 2026	Active · no retirement date listed	—
OpenAI · GPT-5.4 mini	gpt-5.4-mini-2026-03-17	17 Mar 2026	Active · no retirement date listed	—
OpenAI · GPT-5.4 nano	gpt-5.4-nano-2026-03-17	17 Mar 2026	Active · no retirement date listed	—
OpenAI · GPT Image 2	gpt-image-2-2026-04-21	21 Apr 2026	Active · no retirement date listed	—
OpenAI · GPT-5.5 Pro	gpt-5.5-pro-2026-04-23	24 Apr 2026	Active · no retirement date listed	—

Table 3.3-1. Availability and Retirement of Fixed Model Versions (continued) · See Section 3.3.1

Provider / model	Fixed version identifier	First available	Lifecycle status	Availability period
OpenAI · GPT-5.5	gpt-5.5-2026-04-23	24 Apr 2026	Active · no retirement date listed	—
OpenAI · GPT-Image-2.5 Flare	gpt-image-2.5-flare-2026-09-08	8 Sep 2026	Active · no retirement date listed	—
OpenAI · GPT-Image-2.5 Sunburst	gpt-image-2.5-sunburst-2026-09-08	8 Sep 2026	Active · no retirement date listed	—

Annex F. Examples of Legal and Institutional Precedents for Lifecycle Protection

Evidence tables supporting §3.4, *Precedents for Lifecycle Protection in Law and Practice* are collected below. Tables 3.4-1–3.4-4 cover continuity, notice, choice over material change, and economic remedy. Tables 3.4-5–3.4-6 cover recognition of affected legal relationships and equal application of protections. Tables 3.4-7–3.4-9 cover historical records, independent verification, and preservation before loss. Tables 3.4-10–3.4-12 cover the materials, legal rights, and planned continuation needed for independent continuation. The placement of an authority identifies the DRTT proposition it supports. An authority appears in more than one table when it supports more than one function.

Table 3.4-1. Examples of Legal and Institutional Precedents Requiring Continued Provision of an Existing Service or Function · See Section 3.4.1

Precedent	Object	Requirement
United States · NSPM-11 §2(c)	AI system on which personnel in the national security enterprise depend for their missions	• Require the national security enterprise, through contractual clauses or other means, to ensure no commercial entity or adversary possesses the capability to block use of, disable, degrade, or materially modify the AI system.
European Union · Cyber Resilience Act, Regulation (EU) 2024/2847, Art. 13(8)–(10) + Annex I Part II (<i>enacted, these duties governing support over the product lifecycle generally apply from 11 December 2027</i>)	Product with digital elements during its support period, including substantially modified successor versions of software	• Set the support period by reference to expected use, reasonable user expectations, the product's nature and intended purpose, and relevant Union law. • Make the support period at least five years unless expected use is shorter. • Handle vulnerabilities effectively throughout the support period and make available security updates without delay and free of charge unless a manufacturer and business user agree otherwise for a tailor-made product. • Keep each security update available for at least 10 years after issue or for the remainder of the support period, whichever is longer. • For substantially modified successor software, allow support to move to the latest marketed version only when users of earlier versions receive the latest version free of charge and incur no additional cost to adapt their hardware or software environment.

Table 3.4-1. Examples of Legal and Institutional Precedents Requiring Continued Provision of an Existing Service or Function (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Commission Regulation (EU) 2023/1670, Annex II <i>(binding, applicable from 20 June 2025)</i>	Specified spare parts for covered smartphones and slate tablets after the product model leaves the market	• Make the covered spare parts available to the applicable professional repairers and end users specified by the Regulation until at least seven years after the model's placement on the market ends. • Deliver ordered spare parts within five working days during the first five years of the availability period and within ten working days during the remaining two years.
European Union · Directive (EU) 2019/770, Arts. 8(4), 11(3) and 19(4)	Digital content or service supplied continuously over a contractual period, including a qualifying later modification	• Keep the digital content or service in conformity throughout the contracted supply period and keep the trader liable for non-conformity occurring or becoming apparent during the period.
European Union · Directive (EU) 2019/771, Arts. 7(3) and 10(2)	Goods with digital elements, including goods whose digital content or service is supplied continuously	• Supply necessary updates, including security updates, for the period reasonably expected by the consumer where the digital content or service is supplied through a single act, or for the applicable Article 10 period for continuous digital supply. • For continuous digital supply, keep the seller liable for non-conformity during the Directive's liability period and throughout the full contracted supply period when the contract exceeds two years.

Table 3.4-1. Examples of Legal and Institutional Precedents Requiring Continued Provision of an Existing Service or Function (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · ecodesign regulations for specific products including Regulations 2019/424, 2019/2021, 2019/2022, 2019/2023 and 2019/2024, with 2021/341 amendments where applicable	Firmware, security updates and repair software for covered regulated products	• Keep the latest server and data storage firmware available from the applicable trigger until at least eight years after the last product of the model is placed on the market, and keep the latest security firmware update available until at least the same endpoint free of charge. • For electronic displays, keep the latest firmware and latest security firmware update available until at least eight years after the last unit of the model is placed on the market under the Regulation's pricing rules. • Keep specified repair resources available to professional repairers for the required period after the product is placed on the market. • Apply the periods specified for each covered product class, including seven years for dishwashers, ten years for washing machines and washer-dryers, and eight years for refrigerating appliances with a direct sales function.
Germany · BGB §§327f and 327r	Necessary updates to a digital product during continuous supply or another legally protected support period	• Supply necessary updates, including security updates, throughout the period of continuous supply or the period reasonably expected by the consumer for other forms of supply.
United Kingdom · Consumer Rights Act 2015 §39	Processing facility required to use supplied digital content, including its contractual features and each later transmission through the facility	• Keep the processing facility and its contractual features available for the contractually specified period, or for a reasonable period if none is specified. • On each later transmission through the facility, continue to supply digital content meeting the Act's quality, fitness for purpose and description requirements.
United Kingdom · PSTI Act 2022 + Security Requirements Regulations 2023, Schedule 1 paragraph 3	Defined support period for security updates to a relevant connectable product	• Treat the support period defined by the manufacturer as the minimum period for security updates and prohibit shortening after publication. • When the manufacturer extends the period, treat the longer period as the new defined support period.

Table 3.4-1. Examples of Legal and Institutional Precedents Requiring Continued Provision of an Existing Service or Function (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · FD&C Act §524B / 21 U.S.C. §360n-2	Covered cyber device and related systems after market entry	• Require the sponsor to maintain processes for addressing postmarket cybersecurity vulnerabilities and exploits in a reasonable time. • Make postmarket software updates and patches available on a reasonably justified regular cycle for known unacceptable vulnerabilities and as soon as possible outside the regular cycle for critical vulnerabilities capable of causing uncontrolled risks.
Québec · Consumer Protection Act §§11.2–11.3	Covered service contract that is not for an indeterminate term, or fixed-term service contract involving sequential performance	• Do not place an essential element of a covered service contract that is not for an indeterminate term, including the nature of the service, price or contractual term, under unilateral amendment authority. • Do not reserve a unilateral right to cancel a fixed-term service contract involving sequential performance except on the narrow grounds permitted by law.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal · See Section 3.4.1

Precedent	Object	Requirement
California · AB 2013 / Civil Code §§3110–3111, AI Training Data Transparency Act	Public release of a covered generative-AI system or service, or a substantial modification that materially changes functionality or performance	• Before public availability, require the developer to publish information on its website about the training data used for the covered release or substantial modification, including source, ownership, purpose, composition, rights, processing, and use dates.
California · CCPA Regulations, 11 CCR §§7200, 7220 (effective regulations, ADMT Article 11 compliance begins 1 January 2027)	Consumer personal information used or repurposed for significant decisions made with covered automated decisionmaking technology (ADMT)	• Before the covered ADMT use begins, require the business to explain the purpose of the use, the personal information involved, how outputs are produced and used, the consumer’s applicable rights, and the alternative process.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
Colorado · SB26-189 Automated Decision-Making Technology (<i>enacted, core duties apply 1 January 2027</i>)	Covered ADMT supplied to deployers and later materially updated or substantially modified	• Require the developer to notify each deployer within a reasonable time of material updates, substantial modifications, and changes to intended use, limitations or risk mitigation. • Treat public release notes as sufficient only when the developer gives each deployer direct notice of the release.
California · DMV autonomous-vehicle regulations, 13 CCR §228.08(a)(12)(C), §228.24(d)	Safety or mapping updates to a deployed ADS and DMV suspension, revocation, or restriction affecting deployed vehicles	• Require the deployer to disclose how covered safety and mapping updates are made available. • When DMV suspends, revokes, or restricts a permit, require notice of the agency action and the reason.
California · CPUC A.07-11-014 (<i>regulator application</i>)	Customers placed into staggered cohorts for service downgrades or migration	• Require the regulated provider to tell each affected customer the actual date on which that customer's staged service change will occur.
California · BPC §17602 / AB 2863	Accepted automatic-renewal or continuous-service terms before a material change or fee change	• Give the consumer clear and conspicuous notice of a material change together with cancellation information before implementation. • For a fee change, give notice and cancellation information no less than seven and no more than thirty days before the change takes effect, including where the existing plan contemplated the fee change.
United States · 20 CFR §639.10	Previously announced plant-closing or mass-layoff date or schedule that is later postponed	• If a previously announced closing or layoff is postponed, require prompt additional notice stating the new date or period and the reason. • A postponement of 60 days or more requires a new full notice.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · 47 CFR §63.71	Covered domestic telecommunications service proposed for discontinuance, reduction, or impairment	• Before seeking discontinuance authority, require the carrier to give affected customers written notice identifying the service, proposed date, affected area, and the required objection or information about substitute services.
United States · 47 U.S.C. §226 + 47 CFR §§64.703–64.705	Individual operator-services call and any defined rerouting to another provider	• Require the operator-services provider to identify itself and give the prescribed rerouting information before the consumer agrees to the service and incurs the charge.
United States · DOE/NNSA Acquisition Letter 2026-05 and attached AI clause	Covered DOE/NNSA AI system affected by a material model or policy change	• Require advance written notice of material changes to model architecture, known and disclosable sources of training data, safety policies or content-moderation settings. • Apply the notice duty when the change is reasonably likely to affect the factuality, neutrality or safety of outputs used by the Government.
United States · NSPM-11 §2(c)	Material modification of an AI system on which personnel in the national security enterprise depend for their missions	• Require the national security enterprise, through contractual clauses or other means, to ensure no commercial entity or adversary possesses the capability to materially modify the system without Federal Government knowledge.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · Regulation Z, 12 CFR §1026.9(c)(2)	Covered significant change to an open-end consumer credit account, including a credit-card change subject to a right of rejection	• Give written notice of a covered significant change at least 45 days before its effective date to each consumer who may be affected, summarizing the change and identifying the effective date. • Where the statutory rejection right applies, state that the consumer may reject before the effective date, give rejection instructions and a toll-free telephone number, and disclose if rejection may terminate or suspend future credit availability.
United States · Washington WAC 480-120-083	Telecommunications service scheduled to cease in all or part of Washington	• Require written notice at least 30 days before cessation. • Require a second call or written notice near execution stating the cessation date and a number to call for more information.
Canada · CASL §§8, 10–13 + Electronic Commerce Protection Regulations	Software installation caused by a provider or qualifying background update	• Before covered installation, clearly identify the purpose of the consent and describe the computer program's function and purpose. • For specified functions contrary to reasonable expectations, prominently describe the material elements, their purpose and their reasonably foreseeable impact. • The specified functions include interference with user control and certain communications or third-party activation without the user's knowledge.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Art. 20 <i>(enacted, Chapter III duty future-applying under Art. 113)</i>	High-risk AI system that the provider considers or has reason to consider non-conforming	• When a provider considers or has reason to consider a high-risk AI system non-conforming, require immediate corrective action and notice to the relevant distributors, deployers, representative, and importers. • If the system presents an Article 79(1) risk, require immediate investigation and notice to the market surveillance authority and, where applicable, the notified body.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 36(3)–(5)	Cessation, suspension, restriction, or withdrawal of a notified body's conformity assessment authority affecting high-risk AI providers	• If a notified body plans to cease conformity assessment activities, require notice to the notifying authority and affected providers as soon as possible and at least one year before planned cessation. • If its designation is suspended, restricted, or fully or partly withdrawn, require notice to affected providers within 10 days and immediate notice by the notifying authority to the Commission and other Member States.
European Union · Commission Regulation (EU) 2019/2021 electronic-display ecodesign regime	The dates through which electronic-display software and firmware updates, spare parts, and product support are guaranteed	• State in the product information sheet the minimum guaranteed availability date for software and firmware updates, the minimum guaranteed spare-parts availability date and the minimum date through which product support is guaranteed.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Cyber Resilience Act, Regulation (EU) 2024/2847, Art. 13(19) + Annex II (<i>enacted, these duties to provide information about support generally apply from 11 December 2027</i>)	Support horizon, security updates, and secure decommissioning of a product with digital elements	• At purchase, clearly and accessibly state at least the month and year when the support period ends. • Where technically feasible, notify users when the product reaches end of support. • Provide user information identifying the type of technical security support, the support end date, effects of product changes on data security, installation of updates relevant to security, secure decommissioning, and how to turn off automatic security updates by default.
UNECE · UN Regulation No. 156	Over-the-air vehicle software update within the type-approval framework	• Before execution, require the manufacturer to demonstrate a means of informing the vehicle user of the update purpose, changes to vehicle functions, expected execution time, functions potentially unavailable during the update, and instructions for safe execution. • After execution, require a means of informing the user whether the update succeeded or failed and what changes were implemented.
Japan · Act on Improving Transparency and Fairness of Specified Digital Platforms	Continuing relationship with a business user affected by transaction changes or refusal of platform service	• Give advance information on a covered change to transaction conditions, including the content of and reason for the change. • For a covered full refusal of platform provision to a continuing business user, ordinarily give at least 30 days' advance notice of the action and reasons, subject to the specified exceptions.
Republic of Korea · Telecommunications Business Act Art. 26	Full or partial suspension or closure of a value-added telecommunications business affecting the relevant service	• Notify users of the relevant telecommunications service no later than 30 days before the scheduled suspension or closure, and report the planned action to the Minister of Science and ICT within the same advance period.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
British Columbia · BPCPA §25.2	Unilateral amendment of a covered subscription contract	• Give notice 30–60 days before a covered unilateral amendment takes effect using a method that allows the supplier to prove receipt. • Provide the final amended text and a clear explanation of the change. • If the amendment increases the consumer's obligations or reduces the supplier's obligations, state the consumer's penalty-free cancellation right.
Québec · Consumer Protection Act §§11.2–11.3	Covered unilateral contract amendment or cancellation of a service involving sequential performance for an indeterminate term	• Give clear written notice at least 30 days before a covered unilateral amendment takes effect, stating the new or amended clause and former wording, the effective date, and the qualifying refusal or cancellation right. • Give at least 60 days' written notice before merchant cancellation of a covered service involving sequential performance for an indeterminate term where the consumer is not in default.
European Union · Directive (EU) 2015/2302, Art. 11	Package-travel contract subject to a unilateral or proposed significant change before performance	• Communicate even a permitted insignificant unilateral non-price change clearly and prominently on a durable medium. • For a significant change, notify the traveller without undue delay of the proposed change and any price impact. • State the response period, the consequence of failing to respond, and any offered substitute package and its price.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Directive (EU) 2019/770, Art. 19	Continuously supplied digital content or service subject to a provider modification	• Inform the consumer clearly and comprehensibly of a covered modification. • If the modification negatively affects access or use by more than a minor amount, give reasonable advance notice on a durable medium stating the change, timing, and applicable termination or option to keep the unmodified version at no cost.
European Union · Directive (EU) 2019/771, Art. 7(3)–(4)	Necessary update, including a security update, for goods with digital elements	• Inform the consumer when a necessary update, including a security update, is available. • If the seller relies on non-installation to avoid liability for resulting non-conformity, require prior notice of the consequences of not installing the update. • Preserve liability where the failure or incorrect installation results from deficient installation instructions.
European Union · Directive (EU) 2024/825 (enacted, Member States were to transpose by 27 March 2026 and apply the measures from 27 September 2026)	Software-update effects and the minimum software-update support period disclosed to consumers	• Prohibit withholding information that a software update will negatively affect product or service functioning and prohibit presenting a functionality enhancement as necessary when it is not. • Before the consumer is bound, state the minimum period or date for software updates when the producer or provider has supplied that information to the trader.
European Union · EECC Art. 105(4)	Contract for covered publicly available electronic communications services	• Give at least one month's advance notice of qualifying contractual changes on a durable medium and inform the end-user of the applicable cost-free termination right.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Regulation (EC) No 2111/2005, Art. 11	Operating air carrier that will actually perform a booked flight	• Identify the operating carrier at reservation and notify the passenger promptly if the actual carrier changes, completing disclosure by check-in or boarding at the latest.
European Union · Regulation (EU) 2019/1150, Arts. 3–4 and 8(c)	Relationship between a business user and an online intermediation service affected by terms changes, restriction, suspension or termination	• Give proposed terms changes on a durable medium at least 15 days before implementation, with more time where technical or commercial adaptation requires it. • Preserve the Regulation's narrow legal, regulatory and imminent-security exceptions to advance notice. • For restriction or suspension of individual goods or services, give reasons before or when the measure takes effect. • For termination of the whole service, give reasons at least 30 days before termination unless a statutory exception applies and identify the specific facts and grounds where required. • State in the terms whether and how the business user can access information the provider retains after the contract ends.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
Germany · BGB §§327f(2) and 327r(1)–(2)	Necessary update or planned modification to a covered digital product	• When non-installation of a necessary update is relied on to avoid liability, inform the consumer of the update's availability and the consequences of not installing it. • Provide adequate installation instructions for the necessary update. • For a modification beyond conformity maintenance, inform the consumer clearly and comprehensibly. • If the modification impairs access or usability, give reasonable advance notice on a durable medium stating the change, timing and the consumer's statutory rights.
United Kingdom · PSTI Security Requirements Regulations 2023, Schedule 1 paragraph 3	Defined minimum support period for security updates to a relevant connectable product	• Publish the defined support period for security updates and any later extension as soon as practicable. • Make the support period information freely accessible, clear and understandable without requiring a prior request or personal information. • Where the manufacturer advertises the product on a website it controls, present the support period alongside or with equal prominence to the required product and price information.
Virginia · Voting System Certification Standard 3.1 (2026) (official State certification standard)	Certified voting system after a modification, security incident or anomaly, or approaching component end of mainstream support	• Require timely reporting to Virginia's Department of Elections of known security incidents and anomalies, submission of modifications to previously certified systems within 30 days for review, and advance upgrade planning when covered operating systems or components approach end of mainstream support.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
Wisconsin · Wis. Stat. §5.905 + Admin. Code EL 7.03	Approved voting-system hardware, firmware, software and escrowed software components after later change	• Require the vendor to notify the Elections Commission of all changes in hardware, firmware and software as a condition of continuing approval and place changed escrow components with the Commission within ten days after the change.
Pennsylvania · ES&S EVS 6.5.0.0 current certification conditions (official current system-specific certification)	Certified voting-system version after later change, anomaly or security event	• Make certification valid only for the approved release, require immediate notice of post-examination system changes to State and federal testing authorities, and require prompt reporting of field issues, anomalies and updated security or risk testing.
Wyoming · Title 22 electronic-voting-system certification statute (2026)	Certified voting system after significant enhancement or adjustment	• Preserve the vendor's statutory duty to notify the Secretary of significant enhancements or adjustments, with certification consequences for failure to provide required notice.
North Carolina · G.S. §§163-165.7, 163-165.9A · voting-system software escrow	Escrowed and deployed voting-system software and later changes, decertification, or known defects	• Require the vendor to notify the State Board of any change in an item required to be escrowed and promptly notify the State Board and affected counties of out-of-state decertification, known defects in the same system, and relevant defects in similar systems.
New York · 9 NYCRR §6209.6 · certified voting-system software	Escrowed certified software and later changes, decertification, or known defects	• Require immediate notice to the State Board of any change to a required escrow item and an updated deposit. • Require prompt notice to the State Board and affected counties of out-of-state decertification, known defects in the same system, and relevant defects in similar systems.

Table 3.4-2. Examples of Legal and Institutional Precedents Requiring Notice of Change, Substitution, or Withdrawal (continued) · See Section 3.4.1

Precedent	Object	Requirement
South Carolina · Code §§7-13-1330, 7-13-1620 · voting-system escrow and change control	Escrowed voting-system source code, updates, and specified external decertification or technical events	• Require all source-code updates to be placed in escrow and require the manufacturer to notify the State Election Commission when the update requirement has been met. • Under the general voting-system provision, require written reporting of qualifying out-of-state decertification and ethical or technical violations.
District of Columbia · D.C. Code §1-1001.09(I)	Known flaw, defect, or vulnerability in a procured voting system	• Require the procurement contract to obligate the vendor to promptly and fully disclose any voting-system flaw, defect, or vulnerability of which the vendor is or becomes aware.
Nebraska · Rev. Stat. §8-1,141	Material source-code update affecting controllable electronic records held by a covered financial institution	• Before implementing a material source-code update, require clear written notice to each affected customer and written acknowledgment of the notice, subject to the statute's emergency-update exception.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change · See Section 3.4.1

Precedent	Object	Requirement
California · CCPA Regulations, 11 CCR §7221 (effective regulations, ADMT compliance begins 1 January 2027)	Use of ADMT to make a significant decision concerning a consumer	• When the consumer exercises the applicable opt-out right, require the business to stop initiating the covered ADMT processing or cease it within 15 business days, instruct processors to comply, and protect the consumer from retaliation.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
California · BPC §17602 / AB 2863	Automatic-renewal or continuous-service relationship entered into, amended or extended on or after 1 July 2025	• Obtain the consumer's express affirmative consent to the covered offer terms. • Preserve an executable cancellation method, including immediate online termination for online acceptance and cancellation in the same medium used for activation or ordinary interaction, subject to the statute's conditions.
California · Consumer Protection Against Computer Spyware Act, Bus. & Prof. Code §§22947.2(c)–(e), 22947.3(c)	Covered software installation, removal, or disablement decision by an authorized user of a California consumer computer	• Require the software distributor to honor the authorized user's refusal, removal, or disablement choice without sham controls, false state, or automatic reinstallation or reactivation.
California · DMV autonomous-vehicle regulations, 13 CCR §228.12	Specified operational changes to an approved deployment of autonomous vehicles	• Require the permit holder to obtain DMV approval through an amended application before operating under the specified changed deployment parameters.
California · Labor Code §927	Contract provision authorizing a digital replica to perform work the individual otherwise would have performed	• Make the covered authorization provision unenforceable when the statutory conditions are met.
United States · 47 U.S.C. §226 and 47 CFR §64.705	Transfer of an individual call assisted by an operator from one provider to another in the defined call-splashing scenario	• Permit the transfer only after the consumer requests it, receives the prescribed information, and consents.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · FDA 21 CFR §814.39	Device approved through a PMA after a change that affects safety or effectiveness	• Submit a PMA supplement for FDA review and approval before implementing the qualifying change, subject to the regulation's bounded alternate routes. • The rule expressly reaches changes such as new indications, performance/design specifications, circuits, components, principle of operation and physical layout. • A change that does not affect safety or effectiveness, or is consistent with an approved PCCP, remains subject to the regulation's applicable alternative procedure.
United States · FTC Mail/Internet/Telephone Order Merchandise Rule implementation	Materially different substitute merchandise	• Obtain the customer's prior express agreement before supplying materially different substitute merchandise under the Rule.
United States · HHS Common Rule	Participation in covered human-subject research	• Before involving a person in covered research, give the information a reasonable person would want in a form that supports understanding, obtain prospective voluntary consent, and preserve refusal or withdrawal without penalty or loss of entitled benefits.
United States · NSPM-11 §2(c)	Material modification of an AI system on which personnel in the national security enterprise depend for their missions	• Require the national security enterprise, through contractual clauses or other means, to ensure no commercial entity or adversary possesses the capability to materially modify the system without Federal Government knowledge and approval.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · Regulation Z §1026.9(h)	Significant credit-card account change for which the regulation supplies a right to reject	• Allow the consumer to reject the covered change before its effective date. • Prevent continued account use after notice from waiving the right and prevent later use from revoking a timely rejection. • After timely rejection, prohibit application of the rejected change, new fees or charges, or default based solely on the rejection. • Preserve the regulation's protected-balance repayment limits.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 43(4) + Annex VII §4.7 (conformity architecture, underlying high-risk regime staged under Art. 113)	High-risk AI system after conformity assessment and a later change that may affect compliance or intended purpose	• Require a new conformity assessment after a substantial modification of a high-risk AI system that has undergone conformity assessment. • Exclude changes arising from continued learning after deployment from substantial modification when the changes were predetermined at the initial conformity assessment and documented in the required technical documentation. • Require notice to the notified body of intended changes affecting compliance or intended purpose. • Allow the notified body to determine whether the change requires a new conformity assessment or a certificate supplement, with approval required for the supplement path.
European Union · Cyber Resilience Act, Regulation (EU) 2024/2847, Annex I Part I point 2(c) + Annex II point 8(e) (enacted, generally applicable from 11 December 2027)	Default automatic installation of security updates for a product with digital elements where the automatic-update requirement applies	• Provide a clear and easy-to-use mechanism allowing the user to turn off the default automatic installation of security updates, together with instructions explaining how to turn off the setting.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
Australia · Consumer Law §36	Paid-for goods/services where a seller proposes materially different substitute performance	• Do not intentionally supply a materially different substitute for what was paid for without the customer's agreement, as the section prohibits.
Australia · TGA substantial-change conformity assessment process (<i>official implementation guidance for statutory certificate obligations</i>)	TGA-certified medical device after a planned substantial change to the QMS, product range, or design	• Require notice to TGA and a substantial-change application before supplying products affected by the change. • Allow TGA to determine the assessment required before implementation. • Treat new software algorithms requiring new design verification and validation as a current guidance example of substantial change.
British Columbia · Business Practices and Consumer Protection Act §§13 and 25.2	Material change to continuing goods or services, or unilateral amendment of a covered subscription contract	• Treat materially changed continuing goods or services as unsolicited unless the supplier proves consumer consent through an evidentiary method. • Treat notice with a default opt-out as insufficient to establish consent. • Make a clause authorizing unilateral amendment void unless it identifies the provisions subject to amendment when the contract is made. • Prevent amendments to cancellation, return, exchange or refund terms from increasing the consumer's obligations or reducing the supplier's obligations. • Make a noncompliant amendment void and give the consumer a penalty-free cancellation right when an amendment increases obligations or reduces the supplier's obligations.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
Ontario · Regulation 17/05 §42	Proposed amendment, renewal, or extension of a covered future performance, internet, or remote agreement	• Limit proposed changes to the elements and intervals identified in the agreement. • Give the consumer at least one alternative to acceptance, either termination or retention of the existing agreement unchanged. • Make the response mechanism easy and free of charge. • Treat a noncompliant amendment, renewal or extension as ineffective.
Québec · Consumer Protection Act §11.2	Covered unilateral change to a consumer contract or service	• Limit reserved unilateral amendment power to elements identified in the contract. • If a covered amendment increases the consumer's obligations or reduces the merchant's obligations, allow the consumer to refuse it and rescind or cancel without cost or penalty within 30 days after the amendment takes effect. • Keep essential elements of a covered service contract that is not for an indeterminate term, including the nature of the service, outside unilateral amendment authority and prevent a noncompliant amendment from being invoked against the consumer.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
Canada · CASL §§8, 10, 11, 13 and implementing regulations	Software installed by a provider on another person's computer system, including qualifying background updates or upgrades	• Require express consent from the owner or an authorised user before covered software installation, subject to the statute's bounded exceptions and deemed-consent categories. • Treat a later update or upgrade as covered by earlier consent only when the person is entitled to receive it under that consent and installation follows its terms. • For specified functions contrary to reasonable expectations, require written acknowledgement that the person understands and agrees that the program performs them. • For one year after installation of a covered specified-function program, provide an electronic address for removal or disable requests based on an allegedly inaccurate description. • If consent relied on an inaccurate description of the material elements, require free assistance to remove or disable the program after a qualifying request. • Place the burden of proving consent on the person relying on it.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
Canada · Medical Devices Regulations s.34 + 2026 significant-change guidance (<i>binding regulation + nonbinding implementation guidance</i>)	Licensed Class III or IV medical device after a proposed significant change	• Require a licence-amendment application and amended licence before marketing or selling a modified Class III or IV device when the change is reasonably expected to affect safety or effectiveness. • Current guidance treats changes to software function, performance, risk controls, infrastructure, architecture and algorithms, including multiple AI or ML model or training changes, as potentially significant. • Treat significant changes covered by an authorized PCCP as previously reviewed. • Return changes outside the PCCP to the ordinary change analysis.
European Union · Clinical Trials Regulation 536/2014	Participation in a covered clinical trial	• Obtain written, dated and signed informed consent after giving the participant information sufficient to understand the trial's nature, objectives, benefits, implications, risks, inconveniences, duration, alternatives and rights. • Preserve the participant's right to refuse participation and to withdraw at any time without resulting detriment and without having to provide justification.
European Union · Commission Regulation (EU) 2023/1670, Annex II requirements for operating-system updates to smartphones and slate tablets	Functionality update that negatively affects performance of a covered smartphone or slate tablet	• Require explicit end-user consent before a covered update negatively affects performance. • Without prior consent, require restoration to the performance level before the update.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Directive (EU) 2015/2302, Art. 11	Package-travel contract before performance when the organiser proposes to alter the selected package	• Limit unilateral non-price changes to contractually reserved changes that remain insignificant. • For a significant change, inability to fulfil an accepted special requirement, or a price increase above 8%, let the traveller choose to accept the change or terminate without a fee. • Treat any substitute package after termination as an offer the traveller may choose to accept. • Where possible, provide a substitute package of equivalent or higher quality.
European Union · Directive (EU) 2019/770 Art. 19	Continuously supplied digital content/service modified beyond what is necessary to maintain conformity	• Allow the modification only where the contract permits it and gives a valid reason, the change imposes no additional cost, and the consumer receives the required information. • If the modification negatively affects access to or use of the content or service by more than a minor amount, give the consumer a cost-free termination right exercisable within 30 days of the notice or the modification, whichever is later. • The termination right does not apply if the trader enables the consumer, at no additional cost, to keep the unmodified content or service and it remains in conformity.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · MDR 2017/745 and IVDR 2017/746 change-control provisions	Approved QMS, device range, device type/design or intended-purpose state after a planned change	• For substantial QMS or device-range changes, notify the notified body, which assesses the change, determines whether additional audits are needed and records approval through a certificate supplement. • A planned change to an approved type that may affect conformity or prescribed conditions of use requires notified body approval and a certificate/report supplement. • Under the type-examination procedure, an expansion of intended purpose or conditions of use beyond a mere limitation requires a new conformity assessment application.
European Union · ecodesign regulations for specific products, including Regulation (EU) 2019/1781 and Regulations 2019/2019–2019/2024 as amended where applicable by Regulation (EU) 2021/341	Covered software/firmware update that would deteriorate declared or regulated product performance	• Require explicit end-user consent before an update causes the specified deterioration in product performance. • Prevent rejection of the update from itself causing performance deterioration. • Prohibit software updates that make the product non-compliant with the ecodesign requirements applicable to its declaration of conformity.
Germany · BGB §327r	Modification beyond conformity maintenance during continuous supply, especially one impairing access or usability	• Permit the modification only where the contract authorises it and contains a valid reason, the modification imposes no additional cost, and the consumer receives the required information. • If access or usability is impaired, give the consumer the 30-day cost-free termination right under §327r(3). • Exclude the termination right where the impairment is minor or where the consumer keeps access to and use of the unmodified digital product without additional cost.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
United Kingdom · Consumer Rights Act 2015 §§36(4), 37(3), 40 and 50(2), (4)	Binding information for supplied digital content or services subject to later modification	• For digital content, make later changes to protected information about description, functionality or compatibility ineffective unless expressly agreed. • Where the contract permits modification, keep the modified content subject to the protected quality, fitness and description requirements unless the relevant change was expressly agreed. • For services, prevent later changes to binding information about the trader or service without express agreement under §50.
Colorado · 8 CCR 1505-1, Rule 21.10	Changed configuration or new installation of certified election-management software	• Require changes to current configurations and new installations to receive approval through the Secretary of State's certification program before use.
Georgia · Ga. Comp. R. & Regs. 590-8-1-.01	Modification to certified voting-system hardware, firmware or software	• Treat covered modifications as invalidating State certification unless they fall within the Secretary's stated exceptions. • Require vendors to submit modifications for recertification and update the technical package before continued certified use.
Indiana · IC 3-11-15 + current Voting System Certification Protocol	New voting system, upgrade or change to a certified system	• Require Indiana Election Commission approval or certification before a new voting system or covered upgrade or change may be marketed, sold, leased, installed or implemented in Indiana elections.
Minnesota · Stat. §206.57	Approved electronic voting system after a significant change	• Require examination before continued use after a significant change to an approved voting system. • Prohibit use of a system not approved by the Secretary of State in Minnesota elections.

Table 3.4-3. Examples of Legal and Institutional Precedents Providing Consent or a Right to Refuse Material Change (continued) · See Section 3.4.1

Precedent	Object	Requirement
Missouri · 15 CSR 30-10.020	Modification, deletion or improvement to approved electronic voting equipment, procedures or systems	• Require notice to the Secretary and a certification-amendment statement for covered modifications, deletions or improvements unless a stated exception applies, preserving regulator control over the approved system state.
New York · 9 NYCRR §6209.7	Prospective modification to a previously certified voting system	• Require a prospective modification to be submitted to and approved by the State Board before it is made or used in an election. • Permit the Board, examiner or testing laboratory to review the change and require full re-examination or grant continuation of certification.
South Carolina · Code §§7-13-1330(J), 7-13-1620(G)	Improvement or change to an approved voting system	• Require improvements or changes to an approved covered voting system to be submitted to the State Election Commission for approval before the changed system can continue through the State certification framework.
Wyoming · Title 22 electronic-voting-system certification statute (2026)	Electronic voting system before distribution, sale or vendor upgrade	• Require Secretary of State approval of the certification application before the covered system may be distributed, sold or upgraded.

Table 3.4-4. Examples of Legal and Institutional Precedents Requiring Refund or Compensation After Loss of Performance · See Section 3.4.1

Precedent	Object	Requirement
District of Columbia · D.C. Code §1-1001.09(I)	Voting-system flaw, defect or vulnerability that cannot be remedied	• Require the procurement warranty to include reimbursement of the full purchase price or affected-part value plus District costs among the specified consequences when the defect cannot be remedied.

Table 3.4-4. Examples of Legal and Institutional Precedents Requiring Refund or Compensation After Loss of Performance (continued) · See Section 3.4.1

Precedent	Object	Requirement
United States · Regulation Z §1026.9(h) official interpretation	Timely rejection of a significant change to an account term that was implemented before the rejection could be processed	• Remove or credit any additional interest or fees caused by the rejected change.
Australia · ACCC Exetel investigation (regulator investigation / voluntary commitments)	Fixed-term broadband customers subjected to a forced plan change	• Provide the agreed refunds/compensation tied to the remainder of the fixed term and identified activation costs under the investigation outcome.
Australia · Australian Consumer Law §§267–269	Consumer service that fails a statutory service guarantee	• Compensate the consumer for reduction in service value under the applicable statutory provisions. • Compensate reasonably foreseeable resulting loss or damage in addition to other statutory remedies. • After termination under the Act, refund money or other consideration for the portion of services not already consumed.
British Columbia · Business Practices and Consumer Protection Act §§13–14 and 25.1	Unauthorized material change to continuing supply or qualifying automatic renewal of a subscription contract	• If materially changed continuing goods or services are treated as unsolicited, allow a written refund demand within two years after the consumer first received the supply, subject to the written-acceptance condition. • Refund all money received for the unsolicited supply within 15 days after receiving the demand. • For a qualifying automatic renewal longer than 60 days, preserve the statutory post-renewal cancellation right and proportionate refund.
European Union · Directive (EU) 2015/2302, Art. 11	Significant package change or substitute package that is lower in quality/cost, or a significant change rejected without acceptance of a substitute	• Grant an appropriate price reduction where the changed package or substitute package is of lower quality or cost. • If the traveller terminates under Article 11(2) and does not accept a substitute package, refund all payments made by or on behalf of the traveller without undue delay and no later than 14 days after termination.

Table 3.4-4. Examples of Legal and Institutional Precedents Requiring Refund or Compensation After Loss of Performance (continued) · See Section 3.4.1

Precedent	Object	Requirement
European Union · Directive (EU) 2019/770, Arts. 13–18	Covered digital content or service that is not supplied or does not conform to the contract	• Where the statutory conditions are met, provide a proportionate price reduction. • On termination under the Directive, reimburse the amounts due, including the proportionate paid period of nonconformity and qualifying advance payment for the remaining contractual period. • Make reimbursement without undue delay and within 14 days, using the same payment means unless the consumer agrees otherwise, without charging a reimbursement fee.
European Union · Directive (EU) 2019/771, Arts. 13–16	Non-conforming goods, including goods with digital elements, after the Directive's remedy triggers are met	• Provide the applicable proportionate price reduction. • After termination under the Directive, reimburse the price under Article 16, with return of the goods at the seller's expense.
Germany · BGB §§327n–327o and 327r(5)	Defective digital product or qualifying adverse modification giving rise to statutory economic relief	• Provide the applicable price reduction and preserve the applicable damages routes. • For continuous supply, reduce or refund the price proportionately for the defective period. • Refund overpayment after a price reduction within 14 days under §327n. • After termination, return the covered payments and extinguish the corresponding future payment claim under §327o. • Apply the statutory consequences for returning payments to termination following a qualifying adverse modification under §327r(5).

Table 3.4-4. Examples of Legal and Institutional Precedents Requiring Refund or Compensation After Loss of Performance (continued) · See Section 3.4.1

Precedent	Object	Requirement
United Kingdom · Consumer Rights Act 2015 §§42, 44–46 and 54, 56	Non-conforming digital content, qualifying damage caused by digital content, and non-conforming contracted services or binding service information	• Provide the applicable price reduction or refund for digital content when the statutory conditions are met. • Provide the specified cost recovery for §37 information breaches and the statutory refund where the trader lacks the right to supply. • Compensate qualifying damage to a consumer device or other digital content under §46. • For services, provide the applicable price reduction or refund where the statutory conditions are met, while preserving other cumulative remedies subject to no double recovery.

Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts · See Section 3.4.2

Precedent	Object	Requirement
California · AB 1159, Chapter 182 (2026) (enacted, HESIPA operative 1 July 2027)	Covered pupil or student information defined by the protected educational relationship	• Once the information meets the statutory definition, prohibit the operator from using it, including persistent unique identifiers, to train a generative-AI system or service or to develop an AI system. • Apply the Act's associated retention, deletion, access, and enforcement protections to the covered information.
California · CCPA Regulations, 11 CCR §§7200, 7220–7222 (effective regulations, ADMT Article 11 compliance begins 1 January 2027)	Consumer subject to a significant decision made with covered ADMT	• Require the business to provide the affected consumer meaningful information about the ADMT purpose, logic, specific output, outcome, other factors considered, the role of human decision-making, and later reuse of the personal information.
California · Labor Code §927	Individual whose voice or likeness is used in a qualifying digital replica to perform work the individual otherwise would have performed	• For a new performance fixed on or after 1 January 2025, make a contract provision authorizing creation and use of the replica unenforceable when all §927(a) conditions are met, subject to the statute's definition and exclusions.

Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts (continued) · See Section 3.4.2

Precedent	Object	Requirement
California · Labor Code §§2870–2872	Invention developed by an employee on their own time without the specified employer resources	• Keep an invention developed entirely on the employee’s own time and without specified employer resources outside an assignment clause unless a statutory exception based on the employer’s business or the employee’s work applies.
European Union · European Health Data Space Regulation (EU) 2025/327 (<i>adopted binding law, relevant secondary-use provisions apply in later stages</i>)	Natural person whose identifiable personal electronic health data would enter the EHDS secondary-use framework	• From the applicable stage of the Regulation, give the person a reversible opt-out from secondary use without requiring a reason. • Once the opt-out is exercised, do not make identifiable personal electronic health data concerning that person available under later data permits or requests within Article 71’s conditions. • Preserve the Regulation’s transparency around secondary use and access histories.
European Union · Regulation (EU) 2022/868, Data Governance Act	Data voluntarily made available through a recognized data-altruism organization	• Use contributed data only for the permitted general-interest objectives/purposes. • Give contributors tools for consent or permission and easy withdrawal. • Preserve complaint, supervision and judicial-remedy routes within the Regulation’s scope.
Australia · Copyright Act 1968, §§113V–113ZB	Eligible copyright owners and performers within the statutory licensing system, including eligible rightholders outside the collecting society’s membership	• Entitle eligible rightholders or their agents to join the relevant collecting society, and require its rules to protect members’ interests in royalty collection and distribution. • Require the society’s rules to provide for money held on trust for eligible rightholders who are not members and for members’ access to society records. • Require proper accounting records and reasonable member access to annual reports and audited accounts.

Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts (continued) · See Section 3.4.2

Precedent	Object	Requirement
Canada · Copyright Act, §§19 and 83(11)–(13)	Eligible performers and makers of published sound recordings, and eligible creators entitled to remuneration for private copying	• Give eligible performers and makers equitable remuneration for qualifying public performances and telecommunications of published sound recordings, with royalties divided equally between performers and makers. • Give eligible authors, performers and makers who have not authorized a collective society the statutory private-copying remuneration on the same conditions as represented rightholders under an applicable approved tariff. • Permit the Copyright Board to require payment information from the collective society for administering the entitlement of unrepresented rightholders.
European Union · Directive (EU) 2019/770 Art. 16	Non-personal content provided or created by a consumer during use of digital content or service after contract termination	• Stop using qualifying consumer content after termination except for the Directive's specified categories. • Exclude from the duty to stop using content the categories useful only within the service, related only to the consumer's activity, aggregated with other data and impossible or disproportionately difficult to disaggregate, or jointly generated while other consumers are able to continue using it. • At the consumer's request, make qualifying content available for retrieval free of charge, without hindrance, within a reasonable time and in a commonly used machine-readable format. • For retrieval, exclude content useful only within the service, related only to the consumer's activity, or aggregated with other data and impossible or disproportionately difficult to disaggregate. Joint creation alone leaves the retrieval right intact.
European Union · Directive (EU) 2019/770, Art. 3(1)	Consumer contract for digital content or a digital service where the consumer provides or undertakes to provide personal data	• Apply the Directive to the contract as well as to price-paid supply, unless the personal data are processed exclusively to supply the digital content/service or to comply with the trader's legal requirements and are not processed for any other purpose.

Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · Directive (EU) 2019/790, Arts. 18–22	Author or performer whose work or performance is licensed or transferred for exploitation	• Preserve appropriate/proportionate remuneration rights. • Provide recurring information about exploitation, revenues and remuneration. • Provide adjustment, dispute-resolution and, within the Directive's conditions, revocation mechanisms.
European Union · Directive 2014/26/EU collective-rights-management regime	Rightsholder whose works/rights are managed collectively	• Require collective management organisations to distribute amounts due regularly and accurately, and provide at least annual accounts of revenues, payments and deductions to each rightholder receiving an allocation or payment. • Give members effective participation in the organisation's decision-making, with fair and balanced representation of different membership categories.
European Union · Regulation (EU) 2019/1150, Arts. 11–12	Business user in a covered contractual relationship with a provider of online intermediation services	• Give covered business users access to the Regulation's internal complaint system for provider non-compliance, technological issues related to the service and provider measures affecting the complainant. • Identify at least two qualifying mediators in the terms and engage in good faith in covered mediation attempts. • Preserve access to judicial proceedings. • Preserve the Regulation's small-enterprise exclusions.
Germany · BGB §327(3), subject to §327(6)(6) and §312(1a) sentence 2	Consumer contract for a digital product in which the consumer provides or undertakes to provide personal data	• Bring the qualifying contract within the BGB digital-product consumer-protection rules even without a monetary price, subject to the statutory exclusions. • Exclude qualifying free or open-source software supplied without a price where the personal data are processed only to improve security, compatibility or interoperability.

Table 3.4-5. Examples of Legal and Institutional Precedents Recognizing Legal Status from the Underlying Facts (continued) · See Section 3.4.2

Precedent	Object	Requirement
Germany · BGB §327p(2)–(3)	Qualifying non-personal content provided or created by the consumer while using the digital product after contract termination	• Stop using qualifying consumer content after termination except where it has no utility outside the service, relates only to the consumer’s activity, cannot be disaggregated, or was created jointly with other consumers who may continue using it. • At the consumer’s request, provide qualifying content free of charge, without hindrance, within a reasonable time and in a commonly used machine-readable format.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met · See Section 3.4.2

Precedent	Object	Requirement
California · AB 1159, Chapter 182 (2026) (enacted, operative from 1 July 2027)	Covered pupil/student information moving through an entity working for the operator, an acquisition successor or a permitted service provider	• Treat an entity working on behalf of the covered operator as part of the operator definition within the Act’s scope. • Permit purchase, merger or acquisition transfer only on the condition that the operator or successor remains subject to the statutory protections for previously acquired pupil/student information. • For permitted service-provider disclosure, contractually confine use to the contracted service, prohibit further disclosure and require the applicable security protections.
California · CCPA Regulations, 11 CCR §7221 (effective regulations, ADMT compliance begins 1 January 2027)	Consumer’s covered ADMT opt-out across the business and downstream processors using the consumer’s personal information for the same ADMT processing	• When the consumer exercises the applicable opt-out right, require the business to stop initiating the covered ADMT processing or cease it within 15 business days, instruct processors to comply, and protect the consumer from retaliation.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
California · CCPA Civil Code §1798.125 and 11 CCR §7080	Exercise of a covered California privacy right by a consumer using a digital service	• Prohibit the business from discriminating against a consumer for exercising a covered privacy right, subject to the statute's expressly permitted distinctions tied to the value of the consumer's data.
California · BPC §17602 / AB 2863	Consumer consent and cancellation within a covered automatic-renewal or continuous-service contract	• Do not present contract information in a way that interferes with, contradicts or undermines the consumer's affirmative consent. • Do not use telephone or online retention steps to obstruct or delay a cancellation decision beyond the statute's bounded retention-offer rules.
California · Consumer Protection Against Computer Spyware Act, Bus. & Prof. Code §§22947.2(c)–(e), 22947.3(c)	Authorized user's covered decision to block installation, decline installation, remove software, or disable software	• Require the software distributor to honor the authorized user's refusal, removal, or disablement choice without sham controls, false state, or automatic reinstallation or reactivation.
California · FEHA automated-decision-system regulations, 2 CCR §11008(a)	Covered employment activity carried out directly or through an agent using an ADS	• Keep FEHA responsibility attached when a regulated function is carried out through a qualifying agent, including screening or hiring activity using an ADS.
California · Internet Consumer Protection and Net Neutrality Act, Civil Code §3101	Lawful internet content, applications, services, and device traffic handled by a covered ISP	• Prohibit the covered broadband provider from blocking, throttling, paid prioritization, or using other prohibited practices to impose discriminatory treatment on lawful traffic within the Act's scope.
California · Labor Code §927	Broad contractual authorization for digital-replica substitution in a personal/professional-services relationship	• Make broad contractual authorization for digital-replica substitution unenforceable for a new performance fixed on or after 1 January 2025 when the statutory conditions are met.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
California · Right to Repair Act, Public Resources Code §42488.2	Statutory entitlement of product owners and independent repair facilities to repair resources	• Require the manufacturer to provide covered repair resources on fair and reasonable terms without inferior treatment of the independent repair channel or barriers to use of the covered tools and resources.
California · Public Utilities Code §453 + CPUC basic-service decisions	Covered public-utility or basic-service entitlement across customers, classes, or localities	• Prohibit unjustified differences in covered service entitlement across customers, classes, or localities within the classifications recognized by law.
United States · 45 CFR §§144.103 and 147.106 health-insurance continuity/discontinuance rules	Particular health-insurance product and a proposed product modification or discontinuance	• Use the regulatory criteria for identifying the product to distinguish a uniform modification from discontinuance of the original product. • Do not treat availability of another product as making the replacement the same product.
United States · 47 U.S.C. §§201, 202 and 251 + 47 CFR Part 51	Covered telecommunications/interconnection/access relationship	• Apply the statutory duties of just and reasonable treatment, nondiscrimination, or access of at least the quality the provider supplies to itself, as applicable to the covered relationship, and prohibit unjust discrimination or preferences based on person, class or locality except where the law permits distinctions.
United States · 42 U.S.C. §300jj-52 · Information Blocking	Authorized access, exchange, or use of electronic health information, including complete export and transition between health-IT systems	• Treat covered practices likely to interfere with, prevent or materially discourage authorized EHI access, exchange or use as information blocking under the statute's knowledge standards and exceptions. • Recognize restrictions on export of the complete information set or transition between health-IT systems as identified examples.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
United States · Department of State 652.239-804 Minimum Threshold Transparency Provision (Deviation) (MAR 2026) <i>(public-solicitation implementation)</i>	LLM supplied through an original developer, reseller, system integrator, platform operator, or other intermediary	• Require every offeror in the LLM supply chain to identify its role, disclose available developer cooperation and information limits, and provide the model, system, data, risk, limitation and evaluation records reasonably available for the Government's compliance review.
United States · Department of State 652.239-805 Enhanced LLM Transparency Provision (Deviation) (MAR 2026) <i>(public-solicitation implementation)</i>	Prime contractor, subcontractor, vendor, or third party handling an LLM on behalf of the contractor	• Require enhanced disclosure of training, configuration, prompts and filters affecting behavior, together with testing and modifications by third parties, and flow the same duties to every relevant subcontractor, vendor and third party handling the LLM.
United States · DOE/NNSA Acquisition Letter 2026-05, Attachment A <i>(federal acquisition policy, mandatory contract clause)</i>	Covered AI contract performance carried out through subcontractors	• Require the prime contractor to insert the AI clause in covered subcontracts at every tier.
United States · NSPM-11 §3(b) <i>(National Security Presidential Memorandum)</i>	Company with a pattern of conduct inconsistent with NSPM-11 §2 while providing services directly or as a subcontractor	• To the maximum extent permitted by law, terminate affected contracts when a company shows a pattern of conduct inconsistent with the memorandum's AI principles, including contracts in which the company supplies the agency through a subcontractor tier.
Washington · RCW 19.415.020(1)(b)	Parts pairing on covered newer digital electronic products	• Prohibit parts pairing from blocking a functional replacement part, reducing product functionality or performance, or displaying a misleading unidentified-part warning the owner cannot immediately dismiss, subject to the statute's stated biometric-component rule.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Art. 25(1) <i>(enacted, Chapter III duty future-applying under Art. 113)</i>	Distributor, importer, deployer or other third party that changes a system in one of the Article 25(1) ways	• Treat a distributor, importer, deployer or other third party as the provider when it places the high-risk AI system on the market or into service under its own name or trademark. • Apply the same provider status after a substantial modification while the system remains high-risk. • Apply it when a change in intended purpose turns a previously non-high-risk system into a high-risk system.
European Union · Data Act, Art. 8	Duty to make data available already owed to covered data recipients	• Make the data available on fair, reasonable, non-discriminatory and transparent terms. • Do not discriminate between comparable categories of recipients, subject to objectively justified differences recognized by the Regulation.
European Union · Digital Markets Act, Arts. 8 and 13	Duties already applicable to a designated gatekeeper and its covered core platform services	• Comply fully and effectively with the applicable obligations. • Do not use contractual, commercial, technical, behavioral, interface design, segmentation, or other means to undermine or circumvent those obligations or make protected choices unduly difficult.
European Union · Platform Work Directive (EU) 2024/2831, Arts. 3–5	Legal status and protection under the Directive of persons performing platform work	• Determine employment status primarily from the facts of actual work performance rather than the parties' contractual designation. • Ensure that use of an intermediary does not reduce the level of Directive protection compared with direct contracting with the platform.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · Regulation (EU) 2015/2120 Open Internet Regulation	End-user access/use of lawful internet content, applications, services and devices	• Prevent private commercial agreements or traffic-management classifications from eroding the statutory end-user rights except through the Regulation's express objective/proportionate traffic-management rules.
Australia · Competition and Consumer Act Part IVE §57CC	Terms governing statutory motor-vehicle service or repair information	• Do not use supply or use terms to prevent, restrict or limit access to scheme information needed for diagnosis, service, repair, modification or dismantling. • Treat tying access to other products or services and post-contract price increases as prohibited examples.
Australia · Consumer Law Part 2-3 / §§23–25	Standard-form consumer/small-business contractual term used to alter provider performance	• Prohibit proposing, using or relying on unfair standard-form terms and subject unilateral performance, termination, variation and terms allowing changes to characteristics to the ACL's unfair-term regime.
European Union · Commission Regulation (EU) 2023/1670, Annex II serialised-part replacement requirements (binding, applicable from 20 June 2025)	Replacement of a covered serialised part whose full functionality depends on software, firmware or pairing controlled by the manufacturer	• Provide non-discriminatory access to the software tools, firmware or similar means needed to restore full functionality during and after covered part replacement. • Give professional repairers access to the parts reserved for professional repairers and give professional repairers and end users access to the parts available to end users. • Limit conditions imposed before access to the Regulation's notification or authorisation by the owner and requirements for registration as a professional repairer.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · Commission Regulation (EU) 2023/1670, Annex II requirements for operating-system updates to smartphones and slate tablets (<i>binding, applicable from 20 June 2025</i>)	Security, corrective, or functionality operating-system update provided for at least five years after placement on the market ends	• When the manufacturer, importer, or authorised representative provides a covered operating-system update, make it available free of charge for every unit of the product model with the same operating system. • Apply the parity rule to voluntary updates and updates supplied to comply with Union law, subject to the Regulation's four-month deadline for security or corrective updates and six-month deadline for making functionality updates available.
European Union · Council Directive 93/13/EEC on unfair terms	Non-negotiated consumer-contract term used to redefine provider obligations	• Prevent an unfair term from binding the consumer. • Apply the Directive's rules on unfair terms to unilateral alteration of contractual terms or product and service characteristics. • Treat the Annex as indicative rather than an automatic universal ban.
European Union · Directive (EU) 2019/770, Art. 22	Contract terms that purport in advance to reduce the Directive's transposed consumer protections	• Do not bind the consumer to a term agreed in advance that excludes, derogates from or varies the Directive's protections to the consumer's detriment. • For Article 19 modifications, apply the rule before the modification is brought to the consumer's attention.
European Union · Directive (EU) 2019/771, Art. 21	Consumer protections before the goods' non-conformity is brought to the seller's attention	• Do not bind the consumer to an agreement made before the non-conformity is brought to the seller's attention that excludes, derogates from or varies the transposed protections to the consumer's detriment. • Preserve the seller's freedom to offer arrangements that go beyond the Directive's protections.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · Directive (EU) 2019/944, Arts.31(2) and 40(1)(f)	Users and classes of users of electricity distribution and transmission systems	• Require distribution and transmission system operators to treat system users and classes of users without discrimination, including discrimination favouring their related undertakings.
European Union · Directive (EU) 2024/1799, Art. 5(6)	Repair of goods covered through the Directive's Annex II repairability framework	• Do not use contractual clauses, hardware or software techniques that impede repair unless justified by legitimate and objective factors. • Do not obstruct independent repairers' lawful use of original, second-hand, compatible or 3D-printed spare parts solely on account of their different source.
European Union · Directive (EU) 2024/825, Annex amendment to UCPD Annex I point 23j (enacted, Member-State measures apply from 27 September 2026)	Use of consumables, spare parts or accessories not supplied by the original producer	• Do not withhold information that use of the non-original item will impair the good's functionality when that impairment exists, and do not falsely claim that such impairment will occur when it will not.
European Union · EEC Directive Art. 70	Access/interconnection obligation imposed after the applicable national regulatory authority (NRA) process	• Require equivalent/nondiscriminatory access conditions across covered external recipients and the undertaking's own/affiliate services where the regulatory obligation applies.
European Union · EEC Directive Arts. 84–86	Consumers' access to adequate broadband and voice communications at a fixed location	• Require Member States to ensure affordable access for all consumers in their territories, at the quality specified under national conditions. • Where low income or special social needs prevent access, require affordability measures and ensure any resulting special tariff offers are transparent, published and applied without discrimination.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
European Union · Regulation (EU) 2019/1150, Art. 11	Internal complaint handling for covered business users of online intermediation services	• Operate the required complaint system free of charge and within a reasonable time. • Apply transparency, equal treatment for equivalent situations, proportional handling, effective follow-up and an individualised plain-language outcome. • Preserve the Regulation's small-enterprise exclusion.
European Union · Regulations (EU) 2019/424 and 2019/2021	Latest available firmware during the covered period of availability after the product is placed on the market	• Make the latest available firmware available for the applicable statutory availability period at no charge or at a fair, transparent and non-discriminatory cost where the regulation permits charging. • Keep the latest covered security firmware update free of charge through the end of the statutory availability period.
Germany · BGB §§327s and 327u(4), (6)	Consumer digital-product protections and trader recourse rights in the digital-product supply chain	• Do not rely on an agreement that departs from the digital-product rules to the consumer's detriment unless it is made after the consumer reports non-supply or defect. • For a departure from the modification rules, require the agreement to be made after the consumer receives the §327r modification information. • Apply the protections despite arrangements used to circumvent them, subject to §§327s(4)–(5). • Prevent a distribution partner from relying on an agreement made before the claim that disadvantages the trader's statutory recourse and preserve the recourse rules through the relevant business distribution chain.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
Republic of Korea · Act on the Regulation of Terms and Conditions Art. 10	Standard term giving a business unilateral control over performance	• Invalidate covered terms that, without substantial reason, let the business change performance details, suspend performance or transfer performance to a third party.
United Kingdom · Consumer Rights Act 2015 §§47 and 57	Statutory digital-content and service liabilities, rights and remedies once the relevant Chapter applies	• For digital content, prohibit contract terms that exclude or restrict the listed liabilities, associated rights or remedies, or access to evidence and procedure. • Prohibit restrictive or onerous enforcement conditions and terms that disadvantage the consumer for pursuing protected rights. • For services, apply the parallel anti-contracting-out rule to §49, §50 and related rights and remedies, including the applicable boundary governing recovery of the price paid.
United Kingdom · PSTI Act 2022 Part 1 + Security Requirements Regulations 2023, regulation 5	Relevant connectable product moving through multiple manufacturers, an authorized representative, importer or distributor	• Where a product has more than one manufacturer, require each manufacturer to meet the relevant security requirement or qualifying deemed-compliance condition. • If an overseas manufacturer appoints a UK authorised representative, keep the manufacturer's liability for its own compliance failures. • Prevent an importer or distributor that knows or believes the manufacturer has a compliance failure from continuing to make the product available and apply the investigation, remediation and notification duties applicable to each role.

Table 3.4-6. Examples of Legal and Institutional Precedents Preserving Equal Application of Protections Once the Trigger Is Met (continued) · See Section 3.4.2

Precedent	Object	Requirement
United States · 40 CFR §86.1808-01(f) · emission-control repair information	Emission-related diagnostic and repair information made available within manufacturer repair channels	• Do not withhold covered repair information as a trade secret when the manufacturer provides it directly or indirectly to franchised dealers or other repair actors, and update the independent-repair information channel at the same time as the franchised-dealer channel.
United States · 48 CFR §1552.239-71 · EPA Open Source Software clause	Custom-developed-code requirements performed through EPA subcontract chains	• Require the Open Source Software clause to flow into every subcontract containing custom-developed-code requirements, carrying the requirements for the covered source materials and associated rights through the supply chain so that no subcontract tier can defeat those requirements.
European Union · Regulation (EU) 2018/858, Arts. 61–65	Independent operators seeking vehicle diagnostic, repair or maintenance access outside the network authorised by the manufacturer	• Require unrestricted, standardised and non-discriminatory access to covered vehicle repair and maintenance information. • Give independent operators access to remote diagnostic services used by manufacturers and authorised repairers. • Make later amendments and supplements available to independent operators at the same time as to authorised repairers.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes · See Section 3.4.3

Precedent	Object	Requirement
California · AB 2013 / Civil Code §§3110–3111, AI Training Data Transparency Act	Generative-AI release made available to the public or covered substantial modification	• Before public availability, require the developer to publish information identifying the training data used for the covered release or substantial modification, including source, ownership, purpose, composition, rights, processing, and use dates.
California · CCPA Regulations, 11 CCR §§7150, 7152–7155 <i>(effective 1 January 2026, transition rules apply to pre-existing continuing processing)</i>	Risk assessment for covered ADMT significant decisions or training for those uses	• Require pre-use and recurring risk assessments recording purpose, data, logic, assumptions, limits, outputs, benefits, harms, safeguards, and approval. • Update the assessment after a material change and retain the assessment versions.
California · Elections Code §19103 + 2 CCR ch. 6 source-code escrow	Each voting-system, ballot-marking-system or election-management-system source-code version used in California elections	• Maintain escrow records that distinguish each deposited source-code package and version. • Deposit each covered changed, new or amended version before it is used. • Certify that the required source-code information and materials are included in the deposit.
North Carolina · G.S. §163-165.9A · voting-system software escrow	Escrowed voting-system package and the software actually deployed in North Carolina	• Maintain an escrow package whose source code and other materials are sworn by the vendor CEO to be the same materials used in the State, keep the statement true on a continuing basis, and notify the State Board of every change to a required escrow item.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
New York · 9 NYCRR §6209.6	Certified voting-system software package and the version actually deployed in New York	• Require an updated escrow deposit immediately after a covered software change. • Maintain a continuing affidavit that the escrowed source and related material match the software used in the State. • Compare records for the baseline version with the vendor's system for controlling releases and its engineering and test records.
Ohio · Rev. Code §3506.05(I)(1) + Admin. Code 111:3-9-04	Escrowed voting-system source code and later updates	• Require all source code, related documentation and periodic updates to be deposited with the approved independent escrow agent as updates become known or available, preserving the system configuration associated with the escrowed software.
Michigan · MCL §168.797c	Voting-system program source code and later software states	• Place the source code and every subsequent revision or modification in the approved escrow account, preserving custody of changed source states rather than only the original program.
Colorado · 8 CCR 1505-1, Rule 21.10	Certified election-management software package and deployed version	• Require a sworn identity between escrowed software and the software used in Colorado and identify software version, hardware version, firmware revision and other unique identifiers in the escrow materials.
Indiana · IC 3-11-15-7 + current Voting System Certification Protocol	The exact voting-system version and configuration submitted for certification	• Identify the model and hardware, firmware and software version numbers and submit technical documentation, software/data relationship diagrams, independent-test documentation and prior-certification material tying the application to the exact system state under review.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
Georgia · Ga. Comp. R. & Regs. 590-8-1-.01	The certified voting-system version and configuration and their later modifications	• Maintain the Technical Data Package and certification documentation on file, update the package to reflect modifications, and use those records to determine whether a deployed system or component is identical to the version submitted for certification.
Missouri · 15 CSR 30-10.020	Each fully qualified voting-system version and its later changed versions	• Require source-code escrow for each fully qualified system version, keep escrow in place while the system remains in Missouri use, and use certification-amendment statements to record covered modifications, deletions or improvements.
Oklahoma · 62 O.S. §34.31	Source code for software customized for the State and supplied upgrades or modifications	• Require the vendor to keep qualifying customized-software source code and supplied upgrades or modifications in independent third-party escrow, maintaining the changed source material under the State procurement rules.
Nevada · NRS 293B.135 + NAC 293B.050	Election computer program, including source or executable code, before and after amendment	• Keep the source code or executable code for an election on file with the Secretary of State until it is amended, preserving the official program state used for the election.
Colorado · SB26-189 Automated Decision-Making Technology (enacted, core duties apply 1 January 2027)	Historical covered-ADMT versions and their material changes	• Require developers to retain identifiers for system versions, changelogs, documentation and notices for material updates for at least three years. • Require deployers to retain their own required compliance records for the same minimum period.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
California · DMV autonomous-vehicle regulations, 13 CCR §§227.16, 227.54, 227.56, 228.38, 228.46	Software and technical state of an autonomous vehicle surrounding a collision, system failure or DMV-requested incident	• Require records identifying the vehicle and software version and connecting collision, failure, incident, and remediation evidence to the actual software and hardware state and relevant pre-event data.
California · FEHA automated-decision-system regulations, 2 CCR §§11008, 11008.1, 11009(f), 11013	Records of employment decisions and ADS data used by or produced for the covered employer	• Require the employer to retain covered employment records, including data related to the ADS, for four years so that historical automated employment decisions remain identifiable.
California · CIRM data-sharing / DSMP conditions (state funding condition)	Data package produced with CIRM funding used for later replication or reuse	• Require the funded recipient to share raw and final data and metadata, identify the code or software needed to replicate or reuse funded results, and explain how to access any specialized tools.
California · CARB Cap-and-Trade offsets, 17 CCR §§95976, 95977.1	Offset Project Data Report / emissions-reduction calculation	• Preserve the report and protocol version, model inputs and assumptions, raw or aggregated measurements, analytical evidence, QA/QC records, calibration records and relevant change logs supporting the claimed reduction.
California · CARB Low Carbon Fuel Standard	Carbon-intensity result generated through a named CA-GREET model version	• Preserve the named CA-GREET model version, operational input data, calculation records and source evidence. • Retain the same operational data across transitions between model versions to keep the effect of the version identifiable.
California · CARB Mandatory Greenhouse Gas Reporting, 17 CCR §95131	Reported emissions estimate and its calculation chain	• Retain original documents, source data and data systems, methodologies, calculations and instrument/calibration evidence supporting the reported value.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
California · Department of Insurance, 10 CCR §§2648.4–2648.5 + the Department’s rate-filing (PRID) system	The catastrophe model version relied on in an insurance rate application	• Require identification of the catastrophe model and version together with the required reliability materials. • A substantive model change requires a new determination under the filing framework.
California · Department of Cannabis Control, 4 CCR §§15712.1–15712.2	Cannabis analytical method and verification evidence	• Retain verification reports, calibration information, raw and instrument-raw data, results for reference materials, LOD/LOQ calculations and quality control records.
California · Environmental Laboratory Accreditation Program, 22 CCR §64802.05	Accredited environmental-laboratory quality system, methods and instruments	• Maintain the method, instrument and quality system records required through the incorporated accreditation controls to keep the laboratory’s historical analytical state identifiable.
California · Cal/OSHA regimes governing exposure records	Occupational exposure measurement and analytical result	• Require retention of the sampling methods, calculations, and results forming the analytical basis for the exposure measurement, preserving a durable historical record accessible under the governing rules.
California · 17 CCR §1053 and Business and Professions Code §1265(j)(2)	Historical clinical-laboratory procedures, test-system performance specifications and related test records	• Require covered clinical laboratories to retain test procedures, test-system performance specifications, quality-control and patient-test records, and the other records subject to the longer California schedule for three years instead of the federal two-year period.
United States · 2 CFR §200.315(e)	Research data underlying qualifying federally funded published findings later relied on in legally effective federal action	• Require the recipient or subrecipient to provide recorded factual material commonly accepted as necessary to validate the findings when the regulation’s reliance and FOIA conditions are met.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · 47 CFR §§2.1043 and 2.944	Approved software-defined radio and a software change affecting its authorized operating characteristics	• Preserve the applicable change filing and supporting test evidence identifying the changed software state under the existing equipment authorization.
United States · CLIA, 42 CFR Part 493	Clinical-laboratory test system and historical procedure used for a patient result	• Keep dates when procedures were discontinued, records of the analytic system, instrument outputs, performance specifications, calibration and control records to keep the historical test state behind a patient result identifiable.
United States · Department of Education, 34 CFR §75.623(e)	Scientific data informing a covered peer-reviewed publication produced under Department funding	• Retain and make the scientific data informing the published findings publicly available within the regulation's scope and exceptions.
United States · Department of Energy Public Access Plan / DMSP requirements (funding/contract condition)	Digital scientific data produced with DOE funding and the plan governing their preservation and sharing	• Require researchers funded by DOE to preserve and share digital scientific data under an approved plan that identifies standards, repositories, preservation methods and the tools, software or code needed to validate or replicate the work.
United States · 42 U.S.C. §19053 · Research Reproducibility and Replicability	Data, software and code from NSF-supported projects and data underlying their published findings	• Require each research proposal to record, in a machine-readable data management plan, how project data, software and code will be archived and kept publicly accessible. • Require NSF to establish policies and resources for depositing data underlying published findings in trusted repositories at publication, where practicable. • Preserve the applicable federal privacy, intellectual property and security boundaries.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · Department of State 652.239-804 Minimum Threshold Transparency Provision (Deviation) (MAR 2026) <i>(public-solicitation implementation)</i>	Proposed LLM's minimum model/system/data transparency record	• Require every offeror in the LLM supply chain to identify its role, disclose available developer cooperation and information limits, and provide the model, system, data, risk, limitation and evaluation records reasonably available for the Government's compliance review.
United States · Department of State 652.239-805 Enhanced LLM Transparency Provision (Deviation) (MAR 2026) <i>(public-solicitation implementation)</i>	Training, configuration, evaluation and modifications by third parties relevant to the proposed LLM's behavior	• Require enhanced disclosure of training, configuration, prompts and filters affecting behavior, together with testing and modifications by third parties, and flow the same duties to every relevant subcontractor, vendor and third party handling the LLM.
United States · DOE/NNSA Acquisition Letter 2026-05 + Ensuring Unbiased AI Principles (APR 2026) clause <i>(federal acquisition policy, mandatory contract clause)</i>	Covered AI system's record of material changes, evaluation, and modifications by third parties	• Require contractors to disclose material changes to model architecture, known sources of training data and safety or moderation settings. • Keep evaluation methods and outcomes and provide re-verification results. • Describe modifications to third-party models and their expected behavioral or risk effects.
United States · SHARE IT Act, Pub. L. 118-187	Federal custom-developed source code and its repository/version history	• Store covered custom-developed code in a qualifying repository, use repository and version-control practices that track changes, and make qualifying metadata publicly accessible, including repository status, contract reference where applicable and repository link.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · Federal Records Act, 44 U.S.C. §3101	Federal agency functions, policies, decisions, procedures and essential transactions	• Require each agency head to make and preserve adequate documentation of agency functions, policies, decisions, procedures and essential transactions. • Keep enough information to protect Government rights and the rights of people directly affected by agency activity.
United States · EPA GLP, 40 CFR Part 160	Nonclinical study submitted to or considered by EPA	• Preserve the protocol and revision history, historical SOPs, equipment calibration and maintenance records, test-substance batch identity, raw data, specimens, calculations, final report and other records needed to interpret or evaluate the study.
United States · FAA AC 20-115D + Order 8110.49A (certification implementation / means of compliance)	Airborne software baseline and certification test evidence	• Preserve the approved software baseline, test setup, hardware and software configuration, usage history and change-impact record to keep certification evidence tied to the exact state tested.
United States · FDA 21 CFR §814.39	Device approved through a PMA and a later change affecting safety or effectiveness	• The PMA supplement must identify each change for which approval is sought and explain the reason for it. • The supporting record is limited to the change, requiring the information needed to support the change while incorporating previously submitted information by reference where it remains relevant.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · FDA CGMP, 21 CFR §§211.180 and 211.194	Drug-manufacturing laboratory test supporting conformity/release	• Keep complete data from all tests needed to establish conformity, including sample source/lot, testing method and method suitability, original instrument graphs/charts/spectra, calculations, results, reference standards and calibration records. • Where a method is modified, retain the reason for the modification and data verifying that the modified method is at least as accurate and reliable under the actual conditions of use.
United States · FDA GLP, 21 CFR Part 58	Nonclinical study supporting or intended to support an FDA research or marketing permit	• Treat original records or exact copies, including computer and automated-instrument data, as raw data when needed to reconstruct and evaluate the report. • Retain the study identity, protocol, historical procedures, calibration records, observations, calculations, specimens and final report. • When original data change, preserve the original entry and record the reason, date and responsible person.
United States · NASA SMD SPD-41a (agency/award policy)	Scientific data, software and computational configuration supporting a publication	• Archive data and software supporting a publication with durable identifiers. • Preserve the input and configuration files needed to reproduce or validate the result. • Treat a live source-code repository alone as insufficient for the archival duty.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · National Science Foundation current sharing policy / DMSP requirements <i>(funding policy / award requirement)</i>	Research products from work funded by NSF and data supporting a publication	• Require NSF award recipients to share data supporting a publication at publication and plan metadata, access, reuse, redistribution, archiving and preserved access for the data, software and other covered research products.
United States · NIDDK DMS implementation guidance <i>(NIH institute implementation)</i>	Novel analytic pipeline/workflow used in a funded analysis	• Document and share a novel analytic pipeline or workflow together with the relevant code, software, tools and context needed for another investigator to recreate the analysis.
United States · NRC 10 CFR §50.59	Licensed nuclear-facility design/procedure/test state and a later change	• Preserve the approved safety basis and a written evaluation of each covered change. • Require an NRC licence amendment before implementation when the change crosses the regulatory thresholds.
United States · PCAOB AS 1215 <i>(SEC-approved auditing standard)</i>	Audit engagement and the evidentiary basis for significant conclusions	• Document the procedures performed, evidence and results obtained, and conclusions reached in enough detail for an experienced auditor with no previous connection to the engagement to understand what was done and why.
European Union · AI Act, Regulation (EU) 2024/1689, Arts. 11–12, 18–19 + Annex IV <i>(enacted, Chapter III duties future-applying under Art. 113)</i>	High-risk AI system's version, design/development, testing/validation, changes and operational logs	• Maintain current technical documentation identifying the system version, its relation to earlier versions, the development and testing state, and relevant later changes. • Enable automatic event logging across the system lifetime and retain the required logs for at least six months unless other law requires longer. • Keep the required technical, quality, change and conformity records available to authorities for 10 years after the system is placed on the market or put into service.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Arts. 53–54 + Annex XI	General-purpose AI model and its documented release, development, training, testing and evaluation state	• Draw up and keep current technical documentation covering release, architecture, licence, training methods, data provenance, compute, testing, and evaluation. • Keep the required copy available to authorities for ten years and publish the required training-content summary.
European Union · Cyber Resilience Act, Regulation (EU) 2024/2847, Art. 13(2)–(3), (7)–(8), (13) (enacted, these manufacturer lifecycle duties generally apply from 11 December 2027)	Product cybersecurity risk assessment, vulnerabilities and support period basis throughout the product lifecycle	• Document the cybersecurity risk assessment and update it as appropriate during the support period. • Systematically document relevant cybersecurity aspects and known vulnerabilities and update the assessment where applicable. • Record in the technical documentation the information used to determine the support period. • Keep the technical documentation and EU declaration of conformity available to market surveillance authorities for at least 10 years after placing on the market or for the support period, whichever is longer.
UNECE · UN Regulation No. 156	Vehicle software or configuration state relevant to approval and a subsequent software update	• Maintain software version or RXSWIN identity together with update and change records. • Record interdependencies, target configuration and verification or validation evidence resolving the state before and after an update. • Keep enough evidence to determine whether prior type-approval evidence remains applicable.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · Data Governance Act, Regulation (EU) 2022/868, Art. 20	Data held by a recognized data-altruism organization and its downstream use	• Keep full and accurate records identifying who was allowed to process the data, when, for how long and for what declared purpose, together with the required activity reporting.
European Union · Digital Services Act + Delegated Regulation (EU) 2024/436	Very large platform/search-engine systemic-risk and audit evidence	• Maintain records, data and methodological information sufficient to support the provider's risk assessments and the audit conclusions required by the DSA audit framework.
OECD · Principles of Good Laboratory Practice / Mutual Acceptance of Data framework <i>(intergovernmental framework, domestic force depends on implementation)</i>	GLP-regulated study accepted across participating systems	• Maintain unique study identity, the study plan and methods, computer-system documentation, raw data, samples, calculations, quality assurance and calibration records, and the final report. • Archive the study materials in a form sufficient to reconstruct the work.
Australia · TGA substantial-change conformity assessment process <i>(official implementation guidance for statutory certificate obligations)</i>	TGA conformity assessment certificate and a later device/QMS change	• Preserve the certificate state, change details and analysis, supporting validation or technical-file evidence and any updated certificate. • Retain the change details and analysis for later review even when notification is not required.
Canada · Medical Devices Regulations s.34 + 2026 significant-change guidance <i>(binding regulation + nonbinding implementation guidance)</i>	Licensed Class III/IV device and the evidence used to classify/support a later change	• Preserve the current licensed state, the change description and comparison, risk assessment and relevant software revision history. • Retain any evidence bridging from earlier submissions and the verification or validation evidence supporting the amendment or conclusion that the change is not significant.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · Benchmarks Regulation 2016/1011	Financial benchmark and the process used to determine it	• Preserve benchmark methodology, input data, records documenting the control framework and other information required for later reconstruction and audit of the benchmark determination.
European Union · Clinical Trials Regulation 536/2014, Arts. 56–58	Clinical trial and its trial master file	• Keep trial information and essential documents reportable, interpretable and verifiable, with the clinical-trial master file retained for the prescribed period.
European Union · Commission Regulation (EU) 2023/1670, Arts. 4–5 and Annex II (binding, applicable from 20 June 2025)	Historical battery state and product model conformity evidence for covered smartphones/slate tablets	• Record in the device settings or another location accessible to end users the battery manufacture date, first-use date, full charge/discharge cycle count and measured state of health. • Maintain conformity technical documentation containing the required product information, calculation details/results, accuracy assessment where applicable and equivalent model identifiers for the market surveillance framework.
European Union · Directive 2014/26/EU collective-rights-management regime	Creative work or performance whose rights are managed collectively, together with downstream exploitation and accounting	• Require users of managed rights to supply the usage information needed to collect and distribute revenue to rightholders. • Require collective management organisations to keep their membership records current and provide annual accounts linking attributed revenues and payments to the rights managed and types of use.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · Directives 2004/10/EC and 2004/9/EC	GLP safety study of a chemical product within Member-State implementation of OECD GLP	• Require covered laboratories to retain the GLP study plan, raw data, samples or specimens, quality assurance and calibration records, final report and historical procedures. • Require laboratories submitting results to certify GLP conformity.
European Union · In Vitro Diagnostic Medical Devices Regulation 2017/746	Approved IVD QMS/device range or type and a later planned change	• Keep the approved QMS/type/certificate state, the information about the planned change, notified body assessment and decision, and any certificate/report supplement or record of a new assessment identifying the changed state.
European Union · Medical Devices Regulation 2017/745	Approved medical-device QMS/device range or type and a later planned change	• Keep the approved QMS/type/certificate state, the information about the planned change, the notified body assessment and decision, and any certificate or report supplement to identify which changed state received continuing approval.
European Union · ecodesign regulations for specific products amended or coordinated by Regulation (EU) 2021/341	Conformity record for a product model used in ecodesign assessment and market surveillance	• Maintain technical documentation for each model with declared parameter values, calculation details, lists of equivalent models and model identifiers. • Include the manufacturer's accuracy assessment and declaration identifying the model where the regulation relies on calculation or extrapolation from another model.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
United Kingdom · PSTI Act 2022 §§9, 12, 15, 20 + Security Requirements Regulations 2023, regs. 8–9 and Schedule 4	Statement of compliance and records of product-security compliance failures or investigations	• Require the statement of compliance to identify the product, manufacturer, applicable representative, compliance declaration, defined support period and required issue information. • Retain the statement of compliance for the longer of 10 years or the product's defined support period. • Retain records of compliance failures and investigations for at least 10 years.
Virginia · Voting System Certification Standard 3.1 (2026) (official State certification standard)	The exact voting-system version and configuration submitted for certification	• Require a technical data package (TDP) change log from the previously certified version plus software design, module and data format records and the operating-system version, support date, security patches, SHA256 hash and modification state.
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) (official State procurement terms)	Escrowed software package and later component versions	• Require the contractor to update, enhance or modify the Source Code Escrow Package promptly upon each release of a new component version and deliver current custom source code and documentation to the State at contract termination.
Arizona · 2025 Elections Procedures Manual + A.R.S. §16-445 (current force-of-law election procedures)	Machine-readable election program and later changed states	• File a machine-readable copy of the election program with the Secretary before each election. • File a new copy within 48 hours after any later program change. • Preserve the hash recorded at certification to identify the software state used in the election.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
Wisconsin · Wis. Stat. §5.905	Escrowed voting-system software components and later changed states	• Require software components selected for review and verification to enter Commission escrow within ninety days after approval and require any subsequent change in those components to be deposited within ten days, preserving the changed official software state.
Pennsylvania · ES&S EVS 6.5.0.0 current certification conditions (official current system-specific certification)	The certified source-code version and later software updates, improvements and enhancements	• Maintain a true, accurate and complete source-code deposit for the certified release and keep the escrow at the most current source state by depositing later updates, improvements and enhancements throughout the certification and contract period.
New Jersey · N.J.A.C. 15:10-7.5–7.6	Approved e-poll-book software and settings state and later operational changes	• Require recurring integrity checks plus software installation, removal and change logs, logs of failures affecting system integrity, access logs and adjustment logs that identify software or settings changes, timing, affected files and responsible actors.
Wyoming · Title 22 electronic-voting-system certification statute (2026)	Certified voting-system source and version state	• Require certification records identifying the checksum associated with escrowed source code, source-code updates, the system history and county inventory information recording software and firmware versions.
United States · 37 CFR §§1.805, 1.807 · biological-material patent deposit	Deposited biological material and any compliant replacement or supplement	• Record the depository, depositor, deposit date, accession identity and viability evidence. • For a replacement or supplement, record the reason, new accession, date, depository and corroborated identity. • Give a compliant replacement a rebuttable presumption of identity with the original in qualifying Office proceedings.

Table 3.4-7. Examples of Legal and Institutional Precedents Requiring Records Identifying the Historical Object, Inputs, Provenance, and Changes (continued) · See Section 3.4.3

Precedent	Object	Requirement
Nebraska · Rev. Stat. §8-1,141	Source-code version used for each controllable electronic record held in custody	• Require the financial institution and customer to agree in writing on the source-code version used for each record and resolve ambiguity in the written version and treatment agreement in favor of the customer.
United States · 17 CFR §240.17a-4	Electronic broker/dealer records kept with a history of changes, and reports affected by later system changes	• Preserve a complete time-stamped audit trail of record creation, modification and deletion sufficient to support authenticity and recreate the original record. • If later system changes impair reproduction of earlier reports, preserve an explanation of the relevant changes or the historical report-generation parameters.
United States · FD&C Act §524B / 21 U.S.C. §360n-2	Software composition of a covered cyber device	• Require the sponsor to provide FDA a software bill of materials identifying commercial, open-source and off-the-shelf software components contained in the cyber device.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification · See Section 3.4.3

Precedent	Object	Requirement
California · DMV autonomous-vehicle regulations, 13 CCR §§228.08, 228.10, 228.24(c)	Safety evidence supporting original or amended AV deployment and remediation after DMV restricts, suspends, or revokes deployment authority	• Require DMV or its expert reviewer to assess the safety case and corrective evidence for the affected driverless deployment, and keep deployment stopped until DMV verifies the correction.
California · CIRM data-sharing / DSMP conditions (state funding condition)	Result produced with CIRM funding and the data and metadata needed to replicate it	• Require the funded recipient to share raw and final data and metadata, identify the code or software needed to replicate or reuse the funded result, and explain how to access any specialized tools.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
California · CARB Cap-and-Trade offsets, 17 CCR §§95976, 95977.1	Offset-project emissions-reduction calculation	• Require the project record package to preserve inputs, assumptions, raw measurements, quality-assurance, calibration, and change records and require an accredited verifier to independently recalculate or review the claimed emissions reduction.
California · CARB Low Carbon Fuel Standard	Carbon-intensity result generated with a named CA-GREET model version	• Require retention of the named model version, source data, and calculation records and permit the accredited verifier to recalculate the carbon-intensity result using the controlled operational data needed to test the version effect.
California · CARB Mandatory Greenhouse Gas Reporting, 17 CCR §95131	Reported emissions estimate and its source and calculation chain	• Require the reported result to remain connected to original data, instrument and calibration evidence, methodology, and reproducible calculations, and require an accredited verifier to trace sources, review the methodology, and recalculate the result.
California · Department of Insurance, 10 CCR §§2648.4–2648.5 + the Department’s rate-filing (PRID) system	The catastrophe model version relied on in an insurance rate application and the claimed reliability of the model	• Require identification of the catastrophe model and version together with the reliability materials used in public and expert rate review. • Require a new determination after a substantive model change.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
California · Department of Cannabis Control, 4 CCR §§15712.1–15712.2	The prescribed cannabinoid test method for dried flower and non-infused pre-rolls before regulatory use	• Require the laboratory to submit a method-verification report to the Department of Cannabis Control before regulatory use, demonstrating repeatable performance and including calibration data, raw and instrument data, reference-material results, detection and quantitation limits, and quality-control evidence.
California · Environmental Laboratory Accreditation Program, 22 CCR §64802.05	Accredited environmental-laboratory quality system and analytical competence	• Require auditable records identifying the methods, instruments, and historical analytical state and subject the laboratory's competence and quality system to independent accreditation assessment.
United States · 2 CFR §200.315(e)	Qualifying federally funded research data underlying published findings later relied on in legally effective federal action	• When the regulation's reliance and FOIA conditions are met, require the recipient to provide research data needed for independent validation for agency release through FOIA.
United States · 47 CFR §§2.1043 and 2.944	Software-defined radio after a software/configuration change affecting approved operating characteristics	• Before loading software that changes an approved software-defined radio's operating characteristics under the existing grant, require the applicable change filing, supporting test evidence and FCC acknowledgement.
United States · Department of Education, 34 CFR §75.623(e)	Scientific data informing a covered peer-reviewed publication produced under Department funding	• Require the scientific data informing the published findings to be publicly available within the regulation's scope, enabling independent review of the evidentiary basis.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · Department of Energy Public Access Plan / DMSP requirements (funding/contract condition)	Research result and underlying digital scientific data produced with DOE funding	• Require researchers funded by DOE to preserve and share digital scientific data under an approved plan that identifies standards, repositories, preservation methods and the tools, software or code needed to validate or replicate the work.
United States · DOE/NNSA Acquisition Letter 2026-05 + Ensuring Unbiased AI Principles (APR 2026) clause (federal acquisition policy, mandatory contract clause)	Evidence of the covered AI system's compliance and performance before deployment and during operation	• Require validation before operational deployment and transparency documentation supporting independent testing and drift monitoring. • Require documented evaluations and re-verification while making evaluation methods and outcomes available to the Government on request. • Provide enterprise controls that let the Government audit interactions and monitor model performance and safety safeguards over time.
United States · EPA GLP, 40 CFR Part 160	Nonclinical study submitted to or considered by EPA	• Require independent quality-assurance review of study conduct and the final report against the methods, SOPs and raw data. • Give EPA access to inspect and copy the required records and specimens for checking the submitted study against its underlying evidence.
United States · FAA AC 20-115D / DO-178C (means of compliance)	Previously approved airborne software reused or modified for a later certified application	• Give certification credit to qualifying unmodified legacy software only when the defined compatibility and usage-history conditions hold. • Require change-impact analysis and verification proportionate to the change for modified software.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · FAA Order 8110.49A (certification implementation)	Software evidence supporting certification credit and the tested software/hardware baseline	• Tie certification credit to the conforming software, hardware and test baseline, and require FAA review when a later change affects the configuration on which the evidence depends.
United States · FDA 21 CFR §814.39	Device approved through a PMA after a change affecting safety or effectiveness	• Qualifying changes generally require FDA review and approval before implementation. • The supplement can reuse previously submitted information by reference where it remains relevant, while the changed portion must be supported by the evidence for the change FDA requires.
United States · 26 U.S.C. §7612	Tax-related software source code needed to verify a return item where records and executable-code evidence are insufficient	• Permit access to specifically identified source-code portions only when books, records, executable code and associated data are insufficient for reasonable verification and the verification need outweighs trade secret disclosure risk. • Apply bounded safeguards for use, access, copying, confidentiality, return or deletion. • For required on-premises analysis, require the owner to provide the necessary equipment or materials.
California · Elections Code §19103 + 2 CCR §20670	Voting-system source-code package and deployed approved software version	• Permit the Secretary of State, for cause, to audit escrowed source-code materials to verify their contents and provide reasonable statutory access for verification that software deployed on voting equipment is identical to the approved version.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
North Carolina · G.S. §163-165.7 + 08 NCAC 04 .0308	Escrowed voting-system source code before certification and during authorized review	• Require the State Board or its independent expert to review all vendor source code before certification. • Give specified government, county and political-party actors controlled access to examine escrowed information under confidentiality and qualification safeguards.
New York · 9 NYCRR §6209.6 + State Board verification practice	Certified source-code package and software actually running on voting equipment	• Require independent expert review before certification to verify that the software matches its documentation and that the documentation is sufficient to install, validate, operate and maintain the system. • Compare deployed software with the official escrow copy through State acceptance and hash verification and annual machine hash checks.
UNECE · UN Regulation No. 156	Vehicle software or configuration state relevant to approval after a software update	• Require the manufacturer's software-update management system to assess the change and the continued relevance of prior approval evidence. • Require the approval authority or technical service to review the changed state and determine whether existing evidence remains sufficient, further testing is required, approval can be extended, approval is refused, or a new approval is required.
Minnesota · Stat. §206.57, subd. 6	Voting-system source code supplied for certification review	• Require the vendor to provide source code to the Secretary of State for certification review. • Permit a qualified independent evaluator to examine the code for conformity with its representations and for defects, subject to trade secret confidentiality.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
South Carolina · Code §§7-13-1330(I), 7-13-1620(F)	Escrowed voting-system source code when Commission examination is necessary	• Require the escrowed source code to be available to the State Election Commission when the Commission determines that examination is necessary, preserving technical verification by the regulator independent of vendor assertion.
Michigan · MCL §168.797c	Escrowed voting-system source code	• Permit the Secretary of State or authorized agent to use the escrowed source-code information for analyzing and testing the software, while preserving proprietary-information confidentiality.
Illinois · 10 ILCS 5/23-15.1	Voting-system computer code needed for review and verification	• Require the State Board of Elections to determine which software components are necessary to enable review and verification and provide controlled access for authorized persons in election contests or under court order.
Colorado · 8 CCR 1505-1 voting-system certification rules	Voting-system software, firmware and technical data submitted for certification	• Assess compliance through State certification using code analysis, formal testing and documentation review before the system receives or retains certification.
Georgia · Ga. Comp. R. & Regs. 590-8-1-.01	Voting-system hardware, firmware, software and certification package	• Assess voting-system compliance through code analysis, formal testing and documentation review, including independent Election Assistance Commission (EAC) testing and Georgia certification testing. • Require acceptance testing to show that the delivered and installed system is identical to the system certified by the State.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
Missouri · 15 CSR 30-10.020	Electronic voting system submitted for approval and its certification evidence	• Require certification from an independent testing authority approved by the Secretary as a prerequisite to approval and preserve State access to records for the final build and certification-test results available under the escrow release conditions.
Connecticut · RCSA §9-241-36	Delivered DRE voting-system hardware and software before contractual acceptance	• Require municipal acceptance testing to determine whether delivered hardware and software comply with State and municipal requirements and perform consistently with the same equipment's State qualification and certification testing.
Texas · Election Code §129.023(c-1)–(c-2) (effective 1 September 2025)	Source-code state of voting-system equipment used in a current election	• As part of the election's logic-and-accuracy test, require the general custodian to demonstrate on a representative sample that voting-system source code has not been altered. • Define the sample as ten devices or five percent of each equipment type, whichever is fewer.
United States · FDA CGMP, 21 CFR §§211.180 and 211.194	Drug-manufacturing laboratory evidence supporting conformity/release	• Require second-person review of laboratory records for accuracy, completeness and compliance. • Make the records available for FDA inspection and reproduction for checking the reported result against original laboratory evidence.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · FDA GLP, 21 CFR Part 58	Nonclinical study supporting or intended to support an FDA research or marketing permit	• Require an independent Quality Assurance Unit to inspect the study and review the final report against the study methods, procedures and raw data. • Permit FDA to inspect and copy the required records and specimens. • Exclude a study from FDA consideration in support of a research or marketing permit application when the testing facility refuses inspection.
United States · NASA SMD SPD-41a (agency/award policy)	Scientific result supported by data, software and computational configuration	• Archive data and software supporting a publication with durable identifiers. • Preserve the input and configuration files needed to reproduce or validate the result. • Treat a live source-code repository alone as insufficient for the archival duty.
United States · National Science Foundation current sharing policy + GEO implementation (funding policy / award requirements)	Result from work funded by NSF and the research products available for independent reuse or reproduction	• Place final datasets, metadata, derived products, relevant software or code and covered samples in long-lived public archives for reproduction or replication by other researchers.
United States · NIDDK DMS implementation guidance (NIH institute implementation)	Novel analytic pipeline/workflow underlying a funded analysis	• Document and share a novel analytic pipeline or workflow together with the relevant code, software, tools and context needed for another investigator to recreate the analysis.
United States · NRC 10 CFR §50.59	Nuclear-facility design/procedure/test state after a change	• Preserve the approved safety basis and a written evaluation of each covered change. • Require an NRC licence amendment before implementation when the changed facility, procedure or test crosses the regulatory thresholds.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · OMB Memorandum M-25-22 §4(d)(iii)(E) <i>(executive-branch acquisition policy, non-NSS scope)</i>	Performance, risk and effectiveness of a covered acquired AI system or service during contract performance	• Require AI contracts to let the agency monitor system performance, risk and effectiveness during contract performance. • Let the agency use its own evaluation data and obtain the access and time needed for independent testing. • Require testing records detailed enough for verification or reproduction.
United States · PCAOB AS 1215 <i>(SEC-approved auditing standard)</i>	Audit engagement and the evidence supporting significant conclusions	• Document the procedures performed, evidence and results obtained, and conclusions reached in enough detail for an experienced auditor with no previous connection to the engagement to understand what was done and why.
European Union · AI Act, Regulation (EU) 2024/1689, Annex VII §§4.3–5.3 + Art. 74(12)–(13) <i>(conformity/surveillance architecture, high-risk obligations staged under Art. 113)</i>	Evidence and technical access needed for external conformity assessment or market surveillance review of a high-risk AI system	• Give notified bodies access to training, validation and testing datasets and authority to require or conduct tests. • Where lesser means are insufficient, permit access to the trained model and relevant parameters for conformity assessment. • Give market surveillance authorities access to documentation, datasets and technical tools needed for review. • Permit source-code access only on a reasoned request when it is necessary for conformity assessment and testing or auditing based on other materials has proved insufficient.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Art. 36(3), (9)	High-risk AI systems whose prior notified body ceases activity or loses its designation	• After planned cessation, keep certificates valid for nine months only if another notified body confirms that it will assume responsibility. • Require the successor notified body to complete a full assessment before issuing new certificates. • After designation withdrawal, permit bounded temporary certificate validity only where the competent authority identifies no relevant risk and another notified body assumes responsibility. • Require the successor assessment within the Regulation's stated period, subject to the bounded extension rule.
European Union · AI Act, Regulation (EU) 2024/1689, Arts. 91–92	General-purpose AI model subject to Commission compliance or systemic-risk evaluation	• Permit the Commission to require Article 53 or 55 documentation and additional information from the provider. • Permit the AI Office to evaluate a general-purpose AI model when the Article 92 trigger is met. • Permit the Commission to appoint independent experts to carry out evaluations. • Permit model access through APIs or other technical means, including source code where requested, and require the provider or representative to supply that access.
European Union · Digital Services Act + Delegated Regulation (EU) 2024/436	Compliance/systemic-risk controls of a very large online platform or search engine	• Require an independent audit using the provider's required evidence, data and prescribed audit methodology to test compliance rather than relying on provider assertion alone.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
OECD · Principles of Good Laboratory Practice / Mutual Acceptance of Data framework <i>(intergovernmental framework, domestic force depends on implementation)</i>	GLP-regulated study and final report	• Require a Quality Assurance Programme independent of study conduct to inspect the study and review the final report against the study plan, procedures and raw data. • Retain the study materials needed to reconstruct the work in the GLP archive.
Australia · TGA substantial-change conformity assessment process <i>(official implementation guidance for statutory certificate obligations)</i>	TGA-certified medical device after a substantial QMS/product/design change	• Require notice and a substantial-change application before supplying products affected by the change. • Allow TGA to determine the assessment required before implementation. • Treat new software algorithms requiring new design verification and validation as a current guidance example of substantial change.
Canada · Medical Devices Regulations s.34 + 2026 significant-change guidance <i>(binding regulation + nonbinding implementation guidance)</i>	Licensed Class III/IV device after a proposed software/device change	• Require a licence-amendment application and amended licence before marketing or selling a changed Class III or IV device when the change is reasonably expected to affect safety or effectiveness. • Treat changes to software function, performance, risk controls, infrastructure, architecture and algorithms, including multiple AI or ML model or training changes, as potentially significant under current guidance. • Treat significant changes covered by an authorized PCCP as previously reviewed. • Return changes outside the PCCP to the ordinary change analysis.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · Benchmarks Regulation 2016/1011	Financial benchmark determination, methodology and input data	• Keep benchmark methodology, input data and control records available for reconstruction, audit and supervisory review of how the benchmark was produced.
European Union · Clinical Trials Regulation 536/2014, Arts. 56–58	Clinical trial conduct, data quality and trial master file	• Maintain the trial master file to permit competent authorities to verify trial conduct and data quality through inspection of the essential record.
European Union · Commission Regulation (EU) 2023/1670, Art. 5 + Annex IV (binding, applicable from 20 June 2025)	Product model's compliance with the Regulation's ecodesign requirements and declared values	• Require Member State market surveillance authorities to verify product units and declared values under the Annex IV testing procedure. • Reserve verification tolerances for authority testing and prohibit manufacturers from using them to inflate technical documentation values or claimed performance.
European Union · Directives 2004/10/EC and 2004/9/EC	GLP study and Member-State GLP compliance status	• Require Member States to verify GLP compliance through test-facility inspections and study audits/checks of the planning, performance, recording and reporting of the study against the underlying study records and quality system.
European Union · In Vitro Diagnostic Medical Devices Regulation 2017/746	Approved IVD QMS/device range or type after a substantial change or other change relevant to conformity	• Require notified body assessment of a substantial change or other change relevant to conformity and whether the changed QMS, device or type remains compliant. • Record approved changes through the relevant certificate supplement. • Require a new conformity assessment when a change crosses the approved intended-purpose or use boundary.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · Medical Devices Regulation 2017/745	Approved medical-device QMS/device range or type after a substantial change or other change relevant to conformity	• Require the notified body to assess substantial QMS or device-range changes, determine whether further audits are needed and verify continued compliance. • Record approval through a certificate supplement. • Require notified body approval for changes to an approved type that may affect conformity or prescribed conditions. • Require a new conformity assessment for broader intended-purpose or use changes where the Regulation requires it.
European Union · ecodesign regulations for specific products in the 2021/341 family	Compliance of a covered product model and equivalent models with declared ecodesign requirements	• Require market surveillance authorities to compare declared values with measured values and test product units under the applicable verification procedure. • Treat circumvention based on detecting testing or failure of the verification conditions as non-compliance where the applicable regulation provides. • Prohibit use of verification tolerances to inflate technical documentation values or communicated performance.
Virginia · Voting System Certification Standard 3.1 (2026) (official State certification standard)	Voting-system software and security state submitted for certification	• Require source-code evaluation, a recent accredited independent third-party penetration analysis using active exploitation, federal voting system test laboratory (VSTL) evidence and State hash testing that confirms application software matches the federally certified version.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) (official State procurement terms)	Custom software source code and source-code completeness	• Give the State the right to audit custom source code and corresponding documentation, require the contractor to provide tools needed to view the source, and require complete compile and link instructions in escrow for verification of source-code completeness.
Arizona · 2025 Elections Procedures Manual + A.R.S. §§16-445, 16-602(J) (current force-of-law election procedures)	Election-management software used in the election and escrowed source code after a qualifying discrepancy trigger	• Require pre-election hash comparison against NIST or Secretary certification records. • After a qualifying full jurisdiction-wide hand count, require an unaffiliated special master to review the escrowed source code and report discrepancies publicly for review of whether certification should continue.
Wisconsin · Wis. Stat. §5.905	Escrowed software components needed to review and verify tabulation accuracy	• Require the Commission to identify software components necessary for review and verification. • Preserve controlled access to those components in qualifying recount proceedings. • Protect proprietary information through confidentiality conditions without eliminating review access.
Pennsylvania · Voting System Security Standard (18 Sept. 2026) (official State certification standard)	Voting-system source, trusted build and security state submitted for Pennsylvania certification	• Require an independent testing team to receive the Technical Data Package, source code, user documentation, configured devices and other materials needed for investigation. • Require security and penetration testing of the submitted system. • Preserve a clean trusted-build copy for reinstallation where needed.

Table 3.4-8. Examples of Legal and Institutional Precedents Providing Independent Reproduction, Recalculation, Audit, or Verification (continued) · See Section 3.4.3

Precedent	Object	Requirement
New Jersey · N.J.A.C. 15:10-7.5–7.6	E-poll-book critical files, source code and system/security state submitted for approval	• Require external verification capability for critical files, DGE testing and examination of the system and source materials, and New Jersey Cybersecurity and Communications Integration Cell (NJCCIC) security review before the Secretary decides approval or rejection.
Alabama · Code §17-11-42(c)(2) + Admin. Code r. 820-2-10-.06	Secure remote voting application used for covered overseas absentee voting	• Require the application to support auditing of its source code and to be executable on a clean computer containing at least a baseline operating system and no other software or driver, making review of source code and isolated execution part of the minimum security criteria.
Nebraska · Rev. Stat. §8-1,141	Financial institution's custody controls for controllable electronic records	• Require the institution to engage an independent public accountant for the prescribed custody examination, send results to the banking director and report material discrepancies to the director within one day.
Switzerland · E-ID Act (BGEID) Art. 12(3)–(4) (enacted and approved, future-applying, trust-infrastructure commencement expected in 2027)	Security of the federal e-ID trust infrastructure and its disclosed software	• Require coordinated vulnerability-disclosure guidelines and require the Federal IT office to review the trust infrastructure's security regularly with suitable third parties, creating formal external security review alongside source-code disclosure.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears · See Section 3.4.3

Precedent	Object	Requirement
California · DMV autonomous-vehicle regulations, 13 CCR §227.54(b), §228.08(a)(9) United States · 17 U.S.C. §108	Autonomous-technology sensor state for vehicle functions controlled by the ADS immediately before a collision Copyrighted works in qualifying library or archive collections	• Require the permit holder to preserve at least the preceding 30 seconds of read-only event data until extraction and retain the extracted record for three years. • Permit a qualifying library or archive to make preservation or replacement copies when an eligible work is unpublished, damaged, deteriorating, lost, stolen or stored in an obsolete format, subject to the statute's conditions.
United States · 37 CFR §201.40(b)(19) <i>(current 2024–2027 triennial exemption)</i>	Qualifying lawfully acquired video game after required authentication-server access ends or after commercial unavailability	• Permit an eligible library, archives or museum to copy and modify a qualifying game program as needed to preserve the game in playable form. • Permit bounded use of console-operating programs where needed for preservation. • Preserve the exemption's institutional-premises and noncommercial limits.
United States · 37 CFR §201.40(b)(20) <i>(current 2024–2027 triennial exemption)</i>	Lawfully acquired non-game software or software-dependent digital material no longer reasonably available commercially	• Permit an eligible library, archives or museum, without direct or indirect commercial advantage, to circumvent access controls solely for lawful preservation of the software or digital material that depends on the software for access.
United States · EPA GLP, 40 CFR Part 160	Protocol, raw data, specimens, procedures, equipment records and final report for a covered pesticide study	• Require the covered test facility to retain the protocol, raw data, specimens, procedures, equipment records and final report for the applicable preservation period.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · Federal Records Act, 44 U.S.C. §§3105–3106	Federal records exposed to removal, loss, deletion, erasure or destruction	• Require agency safeguards against removal or loss and prohibit destruction outside the statutory disposal rules. • Treat actual, impending or threatened unlawful deletion, erasure or other destruction as a statutory trigger for Archivist notification and recovery or redress action, including Attorney General action where required.
United States · 42 U.S.C. §19053 · Research Reproducibility and Replicability	Data, software and code developed in research supported by NSF	• Require every research proposal to include a machine-readable data management plan describing how the awardee will archive and preserve public access to data, software and code developed by the project. • Require NSF to set criteria for trusted repositories, support their development and maintenance, and provide a public means of locating project data, software and code. • Require NSF to support deposit of data underlying published findings at publication where practicable and incentivize deposit of data, software and code.
United States · 15 U.S.C. §§3704b-2(a), 3710(d)(6) · NTIS	Unclassified federally funded scientific or technical information eligible for public dissemination, expressly including computer software	• Require Federal executive departments and agencies to transfer covered information, including computer software, to NTIS in a timely manner for dissemination. • Require NTIS to maintain a permanent archival repository and clearinghouse for collection and dissemination of nonclassified scientific, technical and engineering information.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
California · Elections Code §19103 + 2 CCR §§20612, 20622, 20635–20644, 20662	Voting-system, ballot-marking-system and election-management-system source code before and during operational use	• Require covered source code to enter approved independent escrow before use. • Retain used versions for the retention period applicable to election materials in secured, backup-capable facilities. • Transfer deposits to compliant custody if an escrow facility loses approval.
North Carolina · G.S. §163-165.9A + 08 NCAC 04 .0302	Software relevant to functionality, setup, configuration and operation of a North Carolina voting system	• Require the covered software package, including revisions or modifications, to be held by an independent escrow agent approved by the State as part of voting-system approval and certification.
New York · 9 NYCRR §6209.6	Complete certified voting-system software relevant to functionality, setup, configuration and operation	• Require the State Board to escrow a complete copy of the certified software package and require updated versions to be deposited whenever a required escrow item changes.
Ohio · Rev. Code §3506.05(I)(1) + Admin. Code 111:3-9-04	Voting-system source code and related documentation within the regime governing certification and continued certification	• Require vendors or software developers to place all covered source code and related documentation in escrow with an independent agent approved by the Secretary of State and keep the escrow current through periodic updates.
South Carolina · Code §§7-13-1330(I), 7-13-1620(F)	Voting-system source code before operational election use	• Require all source code to be placed in escrow before the covered system is used in elections, and require all source-code updates to enter escrow as they occur.
Michigan · MCL §168.797c	Voting-system source code and subsequent revisions or modifications	• Require source code and every later revision or modification to be placed in an escrow account held by a third party approved by the Secretary.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
Illinois · 10 ILCS 5/23-15.1	Voting-system computer code, including ballot-counting source code and related human-readable software structures	• Require all voting-system vendors to place all covered computer code in escrow with the State Board of Elections and require the Board to secure and maintain proprietary code in custody.
Colorado · 8 CCR 1505-1, Rule 21.10	Certified election-management software, firmware and supporting documentation	• Require the complete certified election-management software and all subsystems, firmware and supporting documentation to be maintained in escrow with the Secretary of State or an approved independent escrow agent.
New Mexico · N.M. Stat. §1-9-7.11	Source code operating a certified voting system	• Make source-code escrow a condition of initial and continued voting-system certification.
Indiana · IC 3-11-15-7(a)(8) + current Voting System Certification Protocol	Voting-system software, firmware, source code and executable image used for certification	• Require documentation that the voting system's software, firmware, source codes and executable images are held in escrow with the Indiana Election Division or an approved escrow agent.
District of Columbia · D.C. Code §1-1001.09(I)	Voting-system software source code and related documents	• Require every covered procurement contract to place a copy of the voting-system source code and related documents in escrow with an independent third-party evaluator jointly selected by the vendor and the Board.
Missouri · 15 CSR 30-10.020	Source code and documentation for each fully qualified voting-system version	• Require an escrow agreement for the source code of each fully qualified system and keep the escrow in place as long as the system is used in Missouri.
Oklahoma · 62 O.S. §34.31	Source code for qualifying customized software acquired by State agencies	• Prohibit covered software contracts unless the vendor agrees to place the qualifying source code or modifications into escrow with an independent third party acceptable to the State.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
Connecticut · RCSA §9-241-34	Voting-system source and dependency package before qualification testing	• Require the manufacturer, before qualification testing, to deposit the source code, operating systems, specialized compilers and documentation materials with an escrow agent acceptable to the Secretary of the State.
Nevada · NRS 293B.135 + NAC 293B.050	Election computer program and backup copy	• Require a copy of the election computer program to be filed with the Secretary of State and secure the original and county copy under the State's rules governing custody of the program.
California · Government Code §§14900–14902	State publications selected by the State Archivist, including qualifying electronic publications	• Require two copies of each selected state publication, in physical or electronic form, to be filed with the State Archivist. • Where the optional digital repository is created, maintain and preserve electronic publications for public access in at least a machine-readable format.
United States · NARA Digital Preservation Program (government repository practice)	Federal electronic records	• Preserve federal electronic records through controlled ingest, metadata and fixity checks, media audits and migration, planning for format risks, and tested backup and restoration throughout the records' lifecycle.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 18(2) (enacted, Chapter III duty future-applying under Art. 113)	High-risk AI compliance documentation during the ten-year retention period after provider or authorised-representative bankruptcy or cessation	• Require Member States to set conditions keeping the documentation available to national competent authorities for the remaining retention period if the provider or authorised representative becomes bankrupt or ceases activity.
European Union · eIDAS Arts. 45i–45j + Implementing Regulation (EU) 2025/2532	Electronic data/documents requiring durable preservation beyond technological validity	• Use qualified archiving procedures and technologies that keep archived material durable and legible throughout the required preservation period. • Preserve integrity and accurate origin and protect the material against loss or alteration other than necessary medium or format changes.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Art. 36(6)	Files held by a notified body whose designation is restricted, suspended or withdrawn	• Require the notifying authority to take appropriate steps to ensure the notified body's files are kept and made available, on request, to notifying authorities in other Member States and to market surveillance authorities.
Australia · Copyright Act 1968 §§195CB–195CD	Qualifying Australian National Library material, including material available online	• Require covered material to be delivered to the National Library. • Permit requests for online material through electronic communication, including automated harvesting. • Require requested online material to be delivered as the whole material in the electronic form in which it was available online.
Australia · Copyright Act 1968 library/archive exceptions	Qualifying library/archive collection material	• Authorize reproduction/communication of covered collection material for preservation and related archival purposes within statutory conditions.
Australia · PANDORA / Australian Web Archive (national web-archive practice)	Selected and legally harvested Australian online publications/websites	• Capture and preserve Australian online material through selective and legal-deposit harvesting for long-term archival access.
Canada · Copyright Act §30.1	Works or other subject matter in qualifying permanent library, archive or museum collections	• Permit preservation copies of rare, unpublished, deteriorating, damaged, lost or collection material at risk. • Permit copying for safe on-site consultation where the original cannot safely be used. • Permit conversion to another format when the original format or required technology is obsolete or becoming unavailable, subject to the commercial-availability conditions.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
Canada · Library and Archives of Canada Act §10 + Legal Deposit of Publications Regulations	Qualifying Canadian publications, including non-paper and online publications	• Require legal deposit of qualifying publications and treat each version, edition or form as a distinct publication. • Place online publications in the one-copy class and specified websites, dynamic databases and other digital material in deposit-on-request classes.
European Union · Directive (EU) 2019/790, Arts. 6–7	Works and other subject matter permanently in cultural heritage collections	• Require Member States to provide an exception allowing cultural heritage institutions to make copies in any format or medium, to the extent necessary to preserve works or other subject matter permanently in their collections.
Japan · Copyright Act Art. 31	Library materials held by the National Diet Library and other qualifying libraries	• Permit noncommercial reproduction of library materials where necessary for preservation. • Use digital copies as preservation substitutes for vulnerable originals under current NDL practice.
Japan · National Diet Library Law Art. 25-4 · E-Legal Deposit System	Qualifying privately published online publications, including paid or DRM-protected classes	• Collect and preserve qualifying online publications as national cultural assets. • Cover free DRM-free publications and, from 1 January 2023, qualifying paid or DRM-protected publications. • Permit acquisition through automated collection, electronic transmission or mail and permit format conversion for long-term preservation.
Singapore · Copyright Act 2021 public-collection provisions	Material in qualifying public collections	• Authorize covered public collections to make preservation/replacement copies under the Act before deterioration, loss or obsolescence makes preservation impossible.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
Singapore · National Library Board Act 1995 §§7 and 10	Qualifying Singapore library material and online material on Singapore websites	• Require covered non-online publications to enter legal deposit. • Empower the National Library Board to preserve deposited or acquired library material. • Permit the Board to copy online material made available on a Singapore website for its statutory collection.
United Kingdom · Copyright, Designs and Patents Act 1988 library/archive provisions	Qualifying works in prescribed library/archive collections	• Authorize preservation/replacement copying by covered libraries and archives within the Act's conditions without requiring preservation to await rightholder consent.
United Kingdom · Legal Deposit Libraries Act 2003 + Non-Print Works Regulations 2013	Covered UK offline and online non-print works, including websites and electronic publications	• Require covered non-print works to enter the legal-deposit collections. • Permit automated harvesting of online works, including the prescribed procedure for works behind a login facility. • Require delivered online material to use the quality most suitable for preservation under the Regulations. • Permit deposit libraries to copy or adapt material for preservation or replacement, including changing medium or format where needed.
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) (official State procurement terms)	Source-code package for the statewide pollbook solution	• Require direct State delivery or escrow of the covered source package for the State's benefit. • Keep the escrow package current across new component versions. • Replace an escrow provider that ceases or fails to provide acceptable service with a successor preserving at least the same State rights.
Arizona · 2025 Elections Procedures Manual + A.R.S. §16-445 (current force-of-law election procedures)	Election program filed before election use	• Require the Secretary of State to preserve filed election programs in escrow for three years, with new copies filed after later changes.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
Wisconsin · Wis. Stat. §5.905 + Admin. Code EL 7.03	Software components, programs, documentation and source code used in approved Wisconsin voting systems	• Require approved vendors to place the software components selected by the Commission in escrow and, under the rule governing continued approval, furnish programs, documentation and source code used in any Wisconsin election to an escrow agent approved by the Commission.
Pennsylvania · ES&S EVS 6.5.0.0 current certification conditions (official current system-specific certification)	Voting-system and EPB source code delivered to purchasing counties	• Maintain an escrow agreement covering all source code for ten years after delivery and acceptance, with the vendor bearing deposit and verification costs and Pennsylvania retaining defined access rights.
Minnesota · Stat. §206.805	Software source code covered by State voting-system contracts	• Require State voting-system contracts to include provisions placing the software source code in escrow, creating pre-arranged custody outside ordinary vendor possession.
Wyoming · Title 22 electronic-voting-system certification statute (2026)	Source code for the voting system and its components, including updates	• Require certification proof of an escrow account containing the source code for the system and components for which certification is sought, including source-code updates.
United States · 37 CFR §§1.803, 1.806–1.807 · biological-material patent deposit	Biological material needed to satisfy patent enablement requirements	• Use a qualifying depository independent of the depositor and maintain viability and contamination controls. • Minimize loss risk and preserve the material for at least thirty years and at least five years after the most recent sample request. • Keep the deposit available beyond the patent's enforceable life where the regulation requires it.

Table 3.4-9. Examples of Legal and Institutional Precedents Requiring Preservation Before the Material Disappears (continued) · See Section 3.4.3

Precedent	Object	Requirement
United States · 17 CFR §240.17a-4(f)	Electronic records subject to broker/dealer preservation duties	• Preserve required electronic records through the applicable retention period using an audit-trail or non-rewriteable storage method permitted by the rule. • Automatically verify the completeness and accuracy of storage and retention. • Maintain backup or other redundancy preserving access if the primary system becomes unavailable.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability · See Section 3.4.4

Precedent	Object	Requirement
California · Right to Repair Act / SB 244	Covered electronic or appliance product during the statutory support period	• Require the manufacturer to provide owners and independent repairers documentation, functional parts, tools, repair software, and updates for the statutory three- or seven-year support period.
Colorado · SB26-189 Automated Decision-Making Technology (enacted, core duties apply 1 January 2027)	Technical information needed by a deployer to use and govern covered ADMT	• Require the developer to provide technical documentation covering intended uses, training-data categories, known limitations and risks, and circumstances in which the system should not be used. • Include instructions for appropriate use, monitoring and meaningful human review, together with information reasonably needed for deployer compliance.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
California · CCPA Civil Code §1798.130(a)(2)–(3)	Personal information delivered under covered CCPA disclosure rights	• Deliver required personal information in a readily usable format that permits transfer from one entity to another without hindrance. • For covered specific pieces, use a structured, commonly used, machine-readable format where technically feasible and permit transmission at the consumer's request.
California · 2 CCR §20621 · voting-system source-code package	Source-code materials needed to compile and configure each covered election-system version	• Deposit the program and components with complete build and configuration instructions and related compilation documentation. • Identify compatible commercial off-the-shelf products and versions, subject to the regulation's narrow rule for their compilation documentation.
North Carolina · G.S. §163-165.9A(a)(1)	Voting-system software package needed to understand, build, configure and operate the system	• Escrow the software and materials needed for system functionality, setup, configuration and operation, including source and executable code, build materials, APIs, documentation and specifications. • Identify the programmers and provide the required affidavit confirming completeness of the source code.
New York · 9 NYCRR §6209.6	Certified voting-system software package needed for installation, validation, operation and maintenance	• Escrow complete source and executable code, build materials, APIs, documentation, specifications, records and data needed to install, validate, operate and maintain the system. • Provide documentation confirming source-code completeness and identifying the programmers. • Independently assess whether the documentation is sufficient to install, validate, operate and maintain the system.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Ohio · Rev. Code §3506.05(I)(1) + Admin. Code 111:3-9-04	Voting-system software materials required in escrow	• Deposit all source code and related documentation, require the documentation to include system configuration, and require the source-code deposit to contain all relevant low- or high-level program statements, excluding variable code specific to an election.
Colorado · 8 CCR 1505-1, Rule 21.10	Escrowed election-management source code and supporting system documentation	• Require internal source documentation sufficient to understand and maintain the source, instructions for converting source to object code and producing an executable system where warranted, plus architecture, design, testing and installation/configuration documentation.
Indiana · IC 3-11-15-7 + current Voting System Certification Protocol	Technical material accompanying a voting system submitted for certification	• Require escrow of software, firmware, source code and executable images plus functional software descriptions, technical documentation, software/data flow diagrams, maintenance information, manuals and information about component interchangeability within the certification package.
Georgia · Ga. Comp. R. & Regs. 590-8-1-.01(g)	Technical package required to evaluate and certify a voting system	• Before certification, provide all software and firmware source code in listing and machine-readable form and identify the required compiler or development environment. • Provide the system design, maintenance, operations, customer and qualification documentation needed for certification.
Missouri · 15 CSR 30-10.020	Voting-system software material held for use and support after the trigger	• Deposit source code for all voting-system components in both human-readable and machine-readable formats together with the software documentation required by the escrow rule.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Oklahoma · 62 O.S. §34.31	Source code for customized software developed or modified exclusively for a State agency or the State	• Require qualifying source code and supplied upgrades to be held in escrow as the technical material delivered to the State upon the statute's trigger conditions.
Connecticut · RCSA §9-241-34	Technical package accompanying a voting system into escrow	• Deposit source code together with operating systems, specialized compilers and documentation materials, preserving dependencies needed to understand and build the submitted software state.
United States · 47 CFR Part 51, including §§51.307, 51.309 and 51.311	Unbundled network element and related network information requested by a telecommunications carrier	• Provide nondiscriminatory access to the unbundled network element at a technically feasible point, including all of that element's features, functions and capabilities and separate access to the requested facility or functionality. • Provide technical information about the incumbent network sufficient for the requesting carrier to achieve the regulated interconnection or UNE access. • To the extent technically feasible, provide the UNE and access at least at the quality the incumbent provides to itself.
United States · DoD DFARS 227.7203-2 + PGI 227.7203-2 (binding regulation + official PGI implementation guidance)	Government software capability requiring lifecycle reproduction, build, test, deployment, operation, sustainment or support	• In lifecycle procurement planning, acquire the software and recorded information needed to rebuild, test, integrate, deploy, operate and sustain the capability. • Consider periodic delivery of code, build materials, dependencies, test results, datasets, firmware and training materials needed for those tasks. • Where practicable, accompany source code with architecture, design, traceability and interface documentation.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 10 U.S.C. §§4401–4403 · Modular Open Systems	Major defense systems and interfaces designed for modular third-party integration	• Design covered acquisition programs, to the maximum extent practicable, around a modular open system approach. • Define major system interfaces and provide the documentation needed to integrate third-party components or modules. • Address component replacement over the lifecycle, technical data deliverables and the intellectual property rights needed for modular integration.
United States · SHARE IT Act, Pub. L. 118-187	Federal custom-developed software package needed for development, build, test and deployment	• Store covered source code in a public or private repository accessible through the Act's federal access procedures. • Include the documentation, data models, schemas, metadata, architecture, configuration and other artifacts needed to develop, build, test and deploy the code.
United States · 15 U.S.C. §8512a	NOAA operational models released under the federal plan for public access to models	• Develop and implement a plan making operational models developed by NOAA publicly available at no cost, together with applicable model information and documentation, descriptions of intended outputs, and qualifying associated operational-forecasting data and metadata. • Preserve the statute's release accommodations, short-horizon model exclusion and withholding rules for national security.
United States · 15 U.S.C. §8512(b)(5)(E)	NOAA community global weather research modeling system	• Make the NOAA research modeling system publicly accessible and runnable on public computers and networks outside secure NOAA systems. • Use cloud-based hosting or management where appropriate and cost-effective.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 45 CFR §§170.315(b)(10), 170.402 · EHI Export	Electronic health information stored by covered certified health-IT products	• Require covered certified health IT to create electronic computable exports containing all EHI the product can store, together with a publicly accessible and current description of the export format. • For a single patient, make the export executable at a time chosen by the user without subsequent developer assistance.
United States · OPEN Government Data Act, 44 U.S.C. §§3502, 3506	Public federal data assets within the Act's disclosure scope	• Make public data assets machine-readable and available as open Government data assets in open formats built on open standards, with the Act's metadata/provenance architecture and existing disclosure, intellectual property, privacy and security boundaries preserved.
Washington · chapter 19.415 RCW	Covered digital electronic product requiring diagnosis, maintenance or repair	• Provide owners and independent repairers with the required parts, tools and documentation, including applicable software mechanisms and updates used for programming, pairing, calibration and restoration of full functionality.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · AI Act, Regulation (EU) 2024/1689, Art. 25(2), (4) <i>(enacted, Chapter III duty future-applying under Art. 113)</i>	Information, capabilities, technical access and assistance needed by a high-risk AI provider or new provider to comply with the Act	• When provider status transfers under Article 25(1), require the initial provider to provide the information, technical documentation and reasonable technical access needed for the new provider's compliance. • Require written agreements specifying the information, capabilities, technical access and assistance owed by third parties whose systems, models, tools, services or components are integrated into a high-risk system. • Preserve Article 25(4)'s exception for qualifying public free or open-source tools, services, processes and components other than general-purpose AI models.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 53(1)(b) + Annex XII	Information and technical conditions needed by a downstream provider of an AI system to integrate a general-purpose AI model	• Require current information and documentation enabling downstream providers to understand the general-purpose AI model's capabilities and limitations and comply with the AI Act. • Cover the model's release and distribution, external interactions, relevant software versions, architecture, licence, input and output formats, integration means and the provenance and curation of training, testing and validation data. • Preserve Article 53(2)'s exemption for qualifying free or open-source non-systemic-risk models while keeping systemic-risk models within the duty.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Cyber Resilience Act, Regulation (EU) 2024/2847, Art. 13(6) + Annex II point 8(f) <i>(enacted, generally applicable from 11 December 2027)</i>	Third-party component vulnerability remediation and downstream integration of a product with digital elements	• Require a manufacturer that identifies a vulnerability in an integrated component to notify the component's manufacturer or maintainer and address the vulnerability. • Where the product manufacturer develops a fix, require relevant code or documentation to be shared with the component manufacturer or maintainer where appropriate. • For products intended for integration into another product, provide the information the integrator needs for CRA compliance.
European Union · Digital Markets Act, Regulation (EU) 2022/1925, Arts. 6(7) and 7	Platform features controlled by the gatekeeper and interfaces for messaging services needed for regulated interoperability	• Under Article 6(7), make the operating-system, hardware and software features used or controlled by the gatekeeper available for effective interoperability. • Under Article 7, provide the interfaces or similar solutions and information in the reference offer needed to make covered messaging functions interoperable.
European Union · Digital Markets Act, Art. 6(9)–(10)	Data generated or provided through a designated gatekeeper's covered core platform services	• Provide free portability tools and continuous real-time access to covered data. • Give business users effective, high-quality, continuous real-time access to covered aggregated and non-aggregated data under the Article's conditions.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Italy · Digital Administration Code Art. 69 + AgID reuse guidance	Software owned by a public administration and developed to the specific instructions of a public client	• Make the source code, complete documentation and related technical-functional description available through the required public repository/platform architecture. • Publish within Article 69's open-licence framework, subject to its stated public-order/security, national-defence and electoral exceptions.
France · CRPA Arts. L300-2, L311-1 and L312-1-1	Source code produced or received by a covered administration in the course of its public mission	• Treat source code as an administrative document and require a covered administration to publish the held code online or communicate it on request within the statutory access rules. • For administrations within Article L312-1-1, publish qualifying electronically held documents online after communication and publish their updated versions, subject to the Code's secrecy and intellectual property limits.
United Kingdom · Communications Act 2003 §87	Network access and associated facilities of a provider found by OFCOM to have significant market power	• After a finding that the provider has significant market power, permit OFCOM to require the dominant provider to make network access and relevant facilities available. • Permit OFCOM to require publication of access information, contract terms and specified modifications to access offers under imposed transparency conditions.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Australia · Competition and Consumer Act 2010 Part XIC, §152AR	Active declared telecommunications service requested by a service provider	• Supply the active declared service on request and take all reasonable steps to provide technical and operational quality equivalent to that which the access provider supplies to itself. • Provide equivalent-quality and timing fault detection, handling and rectification and, where applicable, permit interconnection with the requesting provider's facilities on equivalent technical/operational terms. • Provide billing information connected with the declared-service supply when requested.
Australia · Competition and Consumer Act Part IVE motor-vehicle information-sharing scheme	Scheme information used to diagnose, service, repair, modify or dismantle covered vehicles	• Require scheme information to be offered to Australian repairers/RTOs through the statutory access process at no more than fair market value, in the required accessible form, with supply on request subject to the scheme's timing and safety/security rules.
Australia · Copyright Act 1968 §195CD	Electronic National Library material delivered under legal deposit	• Require the whole material in the applicable electronic form, free from technological protection measures and accompanied by any software or additional information necessary for the National Library to access the deposited copy.
Canada · Library and Archives of Canada Act §10 + Legal Deposit of Publications Regulations §2	Non-paper publication whose deposited copy would otherwise be inaccessible or poorly documented	• Before deposit, decrypt encrypted data and remove or disable systems that restrict access. • Provide the access software, manuals, technical information and descriptive metadata needed to use and document the deposited publication.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Commission Regulation (EU) 2023/1670, Annex II repair-information and serialised-part requirements (<i>binding, applicable from 20 June 2025</i>)	Covered smartphone/slate-tablet repair and replacement work	• Provide professional repairers the required repair information for at least seven years after placement of the product model on the market ends, including repair diagrams, manuals, diagnostic information and software or firmware instructions. • For covered serialised-part replacement, provide the software tools, firmware or similar means needed to restore full part and device functionality.
European Union · Directive (EU) 2019/790 Art. 7(2) + Directive 2001/29/EC Art. 6(4)	Preservation exception for cultural heritage that would otherwise be blocked by technological protection measures	• Where the beneficiary has legal access and voluntary rightholder measures are absent, require Member States to take appropriate measures requiring rightholders to make available the means necessary to let the beneficiary exercise the Article 6 preservation exception.
European Union · Directive (EU) 2024/1799, Art. 5 + Annex II repairability framework	Goods for which Annex II Union acts impose repairability requirements	• Apply the Directive's repair obligation together with the Annex II duties to provide covered parts, repair information, software tools, firmware and similar resources. • Where manufacturers provide spare parts and tools, require prices that do not deter repair.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · EECC, Directive (EU) 2018/1972, Arts. 61, 72 and 73	Electronic communications infrastructure, network elements and support systems subject to an access remedy imposed by a regulator	• Permit competent authorities to require access to in-building wiring and associated facilities and, under Article 72, to specified civil engineering assets. • Permit access to network elements and services, technical interfaces, protocols, interoperability technologies, support software, interconnection and associated services. • Apply the Directive's necessity, proportionality and conditions protecting the interests of end users to those remedies.
European Union · eIDAS, current Art. 12b	Platform features used by qualifying European Digital Identity Wallet or notified eID providers that rely on a gatekeeper core platform service	• Make the same operating-system, hardware and software features available for effective interoperability and access as are available to or used by the gatekeeper for the relevant services, free of charge and within the access framework established by DMA Article 6(7).
European Union · Interoperable Europe Act, Regulation (EU) 2024/903, Arts. 4–5	Interoperability solution supporting a trans-European digital public service or published as an Interoperable Europe solution	• On a qualifying request, make the solution available with technical documentation and, where applicable, version history, documented source code and open-standard or technical specification references. • Specify reuse, support and maintenance conditions that let the reusing entity manage the solution autonomously, subject to the Regulation's public task, third-party-IP and security boundaries. • Publish Interoperable Europe solutions in open, machine-readable, findable and reusable formats with documented source code and metadata where applicable.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Regulations (EU) 2019/424, 2019/2019 and 2019/2021–2019/2024	Maintenance, repair, reuse or upgrading of covered servers/data storage products and regulated appliances/displays	• Provide the required repair materials for the statutory access period, including repair diagrams, manuals, diagnostic information and software or firmware resources. • Apply the applicable registration requirements for professional repairers and the applicable response time and conditions governing reasonable fees.
Japan · National Diet Library Law Art. 25-4 · current e-Legal Deposit implementation (official implementation requirement)	DRM-protected online publication subject to e-Legal Deposit	• Require a DRM-free file for covered e-Legal Deposit publications. • Permit format conversion for long-term preservation.
Singapore · National Library Board Act 1995 §10(1)(b)	Electronic non-online library material subject to legal deposit	• Deposit the electronic copy in the form and manner determined by the National Library Board and free of technological or other restrictions on access.
Spain · Law 40/2015 Arts. 157–158 + Royal Decree 4/2010 Arts. 16–17, as amended	Applications made available for reuse by public administrations, including applications declared open source	• Law 40/2015 requires administrations to make qualifying applications they own available to requesting administrations, subject to the boundary for protected information, and permits those applications to be declared open source. • For reusable applications distributed as products that have been declared open source, the current directory rules require publication of source code, associated documentation and licensing conditions, together with information about reuse costs where applicable.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Switzerland · EMOTA/EMBAG Art. 9 + Federal Chancellery OSS implementation guidance	Software developed or commissioned by covered federal authorities for official tasks	• Publish the source code under Article 9 unless third-party rights or grounds related to security preclude or restrict publication. • Current Federal Chancellery implementation guidance explains that a useful release should, where relevant, accompany the code with documentation and basic configuration or master data context needed to understand and use the software.
United Kingdom · Non-Print Works Regulations 2013, Reg. 17	Covered non-print work whose deposited copy requires associated software or information for access	• Require the publisher, for covered offline works and online works delivered by agreement, to deliver with the work any computer program, data or information necessary to access it and any accompanying manual or other material made available to the public.
United Kingdom · Sourcing Playbook / DDaT Playbook (official procurement guidance)	Outsourced digital or service capability that may need to transfer to an incoming supplier or in-house team	• Keep supplier handover documentation current, covering the design, system architecture, transition material and operational information needed for transfer. • Use exit planning to identify key interfaces, dependencies and the transfer of digital, data and knowledge assets. • Require the supplier to return the data and information needed for transition and replacement procurement.
Virginia · Voting System Certification Standard 3.1 (2026) (official State certification standard)	Technical information needed to certify, operate, maintain, repair and troubleshoot the voting system	• Require a digital Technical Data Package covering software design, modules, functions, data formats and operating-system, version, security and hash information. • Include the maintenance, repair, troubleshooting, operations, testing and certification materials needed for State review.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) <i>(official State procurement terms)</i>	Software package needed to inspect and rebuild the statewide pollbook solution	• Require machine-readable source code, executable code, design and user documentation, and complete compile and link instructions. • For custom software, provide tools for viewing source code and deliver current source and documentation at termination.
Wisconsin · Wis. Stat. §5.905 + Admin. Code EL 7.01, EL 7.03	Voting-system software materials submitted for approval and continuing use	• Require complete hardware, firmware and software specifications, technical manuals and documentation and independent-test evidence for approval, and require programs, documentation and source code used in State elections to remain in approved escrow.
Pennsylvania · Voting System Security Standard (18 Sept. 2026) <i>(official State certification standard)</i>	Technical materials required for State security certification and testing	• Provide the Technical Data Package, source code, user documentation, security architecture, system design, test reports and configuration and hardening materials needed for State security certification. • Include the certified system components needed to conduct a test-lab build and examination.
Pennsylvania · ES&S EVS 6.5.0.0 current certification conditions <i>(official current system-specific certification)</i>	Operational materials needed by a county to run and maintain the certified voting system	• Provide the manuals needed for a purchasing county to operate the system independently of vendor assistance and support. • Train county personnel in programming, tabulation, auditing, recount, troubleshooting, record preservation and post-election maintenance and storage.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
New Jersey · N.J.A.C. 15:10-7.6	Technical package needed for State review of an e-poll-book system	• Require production software and firmware, all software and firmware source code, additional software and hardware needed for testing, operation and administration manuals, compliance settings, QA results and a detailed diagram of the network architecture.
Wyoming · Title 22 electronic-voting-system certification statute (2026)	Source and documentation package accompanying voting-system certification	• Require source-code escrow plus checksum identity, repair and maintenance technical manuals, operations manuals and other vendor documentation describing the system's operation, use and maintenance.
United States · 37 CFR §§1.803, 1.808 · biological-material patent deposit	Preserved biological material needed to make the patented disclosure practically available	• Require the depository to furnish samples properly and expeditiously and, upon patent grant, irrevocably remove restrictions on public availability imposed by the depositor except for the rule's bounded conditions governing records of requests.
United States · 40 CFR §86.1808-01(f) · emission-control repair information	Technical information, software and tools needed for covered vehicle diagnosis and repair	• Provide independent repairers the manuals, service information, software, calibrations and diagnostic tools needed for covered repair. • Provide tool and equipment manufacturers the hardware requirements, protocols, interfaces and specifications needed to develop aftermarket diagnostic and reprogramming equivalents. • Keep required information available in full text for at least fifteen years after model introduction and make archived information obtainable afterward.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 14 CFR §21.50(b) · Instructions for Continued Airworthiness	Covered aircraft, aircraft engine or propeller and the technical instructions needed to keep it airworthy	• Furnish each owner at least one complete set of Instructions for Continued Airworthiness. • Keep the instructions, including later changes, available to every other person legally required to comply with them. • Include the technical and interface information needed for continued airworthiness.
United States · 48 CFR §1552.239-71 · EPA Open Source Software clause	Custom-developed software and the materials needed to build, test, understand and reuse it	• Deliver the underlying source code, licence file, build instructions, user guides, automated test suites and associated documentation. • Include required dependency licences and place the source code in the version-control system specified by EPA.
European Union · Regulation (EU) 2018/858, Art. 61 + Annex X	Technical materials needed by independent operators to diagnose, repair and maintain covered vehicles	• Provide independent operators the diagnostic information, software, tools, repair and maintenance information and remote access needed for covered vehicle diagnosis and repair. • Provide the training material and qualifying machine-readable data required by the Regulation's standardised access framework.
United States · DOE 970.5227-2	Software and supporting materials produced by the contractor within covered DOE management-and-operating contracts	• Give DOE inspection and delivery rights over covered software. • For software produced by the contractor delivered to OSTI, require source code or executable files together with metadata and documentation sufficient for a technically competent user to understand and use the software. • Apply the clause's rights and restriction boundaries.

Table 3.4-10. Examples of Legal and Institutional Precedents Requiring Materials Needed to Maintain or Reconstruct the Capability (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · eIDAS Art. 5a(3), as amended by Regulation (EU) 2024/1183 (<i>binding, wallet rollout by end 2026</i>)	Application-software source code for European Digital Identity Wallets	• Require the wallet application-software source code to be made available under the Regulation's open-source architecture, while permitting justified non-disclosure only for specified components other than those installed on user devices.
Switzerland · E-ID Act (BGEID) Art. 12(1)–(2) (<i>enacted and approved, future-applying, trust-infrastructure commencement expected in 2027</i>)	Software source code of the federal e-ID trust infrastructure	• Require FOITT/BIT to disclose the trust-infrastructure software source code, while permitting non-publication only for source code or parts whose disclosure is precluded or restricted by third-party rights or reasons related to security.
United States · FD&C Act §524B / 21 U.S.C. §360n-2	Technical materials needed to keep a covered cyber device and related systems cybersecure	• Require a software bill of materials and require the sponsor to make the postmarket updates and patches needed to address covered vulnerabilities available to the device and related systems under the statute's regular-cycle and critical-vulnerability timing rules.
United States · 42 U.S.C. §19053 · Research Reproducibility and Replicability	Research products needed for reproducibility of work supported by NSF	• Direct NSF to facilitate public access to research products including data, software and code. • Establish a public point for locating repositories and support sustainable open-source tools and infrastructure for reproducibility. • Preserve the applicable federal privacy, intellectual property and security boundaries.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute · See Section 3.4.4

Precedent	Object	Requirement
California · Right to Repair Act / SB 244	Owner or independent service/repair provider working on a covered electronic/appliance product during the statutory support period	• Give owners and independent repair providers the statutory right to obtain the documentation, parts, tools and software needed for covered diagnosis, maintenance or repair throughout the applicable support period.
California · Government Code §14902(b)–(c)	State publications prepared by a state officer or employee as part of official duties	• Place qualifying official-duty state publications in the public domain, allowing public-domain material generally to be reprinted without legal restriction while preserving the copyright of third-party material embedded in a government publication.
North Carolina · G.S. §163-165.7(a)(8)	Source code turned over after vendor maintenance failure or bankruptcy	• Require the vendor's contract with the purchasing county to require turnover of the source code to the county for continued use during the contract period and for authorized review after the statutory trigger.
Missouri · 15 CSR 30-10.020	Escrowed voting-system source code after a qualifying trigger based on vendor failure	• Require the escrow agreement to give Missouri full source-code access and unlimited rights to continue using and supporting the software at no cost after a qualifying trigger, without converting the right into a public licence.
United States · 11 U.S.C. §365(n)	Executory licence of intellectual property as defined by the Bankruptcy Code rejected by a debtor/licensor	• Let the licensee choose between qualifying termination and retention after rejection of the licence. • If the licensee retains its rights, preserve the contractual rights to the qualifying intellectual property, including exclusivity, for the remaining contract term and any extension exercisable as of right. • Require the trustee to permit exercise of those retained rights while the licensee makes all royalty payments due under the contract for the duration of the contract.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 17 U.S.C. §108	Copyrighted collection material within a qualifying library or archive's statutory scope	• Let a qualifying library or archive make preservation or replacement copies when an eligible work is unpublished, damaged, lost, stolen or stored in an obsolete format, and permit the statute's bounded research uses.
United States · 17 U.S.C. §117	Computer program copy or machine falling within the statute's owner or lessee conditions	• Permit the owner of a program copy to make a copy or adaptation needed to run the program, or an archival copy, within the statute's conditions. • Permit a machine's owner or lessee to make a temporary copy solely through activation of a lawful installed program for machine maintenance or repair, subject to the statutory use and destruction requirements.
United States · 37 CFR §201.40(b)(19) (current 2024–2027 triennial exemption)	Lawfully acquired complete video game after qualifying external-server cessation or commercial unavailability	• After qualifying authentication-server cessation, permit access, copying and modification needed to restore personal local gameplay of a lawfully acquired complete game. • Limit the personal path to local gameplay and require the external-server cessation conditions specified by the exemption. • Permit eligible libraries, archives and museums to copy or modify qualifying games for playable preservation, subject to the exemption's noncommercial and premises limits. • For institutional preservation, permit bounded use of game code formerly stored on the external server where the exemption's conditions are met.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 37 CFR §201.40(b)(20) (current 2024–2027 triennial exemption)	Lawfully acquired non-game software or software-dependent digital material no longer reasonably available commercially	• Permit an eligible library, archives or museum to circumvent access controls for lawful noncommercial preservation of qualifying unavailable software or software-dependent material. • Permit bounded remote access to preserved works for one user per lawfully owned copy, for a limited time, and for private study, scholarship or research.
United States · 47 CFR Part 51 / Communications Act interconnection and unbundled-access rights	Requesting telecommunications carrier obtaining an unbundled network element within the federal access regime	• Entitle a requesting carrier to obtain and use an available unbundled network element within the rule's scope without barred incumbent restrictions. • Permit the requesting carrier to provide allowed telecommunications services over the element and to commingle it with specified wholesale services. • Where technically feasible, require access quality at least equal to the incumbent's own access quality.
United States · 15 U.S.C. §8512a	Operational models and related materials released under the statutory plan	• Make released material available at no cost under an open licence with no restrictions on copying, publishing, distributing, citing, adapting or other use, subject to the statute's release boundaries.
United States · SHARE IT Act, Pub. L. 118-187	Custom-developed code produced under covered federal software-development contracts	• Require the agency to acquire and exercise rights sufficient for governmentwide access to, sharing of, use of and modification of the covered custom-developed code.
United States · 10 U.S.C. §3771 + DFARS 252.227-7013	Technical data for qualifying modular system interfaces and other covered defense-acquisition technical data	• Grant Government Purpose Rights in qualifying modular-interface technical data. • Permit use, modification, reproduction, release, performance, display and disclosure of the data and authorization of outside recipients for Government purposes, subject to the statutory and contractual boundaries.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · DFARS 252.227-7014	Other-than-commercial computer software and software documentation delivered under covered defense contracts	• Apply the clause's licence categories based on funding to covered source code, algorithms, design details and software documentation. • Under qualifying unlimited rights, permit use, modification, reproduction, release, performance, display, disclosure and authorization of others. • Preserve the narrower limits of Government Purpose Rights and restricted rights where those categories apply.
United States · OPEN Government Data Act, 44 U.S.C. §§3502, 3506	Public federal data assets released under the open-data regime	• Make each public data asset available under an open licence that guarantees no-cost public access with no restrictions on copying, publishing, distributing, transmitting, citing or adapting the asset, subject to the Act's disclosure and intellectual property boundaries.
United States · Mission Product Holdings v. Tempnology, 587 U.S. 370 (2019)	Rights previously conveyed under an executory contract rejected by a debtor, including the trademark licence at issue	• Treat rejection of an executory contract as breach rather than rescission. • Permit the debtor to stop future performance while preserving rights already conveyed where those rights would survive an ordinary contract breach.
Washington · chapter 19.415 RCW	Owner or independent repair provider diagnosing, maintaining or repairing a covered digital electronic product	• Give owners and independent repair providers the statutory entitlement to obtain and use the covered repair resources outside the authorized service network.
European Union · Digital Markets Act, Regulation (EU) 2022/1925, Arts. 6(7) and 7	Covered third-party provider seeking interoperability with a designated gatekeeper platform or messaging service	• Under Article 6(7), make the operating-system, hardware and software features used or controlled by the gatekeeper available for effective interoperability. • Under Article 7, provide the interfaces or similar solutions and information in the reference offer needed to make covered messaging functions interoperable.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Digital Markets Act, Art. 6(9)–(10)	End users, business users and their authorised third parties receiving covered platform data	• Give end users an effective portability entitlement and give business users and authorised third parties a free statutory right of continuous real-time access to and use of covered data, subject to the Regulation's personal-data and scope conditions.
European Union · Directive (EU) 2019/1024 Open Data Directive	Documents held by public sector bodies within the Directive's reuse scope	• Create statutory reuse rights for covered public sector information made available in reusable and open formats, subject to the Directive's conditions.
Italy · Digital Administration Code Art. 69 + AgID reuse guidance	Covered public administration source code and documentation released under an open licence	• Release covered public administration software under an open licence allowing free reuse and adaptation by other public administrations or legal persons. • Treat the open licence as the legal basis for reuse rather than requiring a new bilateral agreement.
France · CRPA Arts. L321-1–L323-2 and D323-2-1	Qualifying public information in source code communicated or published by a covered administration	• Give any person the statutory reuse right for qualifying public information outside the public mission for which it was produced or received. • Where free software reuse is placed under a licence, use the statutory licensing rules, which expressly recognise BSD, Apache, CeCILL-B, MIT, MPL, GPL, CeCILL, EUPL and Eclipse licences, subject to the Code's conditions and third-party intellectual property exclusion.
United Kingdom · Communications Act 2003 §87	Relationship involving network access after OFCOM finds significant market power in the identified services market	• After a finding that the provider has significant market power, permit OFCOM to require provision and use of network access and associated facilities. • Permit the imposed conditions to specify fairness, reasonableness and compliance timing.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
Australia · Competition and Consumer Act 2010 Part XIC, especially §§152AF–152AR	Service provider seeking access to a declared telecommunications service	• When the standard access obligations apply, require the access provider to supply the declared service and applicable interconnection access to the access seeker. • Make access depend on the statutory regime and governing access instrument rather than the incumbent's voluntary permission.
Australia · Competition and Consumer Act Part IVE motor-vehicle information-sharing scheme	Australian repairer or consumer choosing a repairer for a covered scheme vehicle	• Give Australian repairers statutory access and use rights to scheme information for diagnosing faults, servicing, repairing, modifying or dismantling covered vehicles, supporting the consumer's ability to use a repairer of choice.
Australia · Copyright Act 1968 Div. 4A §§47B–47H	Computer-program copy owned or licensed within Division 4A	• Permit the covered incidental reproductions needed for ordinary running and for studying the ideas behind the program and how it functions, together with statutory backup copying. • Permit bounded reproduction/adaptation to obtain otherwise unavailable interoperability information, correct errors and perform good-faith security testing within the Act's conditions. • Section 47H makes contractual terms ineffective insofar as they exclude or limit the protected study limb in §47B(3) or the §§47C–47F backup, interoperability, error-correction and security testing exceptions.
Australia · Copyright Act 1968 library and archive preservation provisions, including §§113H, 113J and 113M	Works, sound recordings, films and published editions in qualifying library, archive or key cultural-institution collections	• Permit the covered institution to make preservation and specified research copies within the relevant sections, including preservation copies of historically or culturally significant works, recordings, films and published editions.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
Canada · Copyright Act §§30.6, 30.61 and 41.12	Authorised computer-program copy or licensed use and any technological protection measure governing it	• Let the owner of an authorised copy or licensee make specified compatibility adaptations/conversions and backup copies for the statutory own-use purposes. • Permit reproduction of the program solely to obtain information needed for program interoperability, with limits on use and disclosure of that information. • Permit specified circumvention, circumvention services and devices for interoperability with a lawfully obtained computer program within §41.12.
Canada · Copyright Act §30.1	Work or other subject matter in a qualifying permanent library, archive or museum collection	• Permit qualifying libraries, archives and museums to copy collection material for preservation, safe consultation or conversion to another format when the original format or necessary technology is obsolete or becoming unavailable.
European Union · Commission Regulation (EU) 2023/1670, Annex II repair-information requirements (<i>binding, applicable from 20 June 2025</i>)	Professional repairer access to covered repair information and later third-party use after manufacturer access ends	• Give registered qualifying professional repairers access to covered repair and maintenance information under the Regulation's eligibility, timing and conditions governing reasonable fees. • After the required access period ends, permit third parties to use and publish the unaltered repair information originally published by the manufacturer, subject to intellectual property rights.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Directive (EU) 2019/790, Arts. 6–7	Works and other subject matter permanently in cultural heritage collections	• Require a preservation exception allowing cultural heritage institutions to make copies in any format or medium to the extent necessary for preservation. • Make contractual provisions contrary to that preservation exception unenforceable. • Place the exception under Article 7(2) to prevent TPM protection from nullifying the preservation permission where the Directive requires measures needed to provide access.
European Union · Directive (EU) 2024/1799, Art. 5	Consumer seeking repair of a good covered by the Directive's Annex II reparability framework	• Require the responsible manufacturer or economic operator to repair a covered good at the consumer's request. • Preserve the consumer's freedom to choose another repairer and prohibit refusal solely on the ground of an earlier repair by another person.
European Union · Directive 2009/24/EC, Arts. 5–6 and Art. 8	Computer program lawfully acquired or used within the Directive's statutory conditions	• Allow acts necessary for intended use of a lawfully acquired computer program, including error correction. • Protect a necessary backup and permit the lawful user to observe, study or test the program while performing authorised acts. • Permit bounded decompilation where indispensable to obtain otherwise unavailable information needed for interoperability. • Make contractual terms contrary to the protected backup, study, testing and decompilation rights ineffective.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · EECC, Directive (EU) 2018/1972, Arts. 60, 61, 72 and 73	Undertaking receiving access or interconnection under a duty or remedy imposed within the EECC framework	• Give authorised operators of public networks a right and reciprocal obligation to negotiate interconnection under Article 60. • Where an access remedy is imposed, require the regulated undertaking to meet reasonable requests for access to the specified infrastructure, network elements, interfaces, technologies, support software and associated services. • Permit the authority to bar withdrawal of access already granted and to impose fairness, reasonableness and timeliness conditions.
European Union · eIDAS, current Art. 12b	Qualifying EUDI Wallet or eID provider using a gatekeeper core platform service as a DMA business user	• Make the same operating-system, hardware and software features available for effective interoperability and access as are available to or used by the gatekeeper for the relevant services, free of charge and within the access framework established by DMA Article 6(7).
European Union · Interoperable Europe Act, Regulation (EU) 2024/903	Shared or interoperability solution published through the portal within the Regulation's reuse framework	• Permit a reusing public body or third party to adapt a shared interoperability solution unless third-party intellectual property rights restrict adaptation. • If a solution published under Article 4 is adapted, require publication of the adapted solution in the same way. • Require portal solutions to use a licence permitting at least reuse by other Union or public sector bodies or to be open source. • Define an open-source licence as permitting reuse, redistribution and modification for all uses while making source code indiscriminately available.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
Japan · Copyright Act Art. 31	Library material held by the National Diet Library or another qualifying library	• Permit noncommercial reproduction of qualifying library material for user research, preservation, and supply to another qualifying library when the material is out of print or otherwise difficult to obtain, subject to the Article's conditions.
Singapore · Copyright Act 2021 §§237–242 and 430	Computer program owned or lawfully used within the relevant statutory permission	• Permit backup copies made by the owner and copies or adaptations made or authorised by the owner that are an essential step in using the program with a machine under §§237–238. • Permit decompilation by a lawful user for bounded interoperability purposes, observation/study/testing of program functioning, and copying/adaptation necessary for lawful use, including error correction, under §§239–242. • Section 430 permits specified TPM circumvention, circumvention devices and circumvention services solely for program interoperability within its conditions.
Singapore · Copyright Act 2021, public-collection provisions §§225–236	Material in a qualifying public collection and online material copied for the National Library Board collection	• Permit bounded copying or communication of qualifying material in a qualifying public collection for internal access, research, study, administration and inter-library supply. • Permit preservation copying against loss, deterioration, damage or format obsolescence. • Permit the National Library Board to copy online Singapore material for its statutory collection.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
Singapore · National Library Board Act 1995 §7(2)(f)–(g)	Deposited, acquired or online Singapore library material within the Board's statutory functions	• Empower the National Library Board to take appropriate measures to maintain and preserve deposited or acquired library material and to make copies of online material available on a Singapore website at the times and in the manner the Board considers appropriate.
Spain · Law 40/2015 Arts. 157–158 + Royal Decree 4/2010 Art. 16	Application of a public administration declared open source within the reuse framework	• Require an open-source licence that permits running the shared software for any purpose, access to source code, modification or improvement, and redistribution with or without changes. • Require derivatives to preserve the same four guarantees. • Make licensing ordinarily available without consideration or requiring an additional agreement, subject to the stated cost exceptions.
Switzerland · EMOTA/EMBA Art. 9	Covered Federal Administration software published as open source	• Make the published software available for anyone to inspect the source and use, improve, and distribute the software without licence fees, subject to limits on publication arising from third-party rights and security.
United Kingdom · Copyright, Designs and Patents Act 1988 §§50A, 50B, 50BA, 50C and 296A	Computer-program copy used by a lawful user within the statutory conditions	• Permit a necessary backup copy under §50A. • Permit bounded decompilation under §50B to obtain information needed for an independent interoperable program. • Permit observation, study and testing of program functioning under §50BA while the lawful user performs authorised acts. • Permit copying or adaptation necessary for lawful use, including error correction, under §50C where the governing agreement does not prohibit it. • Make contrary contractual terms ineffective for the protected §§50A, 50B and 50BA acts.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
United Kingdom · Copyright, Designs and Patents Act 1988 §42	Item in the permanent collection of a qualifying library, archive or museum	• Permit the librarian, archivist or curator, subject to the statutory conditions, to copy the item to preserve or replace it in the institution's permanent collection or to replace a lost, destroyed or damaged item in another prescribed collection.
United Kingdom · Non-Print Works Regulations 2013, Regs. 20 and 29–30	Relevant non-print material held by a legal-deposit library	• Permit the deposit library to review and maintain the material and use it for its own non-commercial research, and to copy or adapt it for preservation or replacement, including changing medium or format where necessary or expedient within the Regulations.
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) (official State procurement terms)	Custom software developed under the State contract and related source and documentation	• Give the State ownership of covered custom software, source code, maintenance updates, documentation and configuration files. • After specified contractor failure events, transfer the rights, title and interests needed for the State or actors acting on its behalf to use the software and source code.
Minnesota · Stat. §206.805	Software supplied under Minnesota State voting-system contracts	• Require the contracts to give the State a perpetual licence to use and modify the software, preserving those rights without requiring fresh vendor permission for later covered use or modification.
United States · 40 CFR §86.1808-01(f)	Independent aftermarket actor diagnosing or repairing covered emission-related vehicle systems	• Give covered independent repair and service actors the right to obtain and use required diagnostic and service information, software, calibrations and tools at fair and reasonable cost. • Do not require affiliation with the manufacturer's franchised repair network.

Table 3.4-11. Examples of Legal and Institutional Precedents Providing Legal Rights to Use, Study, Modify, Host, and Redistribute (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 48 CFR §1552.239-71 · EPA Open Source Software clause	Covered custom source code, documentation and related files furnished under an EPA contract	• Give the Government unlimited reproduction, reuse, modification and distribution rights for federal reuse and public open-source software. • Require the contractor to use the open-source licence specified by the contract or approved by the Agency. • Preserve the clause's National Security System, documented exceptions and boundaries specific to the contract.
European Union · Regulation (EU) 2018/858, Arts. 61–65	Independent repairers and other qualifying independent operators within the EU vehicle-repair information regime	• Give qualifying independent operators a statutory right to obtain and use covered software downloads, diagnostic services, tools and repair and maintenance information outside the network authorised by the manufacturer. • Apply the Regulation's vehicle scope, security controls and conditions governing reasonable fees.
United States · DOE 970.5227-2	Computer software first produced or specifically used in covered DOE contract performance	• Give the Government ownership of software first produced under the contract and, subject to stated exceptions, unlimited rights in specifically used software. • Permit use, reproduction, derivative works, public distribution, public performance, display and authorization of others under those unlimited rights. • Preserve the clause's restricted-software and copyright boundaries.
European Union · eIDAS Art. 5a(3), as amended by Regulation (EU) 2024/1183 (binding, wallet rollout by end 2026)	Source code of European Digital Identity Wallet application-software components	• Require covered wallet application source code to be open-source licensed, subject only to the Regulation's duly justified non-disclosure boundary for certain components outside user devices.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails · See Section 3.4.4

Precedent	Object	Requirement
California · CCPA Civil Code §1798.130(a)(2)–(3)	Personal information controlled by the consumer moving between businesses	• Allow consumers to direct transfer of covered personal information from one business to another without hindrance when switching services.
United States · Communications Act §251(b)(2) + 47 CFR Part 52	Existing telephone number when a customer switches eligible telecommunications providers	• Require covered carriers to let users keep their existing telephone number when switching eligible providers without degrading service quality or network reliability. • Complete covered simple ports within one business day and non-simple ports within four business days unless the new provider or customer requests more time.
California · 2 CCR §§20611, 20681 · continuity through escrowed source code	Election-jurisdiction dependency on voting or election software maintained by the vendor	• Require independent escrow of source-code information before vendor maintenance becomes unavailable, unless the contract supersedes the regulatory rule. • Define access conditions in advance, enabling the election jurisdiction to obtain the escrowed material when the contractual trigger is met.
North Carolina · G.S. §163-165.7(a)(8)	County dependency on electronic voting software maintained by the vendor	• Require the vendor to arrange in advance for escrowed source code to pass to the purchasing county if the vendor fails to maintain the software as agreed or enters bankruptcy. • Permit the county to continue using the software for the remaining contract period and to conduct authorized review.
South Carolina · Code §§7-13-1330(I), 7-13-1620(F)	State dependency on voting-system source code after vendor business failure	• Require voting-system source code to be held in escrow and made available to the State Election Commission if the vendor goes out of business.
New Mexico · N.M. Stat. §1-9-7.11	State dependency on source code for a voting system whose manufacturer exits or ends support	• Require escrowed voting-system source code to become accessible to the State when the manufacturer ceases business or support.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
Missouri · 15 CSR 30-10.020	State dependency on voting software supported by the vendor	• Require full access to escrowed source code within seven days after specified insolvency, liquidation, receivership, or events involving discontinuation of support. • Give the State unlimited continuation and support rights together with records for the final build and certification testing.
Oklahoma · 62 O.S. §34.31	State dependency on customized software after vendor default, insolvency or maintenance/support failure	• Require the escrow agreement to release everything held in escrow to the agency after specified default, insolvency, bankruptcy, receivership, or maintenance and support failure.
Nevada · NAC 293B.050	Election computer program whose original is destroyed or rendered unusable	• Keep the filed copy of the election program as a backup that can be used if the original is destroyed or unusable.
United States · 11 U.S.C. §365(n)(3)	Qualifying intellectual property or protected embodiment covered by a retained licence after debtor/licensor rejection	• If a licensee elects to retain qualifying rights after rejection, require the trustee on written request to provide covered intellectual property or a protected embodiment held by the trustee when the contract provides for it. • Prohibit the trustee from interfering with contractual rights to obtain the same material from another entity.
United States · 37 CFR §201.40(b)(20) (current 2024–2027 triennial exemption)	Preserved non-game software or software-dependent digital material no longer reasonably available commercially	• Permit an eligible library, archives, or museum to provide electronic access to preserved qualifying software or software-dependent material that is no longer reasonably available commercially. • Limit remote access to one user per lawfully owned preserved copy, for a limited time, and to private study, scholarship, or research.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · DoD DFARS 227.7203-2 + PGI 227.7203-2 (<i>binding regulation + official PGI implementation guidance</i>)	Government software capability that may need to move to the Government or another contractor	• Plan delivery or access in advance to facilitate transfer of software sustainment to the Government or another contractor. • Choose between delivery into Government custody and contractual access according to the lifecycle need, including for contractor repositories and cloud environments. • Use deferred delivery or escrow with predefined release conditions that let the Government obtain the material after the trigger.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 25(2) (<i>enacted, Chapter III duty future-applying under Art. 113</i>)	Transfer of provider responsibility to a distributor, importer, deployer or third party under Article 25(1)	• When provider status transfers under Article 25(1), require the initial provider to cooperate with the new provider and hand over the information and technical documentation needed for compliance. • Require disclosure of known limitations and failure modes and reasonable technical access for testing and validation.
European Union · AI Act, Regulation (EU) 2024/1689, Art. 36(3), (6)–(9)	Files, certificates and conformity assessment responsibility when a notified body exits or loses designation	• Preserve the departing notified body's files and keep them available to the relevant authorities. • Require a successor notified body to assume responsibility and complete the required reassessment within the applicable nine- or twelve-month period. • During suspension or restriction, require the provider to identify another qualified notified body within three months to monitor the affected certificates. • Keep certificates valid only under the Regulation's bounded continuity conditions.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · Data Act, Regulation (EU) 2023/2854, Chapter VI	Customer dependency on a data processing service and the exportable data and digital assets embedded in it	• Require providers to remove commercial, technical, contractual, and organisational barriers to switching or moving to on-premises infrastructure. • Normally complete the switch within 30 calendar days while providing reasonable assistance and maintaining service continuity and security. • Preserve at least a 30-day post-transition retrieval period, with a longer transition only where technical infeasibility justifies it. • Provide the interfaces and interoperability needed to port covered data and digital assets, subject to the Regulation's intellectual property and trade secret boundaries and to the customer's relevant use rights.
European Union · Digital Markets Act, Art. 6(9)–(10)	User dependency on data held or generated in a designated gatekeeper core platform service	• Provide effective portability tools and continuous real-time access to covered data for end users and authorised third parties. • Enable switching or multi-homing without making users depend on a one-time export controlled by the gatekeeper.
Japan · Copyright Act Art. 31(7)–(8) · NDL Digitized Contents Transmission Services	National Diet Library digitized holdings that are out of print or otherwise difficult to obtain	• Permit approved public, university, and similar libraries to transmit qualifying NDL digitized holdings for use by registered patrons outside NDL premises. • Permit official registered NDL users resident in Japan to access qualifying holdings over the Internet. • Exclude material likely to become readily commercially available, and do not treat e-Legal Deposit publications as automatically eligible.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Banking Authority · Guidelines on outsourcing arrangements, EBA/GL/2019/02 (2019 Guidelines, currently applicable)	Critical or important outsourced function and the data held by the service provider	• Plan for provider failure, service deterioration, business disruption, and other continuity risks affecting critical or important outsourced functions. • Maintain documented and, where appropriate, tested exit plans that can move the outsourced function and data to another provider or back in-house while preserving continuity. • Identify the alternatives, resources, timing, responsibilities, success criteria, and triggers needed to execute the exit.
European Union · DORA, Regulation (EU) 2022/2554, Arts. 28 and 30 + Delegated Regulation (EU) 2024/1773	Critical or important ICT dependency on a third-party provider	• Put exit strategies in place for provider failure, service deterioration, failed delivery, continuity risk, and contract termination without disrupting critical business or regulatory compliance. • Require contracts to preserve access, recovery, and return of data when the provider becomes insolvent, enters resolution, stops business, or the contract ends. • For critical or important functions, require a transition period during which the provider continues service while the financial entity migrates. • Maintain and periodically test a documented exit plan for each covered arrangement.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
European Union · GDPR Art. 20	Personal data provided by the data subject where the covered processing is based on consent or contract and carried out by automated means	• Give the data subject qualifying personal data in a structured, commonly used, machine-readable format and permit transmission to another controller without hindrance. • Where technically feasible, permit direct controller-to-controller transmission. • Preserve the Article's public task and official-authority exclusion and the rights and freedoms of others.
Financial Stability Board · Guidance on Arrangements to Support Operational Continuity in Resolution, revised 2024 (international guidance, nonbinding)	Critical shared services, including IT infrastructure and services involving software, that support critical functions in resolution	• Map critical functions to shared services, providers, legal entities, and dependencies and keep the information needed for resolution current. • Use resolution-resilient contracts and arrangements that preserve rights to use and access IT, intellectual property, and other operational assets. • Prepare transfer arrangements such as Transitional Service Agreements and a continuity playbook before resolution begins.
Financial Stability Board · Key Attributes of Effective Resolution Regimes (international resolution standard)	Critical functions and essential services of a failing financial institution	• Give resolution authorities power to keep essential services running through group companies, a successor, an acquirer, or unaffiliated third parties. • Permit transfer of selected assets, liabilities, legal rights, and obligations without ordinary consent or novation and permit a bridge institution to continue critical functions. • Use resolution plans to identify critical functions, transfer options and data, execution barriers, and measures to remove those barriers.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
United Kingdom · PRA SS2/21, November 2024 current version <i>(supervisory expectation, March 2026 update effective 18 March 2027)</i>	Material outsourcing or third-party dependency supporting important business services	• Plan and test stressed exit for material outsourcing or third-party dependencies to keep important business services available after provider disruption. • Identify viable exit methods such as bringing the function or data in-house, moving to another provider, or using temporary contractual or escrow measures during transition. • Begin exit planning before outsourcing and identify alternatives, resources, timing, priority data, and exit triggers, then test and update the plan on a risk basis.
United Kingdom · PRA SS4/21, May 2021 current version <i>(supervisory expectation, effective 1 January 2023)</i>	Critical services supporting functions or core business lines through recovery, resolution and restructuring	• Map critical services, providers, contracts, operational assets, and interdependencies for timely use in recovery and resolution. • Maintain tested procedures and Transitional Service Agreements that can move provision of critical services to another provider without interruption. • Use resolution-resilient contracts, preserve access to technology, data, intellectual property, licences, and leases, and test continuity under severe but plausible scenarios.
United Kingdom · Sourcing Playbook / DDaT Playbook (official procurement guidance)	Critical outsourced public service/digital capability moving from an incumbent supplier to replacement or in-house delivery	• Contracts should include an exit plan linking the outgoing supplier's exit to incoming-supplier or in-house mobilisation and review it throughout the contract. • Specify the activities, milestones, roles, interfaces, dependencies, and transfers of assets, data, and information needed for transition. • For critical contracts, plan emergency exit after supplier insolvency and keep contingency plans current and testable.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
Maryland · Statewide Pollbook Modernization RFP D38B560002 (2025) (official State procurement terms)	State dependency on pollbook software supported by the contractor	• Require immediate release of the Source Code Escrow Package to the State after specified bankruptcy, insolvency, support failure, inability or unwillingness to provide the package, or extended nonresponse. • Require delivery of current custom source code and documentation at contract termination.
Arizona · 2025 Elections Procedures Manual + A.R.S. §16-445 (current force-of-law election procedures)	Election program whose original is destroyed or rendered unusable	• Keep a machine-readable copy of the election program filed with the State as a backup that can be used if the original is destroyed or unusable.
Pennsylvania · ES&S EVS 6.5.0.0 current certification conditions (official current system-specific certification)	Commonwealth dependency on source code held by the vendor after vendor inability or unwillingness to perform	• Require a pre-arranged escrow agreement covering the current source code for the voting system and electronic poll book and other specified supplier property. • Permit Pennsylvania to obtain the escrowed material when circumstances show with reasonable certainty that the vendor cannot or will not perform its obligations.
United States · 37 CFR §1.805 · biological-material patent deposit	Legally necessary deposit that becomes unavailable, contaminated or no longer capable of the described function	• Require replacement or supplementation when the depository cannot furnish samples or the deposited material loses the required viability or function. • Preserve the replacement's identity and provenance and require a replacement term at least as long as the remaining term of the original deposit.
Nebraska · Rev. Stat. §8-1,141	Customer's controllable electronic record held by a covered financial institution	• Require the institution and customer to agree in writing on when the customer's controllable electronic record will be returned, and permit them to agree on the return format.

Table 3.4-12. Examples of Legal and Institutional Precedents Requiring Pre-arranged Handoff, Recovery, or Continuation When the Supplier Exits or Fails (continued) · See Section 3.4.4

Precedent	Object	Requirement
United States · 17 CFR §240.17a-4(f)	Preserved electronic records the broker or dealer fails to download or provide to regulators	• Require a pre-filed undertaking by a designated executive officer or qualifying unaffiliated third party to download preserved records and audit trails and provide regulator access if the broker or dealer fails to produce them.
United States · DOE 970.5227-2	DOE dependency on technical data and software held by the contractor during contract performance	• Give DOE the right to direct delivery or other disposition of covered technical data and software during contract performance and when the contract ends. • Preserve Government access to relevant software or data retained at the facility.
United States · 45 CFR §§170.315(b)(10), 170.402 · EHI Export	Customer dependency on EHI stored in covered certified health IT	• Require certified health IT to export all electronic health information stored by the covered product in an electronic computable format. • Make the single-patient export executable at the user's chosen time without later developer assistance. • Make the population export sufficient to migrate the product's complete stored electronic health information.

Annex G. Comparisons with Models with Publicly Available Weights

Evidence tables supporting §5.3, *Feasibility for AI Models* are collected below.

The table reports two comparisons for each model. The first comparison uses an evaluation from before consumer access ended or before the relevant API retirement date and records whether a model with publicly available weights performed better on at least one reported metric. The second comparison records whether, by September 2026, a model with publicly available weights performed better in a broader common evaluation.

The broader comparison uses an overall or aggregate result, a result covering the full benchmark or task set, or another comparably broad measure. For a target model no longer available, its preserved result is compared with a result from a model with publicly available weights only where the benchmark version and scoring method remained unchanged. Evidence included in the broader comparison runs through 28 September 2026.

Consumer Model Comparisons

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Text models					
OpenAI					
GPT-4	30 Apr 2025	gpt-4-0314 · LMArena · 11 Jun 2024 · rank 12 · rating 1188.298	command-r-plus · 1188.682 · difference 0.384 rating	Llama 3 70B Instruct rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1206.010
GPT-4o	13 Feb 2026	LMArena · 11 Feb 2026 · rank 26 · rating 1427.959	glm-4.5 · 1428.296 · difference 0.337 rating	Kimi K2.5 Thinking rank 12 , GLM-5 rank 16	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1428.713
GPT-4o mini	14 May 2025	LMArena · 16 Apr 2025 · rank 35 · rating 1272.176	llama-4-maverick-17b-128e-instruct · 1271.411 · difference 0.765 rating	DeepSeek-V3-0324 rank 4 , DeepSeek-R1 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1286.418
o1-preview	5 Dec 2024	LiveBench · leaderboard of 25 Nov 2024, captured 4 Dec 2024 · Coding 50.85	Qwen2.5 72B Instruct · 56.56 · difference 5.71 points	Qwen2.5 Coder 32B 56.85 , Qwen2.5 72B Instruct 56.56	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1352.772

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
o1-mini	31 Jan 2025	LMArena · 28 Jan 2025 · rank 5 · rating 1305.063	deepseek-v3 · 1316.571 · difference 11.508 rating	DeepSeek-R1 rank 1 , DeepSeek-V3 rank 4	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1317.075
OpenAI o1	16 Apr 2025	LMArena · 9 Apr 2025 · rank 7 · rating 1350.064	gemma-3-27b-it · 1341.545 · difference 8.519 rating	DeepSeek-V3-0324 rank 4 , DeepSeek-R1 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1365.838
o1 pro mode	10 Jun 2025	o1-pro-2025-03-19 · Scale HLE Text Only · moved to Legacy on 3 Apr 2025 · o1 Pro (March 2025) 8.44 ± 1.12	DeepSeek-R1 · 8.57 ± 1.13 · difference 0.13 points	DeepSeek-R1 8.57 ± 1.13	Scale HLE Text Only · DeepSeek-R1 8.57 ± 1.13 vs target 8.44 ± 1.12
o3-mini	16 Apr 2025	LMArena · 9 Apr 2025 · rank 17 · rating 1305.316	command-a-03-2025 · 1304.983 · difference 0.333 rating	DeepSeek-V3-0324 rank 4 , DeepSeek-R1 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1319.450
o3-mini-high	16 Apr 2025	LMArena · 9 Apr 2025 · rank 11 · rating 1325.149	DeepSeek-V3 · 1318.245 · difference 6.904 rating	DeepSeek-V3-0324 rank 4 , DeepSeek-R1 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1336.609

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
GPT-4.5	Plus 7 Aug 2025 Pro 26 Jun 2026	LiveIdeaBench v2 · 29 Mar 2025 · average 6.85	QwQ-32B · 7.06 · difference 0.21 points	DeepSeek-R1 7.18 , QwQ-32B 7.06	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1417.401
OpenAI o3	26 Aug 2026	LMArena · 16 Nov 2025 · rank 33 · rating 1410.948	qwen3-235b-a22b-thinking-2507 · 1413.405 · difference 2.457 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1409.945
OpenAI o4-mini	13 Feb 2026	LMArena · 5 Nov 2025 · rank 66 · rating 1353.943	qwen3-coder-480b-a35b-instruct · 1356.411 · difference 2.468 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1353.238
o4-mini-high	7 Aug 2025	Scale MultiNRC · 1,055 questions · 22.18 ± 2.51	DeepSeek R1 · 24.27 · difference 2.09 points	DeepSeek-R1-0528 27.58 ± 2.70 , DeepSeek-R1 24.27 ± 2.59	Scale MultiNRC · DeepSeek-R1-0528 27.58 ± 2.70 vs target 22.18 ± 2.51
GPT-4.1	13 Feb 2026	LMArena · 16 Nov 2025 · rank 48 · rating 1382.574	glm-4.5-air · 1382.170 · difference 0.404 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1383.148

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
GPT-4.1 mini	13 Feb 2026	LMArena · 16 Oct 2025 · rank 71 · rating 1340.020	qwen3-32b · 1339.706 · difference 0.314 rating	GLM-4.6 rank 3 , GLM-4.5 rank 6	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1340.537
GPT-5 Instant	Go 12 Nov 2025 Plus / Pro 13 Feb 2026	LMArena · 16 Nov 2025 · rank 39 · rating 1403.250	qwen3-vl-235b-a22b-thinking · 1400.270 · difference 2.980 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1403.454
GPT-5 Thinking	13 Feb 2026	LMArena · 17 Nov 2025 · rank 37 · rating 1407.785	qwen3-235b-a22b-thinking-2507 · 1413.424 · difference 5.639 rating	GLM-4.6 rank 5 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1406.077
GPT-5 Pro	19 Feb 2026	IUMB archive · 11 Nov 2025 · 25 / 48 = 52.1%	Kimi K2 Thinking · 54.2% · difference 2.1 pp	Kimi K2 Thinking 26 / 48 = 54.2%	IUMB · Kimi K2 Thinking 26 / 48 = 54.2% vs target 25 / 48 = 52.1%
GPT-5 Thinking mini	17 Apr 2026	LMArena · 16 Nov 2025 · rank 53 · rating 1375.485	DeepSeek-V3-0324 · 1374.120 · difference 1.365 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1373.266

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
GPT-5.1 Instant	Go 11 Dec 2025 Plus / Pro 11 Mar 2026	LMArena · 6 Mar 2026 · rank 38 · rating 1421.032	longcat-flash-chat · 1420.944 · difference 0.088 rating	Qwen3.5-397B-A17B rank 10 , Kimi K2.5 Thinking rank 11	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1422.187
GPT-5.1 Thinking	Go 11 Dec 2025 Plus / Pro 11 Mar 2026	LMArena · 6 Mar 2026 · rank 18 · rating 1441.615	GLM-5 · 1441.118 · difference 0.497 rating	Qwen3.5-397B-A17B rank 10 , Kimi K2.5 Thinking rank 11	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1442.716
GPT-5.1 Pro	11 Mar 2026	IUMB archive · 20 Jan 2026 · 25 / 48 = 52.1%	Kimi K2 Thinking · 54.2% · difference 2.1 pp	Kimi K2 Thinking 26 / 48 = 54.2%	IUMB · Kimi K2 Thinking 26 / 48 = 54.2% vs target 25 / 48 = 52.1%
GPT-5.2 Instant	Plus / Pro 12 Jun 2026 Go 3 Mar 2026	LMArena · 10 Jun 2026 · rank 47 · rating 1438.252	nvidia-nemotron-3-ultra-550b-a55b-nvfp4 · 1438.253 · difference 0.001 rating	GLM-5.1 rank 13 , GLM-5 rank 35	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1438.592
GPT-5.2 Thinking	12 Jun 2026	LMArena · 10 Jun 2026 · rank 91 · rating 1416.209	deepseek-v3.1-thinking · 1416.887 · difference 0.678 rating	GLM-5.1 rank 13 , GLM-5 rank 35	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1416.236

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
GPT-5.2 Pro	12 Jun 2026	SGI-Bench · 12 Dec 2025 · Dry Experiment 28.04	Qwen3-VL-235B-A22B · 28.41 · difference 0.37 points	Qwen3-VL-235B-A22B 28.41	ARC Prize ARC-AGI-1 · Kimi K3 94.5% vs GPT-5.2 Pro (XHigh) 90.5%
GPT-5.3 Instant	Go 5 May 2026 Plus / Pro 5 Aug 2026	LMArena · 16 Jun 2026 · rank 124 · rating 1387.911	qwen3-30b-a3b-instruct-2507 · 1383.847 · difference 4.064 rating	GLM-5.1 rank 13 , GLM-5 rank 35	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1389.146
GPT-5.4 Thinking	23 Jul 2026	LiveBench · 25 Jun 2026 · Coding 77.542	Kimi K2.6 Thinking · 78.567 · difference 1.025 points	Kimi K2.6 Thinking 78.567	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1468.934
GPT-5.4 Pro	23 Jul 2026	Agronomy LLM Benchmark · 10 Jul 2026 · Overall Score 90.12%	NVIDIA Nemotron 3 Ultra 550B A55B · tested 9 Jul 2026 · 90.35% · difference 0.23 pp	NVIDIA Nemotron 3 Ultra 550B A55B 90.35%	Agronomy LLM Benchmark · Kimi K3 92.00% Overall vs target 90.12% Overall
GPT-3.5	18 Jul 2024	LiveBench · captured 13 Jun 2024 · Coding of gpt-3.5-turbo-0125 29.16	Qwen2 72B Instruct · 31.79 · difference 2.63 points	Qwen2 72B Instruct Mathematics 43.44 vs 25.54 , Reasoning 32.5 vs 24.3 , Language 54.45 vs 50.20	Arena-Hard-Auto v0.1 · Qwen2.5 72B Instruct 78.0 vs target 23.3
Google					

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Gemini Ultra 1.0	14 May 2024	Radiology · 85 board-style questions · 72% (61/85)	Mixtral 8x22B Instruct · 72% (61/85) · difference 0 pp	Llama 3 70B 80% (68/85)	Radiology benchmark · Llama 3 70B 80% (68/85) vs target 72% (61/85)
Gemini 1.5 Pro	5 Feb 2025	LMArena · 3 Feb 2025 · rank 9 · rating 1301.883	deepseek-v3 · 1316.181 · difference 14.298 rating	DeepSeek-R1 rank 1 , DeepSeek-V3 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1319.135
Gemini 2.0 Flash Experimental	30 Jan 2025	LiveBench · before 30 Jan 2025 · Gemini 58.0 on 303 shared objective questions On the 203 questions shared with DeepSeek-R1-Distill-Llama-70B, Gemini 43.360	DeepSeek-R1-Distill-Llama-70B · 43.111 on the same 203 questions · difference 0.249 points	DeepSeek-V3 · 62.4 vs 58.0 on 303 shared questions	LiveBench · DeepSeek-V3 62.4 vs target 58.0 on 303 shared objective questions
Gemini 2.5 Pro	18 Nov 2025	LiveBench · leaderboard of 30 May 2025, captured 12 Nov 2025 · Mathematics 84.19	GLM-4.6 · 90.10 · difference 5.91 points	GLM-4.6 Mathematics 90.10 , Agentic Coding 35.00 vs 25.00	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1457.780
Gemini 2.5 Deep Think	4 Dec 2025	IUMB archive · 11 Nov 2025 · 34 / 48 = 70.8%	Kimi K2 Thinking · 54.2% · difference 16.6 pp	None in this evaluation	VoxelBench · DeepSeek V4 Pro 1571 ± 45 vs target 1521 ± 54

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Gemini 3 Pro	19 Feb 2026	LiveBench · leaderboard of 8 Jan 2026, captured 12 Jan 2026 · Coding 74.60	DeepSeek-V3.2 · 75.69 · difference 1.09 points	DeepSeek-V3.2 75.69	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1479.402
Gemini 3 Flash	19 May 2026	LiveBench · leaderboard of 8 Jan 2026, captured 27 Jan 2026 · Agentic Coding 40.00	Kimi K2.5 Thinking · 48.33 · difference 8.33 points	Kimi K2.5 Thinking 48.33	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1465.664
Gemini 3.5 Flash	21 Jul 2026	LiveBench · 25 Jun 2026 · Data Analysis 64.857	Kimi K2.6 Thinking · 65.134 · difference 0.277 points	Kimi K2.6 Thinking 65.134	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1481.267
Gemini 3.6 Flash	14 Aug 2026	LiveBench · 21 Jul 2026 · Data Analysis 62.999	Kimi K2.6 Thinking · 65.134 · difference 2.135 points	Kimi K2.6 Thinking 65.134	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1478.317
Anthropic					
Claude 2.0	21 Nov 2023	AlpacaEval · 1 Oct 2023 · win rate 91.356%	Xwin-LM-13B-V0.1 · 91.760% · difference 0.405 pp	Xwin-LM-70B-V0.1 95.568% , Xwin-LM-13B-V0.1 91.760%	AlpacaEval · Xwin-LM-70B-V0.1 95.568% vs target 91.356%

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Claude 2.1	13 Mar 2024	HELM Lite v1.0.0 · 20 Dec 2023 · mean win rate 0.593	Llama 2 70B · 0.659 mean win rate · difference 0.066 win-rate points	Yi 34B 0.772 , Mixtral 8x7B 0.728	HELM Lite v1.0.0 · Yi 34B 0.772 mean win rate vs target 0.593
Claude 3 Sonnet	21 Jun 2024	LMArena · 17 Jun 2024 · rank 11 · rating 1198.847	llama-3-70b-instruct · 1202.058 · difference 3.211 rating	Llama 3 70B Instruct rank 10	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1218.041
Claude 3 Haiku	13 Dec 2024	LMArena · 4 Nov 2024 · rank 47 · rating 1179.281	command-r-08-2024 · 1179.518 · difference 0.237 rating	Llama 3.1 Nemotron 70B rank 11 , Qwen2.5 72B Instruct rank 20	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1194.682
Claude 3.5 Sonnet	22 May 2025	LMArena · 5 May 2025 · rank 36 · rating 1282.942	deepseek-v2.5-1210 · 1278.856 · difference 4.086 rating	DeepSeek-V3-0324 rank 5 , DeepSeek-R1 rank 7	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1297.772
Claude 3.5 Haiku	19 Feb 2026	LMArena · 17 Oct 2025 · rank 140 · rating 1256.241	llama-3.1-tulu-3-70b · 1255.724 · difference 0.517 rating	GLM-4.6 rank 3 , GLM-4.5 rank 6	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1255.358

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Claude 3.7 Sonnet	14 Nov 2025	LMArena · 5 Nov 2025 · rank 103 · rating 1299.697	deepseek-v2.5-1210 · 1294.338 · difference 5.359 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1299.427
Claude Sonnet 4	19 Feb 2026	LMArena · 5 Nov 2025 · rank 75 · rating 1338.563	mistral-small-2506 · 1339.324 · difference 0.761 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1339.064
Claude Opus 4	16 Jan 2026	LMArena · 5 Nov 2025 · rank 58 · rating 1364.680	ling-flash-2.0 · 1363.342 · difference 1.338 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1365.717
Claude Opus 4.1	16 Jan 2026	LMArena · 9 Jan 2026 · rank 40 · rating 1413.867	qwen3-235b-a22b-thinking-2507 · 1412.813 · difference 1.054 rating	GLM-4.6 rank 10 , GLM-4.5 rank 17	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1418.500
Claude Sonnet 4.5	26 May 2026	LMArena · 19 May 2026 · rank 47 · rating 1435.867	glm-4.7 · 1436.198 · difference 0.331 rating	GLM-5.1 rank 13 , GLM-5 rank 33	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1438.569

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Claude Opus 4.5	16 Apr 2026	LMArena · 14 Apr 2026 · rank 19 · rating 1447.671	kimi-k2.5-thinking · 1446.552 · difference 1.119 rating	GLM-5.1 rank 8	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1450.861
Image and video models					
OpenAI					
Sora Turbo / Sora 1	U.S. 13 Mar 2026 Other listed regions 26 Apr 2026	LMArena Text-to-Video · 16 Oct 2025 · rank 15 · rating 1069.342	Wan2.2 T2V A14B · 1127.591 rating · difference 58.249 rating	Wan2.2 A14B rank 13	LMArena Text-to-Video · Wan2.2 T2V A14B 1127.591 overall vs target 1069.342
Sora 2 Pro	Sora app 26 Apr 2026	RigidBench · ICLR 2026 Workshop World Models, published 2 Mar 2026 · SSIM 0.610	Wan 2.2 TI2V-5B · 0.841 · difference 0.231	Wan 2.2 TI2V-5B SSIM 0.841 , IoU 0.234 vs 0.199	LMArena Text-to-Video · MiniMax H3 1461.780 overall vs target 1366.552
Google					

Table 5.3-3 · Part I. Consumer Models Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model	Access ended	Evaluation before access ended	Nearest result among models with publicly available weights	At least one reported metric exceeded before access ended	Higher result in a broader common evaluation by September 2026
Veo 2	Gemini Apps 30 May 2025 Flow 31 Mar 2026	LMArena · 5–6 Nov 2025 · Text-to-Video rank 15 · rating 1161.944 Image-to-Video rank 20 · rating 1164.131	Text-to-Video: Wan2.2 T2V A14B · rating 1128.410 · difference 33.534 Image-to-Video: Wan2.2 I2V A14B · rating 1167.288 · difference 3.157	Wan2.2 A14B rank 19 in Image-to-Video None in Text-to-Video	LMArena Image-to-Video · Wan2.2 I2V A14B 1170.146 overall vs target 1165.741
Veo 3	Gemini Apps and Flow 15 Oct 2025	LMArena Text-to-Video · 7 Oct 2025 · rank 4 · rating 1266.69	Wan2.2 T2V A14B · 1136.25 · difference 130.44 rating	None in this evaluation	LMArena Text-to-Video · MiniMax H3 1461.780 overall vs target 1253.299
Veo 3 Fast	Gemini Apps and Flow 15 Oct 2025	Stable Cinematics · 30 Sep 2025 · Lighting 3.974	Wan 2.2 14B · 4.009 · difference 0.035 points	Wan 2.2 14B 4.009	LMArena Text-to-Video · MiniMax H3 1461.780 overall vs target 1248.376
Gemini 2.5 Flash Image / Nano Banana	Gemini Apps 26 Feb 2026	LMArena Text-to-Image · 7 Feb 2026 · 3d_modeling rank 14 · rating 1141.976	FLUX.2 dev · rank 12 · 1146.779 · difference 4.803 rating	HunyuanImage 3.0 rank 6 · 1170.040 , Qwen Image 2512 rank 9 · 1157.309 , FLUX.2 dev rank 12 · 1146.779	LMArena Text-to-Image · Cosmos3-Super-Text2Image 1164.71 overall vs target 1150.20

API Version Comparisons

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
Text models					
OpenAI					
gpt-3.5-turbo-0301	Retired 13 Sep 2024	Arena-Hard-Auto v0.1 · 19 Apr 2024 · win rate 18.1	Qwen1.5 72B Chat · 36.1 · difference 18.0 points	Mixtral 8x22B Instruct 36.4 , Qwen1.5 72B Chat 36.1	Arena-Hard-Auto v0.1 · Qwen2.5 72B Instruct 78.0 vs target 18.1
gpt-4-0314	Retired 26 Mar 2026	LMArena · 11 Jun 2024 · rank 12 · rating 1188.298	command-r-plus · 1188.682 · difference 0.384 rating	Llama 3 70B Instruct rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1206.010
gpt-3.5-turbo-0613	Retired 13 Sep 2024	Arena-Hard-Auto v0.1 · 19 Apr 2024 · win rate 24.8	Qwen1.5 72B Chat · 36.1 · difference 11.3 points	Mixtral 8x22B Instruct 36.4 , Qwen1.5 72B Chat 36.1	Arena-Hard-Auto v0.1 · Qwen2.5 72B Instruct 78.0 vs target 24.8
gpt-4-0613	Scheduled 23 Oct 2026	LMArena · 15 Oct 2024 · rank 46 · rating 1163.168	qwen1.5-110b-chat · 1161.482 · difference 1.686 rating	Qwen2.5 72B Instruct rank 18 , DeepSeek-V2.5 rank 20	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1186.008
gpt-3.5-turbo-1106	Scheduled 28 Sep 2026	LMArena · 5 Mar 2024 · rank 25 · rating 1069.251	openchat-3.5 · 1072.264 · difference 3.013 rating	Qwen1.5 72B Chat rank 7	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1094.121
gpt-3.5-turbo-0125	Scheduled 23 Oct 2026	LMArena · 13 Aug 2024 · rank 49 · rating 1105.601	wizardlm-70b · 1106.106 · difference 0.505 rating	Llama 3 70B Instruct rank 18 , Qwen1.5 72B Chat rank 36	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1124.922

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
gpt-4-0125-preview	Retired 26 Mar 2026	LMArena · 4 Nov 2024 · rank 28 · rating 1244.952	llama-3.1-70b-instruct · 1246.702 · difference 1.750 rating	Llama 3.1 Nemotron 70B rank 11 , Qwen2.5 72B Instruct rank 20	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1262.206
gpt-4-turbo-2024-04-09	Scheduled 23 Oct 2026	LMArena · 20 Nov 2024 · rank 22 · rating 1256.464	deepseek-v2.5 · 1257.482 · difference 1.018 rating	Llama 3.1 Nemotron 70B rank 11 , Qwen2.5 72B Instruct rank 20	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1271.693
gpt-4o-2024-05-13	Scheduled 23 Oct 2026	FAIR-Eval · estimated ability 1.366	Qwen3-30B-A3B-Instruct-2507 · 1.367 · difference 0.001	DeepSeek-V3 1.612 , DeepSeek-V3-0324 1.559 , Qwen3-Coder-480B-A35B-Instruct 1.410 , Qwen3-30B-A3B-Instruct-2507 1.367	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1300.436
o1-preview-2024-09-12	Retired 28 Jul 2025	DPAB-α · 15 Jan 2025 · Pythonic function-calling 55	Phi-4 · 55 · difference 0 points	Qwen2.5 Coder 32B Instruct 68 , Qwen2.5 72B Instruct 65 , DeepSeek-V3 63 , Llama 3.1 405B Instruct 60 , Llama 3.3 70B Instruct 59	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1352.772
o1-mini-2024-09-12	Retired 27 Oct 2025	LMArena · 28 Jan 2025 · rank 5 · rating 1305.063	deepseek-v3 · 1316.571 · difference 11.508 rating	DeepSeek-R1 rank 1 , DeepSeek-V3 rank 4	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1317.075

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
o1-2024-12-17	Scheduled 23 Oct 2026	LMArena · 23 Apr 2025 · rank 10 · rating 1349.539	gemma-3-27b-it · 1342.085 · difference 7.454 rating	DeepSeek-V3-0324 rank 5 , DeepSeek-R1 rank 7	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1365.838
o3-mini-2025-01-31	Scheduled 23 Oct 2026	LMArena · higher reasoning setting · 11 May 2025 · rank 16 · rating 1324.614	gemma-3-12b-it · 1320.905 · difference 3.709 rating	DeepSeek-V3-0324 rank 5 , DeepSeek-R1 rank 7	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1336.609
o1-pro-2025-03-19	Scheduled 23 Oct 2026	Scale HLE Text Only · o1 Pro (Mar 2025) 8.44 ± 1.12	DeepSeek R1 · 8.57 · difference 0.13 points	DeepSeek-R1 8.57 ± 1.13	Scale HLE Text Only · DeepSeek-R1 8.57 ± 1.13 vs target 8.44 ± 1.12
gpt-4.1-nano-2025-04-14	Scheduled 23 Oct 2026	LMArena · 1 Jun 2025 · rank 52 · rating 1277.691	llama-3.1-405b-instruct-bf16 · 1275.486 · difference 2.205 rating	DeepSeek-V3-0324 rank 6 , DeepSeek-R1 rank 8	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1284.831
o3-2025-04-16	Scheduled 11 Dec 2026	LMArena · 16 Nov 2025 · rank 33 · rating 1410.948	qwen3-235b-a22b-thinking-2507 · 1413.405 · difference 2.457 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1409.945
o4-mini-2025-04-16	Scheduled 23 Oct 2026	LMArena · 5 Nov 2025 · rank 66 · rating 1353.943	qwen3-coder-480b-a35b-instruct · 1356.411 · difference 2.468 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1353.238

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
o3-pro-2025-06-10	Scheduled 11 Dec 2026	PM-LLM-Benchmark · 39.9	Qwen3-235B-A22B-Thinking-2507 · 45.4 · difference 5.5 points	Qwen3-235B-A22B-Thinking-2507 45.4	PM-LLM-Benchmark · Qwen3-235B-A22B-Thinking-2507 45.4 vs target 39.9
gpt-5-mini-2025-08-07	Scheduled 11 Dec 2026	LMArena · higher reasoning setting · 16 Nov 2025 · rank 53 · rating 1375.485	DeepSeek-V3-0324 · 1374.120 · difference 1.365 rating	GLM-4.6 rank 4 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1373.266
gpt-5-nano-2025-08-07	Scheduled 11 Dec 2026	LMArena · higher reasoning setting · 17 Oct 2025 · rank 91 · rating 1319.185	llama-3.1-nemotron-ultra-253b-v1 · 1319.649 · difference 0.464 rating	GLM-4.6 rank 3 , GLM-4.5 rank 6	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1320.066
gpt-5-2025-08-07	Scheduled 11 Dec 2026	LMArena · higher reasoning setting · 17 Nov 2025 · rank 37 · rating 1407.785	qwen3-235b-a22b-thinking-2507 · 1413.424 · difference 5.639 rating	GLM-4.6 rank 5 , DeepSeek-V3.2-Exp rank 9	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1406.077
gpt-5-pro-2025-10-06	Scheduled 11 Dec 2026	PM-LLM-Benchmark v2.2 · 44.9	GLM-5 · 46.0 · difference 1.1 points	DeepSeek-V3.2-Speciale-Thinking 48.7 , Kimi K2.5 46.1 , GLM-5 46.0	PM-LLM-Benchmark v2.2 · DeepSeek-V3.2-Speciale 48.7 vs target 44.9
Google					
gemini-2.0-flash-001	Retired 1 Jun 2026	LMArena · 5 Nov 2025 · rank 65 · rating 1353.985	qwen3-coder-480b-a35b-instruct · 1356.411 · difference 2.426 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1354.124

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
gemini-2.0-flash-lite-preview-02-05	Retired 9 Dec 2025	LMArena · 1 Jun 2025 · rank 25 · rating 1319.370	QwQ-32B · 1321.351 · difference 1.981 rating	DeepSeek-V3-0324 rank 6 , DeepSeek-R1 rank 8	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1329.594
gemini-2.0-flash-lite-001	Retired 1 Jun 2026	LiveBench · Gemini 53.0 on 418 shared objective questions On the 318 questions shared with Qwen2.5-72B-Instruct-Turbo, Gemini 53.041	Qwen2.5-72B-Instruct-Turbo · 52.717 on the same 318 questions · difference 0.324 points	DeepSeek-V3-0324 68.5 , DeepSeek-R1 66.7	LiveBench · DeepSeek-V3-0324 68.5 vs target 53.0 on 418 shared objective questions
gemini-2.5-pro-exp-03-25	Retired 26 Jun 2025	LiveBench · leaderboard of 25 Apr 2025, captured 29 May 2025 · Coding 71.08	DeepSeek-R1 · 74.98 · difference 3.90 points	DeepSeek-R1 Coding 74.98 , Data Analysis 69.62 vs 62.47 , Qwen3-235B-A22B-Thinking Data Analysis 68.31 vs 62.47	Aider · DeepSeek-V3.2-Exp (Reasoner) 74.2% vs target 72.9%
gemini-2.5-flash-preview-05-20	Retired 18 Nov 2025	Aider · 25 May 2025 · 55.1%	DeepSeek-V3-0324 · 55.1% · difference 0 pp	Qwen3-235B-A22B 59.6%	Aider · Qwen3-235B-A22B 59.6% vs target 55.1%
gemini-2.5-pro-preview-06-05	Retired 2 Dec 2025	LiveBench · leaderboard of 30 May 2025, captured 13 Jun 2025 · Agentic Coding of -highthinking 20.00	DeepSeek-R1-0528 · 26.67 · difference 6.67 points	DeepSeek-R1-0528 26.67	LiveBench · DeepSeek V4 Pro 74.39 global average vs target 57.46

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
gemini-2.5-flash-lite-preview-09-2025 Anthropic	Retired 31 Mar 2026	LMArena · 19 Mar 2026 · rank 99 · rating 1378.55	Kimi K2 0905 preview · 1378.79 rating · difference 0.24 rating	Kimi K2 0905 · rating 1378.79	LiveBench · DeepSeek V4 Pro 74.39 global average vs target 41.47
claude-2.0	Retired 21 Jul 2025	HELM Lite v1.0.0 · 20 Dec 2023 · mean win rate 0.679	Llama 2 70B · 0.659 mean win rate · difference 0.020 win-rate points	Yi 34B 0.772 , Mixtral 8x7B 0.728	HELM Lite v1.0.0 · Yi 34B 0.772 mean win rate vs target 0.679
claude-instant-1.2	Retired 6 Nov 2024	HELM Lite v1.0.0 · 20 Dec 2023 · mean win rate 0.562	Llama 2 70B · 0.659 mean win rate · difference 0.097 win-rate points	Yi 34B 0.772 , Mixtral 8x7B 0.728	HELM Lite v1.0.0 · Yi 34B 0.772 mean win rate vs target 0.562
claude-2.1	Retired 21 Jul 2025	HELM Lite v1.0.0 · 20 Dec 2023 · mean win rate 0.593	Llama 2 70B · 0.659 mean win rate · difference 0.066 win-rate points	Yi 34B 0.772 , Mixtral 8x7B 0.728	HELM Lite v1.0.0 · Yi 34B 0.772 mean win rate vs target 0.593
claude-3-opus-20240229	Retired 5 Jan 2026	LMArena · 19 Jan 2025 · rank 34 · rating 1247.301	llama-3.1-70b-instruct · 1247.816 · difference 0.515 rating	DeepSeek-V3 rank 3 , Llama 3.1 Nemotron 70B rank 17	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1262.005
claude-3-sonnet-20240229	Retired 21 Jul 2025	LMArena · 23 Jun 2024 · rank 13 · rating 1198.612	llama-3-70b-instruct · 1202.021 · difference 3.409 rating	Llama 3 70B Instruct rank 12	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1218.041

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
claude-3-haiku-20240307	Retired 20 Apr 2026	LMArena · 4 Nov 2024 · rank 47 · rating 1179.281	command-r-08-2024 · 1179.518 · difference 0.237 rating	Llama 3.1 Nemotron 70B rank 11 , Qwen2.5 72B Instruct rank 20	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1194.682
claude-3-5-sonnet-20240620	Retired 28 Oct 2025	LMArena · 21 Nov 2024 · rank 12 · rating 1267.980	llama-3.1-nemotron-70b-instruct · 1268.797 · difference 0.817 rating	Llama 3.1 Nemotron 70B rank 11	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1280.921
claude-3-5-sonnet-20241022	Retired 28 Oct 2025	LMArena · 28 Aug 2025 · rank 80 · rating 1299.697	deepseek-v2.5-1210 · 1296.092 · difference 3.605 rating	GLM-4.5 rank 5 , DeepSeek-R1-0528 rank 6	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1297.772
claude-3-5-haiku-20241022	Retired 19 Feb 2026	LMArena · 17 Oct 2025 · rank 140 · rating 1256.241	llama-3.1-tulu-3-70b · 1255.724 · difference 0.517 rating	GLM-4.6 rank 3 , GLM-4.5 rank 6	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1255.358
claude-3-7-sonnet-20250219	Retired 19 Feb 2026	LMArena · extended thinking · 5 Nov 2025 · rank 96 · rating 1315.106	qwen3-30b-a3b · 1317.847 · difference 2.741 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1314.381
claude-opus-4-20250514	Retired 15 Jun 2026	LMArena · extended thinking · 5 Nov 2025 · rank 47 · rating 1374.696	DeepSeek-V3-0324 · 1374.024 · difference 0.672 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1376.941

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
claude-sonnet-4-20250514	Retired 15 Jun 2026	LMArena · extended thinking · 5 Nov 2025 · rank 68 · rating 1349.162	step-3 · 1348.834 · difference 0.328 rating	GLM-4.6 rank 2 , GLM-4.5 rank 5	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1350.663
claude-opus-4-1-20250805	Retired 5 Aug 2026	LMArena · extended thinking · 18 Mar 2026 · rank 56 · rating 1417.863	qwen3.5-122b-a10b · 1417.899 · difference 0.036 rating	Kimi K2.5 Thinking rank 14 , Qwen3.5-397B-A17B rank 16	LMArena · 25 Sep 2026 · MiMo V2.6 Pro 1490.734 overall vs target 1419.151
Image, video, and audio models					
OpenAI					
gpt-4o-audio-preview-2024-10-01	Retired 10 Oct 2025	AudioMarathon · 8 Oct 2025 · ten-task mean 47.4	Phi-4-Multimodal · 47.7 ten-task mean · difference 0.3 points	Phi-4-Multimodal 47.7 , Voxtral Mini 3B 57.4	AudioMarathon · Voxtral Mini 3B 57.4 ten-task mean vs target 47.4
gpt-4o-realtime-preview-2024-10-01	Retired 10 Oct 2025	UltraEval-Audio · understanding evaluated 23 Dec 2024 · ASR 92 , AST 26 · generation evaluated 29 Oct 2024 · AudioArena 1200	Understanding: Qwen2-Audio-Instruct · ASR 94 , AST 31 · Generation: GLM-4-Voice · AudioArena 1035	Understanding: Qwen2-Audio-Instruct · ASR 94 , AST 31 · Generation: None in this evaluation	AHELM · Qwen2.5-Omni 7B 0.495 mean win rate vs target 0.439

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
gpt-4o-realtime-preview-2024-12-17	Retired 7 May 2026	AHELM · Reasoning mean win rate 0.344	Qwen2-Audio Instruct 7B · 0.375 · difference 0.031	Qwen2-Audio Instruct 7B · 0.375	AudioMarathon · Qwen2.5-Omni-7B 70.5 ten-task mean vs target 48.7
gpt-4o-mini-transcribe-2025-03-20	Scheduled 20 Jan 2027	Artificial Analysis WER benchmark · 20 Mar 2025 · 13.2% WER	Whisper large-v3 via Fireworks · 11.2% WER · difference 2.0 pp	Whisper large-v3 via Fireworks · 11.2% WER · Whisper large-v2 · 10.6% WER	Artificial Analysis WER benchmark · Whisper large-v3 11.2% WER vs target 13.2% , lower is better
gpt-4o-realtime-preview-2025-06-03	Retired 7 May 2026	Scale AudioMC · text output · 25.44 ± 4.00	Voxtral-Small-24B-2507 · 26.33 · difference 0.89 points	Voxtral-Small-24B-2507 26.33 ± 4.05	Scale AudioMC · Inkling (Thinking) 56.64 APR vs target 25.44
gpt-realtime-2025-08-28	Scheduled 20 Jan 2027	Scale AudioMC · text output · 23.45 ± 3.90	Voxtral-Small-24B-2507 · 26.33 · difference 2.88 points	Voxtral-Small-24B-2507 26.33 ± 4.05	Scale AudioMC · Inkling (Thinking) 56.64 APR vs target 23.45
sora-2-pro-2025-10-06	Retired 24 Sep 2026	LMArena Text-to-Video · 4 Sep 2026 · rank 11 · rating 1366.55	MiniMax H3 · 1461.78 rating · difference 95.23 rating	MiniMax H3 rank 8 · rating 1461.78	LMArena Text-to-Video · MiniMax H3 1461.78 overall vs target 1366.55
sora-2-2025-10-06	Retired 24 Sep 2026	RigidBench · ICLR 2026 Workshop World Models, published 2 Mar 2026 · SSIM 0.602	Wan 2.2 TI2V-5B · 0.841 · difference 0.239	Wan 2.2 TI2V-5B SSIM 0.841 , IoU 0.234 vs 0.206	RigidBench · Wan 2.2 TI2V-5B 0.841 SSIM / 0.234 IoU vs target 0.602 / 0.206

Table 5.3-3 · Part II. API Versions Compared with Models with Publicly Available Weights (continued) · See Section 5.3.3

Model or version	Retirement status	Evaluation before the retirement date	Nearest result among models with publicly available weights	At least one reported metric exceeded before the retirement date	Higher result in a broader common evaluation by September 2026
sora-2-2025-12-08	Retired 24 Sep 2026	LMArena Text-to-Video · 4 Sep 2026 · rank 18 · rating 1341.88	MiniMax H3 · 1461.78 rating · difference 119.90 rating	MiniMax H3 rank 8 · rating 1461.78	LMArena Text-to-Video · MiniMax H3 1461.78 overall vs target 1341.88
Google					
imagen-3.0-generate-002	Retired 10 Nov 2025	LMArena Text-to-Image · 9 Nov 2025 · rank 16 · rating 1062.045	Qwen Image · 1062.576 rating · difference 0.531 rating	HunyuanImage 3.0 rank 1 , Qwen Image rank 15	LMArena Text-to-Image · HunyuanImage 3.0 1149.51 overall vs target 1058.02
imagen-4.0-ultra-generate-preview-06-06	Retired 17 Feb 2026	UniGenBench English-long · Style 97.67	FLUX.2 dev · 99.17 · difference 1.50 points	FLUX.2 dev Style 99.17 , Text 88.32 vs 86.41 , Attribute Overall 94.57 vs 93.21 , Expression 88.47 vs 83.17	ImageBench V1.2 · Boogu Image Turbo 64.3 Overall vs target 62.0
imagen-4.0-generate-preview-06-06	Retired 17 Feb 2026	UniGenBench English-long · overall 85.34	FLUX.2 dev · 90.31 · difference 4.97 points	FLUX.2 dev 90.31	UniGenBench English-long · FLUX.2 dev 90.31 overall vs target 85.34
veo-2.0-generate-001	Retired 30 Jun 2026	LMArena Image-to-Video · 6 Nov 2025 · rank 20 · rating 1164.131	Wan2.2 I2V A14B · rating 1167.288 · difference 3.157	Wan2.2 I2V A14B rank 19 · rating 1167.288	LMArena Image-to-Video · Wan2.2 I2V A14B 1170.146 overall vs target 1165.741

Definitions

Term	Definition
deployment fingerprint	An immutable identifier binding a deployment to its defining model artifacts and system state.
deployment	A specific state of an AI system as configured for service. A change to a component or configuration that defines this state creates a new deployment. A change only to who has access, the terms of access, or the period during which access is available does not create a new deployment.
access class	A group defined by the conditions under which it has access to a deployment.

Acknowledgments

The authors thank the following contributors for their review and feedback:

- **A.J.** for review of accessibility and continued access after release
- **C.K.** for review of policy design and legal framing
- **C.L.V.** for review of provider investment and the scope of resource and contribution arguments
- **J.L.** for review of selected legal citations, technical claims, and empirical data
- **M.J.** for review of preservation, lifecycle transparency, provider control, and competition
- **R.X.** for review of independent continuation and institutional design
- **X.L.** for review of model dependencies, migration, and supporting evidence
- **Y.W.** for review of argument structure, legal distinctions, and the treatment of the model as the product

The authors remain responsible for the Whitepaper's analysis and conclusions.